

Chapter 2

Instructions: Language of the Computer

Instruction Set

- The repertoire of instructions of a computer
- Different computers have different instruction sets
 - But with many aspects in common
- Early computers had very simple instruction sets
 - Simplified implementation
- Many modern computers also have simple instruction sets

The MIPS Instruction Set

- Used as the example throughout the book
- Stanford MIPS commercialized by MIPS Technologies (www.mips.com)
- Large share of embedded core market
 - Applications in consumer electronics, network/storage equipment, cameras, printers, ...
- Typical of many modern ISAs
 - See MIPS Reference Data, and Appendixes B and E

MIPS Instruction Set

■ MIPS Instruction Types

- Arithmetic
- Logical operation
- Data transfer
- Flow control

■ Design Principles

- Design Principle 1: Simplicity favors regularity.
- Design Principle 2: Smaller is faster.
- Design Principle 3: Good design demands a compromise
- Design Principle 4: Make common case fast.

Arithmetic Operations

- Add and subtract, three operands
 - Two sources and one destination
add a, b, c # a gets b + c
- All arithmetic operations have this form
 - Long expressions are broken into small ops
 - (e.g., how to calculate $a = b + c + d + e$?)
- **Design Principle 1: Simplicity favors regularity**
 - Regularity makes implementation simpler
 - Simplicity enables higher performance at lower cost



Arithmetic Example

- C code:

$f = (g + h) - (i + j);$


- Compiled MIPS code:

```
add t0, g, h    # temp t0 = g + h
add t1, i, j    # temp t1 = i + j
sub f, t0, t1   # f = t0 - t1
```

- Break down one-line of C into several basic arithmetic operations
 - Each handles only two inputs and one output
 - **t0** and **t1** are temporary variables normally created by the compiler
- Program cannot access variables in memory directly
 - Need to load from memory to CPU registers

Register Operands

- Arithmetic instructions use register operands
- MIPS has a 32×32 -bit register file
 - Use for frequently accessed data
 - Numbered 0 to 31
 - 32-bit data called a “word”



- Why not use memory operands?
 - **Design Principle 2: Smaller is faster**
 - More registers → complex hardware → slower clocks
 - More registers → more bits to represent in ISA format
 - c.f. main memory: millions of locations

Register Operand Example

- C code:

$f = (g + h) - (i + j);$

- g, \dots, j in $\$s1, \dots, \$s4$

← moved from memory to registers by data transfer instructions.

- Compiled MIPS code:

add \$t0, \$s1, \$s2

add \$t1, \$s3, \$s4

sub \$s0, \$t0, \$t1

Assembler names

\$t0, \$t1, ..., \$t9 for temporary values

\$s0, \$s1, ..., \$s7 for saved variables


Memory Operands

- Main memory used for composite data
 - Arrays, structures, dynamic data
- To apply arithmetic operations
 - Load values from memory into registers
 - Store result from register to memory
- Memory is byte addressed
 - Each address identifies an 8-bit byte
- Words are **aligned** in memory
 - Address must be a multiple of 4
- MIPS is **Big Endian**
 - Most-significant byte at least address of a word
 - *c.f.* Little Endian: least-significant byte at least address

Memory Operands

Words alignment

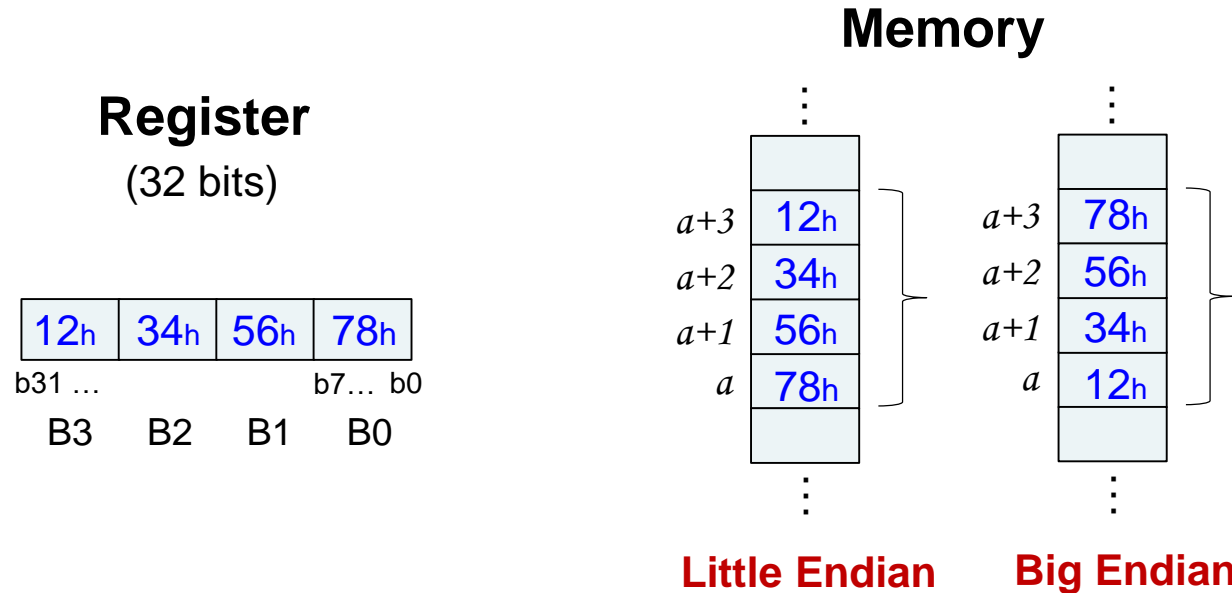
- Address must be a multiple of 4



Memory is
byte addressed

Memory Operands

- MIPS is **Big-Endian**
 - Most significant byte at least address of a word
 - c.f. **Little-Endian**: least significant byte at least address
 - Ex. To store a word (32 bits) : 0x12345678



Data Transfer Operations

- C code:

`g = h + A[8];` // g, h, A[] are type of int (4 bytes)

- g in \$s1, h in \$s2, base address of A in \$s3


- Compiled MIPS code:

- Index 8 requires offset of 32
 - 4 bytes per word

```
lw    $t0, 32($s3) #load word
add   $s1, $s2, $t0
```

offset

base register



Data Transfer Operations

- C code:

`A[12] = h + A[8];`

- `h` in `$s2`, base address of `A` in `$s3`

- Compiled MIPS code:

- Index 8 means offset of 32
- Index 12 means offset of 48

```
lw    $t0, 32($s3)    # load word
add   $t0, $s2, $t0
sw    $t0, 48($s3)    # store word
```

Registers vs. Memory

- Registers are faster to access than memory
- Operating on memory data requires loads and stores
 - More instructions to be executed
- Compiler must use registers for variables as much as possible
 - Only spill to memory for less frequently used variables
 - Register optimization is important!

Immediate Operands

- Constant data specified in an instruction
`addi $s3, $s3, 4`
- No subtract immediate instruction
 - Just use a negative constant
`addi $s2, $s1, -1`
- **Design Principle 3: Make the common case fast**
 - Small constants are common
 - Immediate operand avoids a load instruction

The Constant Zero

- MIPS register 0 (\$zero) is the constant 0
 - The number, 0, is so common in code that MIPS hardwired register 0 (\$0) to 0.
 - Cannot be overwritten
 - e.g., `addi $0, $0, 5` # no effect
- Also useful for other common operations
 - e.g., move between registers
`add $t2, $s1, $zero`
 - So, there is no need to implement a move instruction in hardware → simpler CPU design and implementation

Unsigned Binary Integers

- Given an n-bit number

$$x = x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \dots + x_12^1 + x_02^0$$

- Range: 0 to $+2^n - 1$

- Using 32 bits

- 0 to +4,294,967,295

- Example

- $0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 1011_2$
= $0 + \dots + 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0$
= $0 + \dots + 8 + 0 + 2 + 1 = 11_{10}$

2's-Complement Signed Integers

- Given an n-bit number

$$x = -x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \dots + x_12^1 + x_02^0$$

- Range: -2^{n-1} to $+2^{n-1} - 1$
- Using 32 bits (n=32)
 - $-2,147,483,648$ to $+2,147,483,647$
- Example

$$\begin{aligned} & \text{■ } 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1100_2 \\ & = -1 \times 2^{31} + 1 \times 2^{30} + \dots + 1 \times 2^2 + 0 \times 2^1 + 0 \times 2^0 \\ & = -2,147,483,648 + 2,147,483,644 = -4_{10} \end{aligned}$$

Signed v.s. Unsigned

4 bits

3 bits

(10-based)	(Binary)	(10-based)
0	0 0 0	0
1	0 0 1	1
2	0 1 0	2
3	0 1 1	3
-4	1 0 0	4
-3	1 0 1	5
-2	1 1 0	6
-1	1 1 1	7


**Signed binary
(2's complement)**

Unsigned binary


(10-based)

(Binary)

(10-based)

0	0 0 0 0	0
1	0 0 0 1	1
2	0 0 1 0	2
3	0 0 1 1	3
4	0 1 0 0	4
5	0 1 0 1	5
6	0 1 1 0	6
7	0 1 1 1	7
-8	1 0 0 0	8
-7	1 0 0 1	9
-6	1 0 1 0	10
-5	1 0 1 1	11
-4	1 1 0 0	12
-3	1 1 0 1	13
-2	1 1 1 0	14
-1	1 1 1 1	15


**Signed binary
(2's complement)**

Unsigned binary


2's-Complement Signed Integers

- Bit 31 is sign bit
 - 1 for negative numbers
 - 0 for non-negative numbers

- Some specific numbers

- Most-negative: 1000 0000 ... 0000
- Most-positive: 0111 1111 ... 1111
- 0: 0000 0000 ... 0000
- -1: 1111 1111 ... 1111

- Positive numbers have the same unsigned and 2s-complement representation

How to represent -17 in binary for a 8-bit signed integer?

-2^7 1 1 1 0 1 1 1 1

$$-128 + ? = -17$$

$$? = 111 = 1101111_2$$

$$\begin{array}{r} 011111111111111 \\ + 1 \\ \hline 100000000000000 \end{array}$$

Signed Negation

- Negate a signed integer: *Complement and add 1*
 - **Complement** means $1 \rightarrow 0, 0 \rightarrow 1$

$$x + \bar{x} = 1111 \dots 111_2 = -1$$
$$-x = \bar{x} + 1$$

Ex3

How to represent -11 in 5-bit binary?

+11 binary \rightarrow **01011₂**

Negation \rightarrow **10101₂**

- Ex1: To negate +2
 - $+2 = 0000 \ 0000 \ \dots \ 0010_2$
 - $-2 = 1111 \ 1111 \ \dots \ 1101_2 + 1$
 $= 1111 \ 1111 \ \dots \ 1110_2$
- Ex2: To negate -2

Sign Extension

- Purpose: Representing a number using more bits, but preserve the numeric value
- How: Replicate the sign bit to the left
 - c.f. unsigned values: extend with 0s
 - Examples: 8-bit to 16-bit
 - +2: 0000 0010 => 0000 0000 0000 0010
 - -2: 1111 1110 => 1111 1111 1111 1110
- In MIPS instruction set, sign extension is needed.
 - Example:
 - addi: extend immediate value
 - lb, lh: extend loaded byte/halfword
 - beq, bne: extend the displacement

Representing Instructions

- Instructions are encoded in binary
 - Called machine code
- MIPS instructions
 - Encoded as 32-bit instruction words
 - Small number of formats encoding operation code (opcode), register numbers, ...
 - Regularity!
- Register numbers
 - \$t0 – \$t7 are reg's 8 – 15
 - \$t8 – \$t9 are reg's 24 – 25
 - \$s0 – \$s7 are reg's 16 – 23



MIPS R-format Instructions



■ Instruction fields

- op: operation code (opcode)
- rs: first source register number
- rt: second source register number
- rd: destination register number
- shamt: shift amount (00000 for now)
- funct: function code (extends opcode)

R-format Example

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

add \$t0, \$s1, \$s2

special	\$s1	\$s2	\$t0	0	add
---------	------	------	------	---	-----

0	17	18	8	0	32
---	----	----	---	---	----

000000	10001	10010	01000	00000	100000
--------	-------	-------	-------	-------	--------

$00000010001100100100000000100000_2 = 02324020_{16}$

Hexadecimal

- Base 16
 - Compact representation of bit strings
 - 4 bits per hex digit

0	0000	4	0100	8	1000	c	1100
1	0001	5	0101	9	1001	d	1101
2	0010	6	0110	a	1010	e	1110
3	0011	7	0111	b	1011	f	1111

- Example: eca8 6420 ...h, 0x..
 - 1110 1100 1010 1000 0110 0100 0010 0000

MIPS I-format Instructions



- Immediate arithmetic and load/store instructions
 - rt: destination reg (I-type) or source reg (R-type)
 - Constant: -2^{15} to $+2^{15} - 1$
 - Address: **offset** added to **base address in rs**
- **Design Principle 4: Good design demands good compromises**
 - Different formats complicate decoding, but allow 32-bit instructions uniformly
 - Keep formats as similar as possible

R-format v.s. I-format



Example:

C code : `A[300] = h + A[300];`

Assembly code: `lw $t0, 1200($t1) rs`

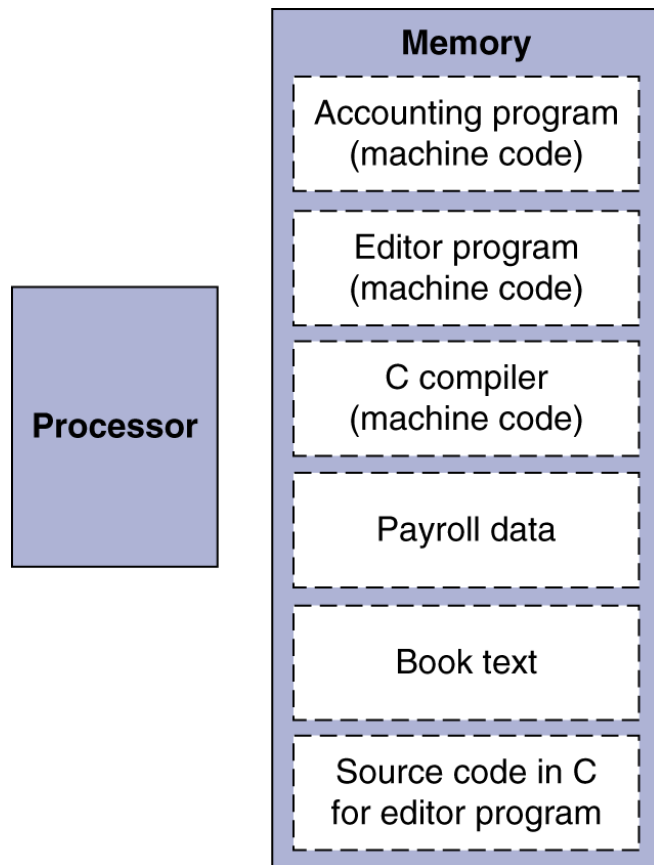
`add $t0, $s2, $t0`

`sw $t0, 1200($t1) rs`

Machine code:

op	rs	rt	rd	Addr/ shamt	funct
35	9	8	1200		
0	18	8	8	0	32
45	9	8	1200		

Stored Program Computers



- Instructions represented in binary, just like data
- Instructions and data are stored in memory
- Programs can operate on programs
 - e.g., compilers, linkers, ...
- Binary compatibility allows compiled programs to work on different computers
 - Standardized ISAs

Logical Operations

■ Instructions for bitwise manipulation

Operation	C	Java	MIPS
Shift left	<<	<<	sll
Shift right	>>	>>	srl
Bitwise AND	&	&	and, andi
Bitwise OR			or, ori
Bitwise NOT	~	~	nor

= NOT OR

■ Useful for extracting and inserting groups of bits in a word

If one operand of NOR is zero, then it is equivalent to NOT

Shift Operations

R-type:

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

- shamt: how many positions to shift
- Shift left logical (sll)
 - Shift left and fill with 0 bits
 - sll by i bits multiplies by 2^i
- Shift right logical (srl)
 - Shift right and fill with 0 bits
 - srl by i bits divides by 2^i (unsigned only)

- sll \$t2, \$s0, 4
 - srl \$t2, \$s0, 6
- rt

op	rs	rt	rd	shamt	funct
0	0	16	10	4	0
0	0	16	10	6	2



AND Operations

- Useful to mask bits in a word
 - Select some bits, clear others to 0
- and \$t0, \$t1, \$t2

\$t2	0000 0000 0000 0000 0000 1101 1100 0000
\$t1	0000 0000 0000 0000 0011 1100 0000 0000
\$t0	0000 0000 0000 0000 0000 1100 0000 0000

OR Operations

- Useful to include bits in a word
 - Set some bits to 1, leave others unchanged
- or \$t0, \$t1, \$t2

\$t2 0000 0000 0000 0000 0000 1101 1100 0000

\$t1 0000 0000 0000 0000 0011 1100 0000 0000

\$t0 0000 0000 0000 0000 0011 1101 1100 0000

NOT Operations

- Useful to invert bits in a word
 - Change 0 to 1, and 1 to 0
- MIPS has NOR 3-operand instruction
 - $a \text{ NOR } b == \text{NOT} (a \text{ OR } b)$

`nor $t0, $t1, $zero` ←
(use NOR to perform NOT)

Register 0: always
read as zero

\$t1 0000 0000 0000 0000 0011 1100 0000 0000


\$t0 1111 1111 1111 1111 1100 0011 1111 1111

Category	Instr	Example	Meaning	Comments
Logical	and	and \$s1, \$s2, \$s3	\$s1 = \$s2 & \$s3	bit-by-bit AND
	or	or \$s1, \$s2, \$s3	\$s1 = \$s2 \$s3	bit-by-bit OR
	nor	nor \$s1, \$s2, \$s3	\$s1 = ~(\$s2 \$s3)	bit-by-bit NOR
	andi	andi \$s1, \$s2, 100	\$s1 = \$s2 & 100	AND with const
	ori	ori \$s1, \$s2, 100	\$s1 = \$s2 100	OR with const.
	sll	sll \$s1, \$s2, 10	\$s1 = \$s2 << 10	shift left by const
	srl	srl \$s1, \$s2, 10	\$s1 = \$s2 >> 10	shift right by const



Assembly code
MIPS
Logical Instruction

Machine code



Name	Format	op	rs	rt	rd	shamt	func
and	R	0	18	19	17	0	36
or	R	0	18	19	17	0	37
nor	R	0	18	19	17	0	39
andi	I	12	18	17	100		
ori	I	13	18	17	100		
sll	R	0	0	18	17	10	0
srl	R	0	0	18	17	10	2

Flow Control Operations

- Branch to a labeled instruction if a condition is true
 - Otherwise, continue sequentially
- `beq rs, rt, L1`
 - if (`rs == rt`) branch to instruction labeled L1;
- `bne rs, rt, L1`
 - if (`rs != rt`) branch to instruction labeled L1;
- `j L1`
 - unconditional jump to instruction labeled L1



Compiling If Statements

- C code:

```
if (i==j) f = g+h;  
else f = g-h;
```

- f, g, ... in \$s0, \$s1, ...

- Compiled MIPS code:


```
bne $s3, $s4, Else  
add $s0, $s1, $s2
```

```
j Exit
```

```
Else: sub $s0, $s1, $s2
```

```
Exit: ...
```

Assembler calculates addresses



Compiling Loop Statements

- C code:

```
while (save[i] == k) i += 1;
```

- i in \$s3, k in \$s5, address of save in \$s6

- MIPS code:

```
Loop:  sll    $t1, $s3, 2
        add   $t1, $t1, $s6
        lw    $t0, 0($t1)
        bne   $t0, $s5, Exit
        addi   $s3, $s3, 1
        j     Loop
Exit:  ...
```

000000**1**0₂

00000**1**00₂

x2

0000**1**000₂

x4

00000**1**10₂

0000**1**100₂


x2

000**1**1000₂

x4

Basic Blocks

- A basic block is a sequence of instructions with
 - No embedded branches (except at end)
 - No branch targets (except at beginning)



- A compiler identifies basic blocks for optimization
- An advanced processor can accelerate execution of basic blocks

Critical thinking...

- We have used beq/bne for structures below

```
if (i==j) f = g+h;  
else f = g-h;
```

```
while (save[i] == k) i += 1;
```

- How about following structures?

```
if (i<j) f = g+h;  
else f = g-h;
```

```
while (save[i] >= k) i += 1;
```


More Conditional Operations

- Set result to 1 if a condition is true
 - Otherwise, set to 0
- `slt rd, rs, rt`
 - if ($rs < rt$) $rd = 1$; else $rd = 0$;
- `slti rt, rs, constant`
 - if ($rs < \text{constant}$) $rt = 1$; else $rt = 0$;
- Use in combination with `beq`, `bne`

```
slt $t0, $s1, $s2  
bne $t0, $zero, L
```

if ($\$s1 < \$s2$)
branch to L

(b1t)

Branch Instruction Design

- Why not `blt`, `bgt`, etc?
- Hardware for `<`, `≥`, ... slower than `=`, `≠`
 - Combining with branch involves more work per instruction, requiring a slower clock
 - All instructions penalized!
- `beq` and `bne` are the common case
- This is a good design compromise

Critical thinking ...

```
slt $t0, $s1, $s2  
bne $t0, $zero, L
```

if ($\$s1 < \$s2$)
branch to L
(b1t)

?

if ($\$s1 > \$s2$)
branch to L
(bgt)

?

if ($\$s1 \leq \$s2$)
branch to L
(ble)

?

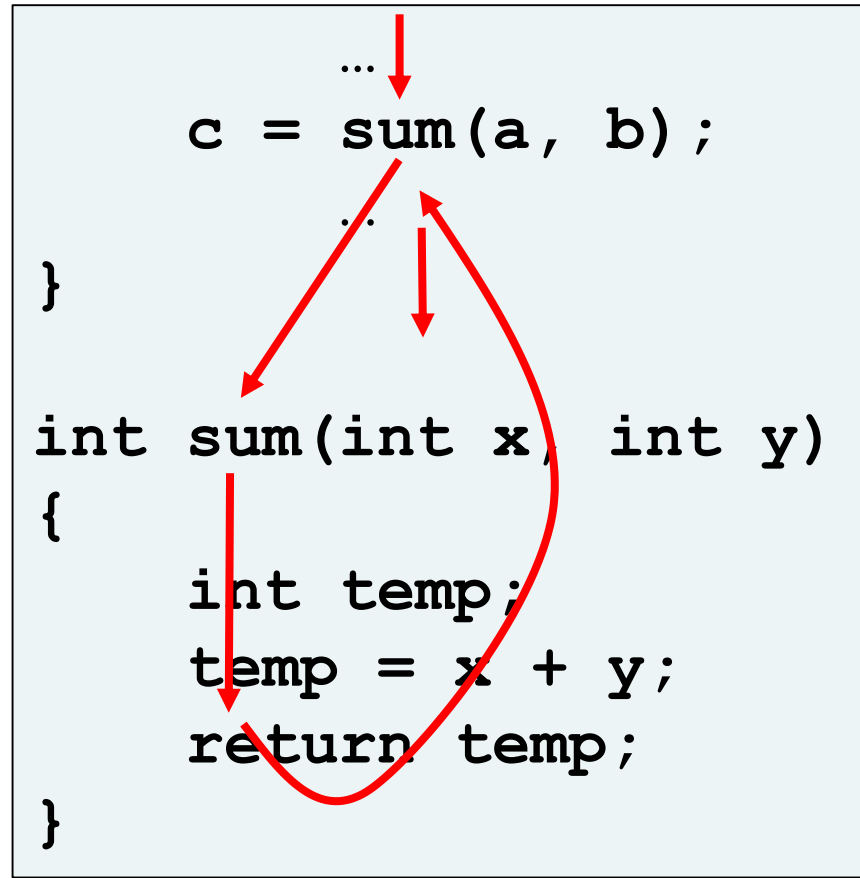
if ($\$s1 \geq \$s2$)
branch to L
(bge)

Signed vs. Unsigned

- Signed comparison: `slt`, `slti`
- Unsigned comparison: `sltu`, `sltui`
- Example
 - `$s0 = 1111 1111 1111 1111 1111 1111 1111 1111`
 - `$s1 = 0000 0000 0000 0000 0000 0000 0000 0001`
 - `slt $t0, $s0, $s1 #signed`
 - $-1 < +1 \Rightarrow \text{\textcolor{blue}{\$t0}} = 1$
 - `sltu $t0, $s0, $s1 #unsigned`
 - $+4,294,967,295 > +1 \Rightarrow \text{\textcolor{blue}{\$t0}} = 0$

Procedure call in HLL

- Procedure/function call changes control flow



Must specify:

- Procedure address
- Arguments (or parameters)
- Local variables
- Return value
- Return address

**How to implement it
in machine code?**

Procedure Calling

Caller:

1. Place arguments in **registers**
2. Jump to beginning of procedure, save next PC

Callee:

3. Acquire storage on program stack
(for *arguments*, *return address* *local variables*, reg backup...)
4. Perform procedure's operations
5. Place result in **register** for caller
6. Return to place of call at caller

Caller:

7. Get return result from **register**

Caller and callee use registers to communicate as much as possible. If registers are insufficient, stack is used.

Register Usage

- \$a0 – \$a3: arguments (reg's 4 – 7)
- \$v0, \$v1: result values (reg's 2 and 3)
- \$t0 – \$t9: temporaries
 - Can be overwritten by callee
- \$s0 – \$s7: saved
 - Must be saved/restored by callee
- \$gp: global pointer for static data (reg 28)
- \$sp: stack pointer (reg 29)
- \$fp: frame pointer (reg 30)
- \$ra: return address (reg 31)

Register 1 (\$at) reserved for assembler, 26-27 for operating system

Procedure Call Instructions

- **Procedure call:** jump and link

`jal ProcedureLabel`

- Address of following instruction put in \$ra
- Jumps to target address

- **Procedure return:** jump register

`jr $ra`

- Copies \$ra to program counter
- Can also be used for computed jumps
 - e.g., for case/switch statements

```
...  
jal Proc1  
xxx
```

Proc1:

```
...  
jr $ra
```


Leaf Procedure Example

■ C code:

```
int leaf_example (int g, h, i, j)
{ int f;
  f = (g + h) - (i + j);
  return f;
}
```

- Arguments g, ..., j in \$a0, ..., \$a3
- f in \$s0 (hence, need to save \$s0 on stack)
- Result in \$v0

Leaf Procedure

- Local/global variables
- Stack's growth direction

■ MIPS code:

leaf_example:

```
addi $sp, $sp, -4  
sw   $s0, 0($sp)
```

Save \$s0
on stack

```
add   $t0, $a0, $a1  
add   $t1, $a2, $a3  
sub   $s0, $t0, $t1  
add   $v0, $s0, $zero
```

Procedure
body

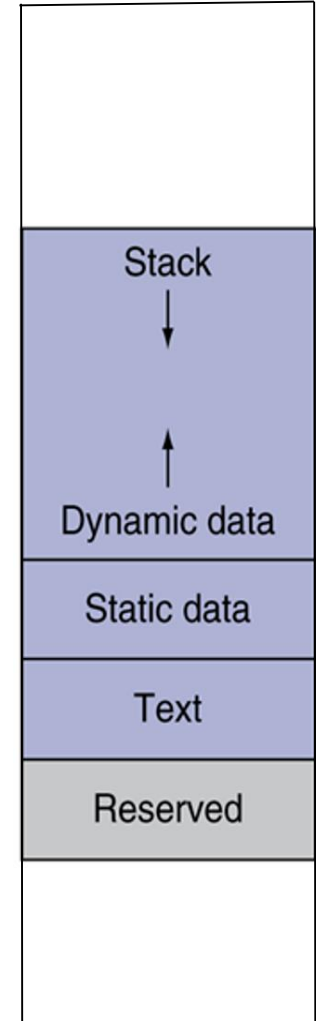
```
lw    $s0, 0($sp)  
addi  $sp, $sp, 4
```

Restore \$s0

```
jr    $ra
```

Return



$2^{32} - 1$



0

Non-Leaf Procedures

- Procedures that call other procedures
- For **nested call**, caller needs to save on the stack:
 - Its return address
 - Any arguments and temporaries needed after the call
- Restore from the stack after the call



Non-Leaf Procedure Example

■ C code:

```
int fact(int n)
{
    if (n < 1) return 1;
    else return n * fact(n - 1);
}
```

- Argument n in \$a0
- Result in \$v0

$f(4) \rightarrow 4 \times f(3)$

$3 \times f(2)$

$2 \times f(1)$

$1 \times f(0)$

1

Non-Leaf Procedure

```
int fact(int n)
{
    if (n < 1) return 1;
    else return n*fact(n-1);
}
```

MIPS code:

fact:		
addi \$sp, \$sp, -8	# adjust stack for 2 items	
sw \$ra, 4(\$sp)	# save return address	
sw \$a0, 0(\$sp)	# save argument	
slti \$t0, \$a0, 1	# test for n < 1	
beq \$t0, \$zero, L1	# branch if n >= 1	
addi \$v0, \$zero, 1	# if so, result is 1	
addi \$sp, \$sp, 8	# pop 2 items from stack	
jr \$ra	# and return	
L1: addi \$a0, \$a0, -1	# else decrement n	
jal fact	# recursive call	
lw \$a0, 0(\$sp)	# restore original n	
lw \$ra, 4(\$sp)	# and return address	
addi \$sp, \$sp, 8	# pop 2 items from stack	
mul \$v0, \$a0, \$v0	# multiply to get result	
jr \$ra	# and return	



Non-Leaf Pro

MIPS code:

fact:

```

addi $sp, $sp, -8
sw    $ra, 4($sp)
sw    $a0, 0($sp)
slti  $t0, $a0, 1
beq   $t0, $zero, L1
addi  $v0, $zero, 1
addi  $sp, $sp, 8
jr    $ra
L1:   addi $a0, $a0, -1
      jal  fact
      lw   $a0, 0($sp)
      lw   $ra, 4($sp)
      addi $sp, $sp, 8
      mul  $v0, $a0, $v0
      jr   $ra
    
```

M →

a0 = 3

jal (ra = jal+4 in main)

fact

...

a0 = 2

jal (ra = M)

fact

...

a0 = 1

jal (ra = M)

fact

...

a0 = 0

jal (ra = M)

fact

v0=1 pop stack

pop stack (a0=1, ra=M)

v0= 1x1 = 1

pop stack (a0=2, ra=M)

v0= 2x1 = 2

pop stack (a0=3, ra:in main)

v0= 3x2 = 6

ra (in main)

a0 = 3

ra (M)

a0 = 2

ra (M)

a0 = 1

ra (M)

a0 = 0

Procedure – cont.

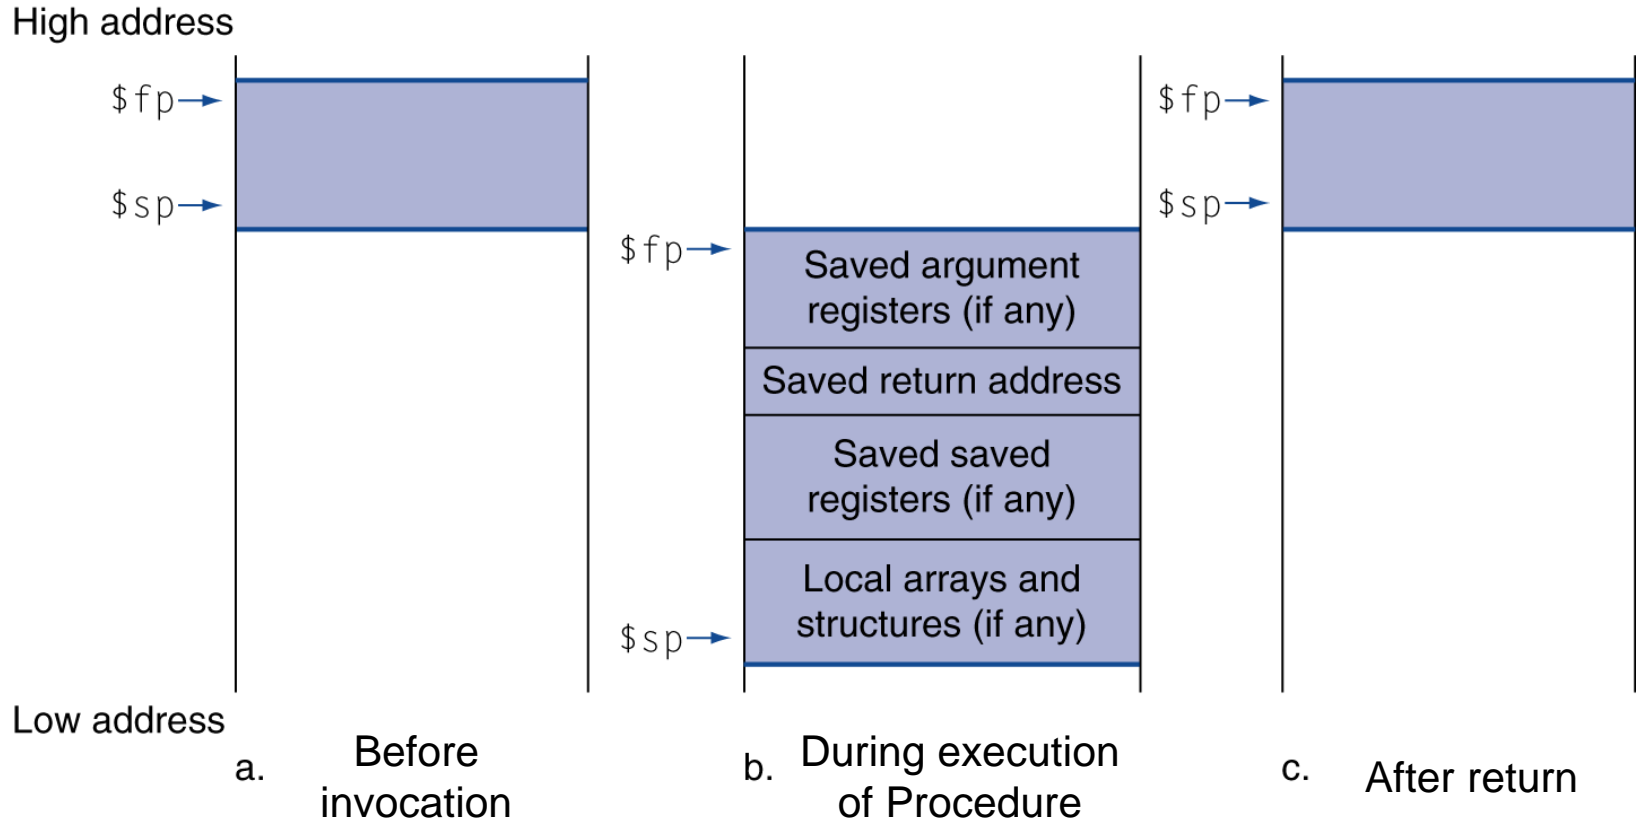
■ Data preservation in MIPS

- To avoid saving and restoring a register whose value is never used, MIPS registers are separated to two groups.

Preserved	Not preserved
Saved registers: \$s0 - \$s7	Temporary registers: \$t0 - \$t9
Stack pointer register: \$sp	Argument registers: \$a0 - \$a3
Return address register: \$ra	Return value registers: \$v0 - \$v1
Stack above the stack pointer	Stack below the stack pointer

- Preserved: The caller assumes them will be preserved on a procedure call. If they are used in the procedure, the callee must save and restore them.
 - In leave procedure, we need to save and restore \$s0


Local Data on the Stack



- Procedure frame (activation record)
 - Space allocated specifically for this invocation of the procedure

Memory Layout

- Text: program code
- Static data: global variables
 - e.g., static variables in C, constant arrays and strings
 - \$gp initialized to address allowing \pm offsets into this segment
- Dynamic data: heap
 - E.g., malloc in C, new in Java
- Stack: automatic storage



Character Data

- Byte-encoded character sets
 - ASCII: 128 characters
 - 95 graphic, 33 control
 - Latin-1: 256 characters
 - ASCII, +96 more graphic characters
- Unicode: 32-bit character set
 - Used in Java, C++ wide characters, ...
 - Most of the world's alphabets, plus symbols
 - UTF-8, UTF-16: variable-length encodings

Byte/Halfword Operations

- MIPS byte/halfword load/store
 - String processing is a common case
 - Sign extend to 32 bits in rt
 - `lb rt, offset(rs)`
 - `lh rt, offset(rs)`
 - Zero extend to 32 bits in rt
 - `lbu rt, offset(rs)`
 - `lhu rt, offset(rs)`
 - Store just rightmost byte/halfword
 - `sb rt, offset(rs)`
 - `sh rt, offset(rs)`

String Copy Example

- C code (naïve):

- Null-terminated string

```
void strcpy (char x[], char y[])  
{  
    int i=0;  
    while ((x[i]=y[i])!='\0')  
        i += 1;  
}
```

- Addresses of x, y in \$a0, \$a1
 - i in \$s0

String Copy Example

- MIPS code:

strcpy:		
addi \$sp, \$sp, -4	# adjust stack for 1 item	
sw \$s0, 0(\$sp)	# save \$s0	
add \$s0, \$zero, \$zero	# i = 0	
L1: add \$t1, \$s0, \$a1	# addr of y[i] in \$t1	
lbu \$t2, 0(\$t1)	# \$t2 = y[i]	
add \$t3, \$s0, \$a0	# addr of x[i] in \$t3	
sb \$t2, 0(\$t3)	# x[i] = y[i]	
beq \$t2, \$zero, L2	# exit loop if y[i] == 0	
addi \$s0, \$s0, 1	# i = i + 1	
j L1	# next iteration of loop	
L2: lw \$s0, 0(\$sp)	# restore saved \$s0	
addi \$sp, \$sp, 4	# pop 1 item from stack	
jr \$ra	# and return	

We don't need to multiply s0 by four because they are char (byte) arrays, instead of integer arrays.

32-bit Constants

- Most constants are small
 - 16-bit immediate is sufficient
- For the occasional 32-bit constant
 - Copies 16-bit constant to left 16 bits of `rt`
 - Clears right 16 bits of `rt` to 0**

$61 = 0000\ 0000\ 0011\ 1101_2$

`lui $s0, 61`

0000 0000 0011 1101	0000 0000 0000 0000
---------------------	---------------------

`ori $s0, $s0, 2304`

0000 0000 0011 1101	0000 1001 0000 0000
---------------------	---------------------

$2304 = 0000\ 1001\ 0000\ 0000_2$



32-bit Constants

- How to load the 32-bit constant below into register \$s0 using MIPS assembly code?

0000 0000 0011 1101 0000 1001 0000 0000

```
lui $s0, 0000000000111101b
```


```
ori $s0, $s0, 0000100100000000b
```

0000 0000 0011 1101	0000 1001 0000 0000
---------------------	---------------------

Branch Addressing

- Branch instructions specify
 - Opcode, two registers, target address
- Most branch targets are near branch
 - Forward or backward


I format



- **PC-relative addressing**
 - Target address = PC + offset \times 4
 - PC already incremented by 4 by this time

Jump Addressing

- Jump targets could be anywhere in text segment
 - Encode full address in instruction



Direct addressing

- Target address = $PC_{31...28} : (\text{address} \times 4)$

(十進位)

10000 L: ...

...

J L

jal L


Jr \$ra

2	2500 (word address)					Go to 2500 (word)
3	2500 (word address)					\$ra = PC (cur+4); go to 2500 (word)
0	31	0	0	0	8	Go to \$ra



Jump Addressing


1. Why the jump addressing is so complex as to concatenate the 26 bits in the instruction and 4 bits from PC?



The space is insufficient to store a full address.

2. Why only needs 4 bits from PC?

Instr. are word alignment \rightarrow target address is a multiple of 4
 \rightarrow rightmost two bits are 00 $\rightarrow 32 - (26+2) = 4$



Jump Addressing

3. How to make sure that the first 4 bits of the target address is the same to the first 4 bits of PC?

To limit the range that j instruction can reach.


4. What if we need to go to the address with the first 4 bits different from the first 4 bits of PC?

With the help of beq

0x2FFFFFFC

PC → 0x22334455

0x20000000



Target Addressing Example

- Loop code from earlier example
 - Assume Loop at location 80000

```
while (save[i] == k)
    i += 1;
```

(十進位)

Loop: sll	\$t1, \$s3, 2	80000	0	0	19	9	2	0
add	\$t1, \$t1, \$s6	80004	0	9	22	9	0	32
lw	\$t0, 0(\$t1)	80008	35	9	8	0		
bne	\$t0, \$s5, Exit	80012	5	8	21	2		
addi	\$s3, \$s3, 1	80016	8	19	19	1		
j	Loop	80020	2	20000				
Exit: ...		80024						

Branching Far Away

- What is the range that *beq* can reach?
 - $\pm 2^{15}$ (word) = $\pm 2^{17}$ (bytes) = $\pm 128\text{k}$ (PC relative)
- What is the range that *j* can reach?
 - 2^{26} (word) = 2^{28} (bytes) = 2^8 MB = 256MB address where the PC supplies the upper 4bits
- If branch target is too far to encode with 16-bit offset, assembler rewrites the code

```
beq $s0,$s1, L1
```



```
bne $s0,$s1, L2
```

```
j L1
```

```
L2: ...
```

Overview of Instruction Formats

- MIPS instr formats:
 - Simple instructions all 32 bits wide
 - Very structured, no unnecessary baggage
 - Three instruction formats:

Name	Fields						Comments
Field size	6 bits	5 bits	5 bits	5 bits	5 bits	6 bits	All MIPS instructions 32 bits
R-format	op	rs	rt	rd	shamt	funct	Arithmetic instructions format
I-format	op	rs	rt	address/immediate			Transfer, branch, imm. format
J-format	op	target address					Jump instruction format

Addressing Mode Summary


1. **Immediate addressing**
 - the operand is a constant inside the instr.
2. **Register addressing**
 - the operand is a register
3. **Base or displacement addressing**
 - the operand is at the memory location whose addr. is the sum of a reg and a constant inside instr. (lw/sw/lh/sh/lb/sb)
4. **PC-relative addressing**
 - the address is a 16-bit constant in instr., shifted left 2 bits (word to byte), and then added to (PC+4). (i.e., next instr.)
5. **Pseudodirect addressing**
 - the address is a 26-bit constant in instruction, shifted left 2 bits (word to byte), and then concatenated with the 4 upper bits of PC. Address boundary: $2^{28} = 256\text{MB}$.

Addressing Mode Summary

1. Immediate addressing




2. Register addressing




Registers

Register


3. Base addressing




Memory




4. PC-relative addressing




Memory



5. Pseudodirect addressing



Memory



Put it all together – Sort

- Illustrates use of assembly instructions for a C bubble sort function
- Swap procedure (leaf)

```
void swap(int v[], int k)
{
    int temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}
```

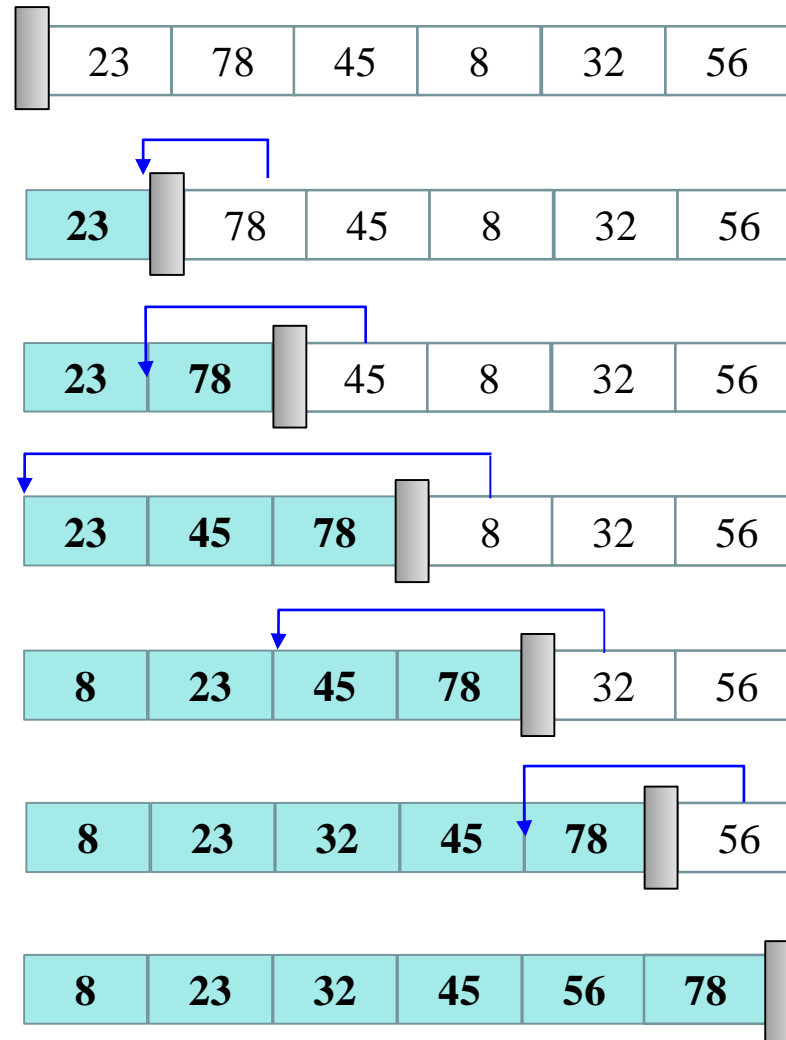
- v in \$a0, k in \$a1, temp in \$t0



The Procedure Swap

swap:	sll \$t1, \$a1, 2	# \$t1 = k * 4
	add \$t1, \$a0, \$t1	# \$t1 = v+(k*4)
		# (address of v[k])
	lw \$t0, 0(\$t1)	# \$t0 (temp) = v[k]
	lw \$t2, 4(\$t1)	# \$t2 = v[k+1]
	sw \$t2, 0(\$t1)	# v[k] = \$t2 (v[k+1])
	sw \$t0, 4(\$t1)	# v[k+1] = \$t0 (temp)
	jr \$ra	# return to calling routine

The Insertion Sort Algorithm



The Sort Procedure in C

- Non-leaf (calls swap)

```
void sort (int v[], int n)
{
    int i, j;
    for (i = 0; i < n; i += 1){
        for(j = i-1; j>=0 && v[j]>v[j+1];j -= 1) {
            swap(v,j);
        }
    }
}
```

- v in \$a0, k in \$a1, i in \$s0, j in \$s1

The Procedure Body

	move \$s2, \$a0	# save \$a0 into \$s2	Move params
	move \$s3, \$a1	# save \$a1 into \$s3	
	move \$s0, \$zero	# i = 0	Outer loop
for1tst:	slt \$t0, \$s0, \$s3	# \$t0 = 0 if \$s0 ≥ \$s3 (i ≥ n)	
	beq \$t0, \$zero, exit1	# go to exit1 if \$s0 ≥ \$s3 (i ≥ n)	Inner loop
	addi \$s1, \$s0, -1	# j = i - 1	
for2tst:	slti \$t0, \$s1, 0	# \$t0 = 1 if \$s1 < 0 (j < 0)	
	bne \$t0, \$zero, exit2	# go to exit2 if \$s1 < 0 (j < 0)	
	sll \$t1, \$s1, 2	# \$t1 = j * 4	
	add \$t2, \$s2, \$t1	# \$t2 = v + (j * 4)	
	lw \$t3, 0(\$t2)	# \$t3 = v[j]	
	lw \$t4, 4(\$t2)	# \$t4 = v[j + 1]	
	slt \$t0, \$t4, \$t3	# \$t0 = 0 if \$t4 ≥ \$t3	
	beq \$t0, \$zero, exit2	# go to exit2 if \$t4 ≥ \$t3	
	move \$a0, \$s2	# 1st param of swap is v (old \$a0)	Pass params & call
	move \$a1, \$s1	# 2nd param of swap is j	
	jal swap	# call swap procedure	Inner loop
	addi \$s1, \$s1, -1	# j -= 1	
	j for2tst	# jump to test of inner loop	Outer loop
exit2:	addi \$s0, \$s0, 1	# i += 1	
	j for1tst	# jump to test of outer loop	

The Full Procedure

sort:	addi \$sp,\$sp, -20	# make room on stack for 5 registers
	sw \$ra, 16(\$sp)	# save \$ra on stack
	sw \$s3,12(\$sp)	# save \$s3 on stack
	sw \$s2, 8(\$sp)	# save \$s2 on stack
	sw \$s1, 4(\$sp)	# save \$s1 on stack
	sw \$s0, 0(\$sp)	# save \$s0 on stack
	...	# procedure body
	...	
exit1:	lw \$s0, 0(\$sp)	# restore \$s0 from stack
	lw \$s1, 4(\$sp)	# restore \$s1 from stack
	lw \$s2, 8(\$sp)	# restore \$s2 from stack
	lw \$s3,12(\$sp)	# restore \$s3 from stack
	lw \$ra,16(\$sp)	# restore \$ra from stack
	addi \$sp,\$sp, 20	# restore stack pointer
	jr \$ra	# return to calling routine

Put it all together — Arrays vs. Pointers

- Array indexing involves
 - Multiplying index by element size
 - Adding to array base address
- Pointers correspond directly to memory addresses
 - Can avoid indexing complexity



Example: Clearing an Array

```
clear1(int array[], int size) {  
    int i;  
    for (i = 0; i < size; i += 1)  
        array[i] = 0;  
}
```

```
loop1: move $t0,$zero    #i = 0  
      sll $t1,$t0,2      # $t1 = i * 4  
      add $t2,$a0,$t1    # $t2 =  
                        # &array[i]  
      sw $zero, 0($t2)   # array[i]=0  
      addi $t0,$t0,1     # i = i + 1  
      slt $t3,$t0,$a1    # $t3 = (i < size)  
      bne $t3,$zero,loop1  
                        # if(...) goto loop1
```

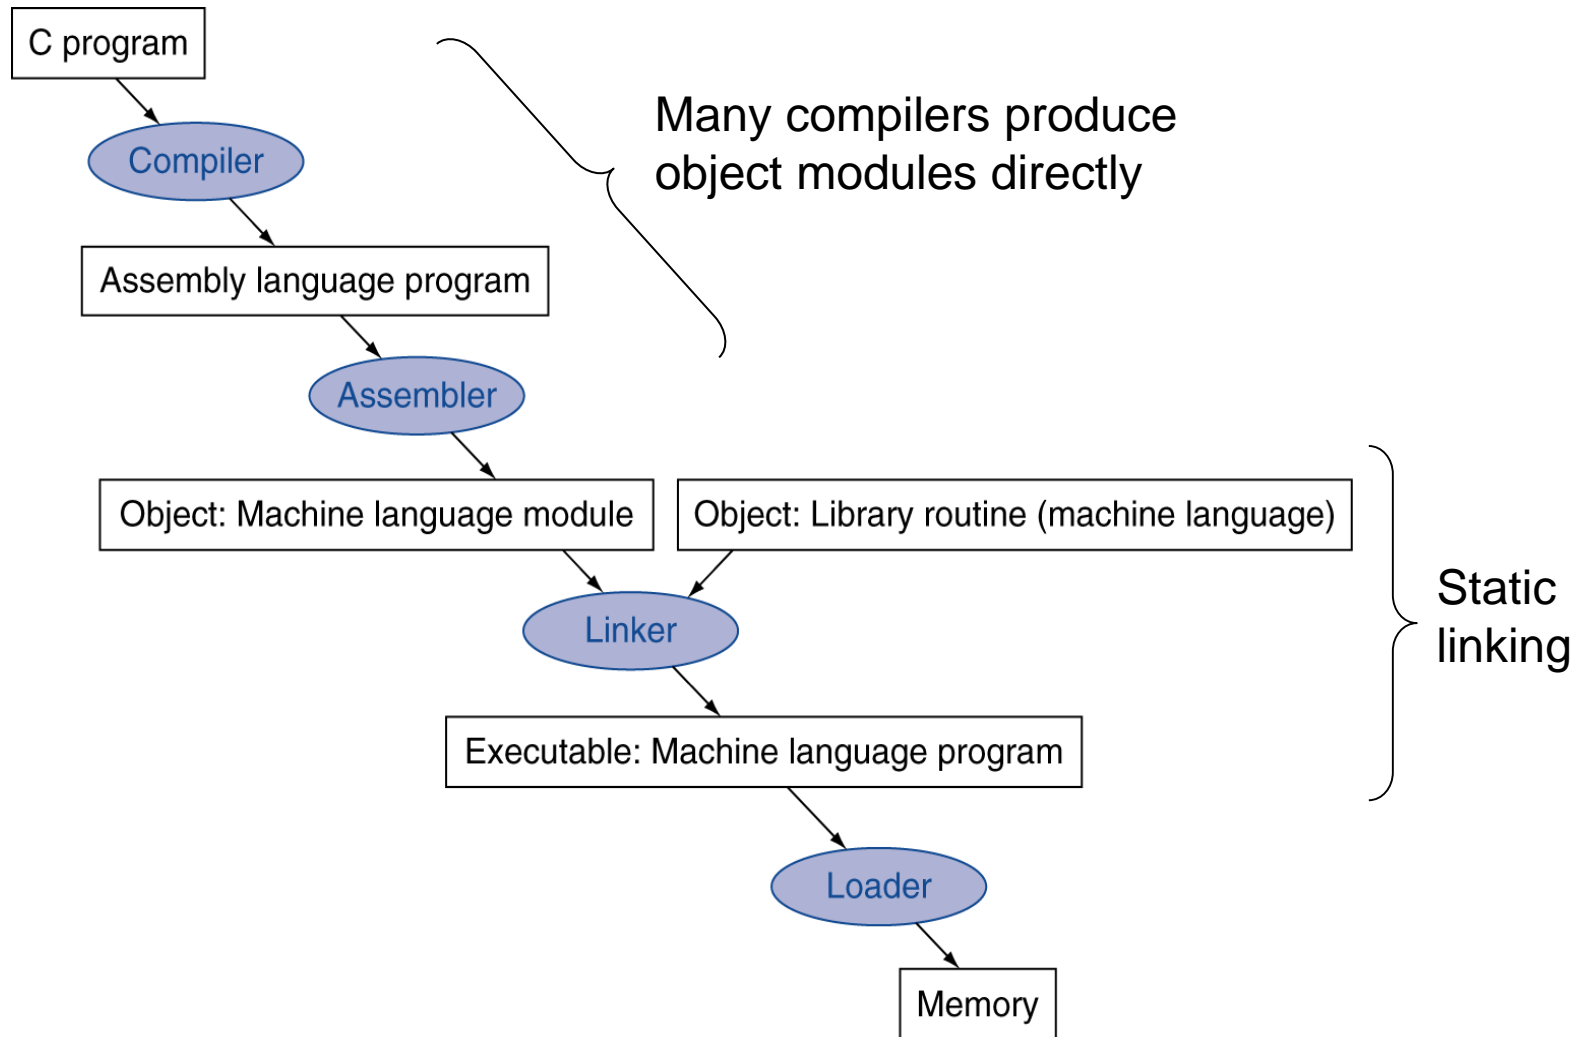
```
clear2(int *array, int size) {  
    int *p;  
    for (p = &array[0]; p < &array[size];  
        p = p + 1)  
        *p = 0;  
}
```

```
      move $t0,$a0      # p=&array[0]  
      sll $t1,$a1,2      # $t1 = size*4  
      add $t2,$a0,$t1    # $t2 =  
                        # &array[size]  
loop2: sw $zero,0($t0)   # Memory[p] = 0  
      addi $t0,$t0,4     # p = p + 4  
      slt $t3,$t0,$t2    # $t3=(p<&array[size])  
      bne $t3,$zero,loop2  
                        # if(...) goto loop2
```


Comparison of Array vs. Ptr

- Multiply “strength reduced” to shift
- Array version requires shift to be inside loop
 - Part of index calculation for incremented i
 - c.f. incrementing pointer
- Compiler can achieve same effect as manual use of pointers
 - Induction variable elimination
 - Better to make program clearer and safer

Translation and Startup



Assembler Pseudoinstructions

- Most assembler instructions represent machine instructions one-to-one
- Pseudoinstructions: figments of the assembler's imagination

`move $t0, $t1` \rightarrow `add $t0, $zero, $t1`

`blt $t0, $t1, L` \rightarrow `slt $at, $t0, $t1`
 `bne $at, $zero, L`

- `$at` (register 1): assembler temporary

Producing an Object Module

- Assembler (or compiler) translates program into machine instructions
- Provides information for building a complete program from the pieces
 - Header: described contents of object module
 - Text segment: translated instructions
 - Static data segment: data allocated for the life of the program
 - Relocation info: for contents that depend on absolute location of loaded program
 - Symbol table: global definitions and external refs
 - Debug info: for associating with source code

Linking Object Modules





- Produces an executable image
 1. Merges segments
 2. Resolve labels (determine their addresses)
 3. Patch location-dependent and external refs
- Could leave location dependencies for fixing by a relocating loader
 - But with virtual memory, no need to do this
 - Program can be loaded into absolute location in virtual memory space

Loading a Program




- Load from image file on disk into memory
 1. Read header to determine segment sizes
 2. Create virtual address space
 3. Copy text and initialized data into memory
 - Or set page table entries so they can be faulted in
 4. Set up arguments on stack
 5. Initialize registers (including \$sp, \$fp, \$gp)
 6. Jump to startup routine
 - Copies arguments to \$a0, ... and calls main
 - When main returns, do exit syscall

Effect of Compiler Optimization

Compiled with gcc for Pentium 4 under Linux



Effect of Language and Algorithm



Lessons Learnt

- Instruction count and CPI are not good performance indicators in isolation
- Compiler optimizations are sensitive to the algorithm
- Java/JIT compiled code is significantly faster than JVM interpreted
 - Comparable to optimized C in some cases
- Nothing can fix a dumb algorithm!

ARM & MIPS Similarities

- ARM: the most popular embedded core
- Similar basic set of instructions to MIPS

	ARM	MIPS
Date announced	1985	1985
Instruction size	32 bits	32 bits
Address space	32-bit flat	32-bit flat
Data alignment	Aligned	Aligned
Data addressing modes	9	3
Registers	15 × 32-bit	31 × 32-bit
Input/output	Memory mapped	Memory mapped

Implementing IA-32


- Complex instruction set makes implementation difficult
 - Hardware translates instructions to simpler microoperations
 - Simple instructions: 1–1
 - Complex instructions: 1–many
 - Microengine similar to RISC
 - Market share makes this economically viable
- Comparable performance to RISC
 - Compilers avoid complex instructions

Fallacies

- Powerful instruction \Rightarrow higher performance
 - Fewer instructions required
 - But complex instructions are hard to implement
 - May slow down all instructions, including simple ones
 - Compilers are good at making fast code from simple instructions
- Use assembly code for high performance
 - But modern compilers are better at dealing with modern processors
 - More lines of code \Rightarrow more errors and less productivity

Fallacies

- Backward compatibility \Rightarrow instruction set doesn't change
 - But they do accrete more instructions



x86 instruction set

Pitfalls

- Sequential words are not at sequential addresses
 - Increment by 4, not by 1!
- Keeping a pointer to an automatic variable after procedure returns
 - e.g., passing pointer back via an argument
 - Pointer becomes invalid when stack popped

Concluding Remarks

- Design principles
 1. Simplicity favors regularity
 2. Smaller is faster
 3. Make the common case fast
 4. Good design demands good compromises
- Layers of software/hardware
 - Compiler, assembler, hardware
- MIPS: typical of RISC ISAs
 - c.f. x86

Concluding Remarks

- Measure MIPS instruction executions in benchmark programs
 - Consider making the common case fast
 - Consider compromises

Instruction class	MIPS examples	SPEC2006 Int	SPEC2006 FP
Arithmetic	add, sub, addi	16%	48%
Data transfer	lw, sw, lb, lbu, lh, lhu, sb, lui	35%	36%
Logical	and, or, nor, andi, ori, sll, srl	12%	4%
Cond. Branch	beq, bne, slt, slti, sltiu	34%	8%
Jump	j, jr, jal	2%	0%