



# Chapter 4

## The Processor

# Introduction

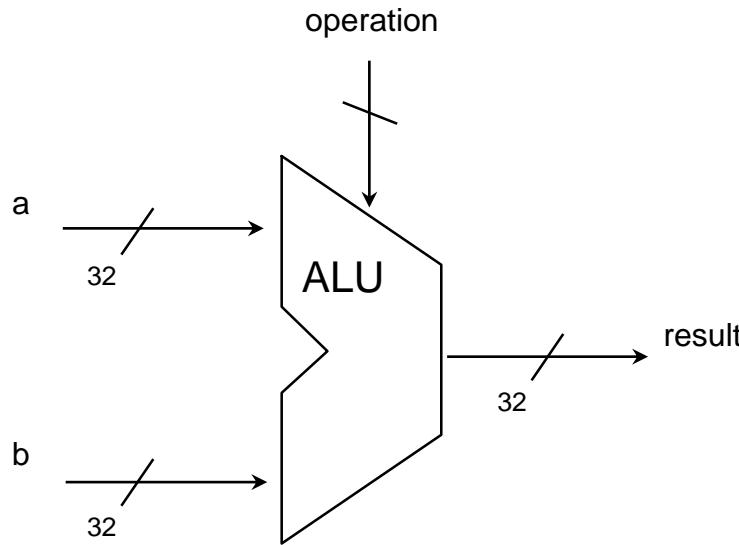
- CPU performance factors
  - Instruction count
    - Determined by ISA and compiler
  - CPI and Cycle time
    - Determined by CPU hardware
- We will examine two MIPS implementations
  - A simplified version (single-cycle CPU)
  - A more realistic version (pipelined CPU)

# Introduction

- The simplified CPU (Single-cycle CPU)
  - ALU (Arithmetic Logic Unit)
  - Datapath
  - Control Signals
    - ALU control
    - Main control
- Instructions supported
  - Memory reference: lw, sw
  - Arithmetic/logical: add, sub, and, or, slt
  - Control transfer: beq, j
- Simple subset, shows most aspects

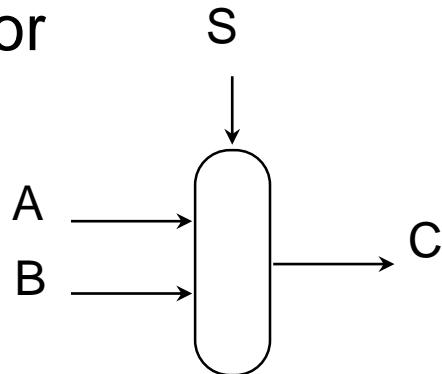
# The ALU (Appendix)

## A 32-bit ALU



- First, let's review Boolean Logic we'll need

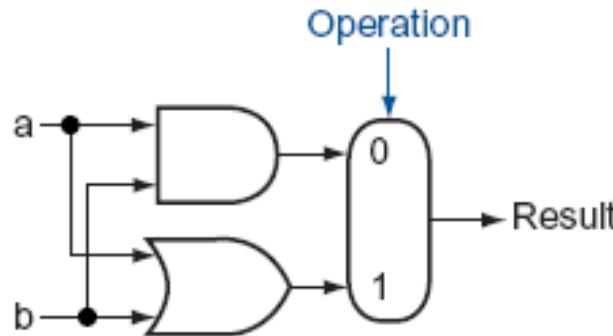
- Multiplexor



*note: we call this a 2-input mux even though it has 3 inputs!*

# The ALU

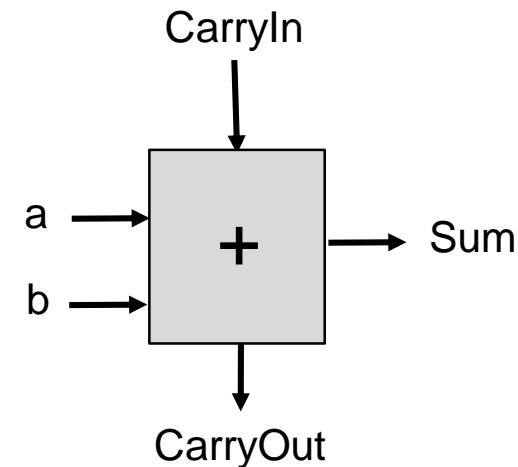
- 1-bit and/or



- 1-bit addition:

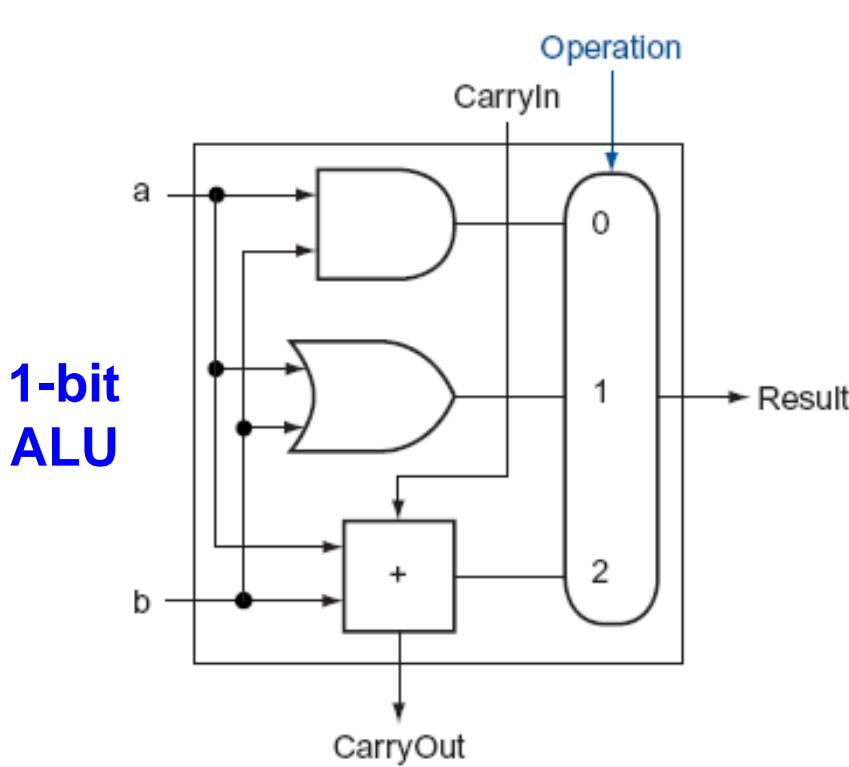
$$C_{out} = a \cdot b + a \cdot C_{in} + b \cdot C_{in}$$

$$\text{sum} = a \oplus b \oplus C_{in}$$



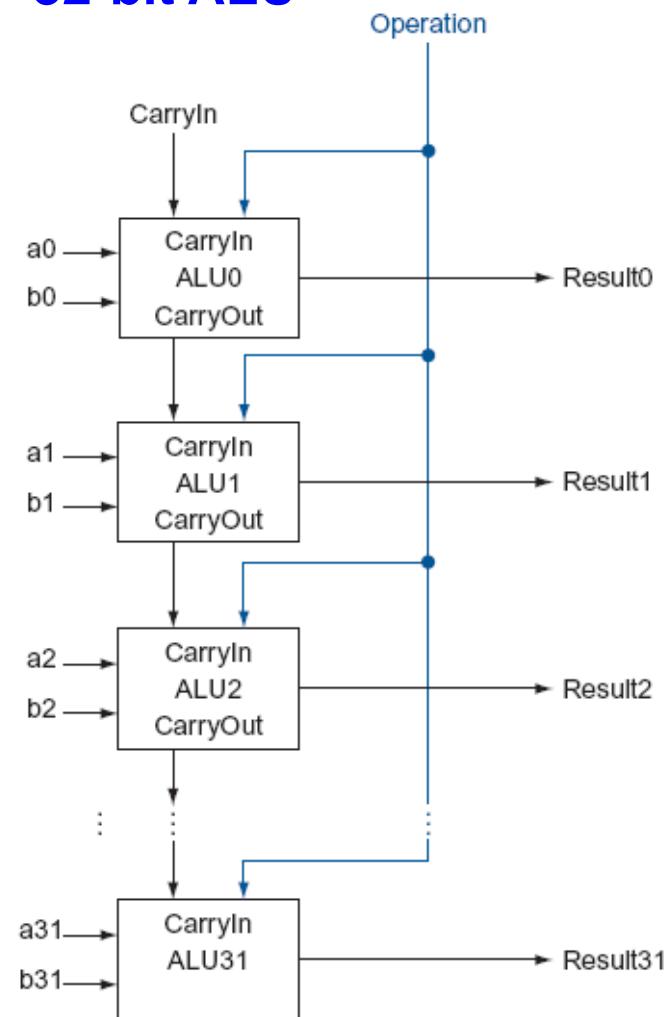
- How could we build a 1-bit ALU for add, and, and or?
- How could we build a 32-bit ALU?

# The ALU – ADD function



| Control lines |          |
|---------------|----------|
| Operation     | Function |
| 00            | and      |
| 01            | or       |
| 10            | add      |

**32-bit ALU**



# The ALU – SUB function

- What about subtraction ( $a-b$ )?  $\rightarrow a + (-b)$

- How to negate b?

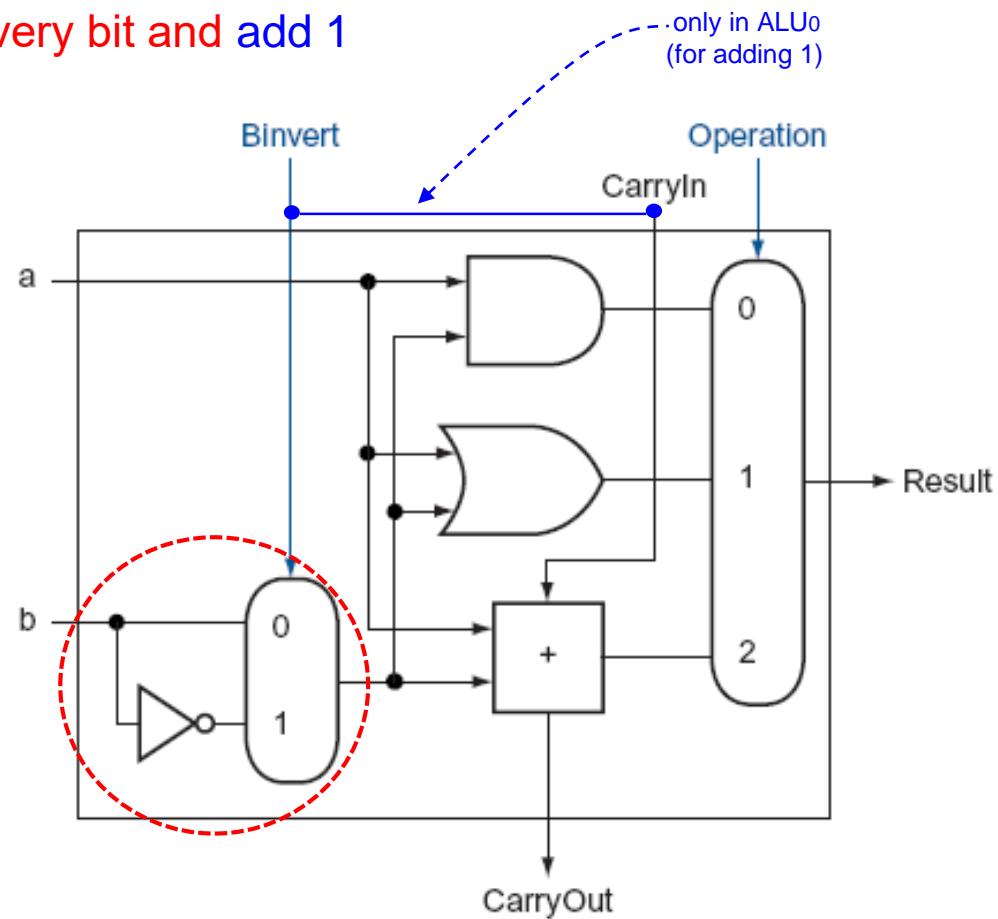
- 2's complement: invert every bit and add 1

- A very clever solution:

| Control lines |           |          |
|---------------|-----------|----------|
| Binvert       | Operation | Function |
| 0             | 00        | and      |
| 0             | 01        | or       |
| 0             | 10        | add      |
| 1             | 10        | sub      |

How to do “add 1” for **sub** ?  
(need to set  $C_{in0} = 1$ )

Combining Binvert and  $C_{in0}$



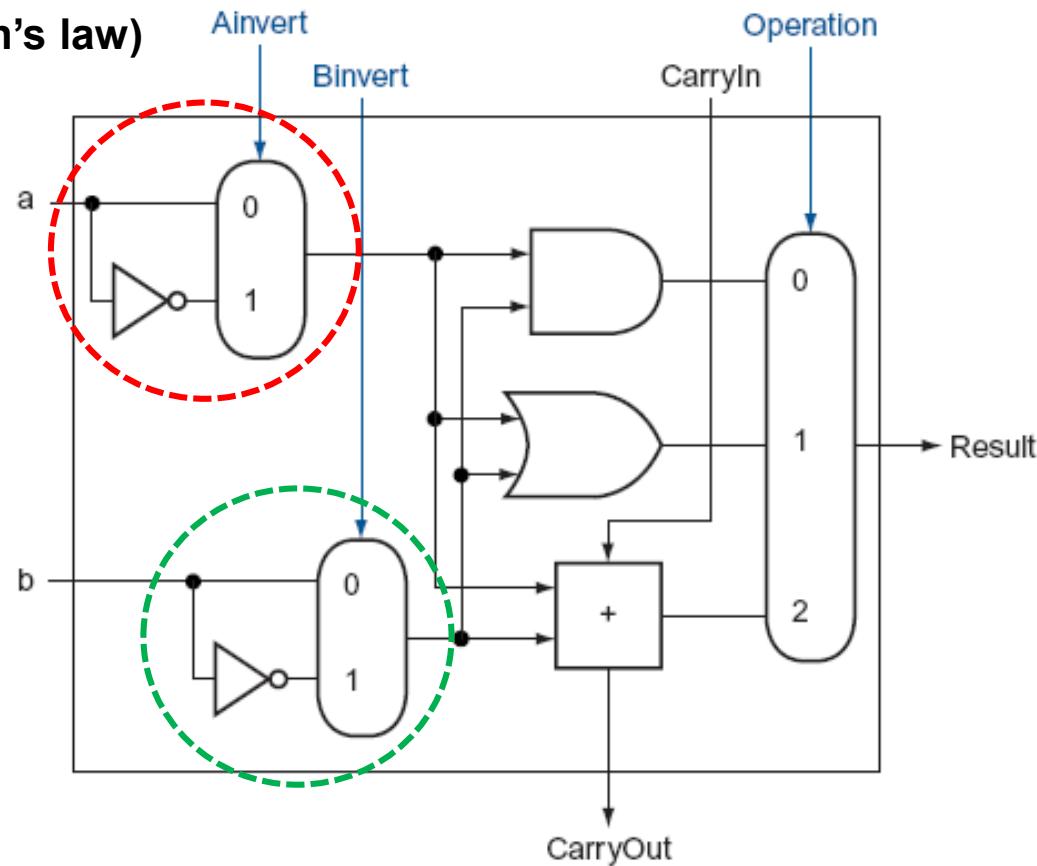
# The ALU - NOR function

## NOR function

- How do we get “a NOR b” ?

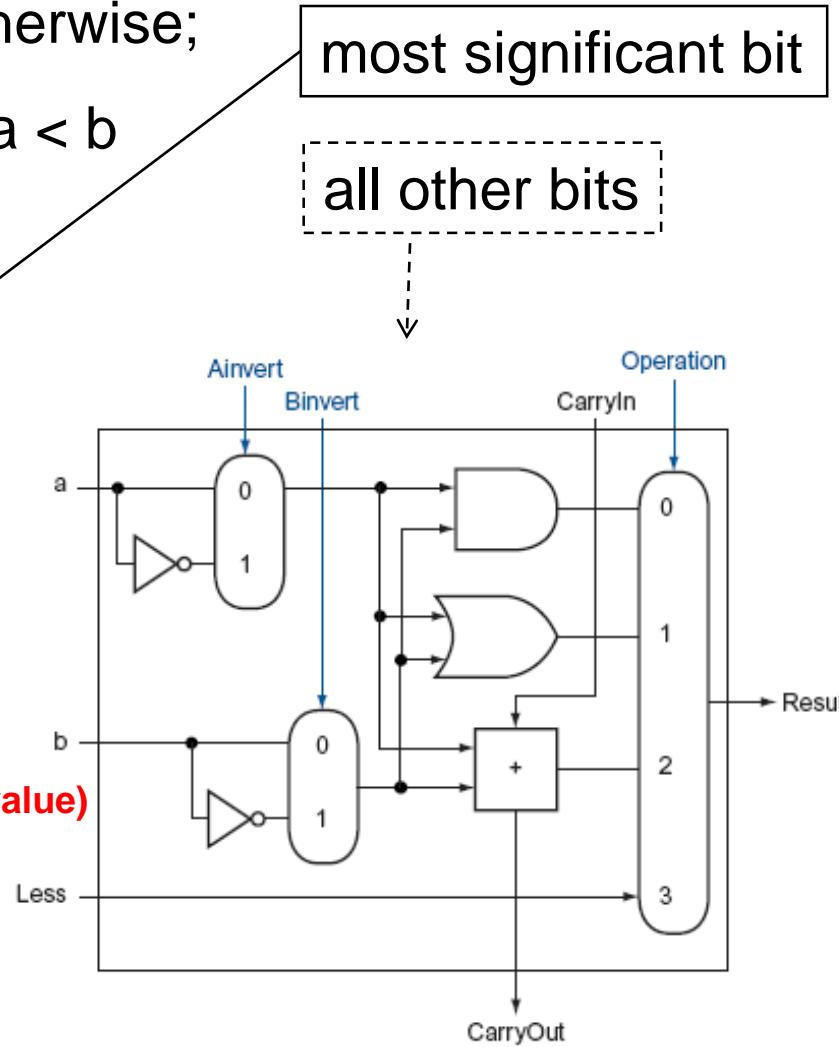
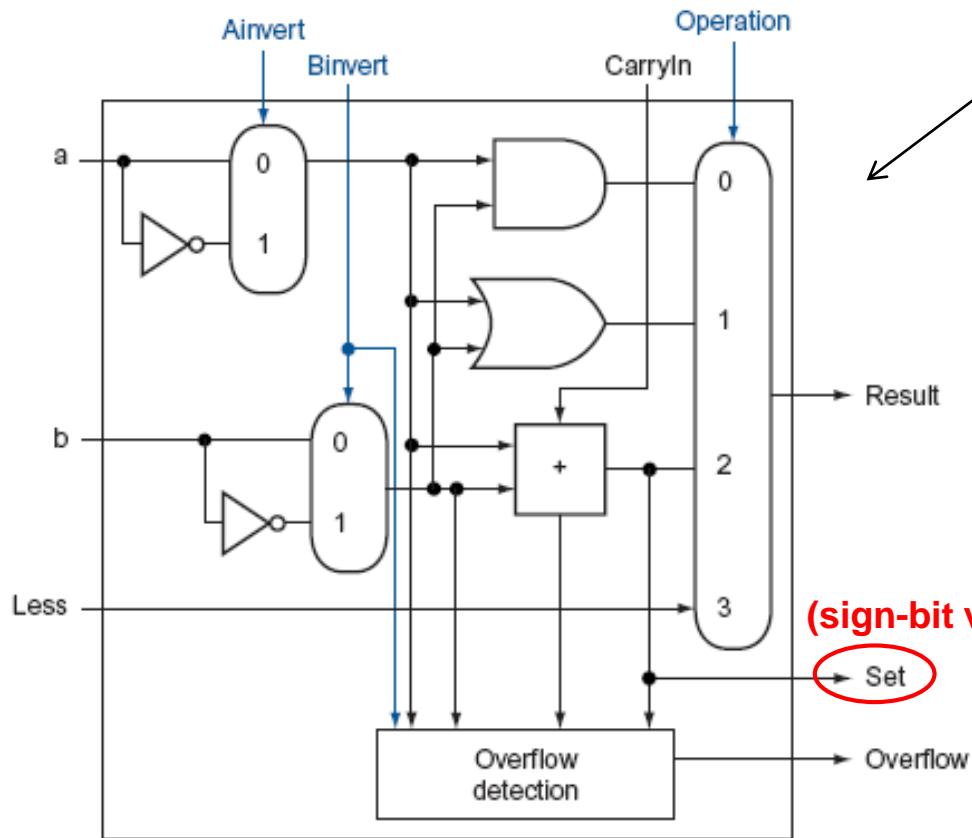
$$(\overline{a + b}) = \overline{a} \cdot \overline{b} \quad (\text{DeMorgan's law})$$

| Control lines |      |           | Function |
|---------------|------|-----------|----------|
| Ainv          | Binv | Operation |          |
| 0             | 0    | 00        | and      |
| 0             | 0    | 01        | or       |
| 0             | 0    | 10        | add      |
| 0             | 1    | 10        | sub      |
| 1             | 1    | 00        | nor      |



# The ALU – SLT function

- SLT produces a 1 if  $rs < rt$  and 0 otherwise;
- use subtraction:  $(a-b) < 0$  implies  $a < b$

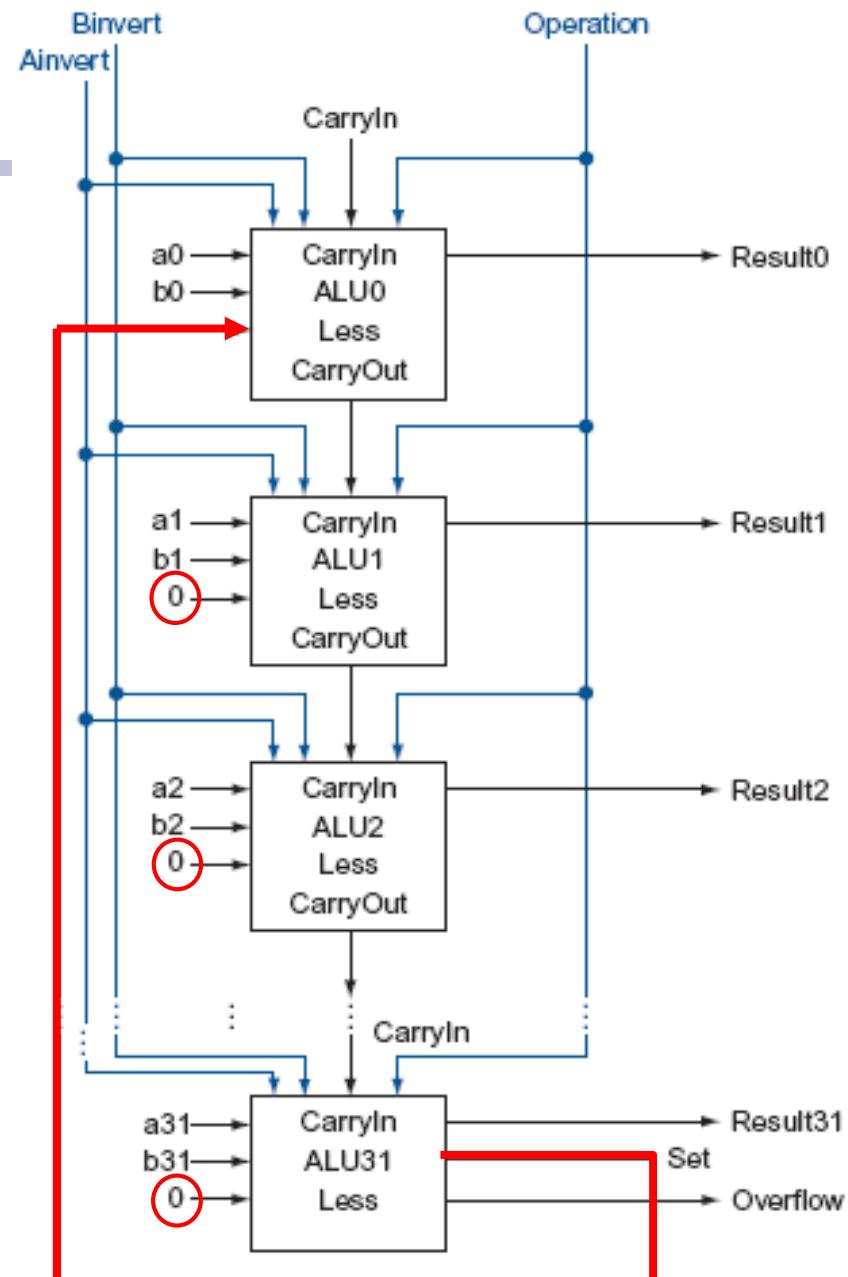


# The ALU - SLT

```
if (a - b) < 0, Result = 0...01;  
else Result = 0...00
```

Control lines

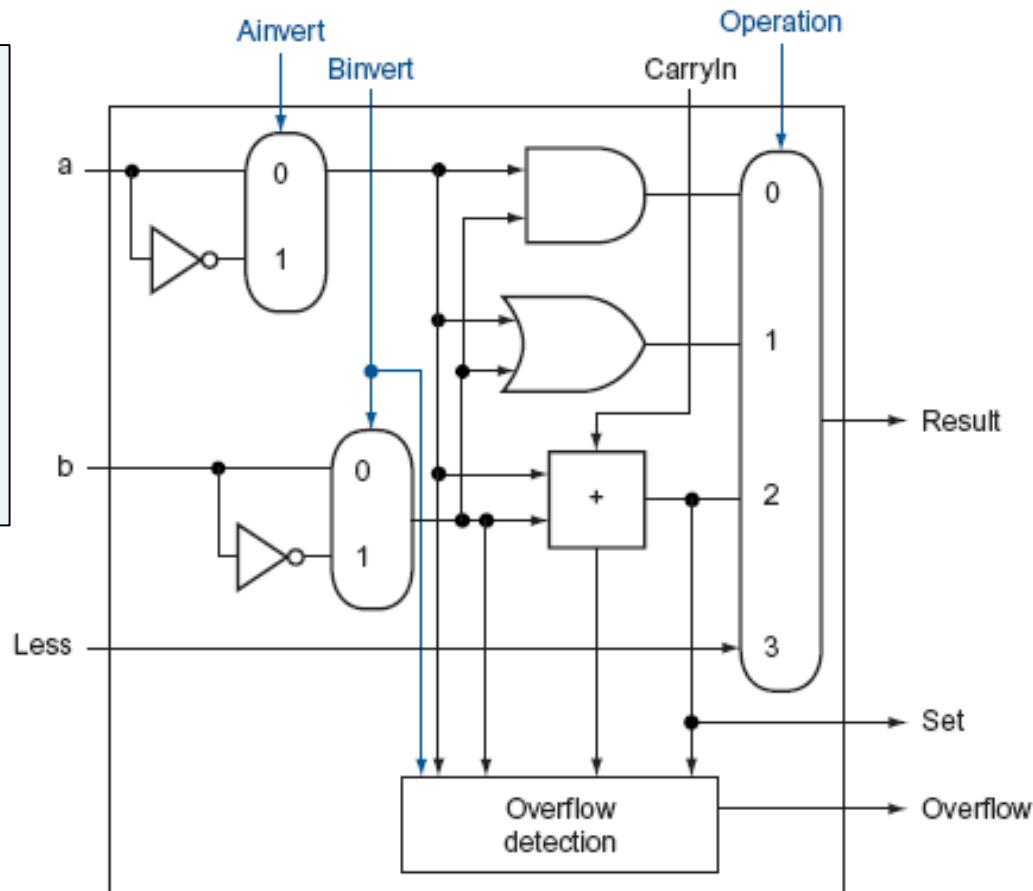
| Ainv | Binv | Operation | Function |
|------|------|-----------|----------|
| 0    | 0    | 0 0       | and      |
| 0    | 0    | 0 1       | or       |
| 0    | 0    | 1 0       | add      |
| 0    | 1    | 1 0       | sub      |
| 0    | 1    | 1 1       | slt      |
| 1    | 1    | 0 0       | nor      |



# The ALU

Control lines

| Ainv | Binv | Operation | Function |
|------|------|-----------|----------|
| 0    | 0    | 0 0       | and      |
| 0    | 0    | 0 1       | or       |
| 0    | 0    | 1 0       | add      |
| 0    | 1    | 1 0       | sub      |
| 0    | 1    | 1 1       | slt      |
| 1    | 1    | 0 0       | nor      |



# ALU- Test for equality (for beq)

- Use subtraction:

$$(a-b) = 0 \text{ implies } a = b$$

- Notice control lines:

0000 = and

0001 = or

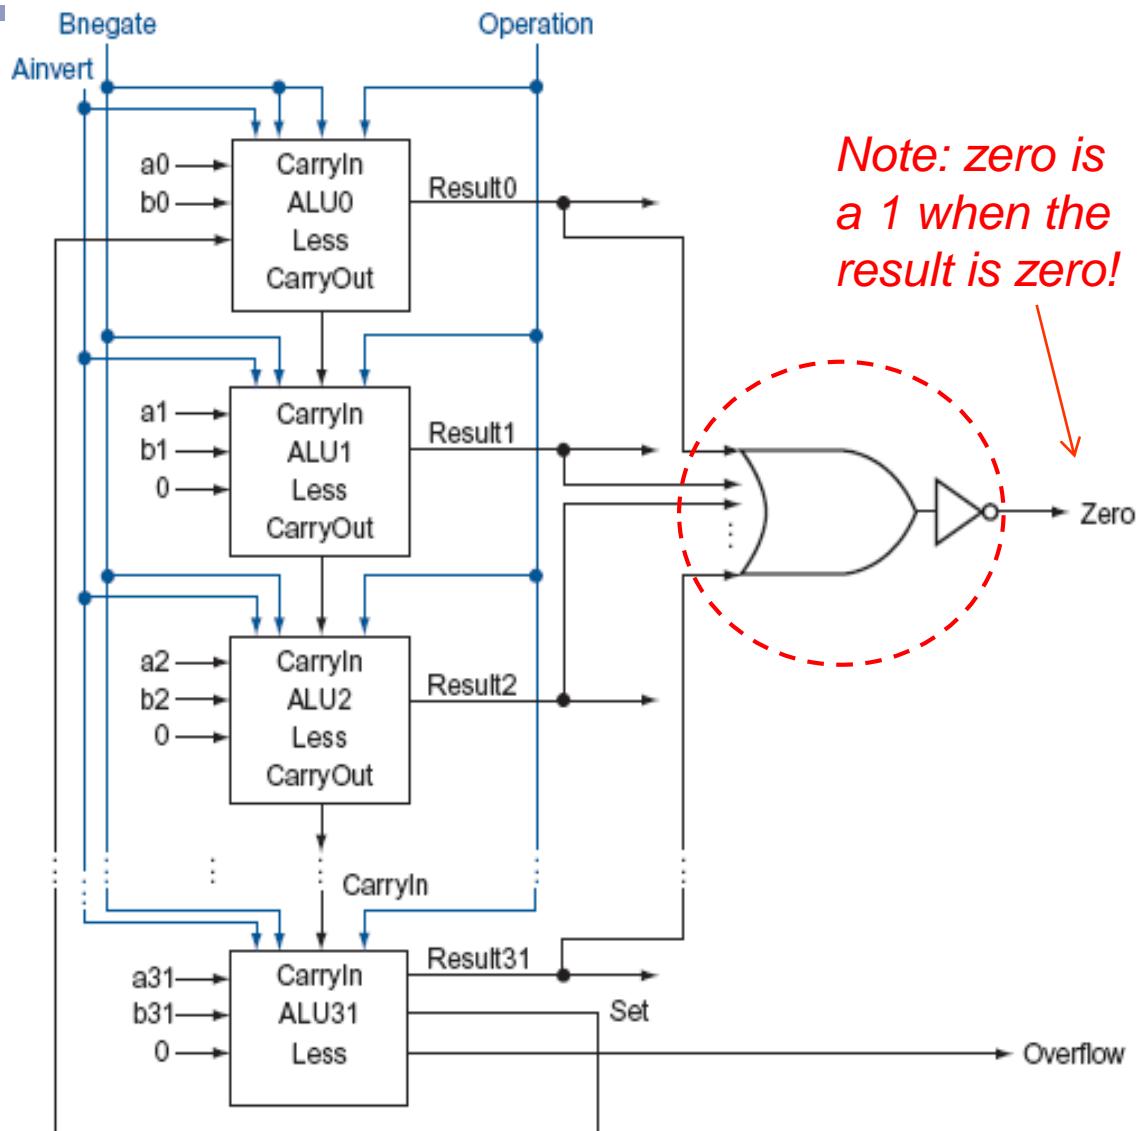
0010 = add

0110 = subtract

0111 = slt

1100 = NOR

Note: we only perform the first five functions in the implementation later



# ALU: Overflow

- 4-bit signed addition

Ex.1  $7 + 3 = 10$ , but ...

$$\begin{array}{r} (0) \quad (1) \quad (1) \quad (1) \\ \boxed{0} \quad 1 \quad 1 \quad 1 \\ + \quad 0 \quad 0 \quad 1 \quad 1 \\ \hline 1 \quad 0 \quad 1 \quad 0 \end{array} \quad \begin{array}{l} 7 \\ 3 \\ -6 \end{array}$$

Ex.2  $(-4) + (-5) = -9$ , but ...

$$\begin{array}{r} (1) \quad (0) \quad (0) \quad (0) \\ \boxed{1} \quad 1 \quad 0 \quad 0 \\ + \quad 1 \quad 0 \quad 1 \quad 1 \\ \hline 0 \quad 1 \quad 1 \quad 1 \end{array} \quad \begin{array}{l} -4 \\ -5 \\ 7 \end{array}$$

Overflow !

| 4 bits     |          |
|------------|----------|
| (10-based) | (Binary) |
| 0          | 0 0 0 0  |
| 1          | 0 0 0 1  |
| 2          | 0 0 1 0  |
| 3          | 0 0 1 1  |
| 4          | 0 1 0 0  |
| 5          | 0 1 0 1  |
| 6          | 0 1 1 0  |
| 7          | 0 1 1 1  |
| -8         | 1 0 0 0  |
| -7         | 1 0 0 1  |
| -6         | 1 0 1 0  |
| -5         | 1 0 1 1  |
| -4         | 1 1 0 0  |
| -3         | 1 1 0 1  |
| -2         | 1 1 1 0  |
| -1         | 1 1 1 1  |

Signed binary  
(2's complement)

# ALU: Overflow

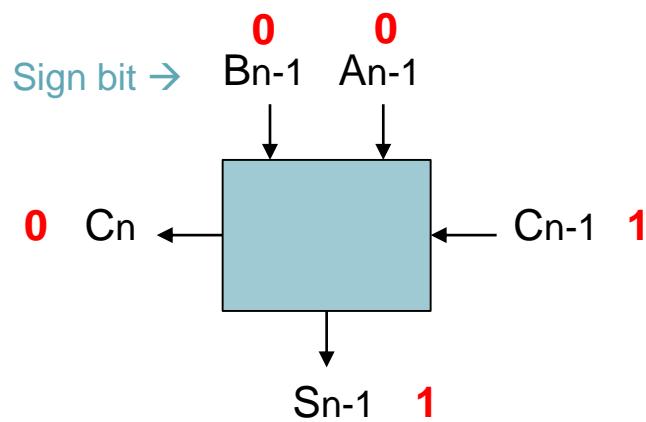
- **Overflow:** result too big/small to represent
  - For addition
    - Adding two +ve operands: overflow if result sign is 1
    - Adding two –ve operands: overflow if result sign is 0
    - Adding +ve and –ve operands: No overflow
  - How about subtraction?

| Operation | Operand A | Operand B | Result indicating overflow |
|-----------|-----------|-----------|----------------------------|
| A+B       | $\geq 0$  | $\geq 0$  | $< 0$                      |
| A+B       | $< 0$     | $< 0$     | $\geq 0$                   |
| A-B       | $\geq 0$  | $< 0$     | $< 0$                      |
| A-B       | $< 0$     | $\geq 0$  | $\geq 0$                   |

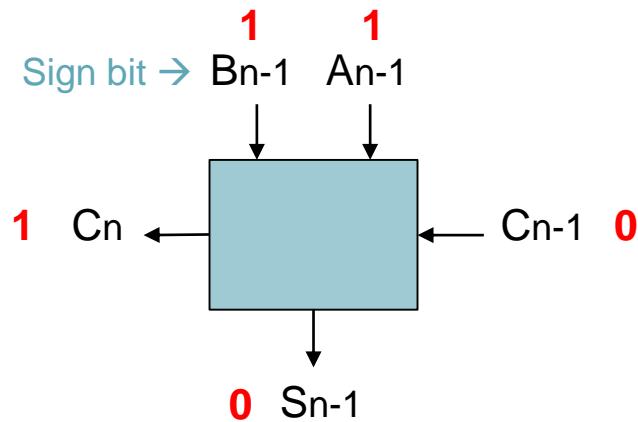
- **Overflow detection:** Carry into MSB  $\neq$  Carry out of MSB

# Overflow Detection

- Signed n-bit:  $A_{n-1} A_{n-2} \dots A_0 \rightarrow$  Sign bit :  $A_{n-1}$
- For n-bit adder, overflow occurs if
  - 1. Add two positive numbers, result is negative
  - 2. Add two negative numbers, result is positive



→  $C_{n-1} = 1$  and  $C_n = 0$



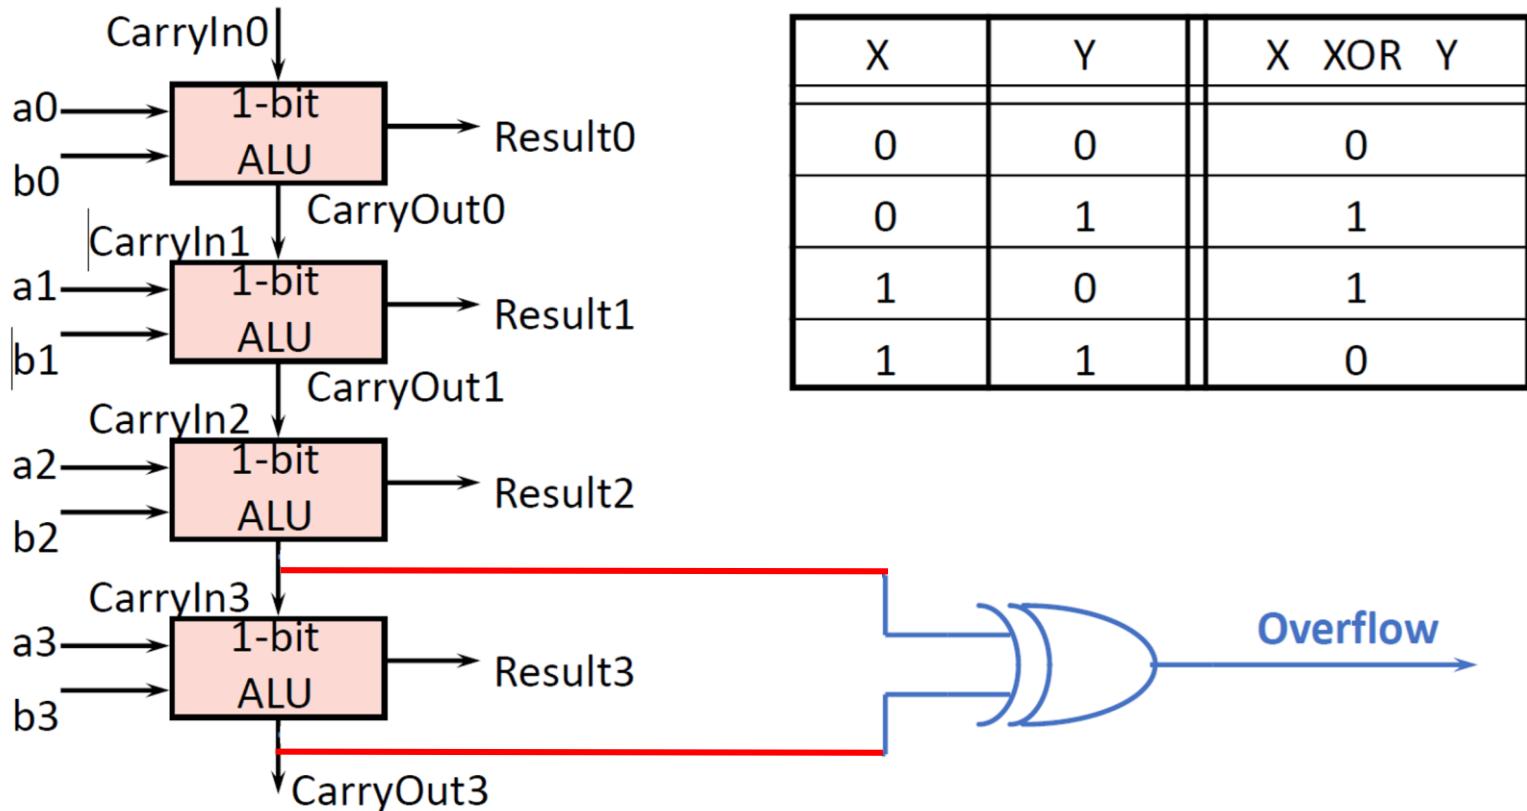
→  $C_{n-1} = 0$  and  $C_n = 1$

$$C_{n-1} \text{ XOR } C_n = 1$$

# Overflow Detection

## Detection logic

- Overflow = CarryIn[N-1] **XOR** CarryOut[N-1]



# Introduction

- The simplified CPU (Single-cycle CPU)
  - ALU (Arithmetic Logic Unit)
  - Datapath
  - Control Signals
- We have seen the design of a simple 32-bit ALU that perform and/or/add/sub/slt/nor
- Now we are ready to design the whole processor: Datapath + Control signals

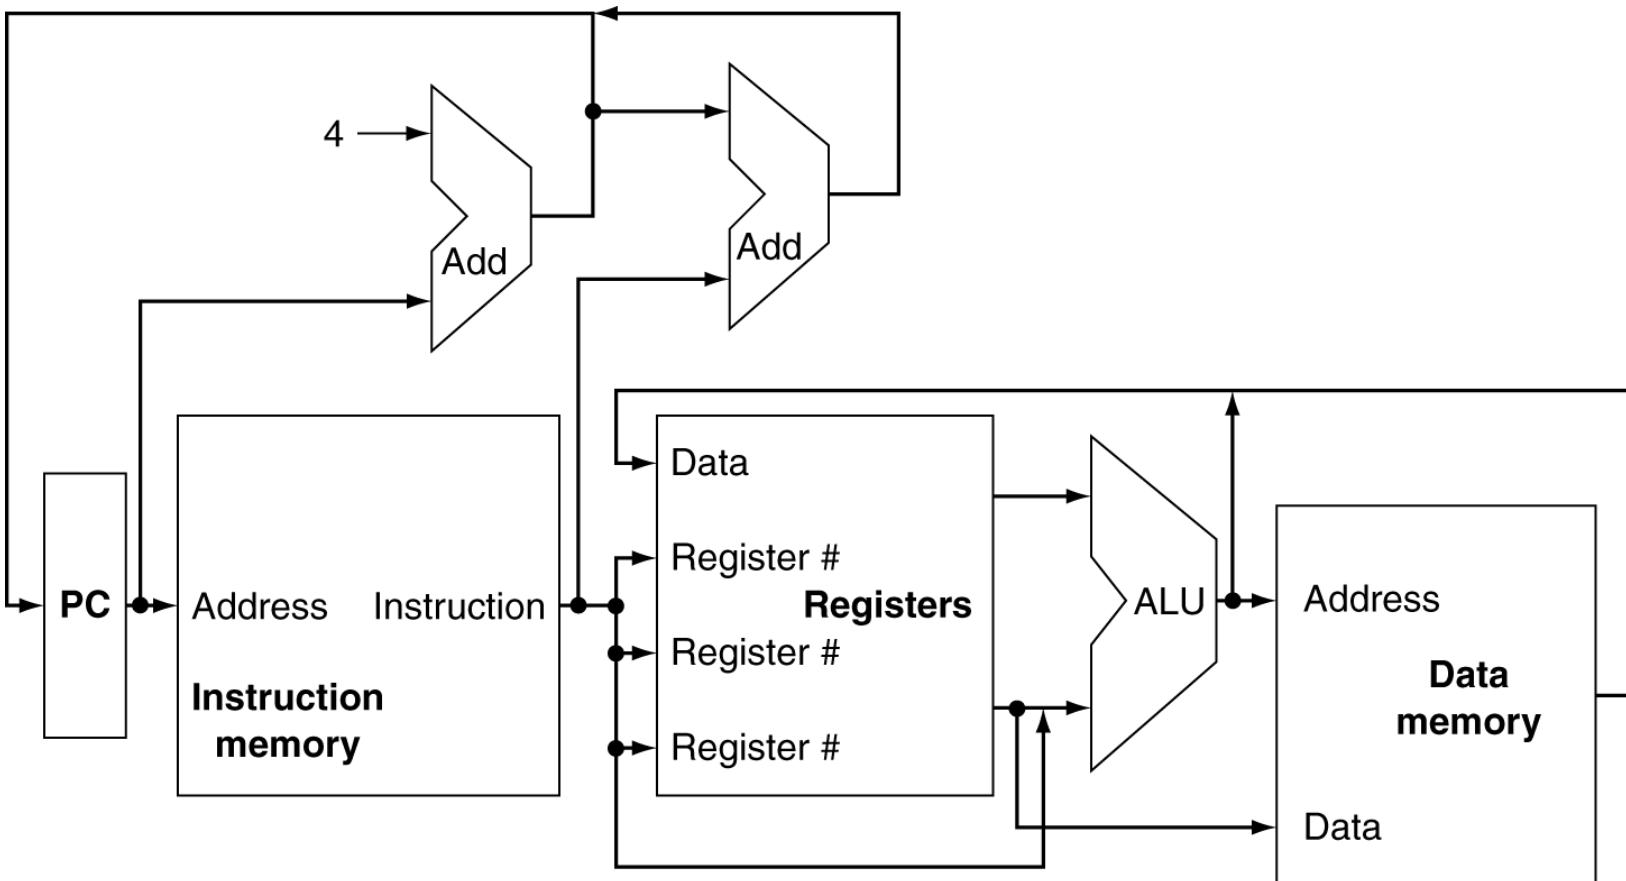
# CPU Overview

## Instruction Execution

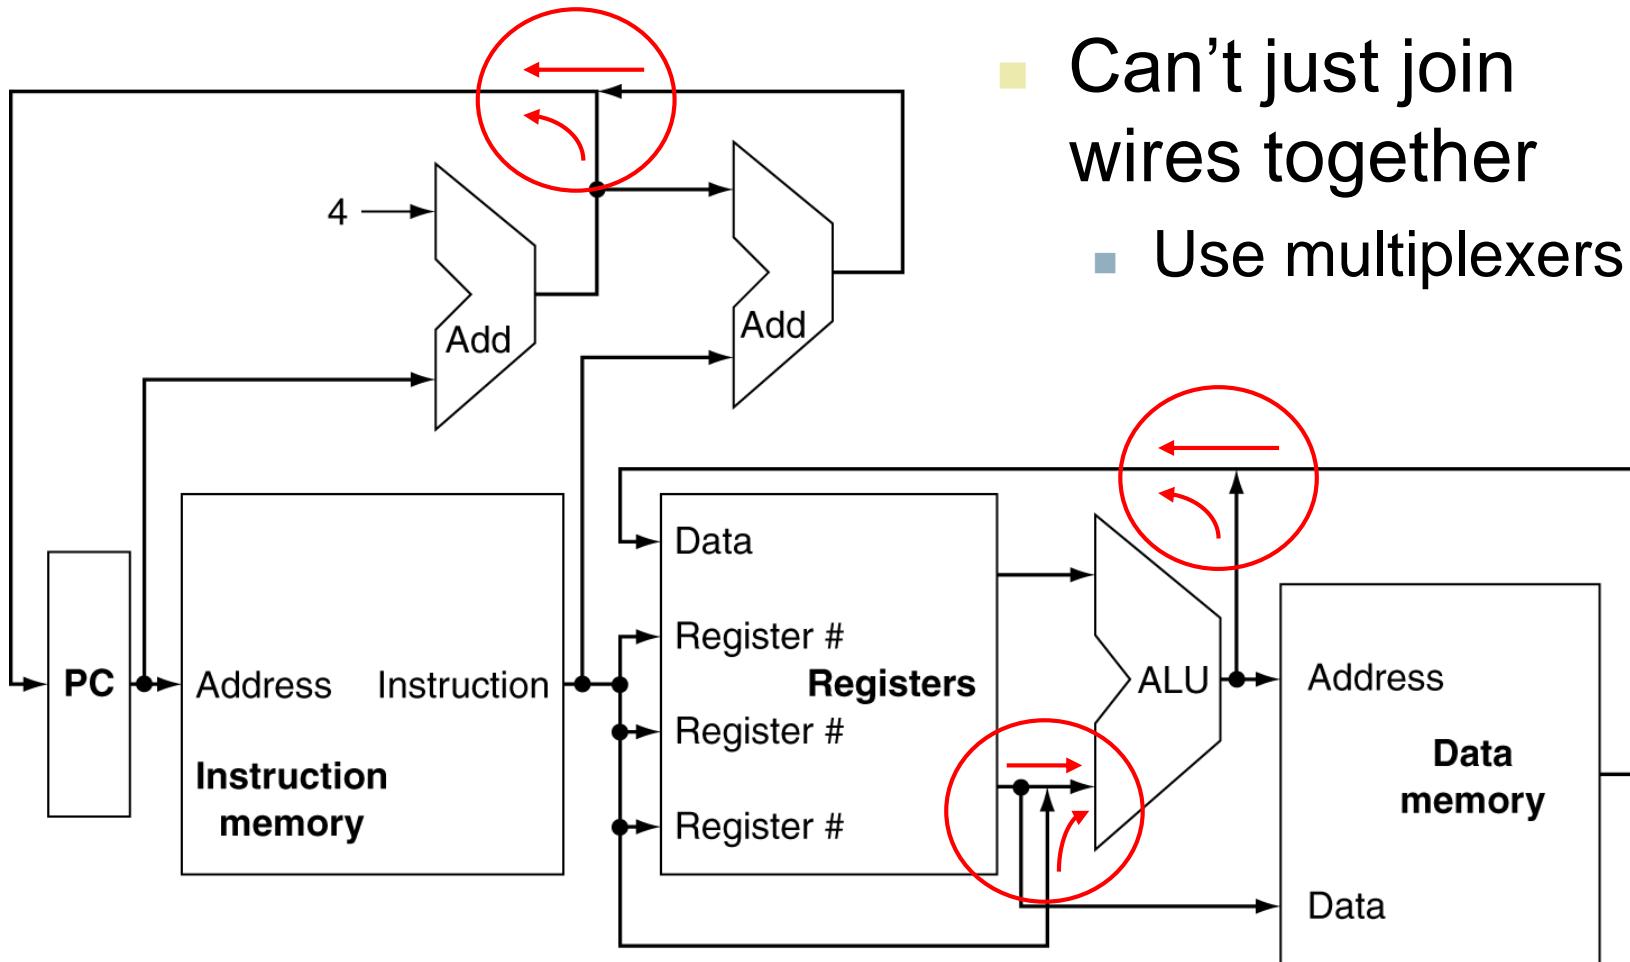
- PC → instruction memory, fetch instruction
- Register numbers → register file, read registers
- Depending on instruction class
  - Use ALU to calculate
    - Arithmetic/logic result
    - Memory address for load/store
    - Branch target address
  - Access data memory for load/store
  - PC ← target address or PC + 4

# CPU Overview

- The datapath

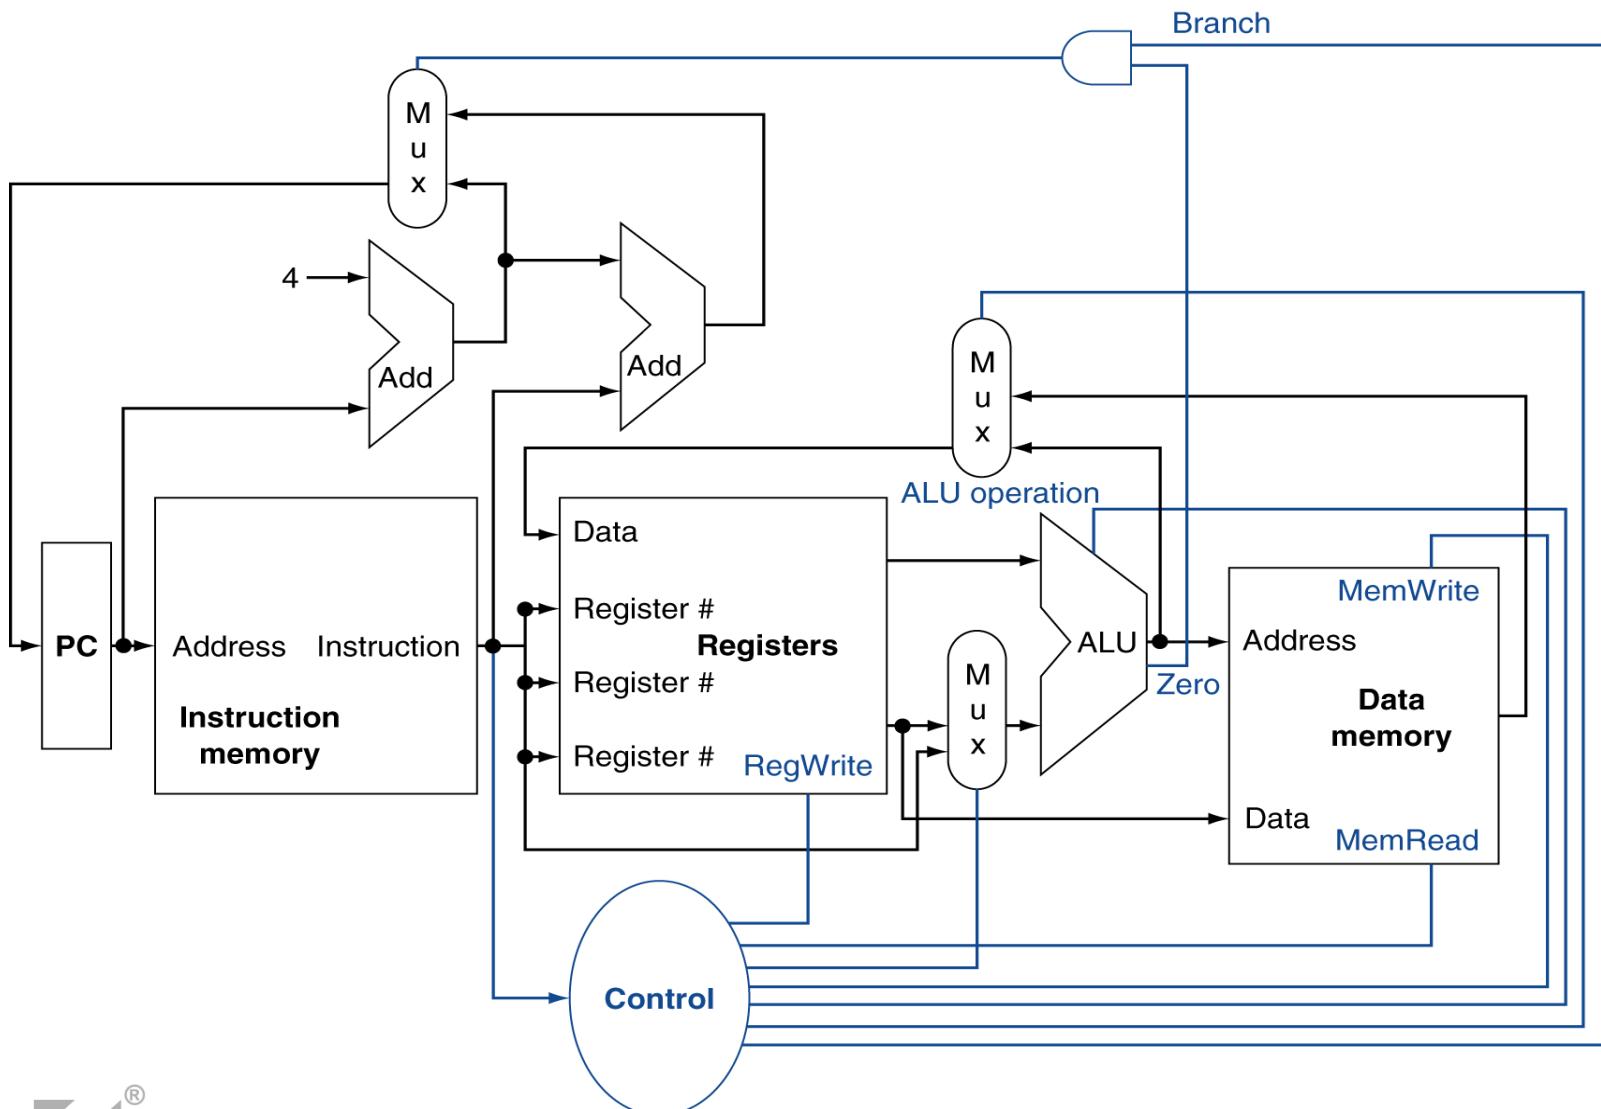


# Multiplexers



- Can't just join wires together
  - Use multiplexers

# Control

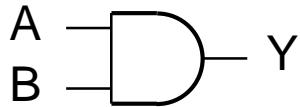


# Logic Design Basics

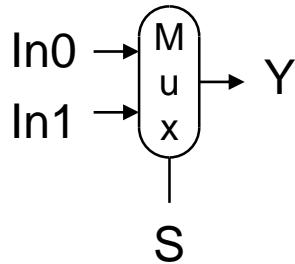
- Information encoded in binary
  - Low voltage = 0, High voltage = 1
  - One wire per bit
  - Multi-bit data encoded on multi-wire buses
- Combinational element
  - Operate on data
  - Output is a function of input
- State (sequential) elements
  - Store information

# Combinational Elements

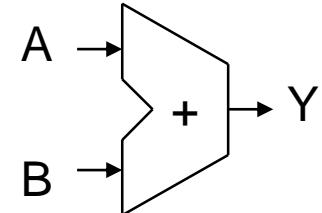
- AND-gate
  - $Y = A \& B$



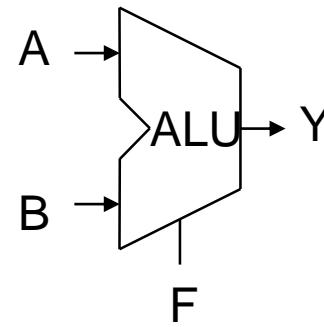
- Multiplexer
  - $Y = S ? In1 : In0$



- Adder
  - $Y = A + B$

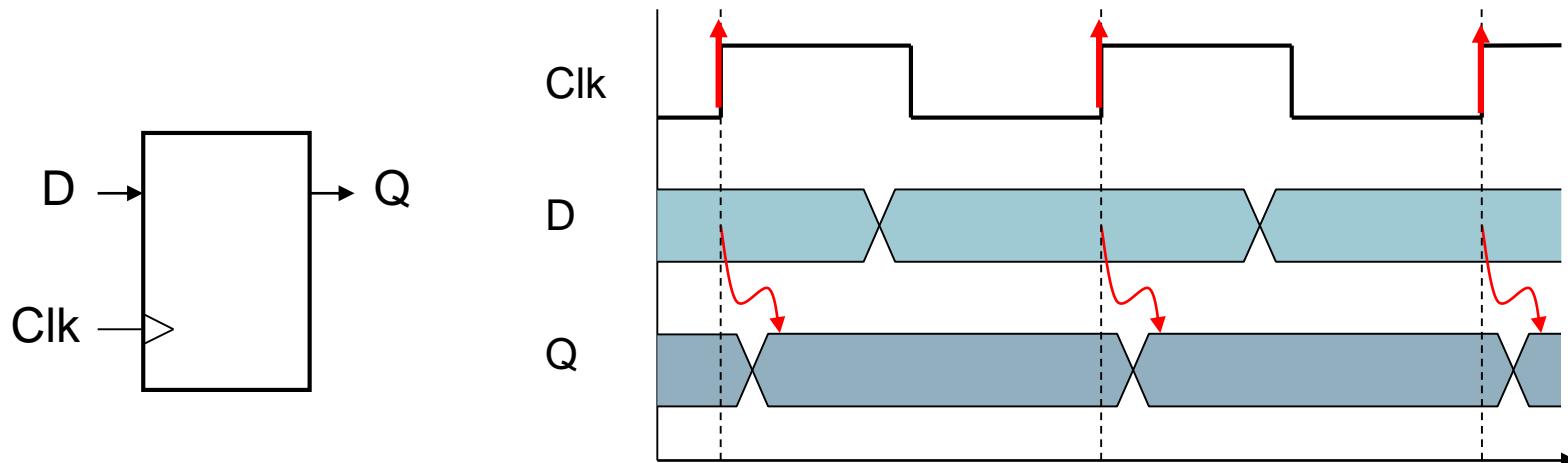


- Arithmetic/Logic Unit
  - $Y = F(A, B)$



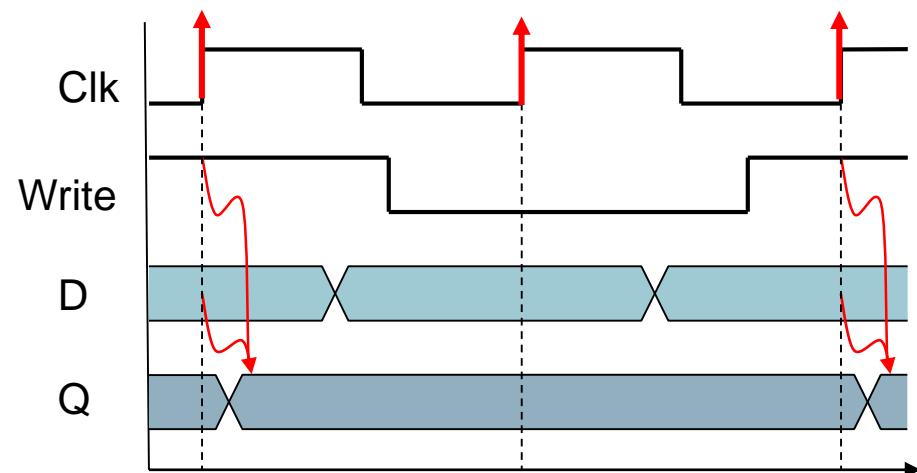
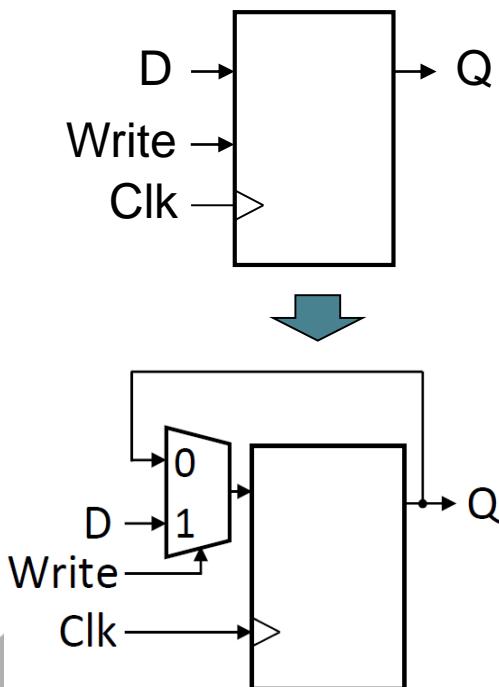
# Sequential Elements

- Register: stores data in a circuit
  - Uses a clock signal to determine when to update the stored value
  - Edge-triggered: update when Clk changes from 0 to 1



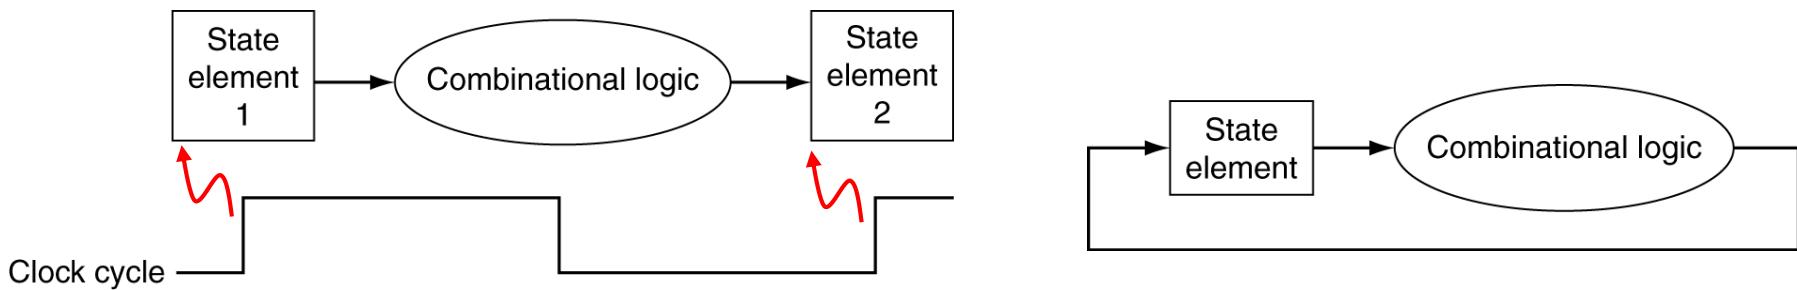
# Sequential Elements

- Register with **write** control
  - Only updates on **clock edge** when **write** control input is 1
  - Used when stored value is required later



# Clocking Methodology

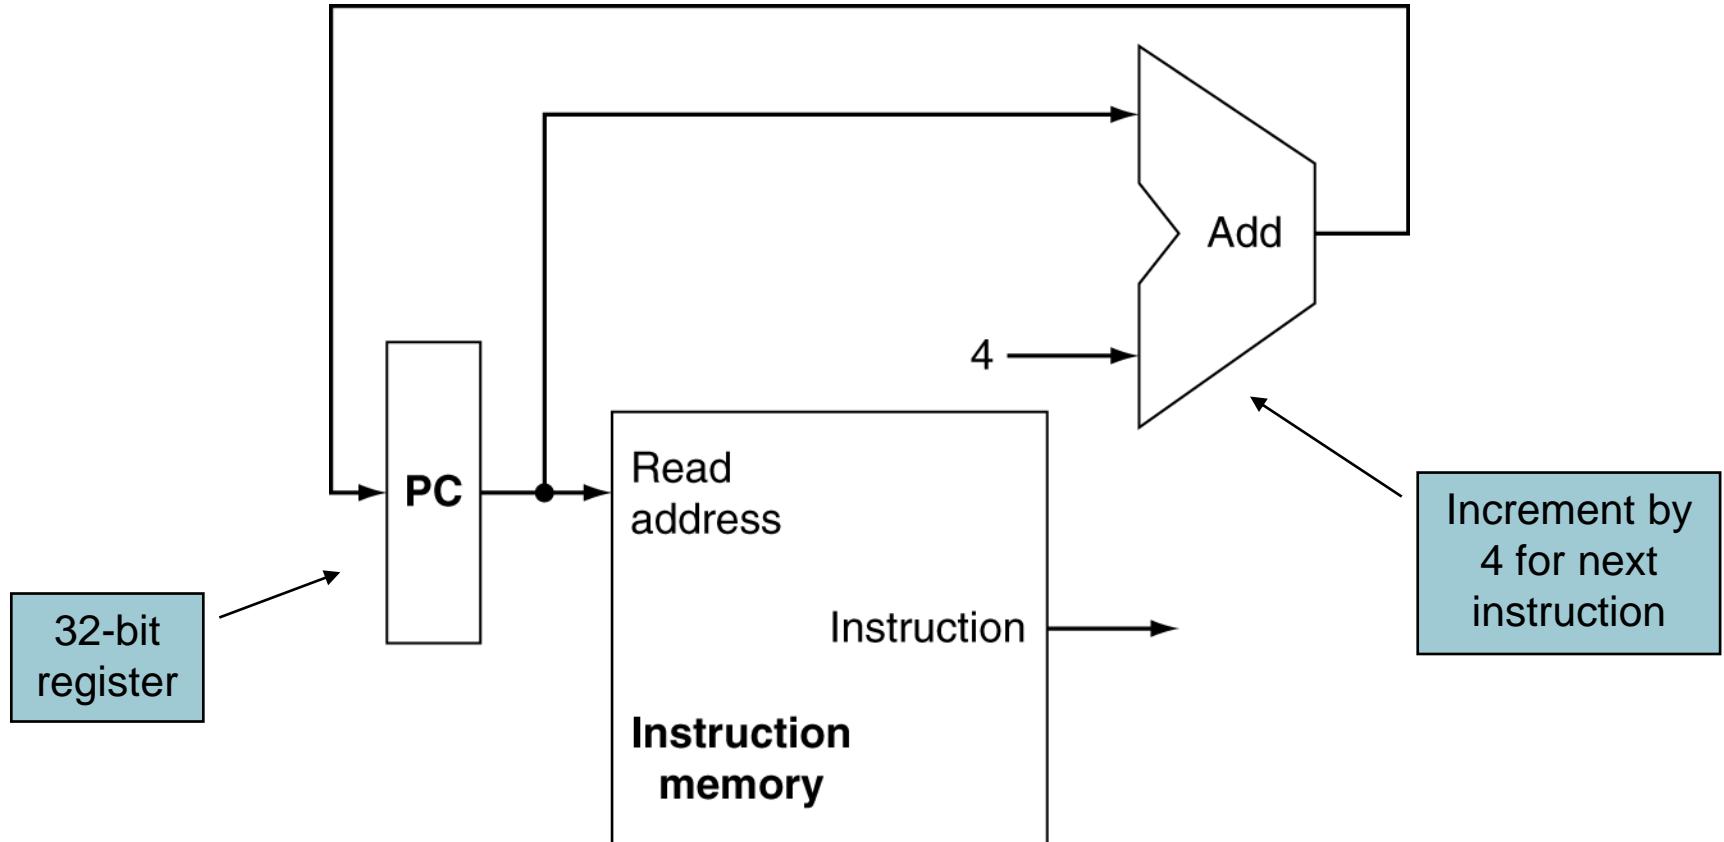
- Combinational logic transforms data during clock cycles
  - Between clock edges
  - Input from state elements, output to state element
  - **Longest delay determines clock period**



# Building a Datapath

- Datapath
  - Elements that process data and addresses in the CPU
    - Registers, ALUs, mux's, memories, ...
- We will build a MIPS datapath incrementally
  - Refining the overview design

# Instruction Fetch

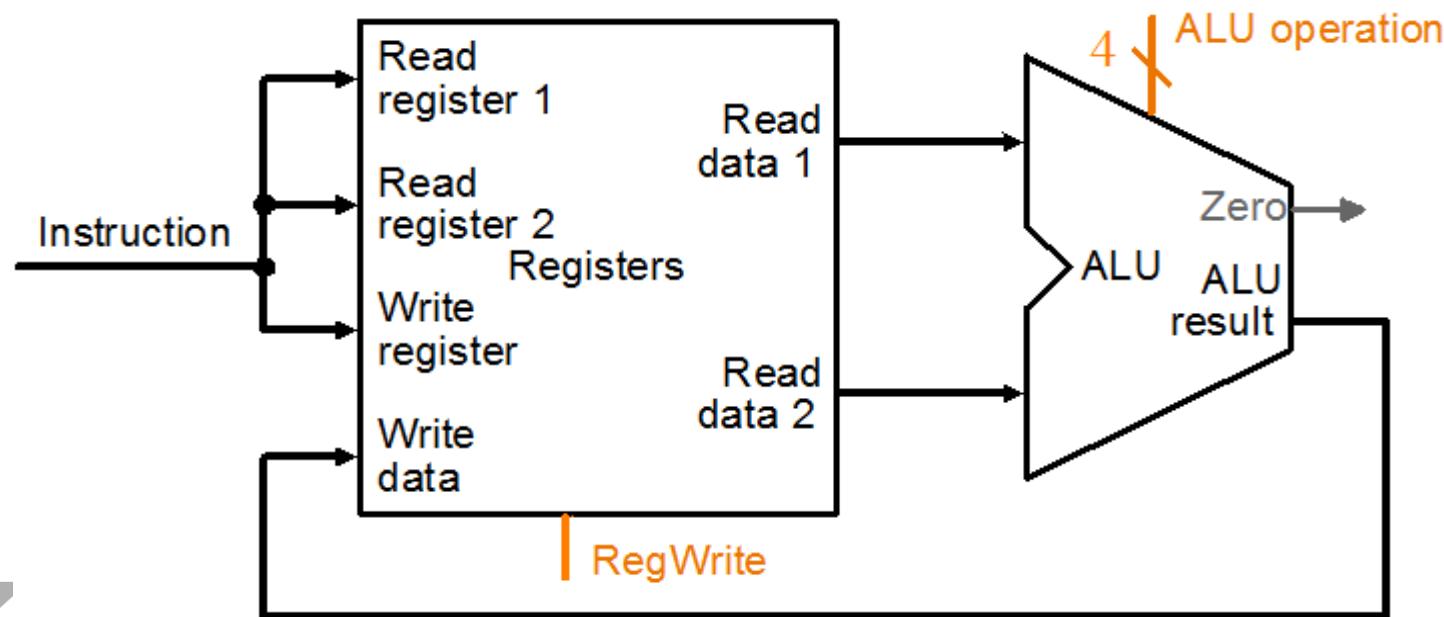


# R-Format Instructions

- Read two register operands
- Perform arithmetic/logical operation
- Write register result

add \$rd, \$rs, \$rt

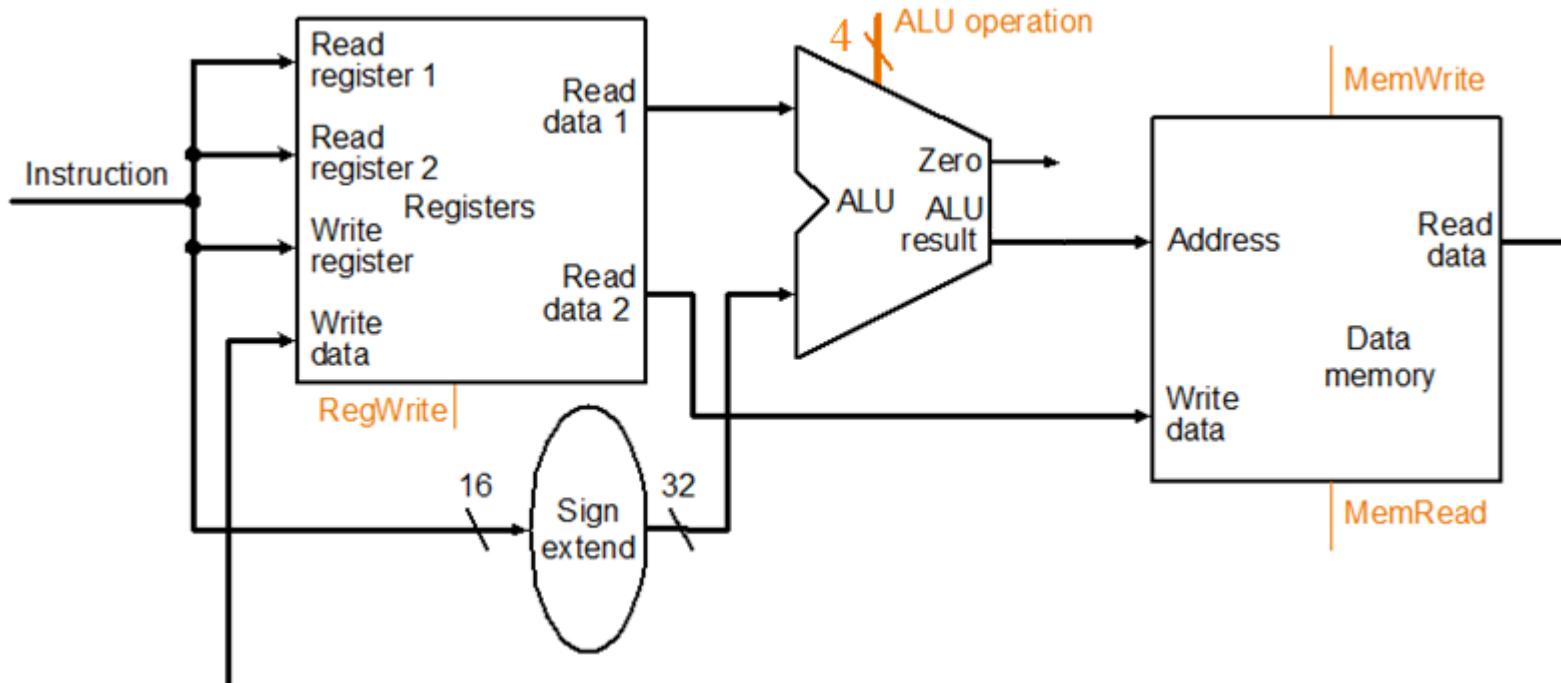
|    |    |    |    |       |       |
|----|----|----|----|-------|-------|
| op | rs | rt | rd | shamt | funct |
|----|----|----|----|-------|-------|



# Load/Store Instructions

- Read register operands
- Calculate address using 16-bit offset
  - Use ALU, but sign-extend offset
- Load: Read memory and update register
- Store: Write register value to memory

**lw \$rt, offset(\$rs)**  
**sw \$rt, offset(\$rs)**

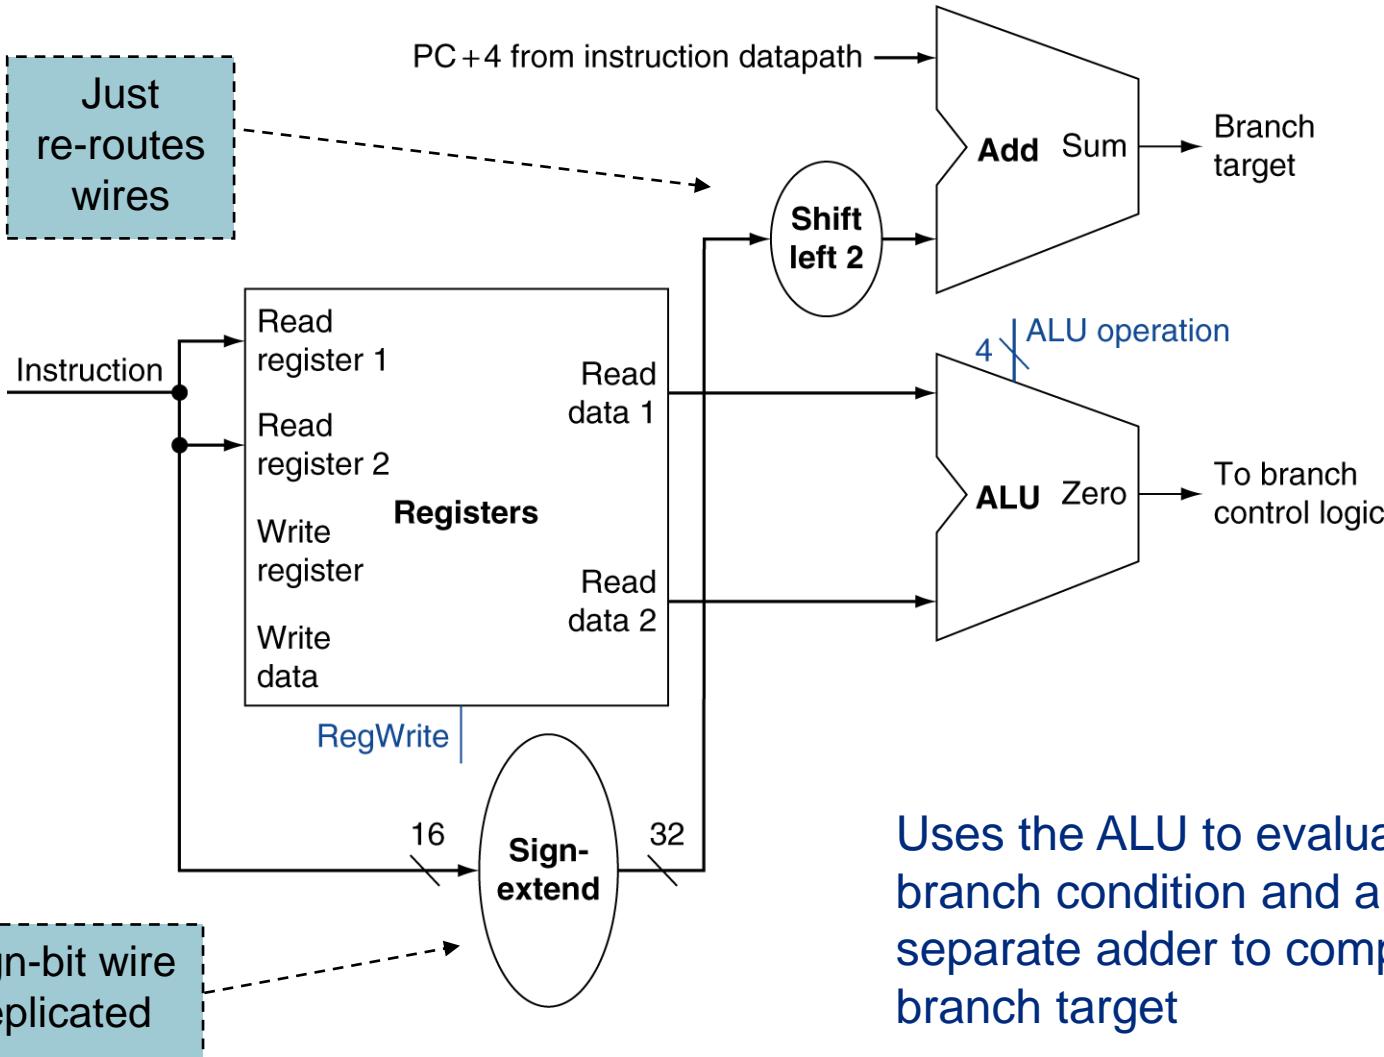


# Branch Instructions

- Read register operands
- Compare operands
  - Use ALU, subtract and check Zero output
- Calculate target address
  - Sign-extend displacement
  - Shift left 2 places (word displacement)
  - Add to PC + 4
    - Already calculated by instruction fetch

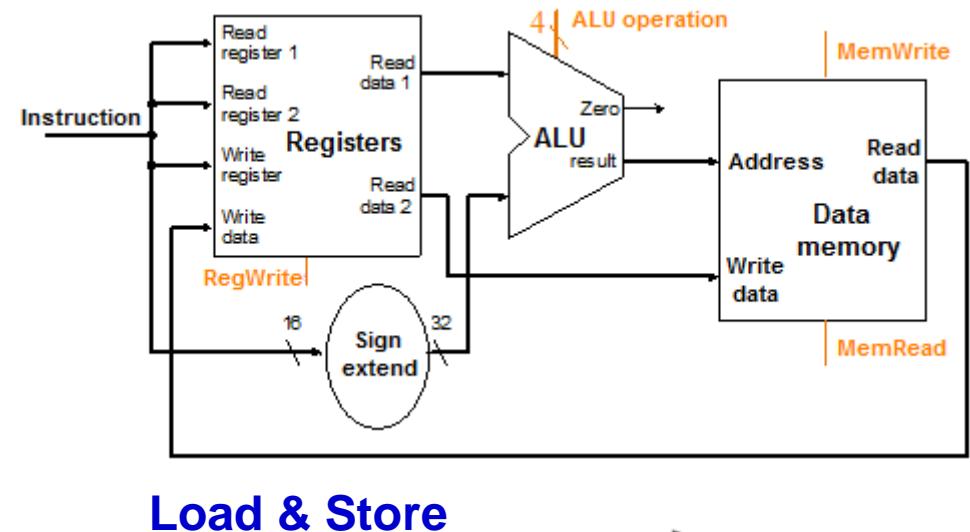
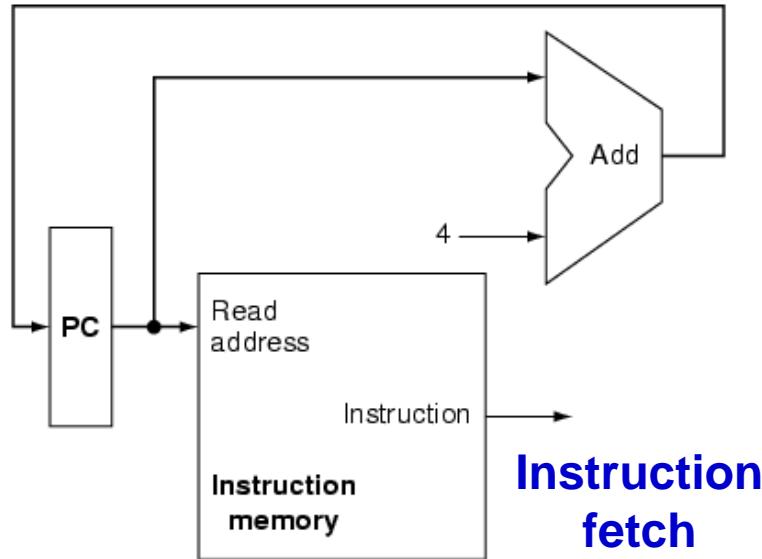
**Beq \$rs, \$rt, L**

# Branch Instructions



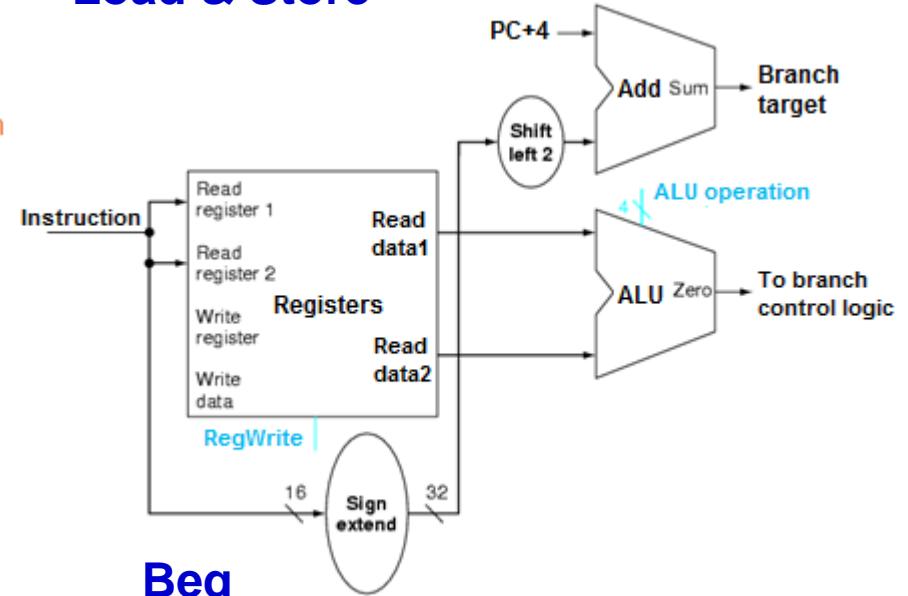
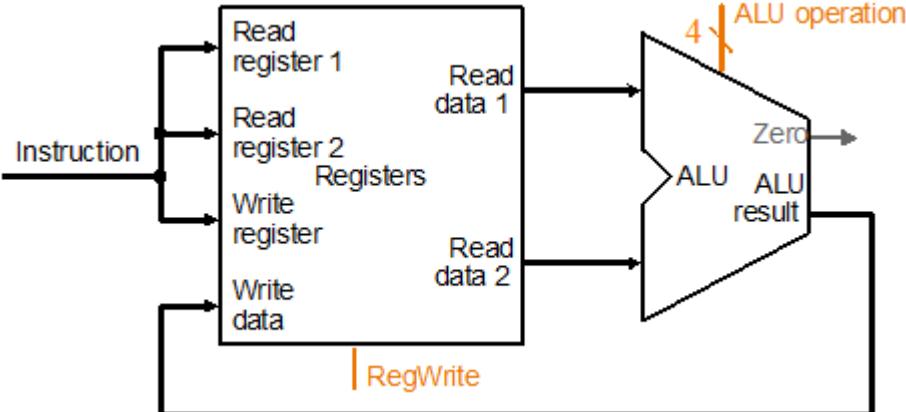
Uses the ALU to evaluate the branch condition and a separate adder to compute the branch target

# How to combines these together?



**Instruction  
fetch**

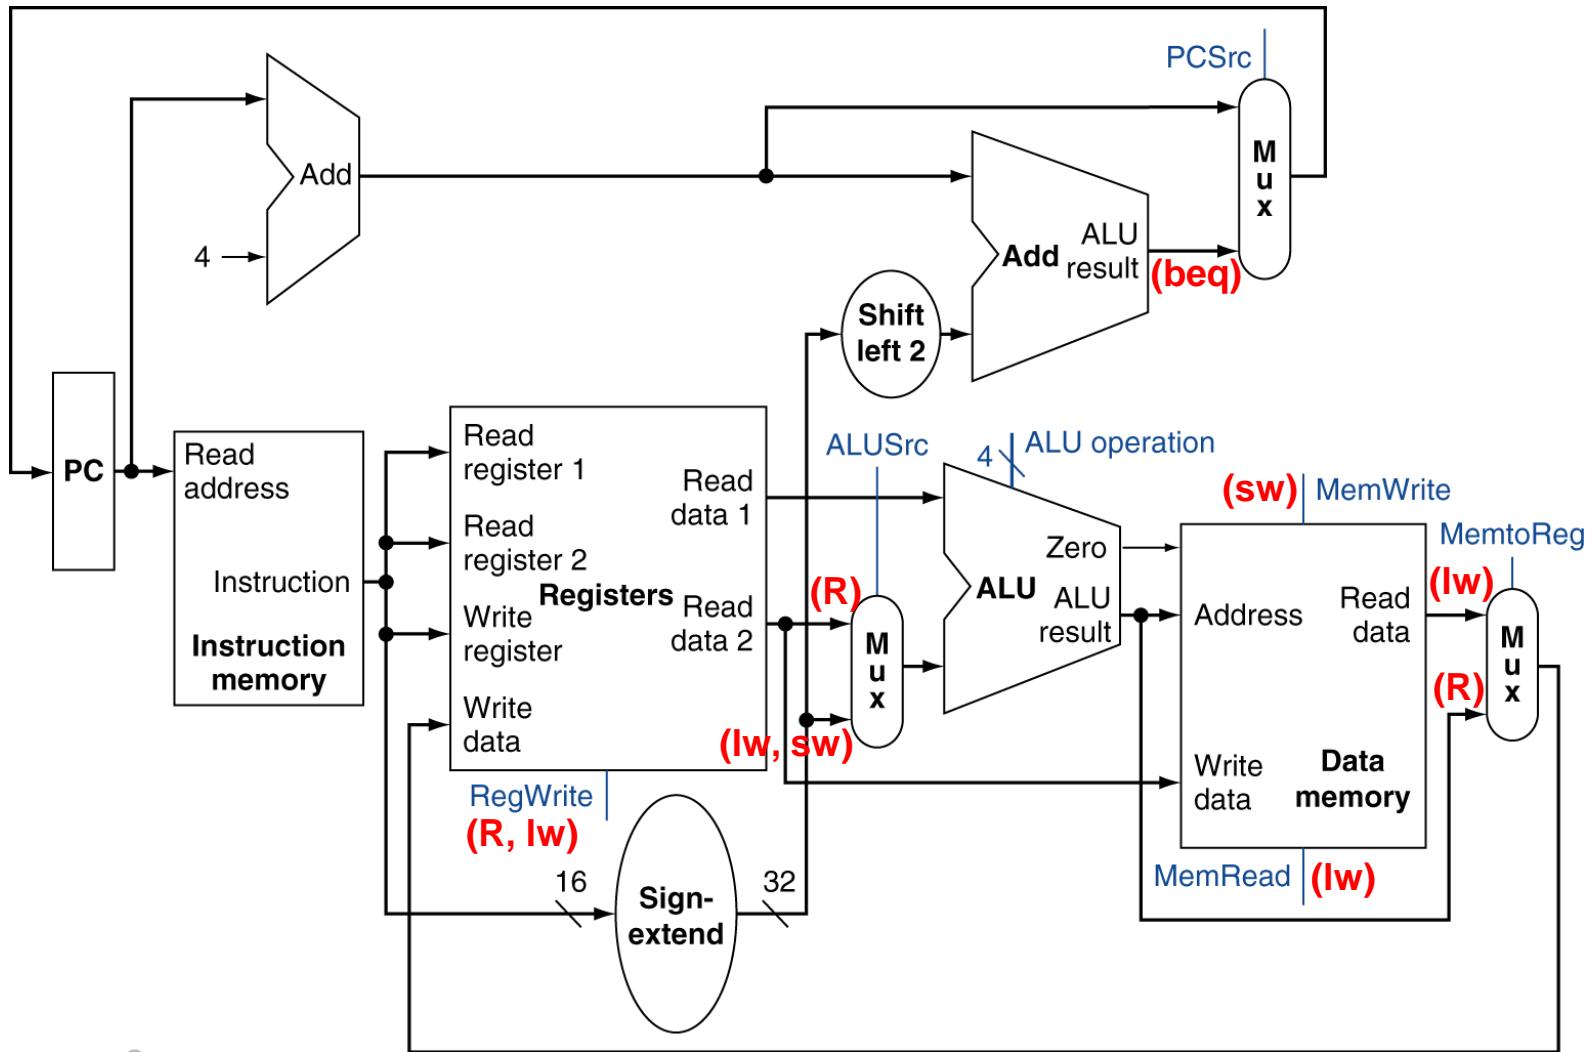
**Load & Store**



# Composing the Elements

- First-cut data path does an instruction in one clock cycle
  - Each datapath element can only do one function at a time
  - Hence, we need separate instruction and data memories
- Use multiplexers where alternate data sources are used for different instructions

# The Datapath with MUXs and Control Signals



# Difference in the write-register field

| R-type | 0     | rs    | rt    | rd    | shamt | funct |
|--------|-------|-------|-------|-------|-------|-------|
|        | 31:26 | 25:21 | 20:16 | 15:11 | 10:6  | 5:0   |

| Load/<br>Store | 35 or 43 | rs    | rt    | address |
|----------------|----------|-------|-------|---------|
|                | 31:26    | 25:21 | 20:16 | 15:0    |

**add \$t0, \$t1, \$t2**

rd    rs    rt

**lw \$t0, 4(\$t1)**

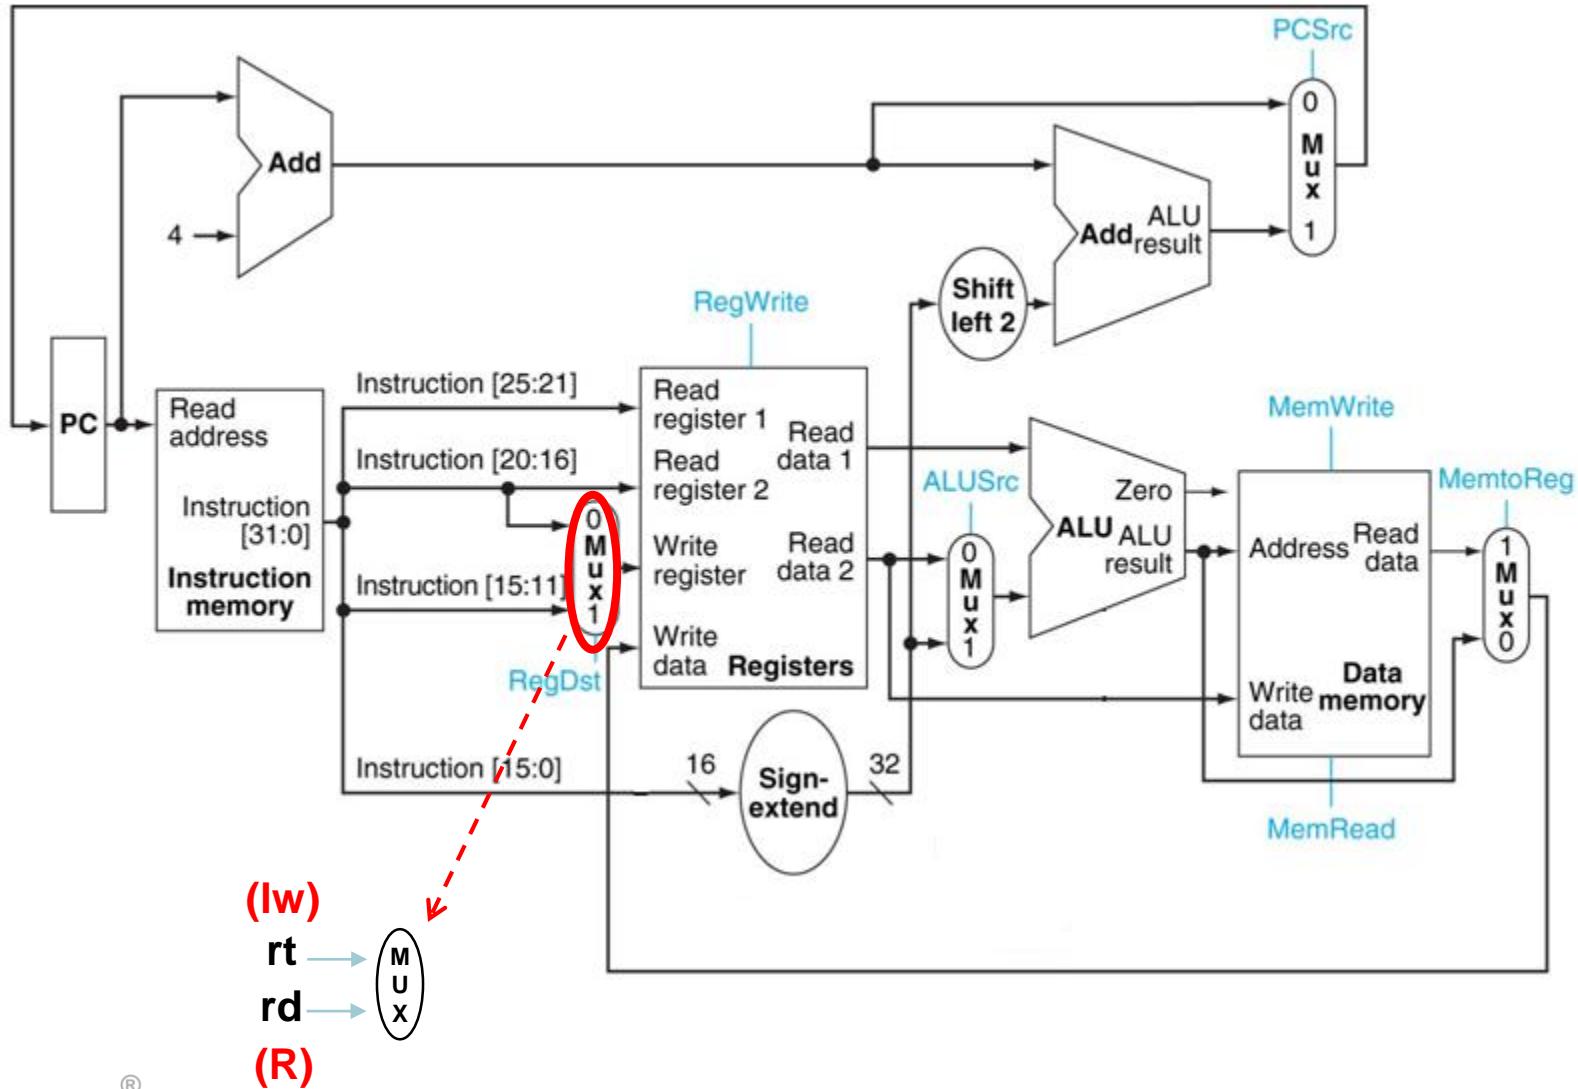
rt    rs

**beq \$t0, \$t1, L**

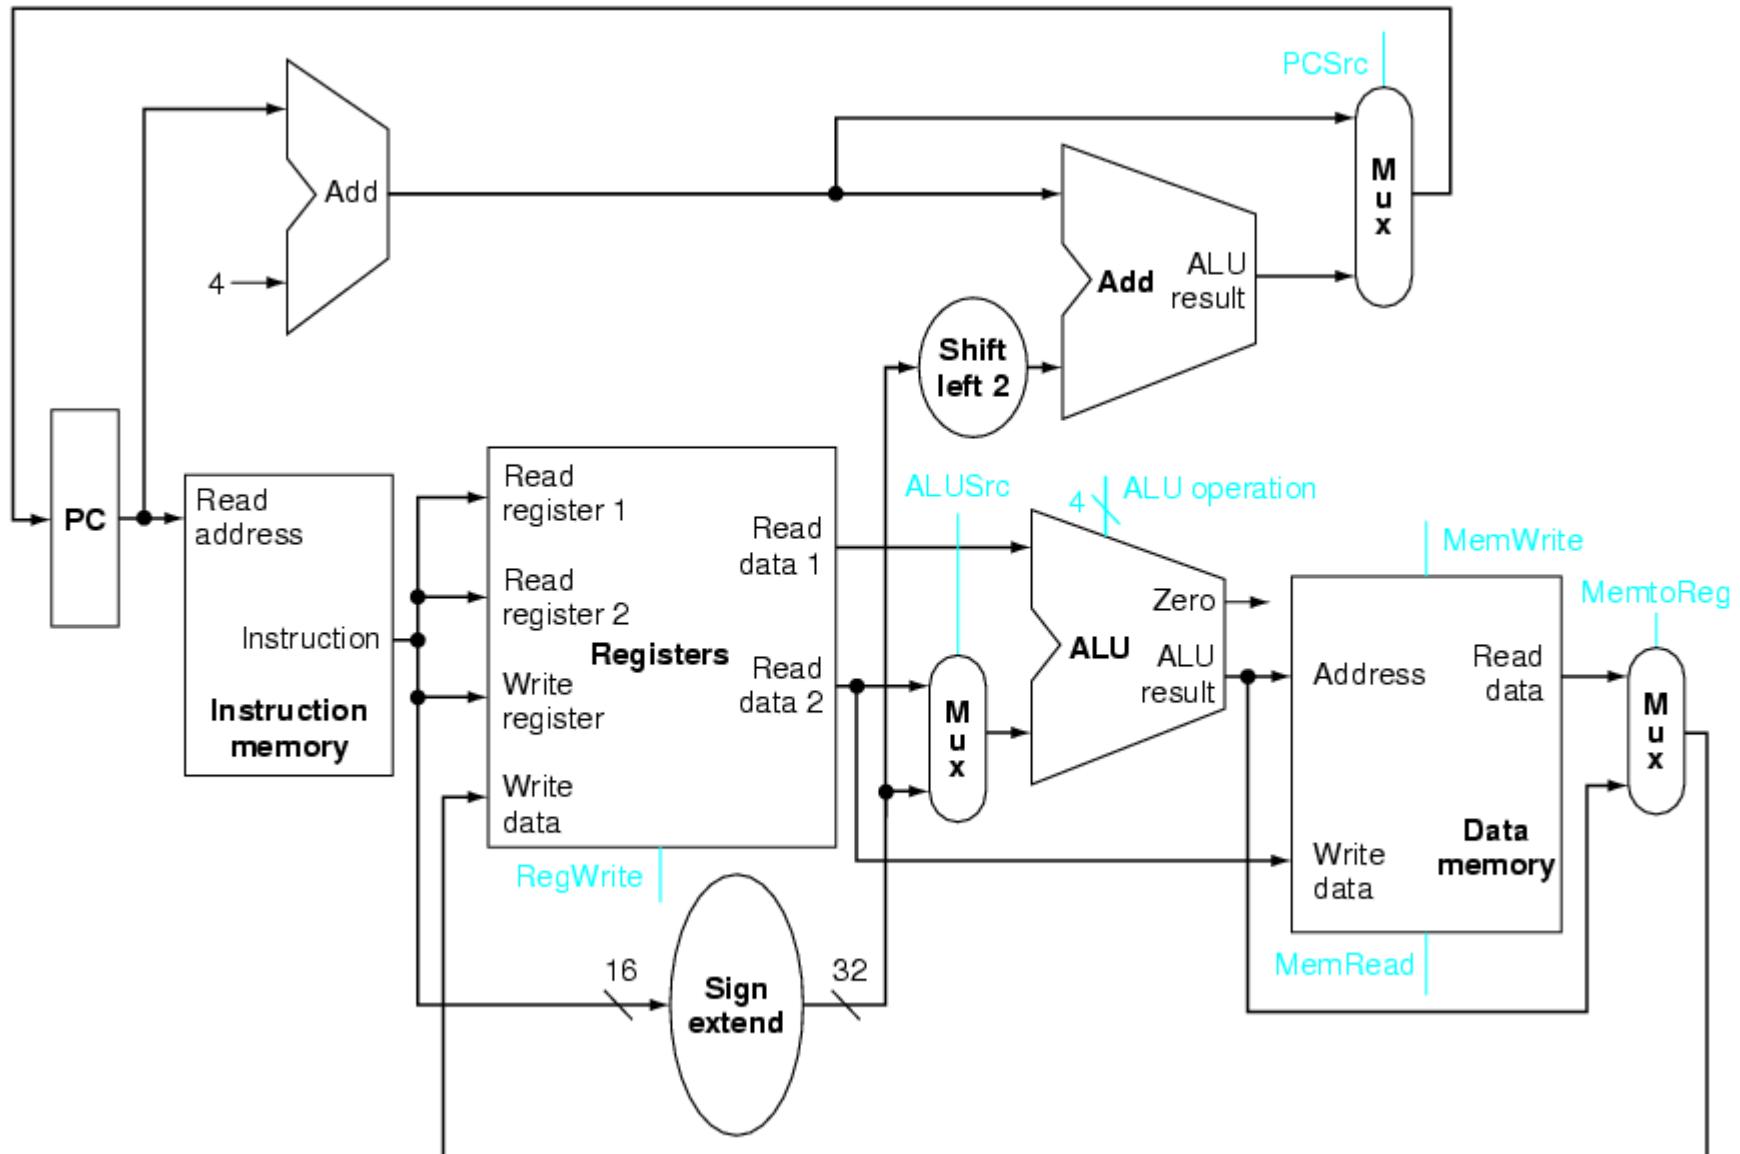
rs    rt

write for  
R-type  
and load

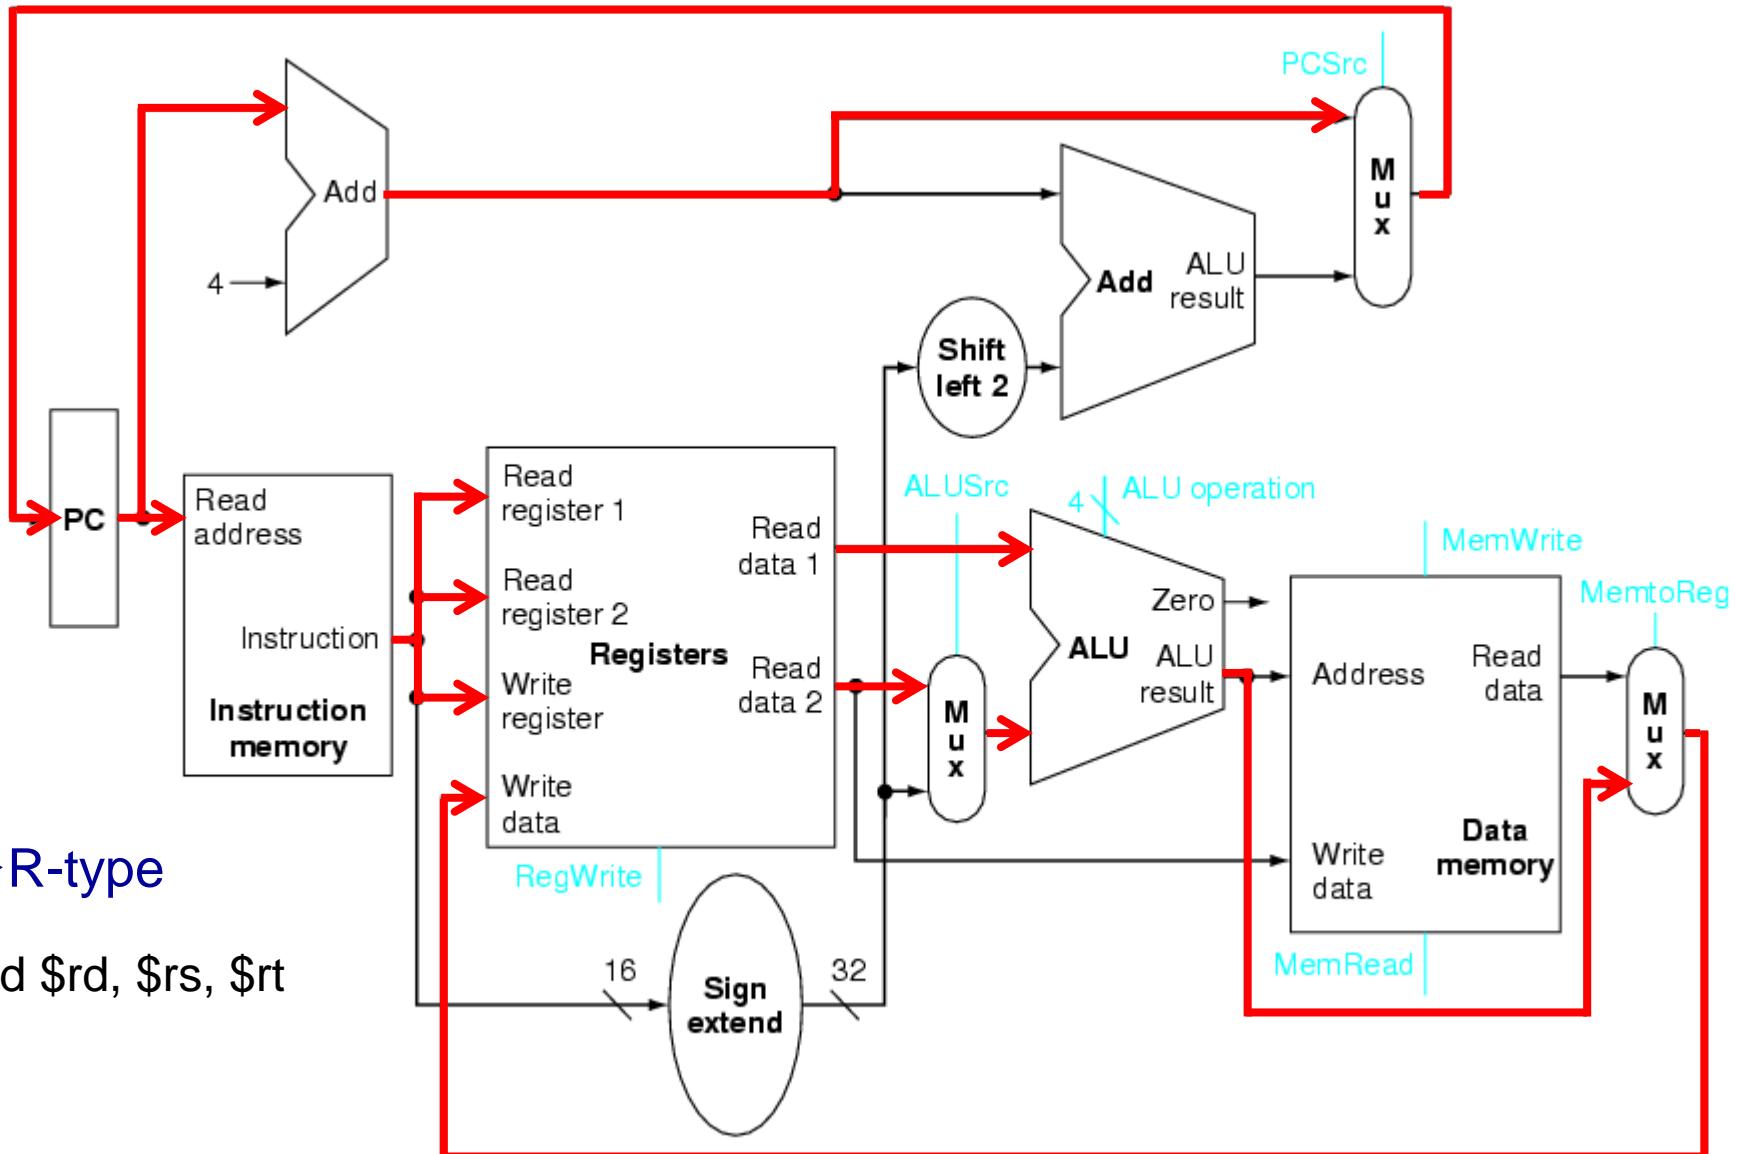
# Full Datapath



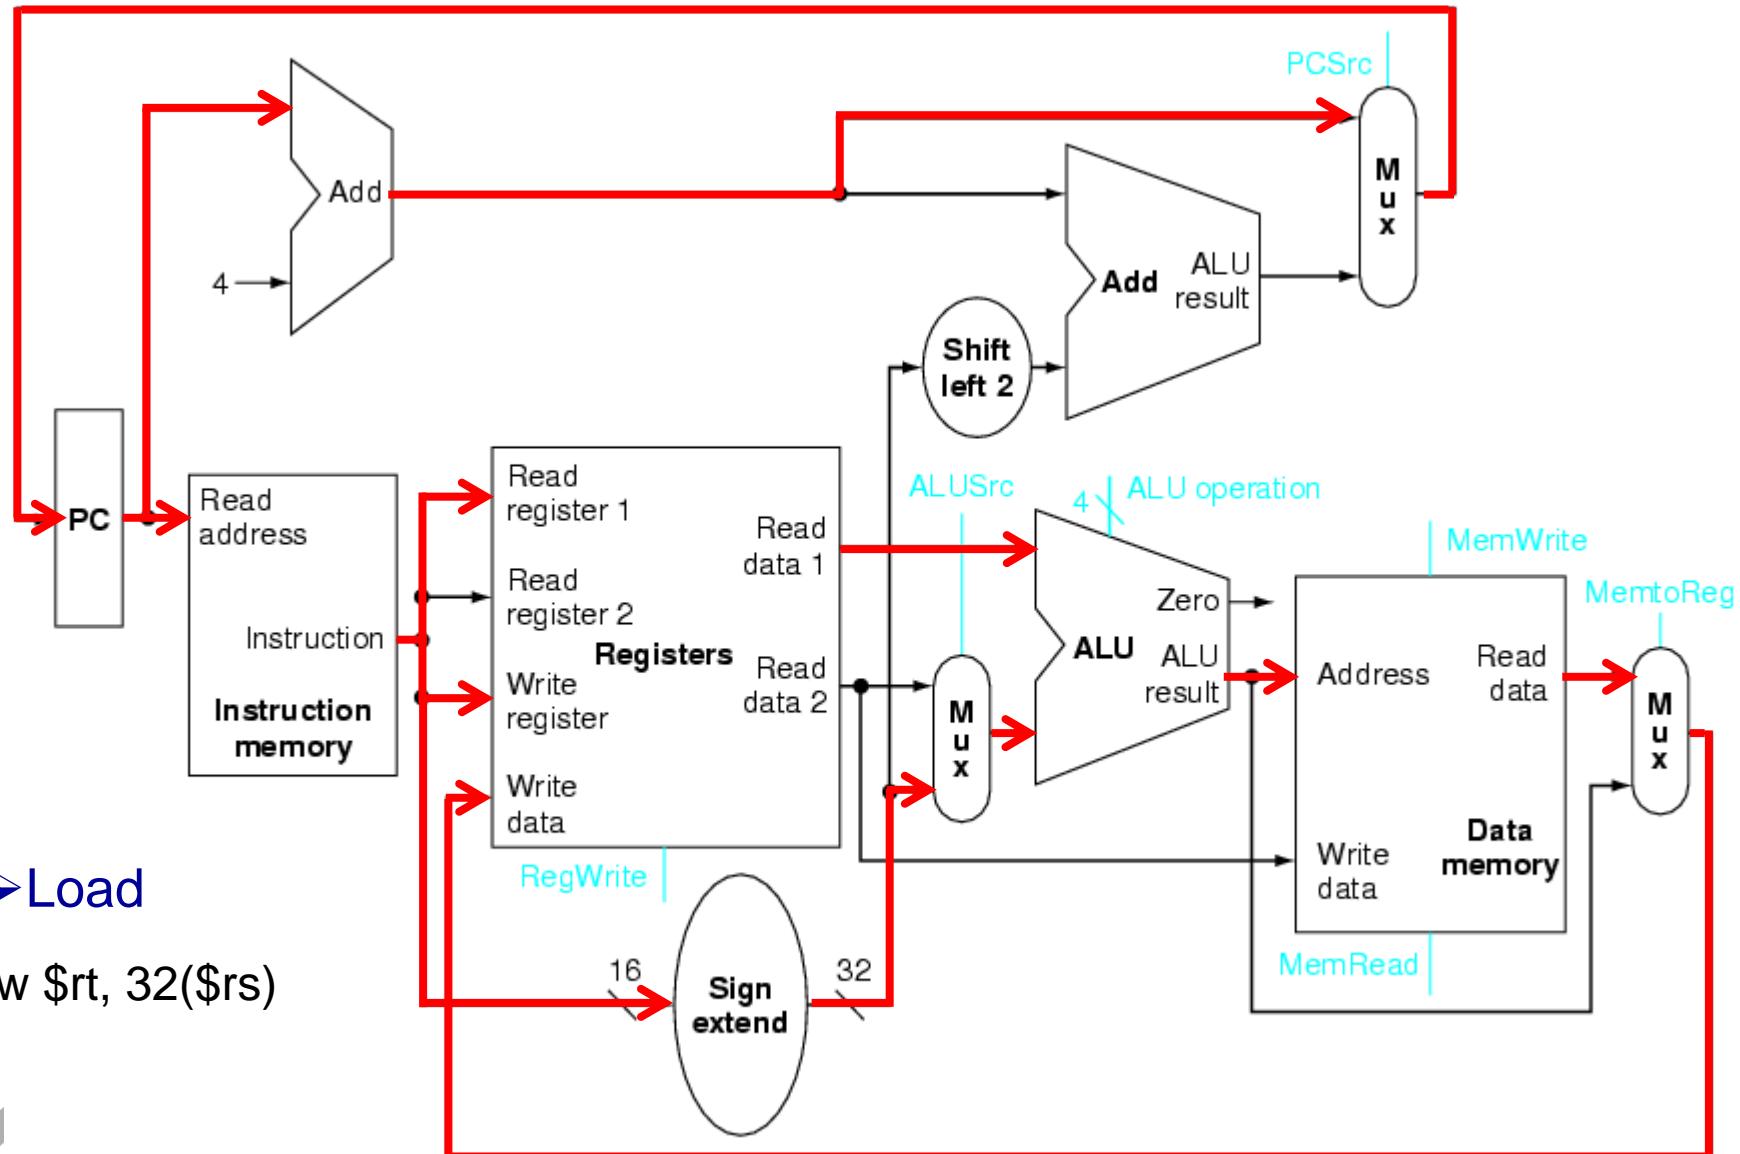
# A Simple Datapath for MIPS Inst. Set



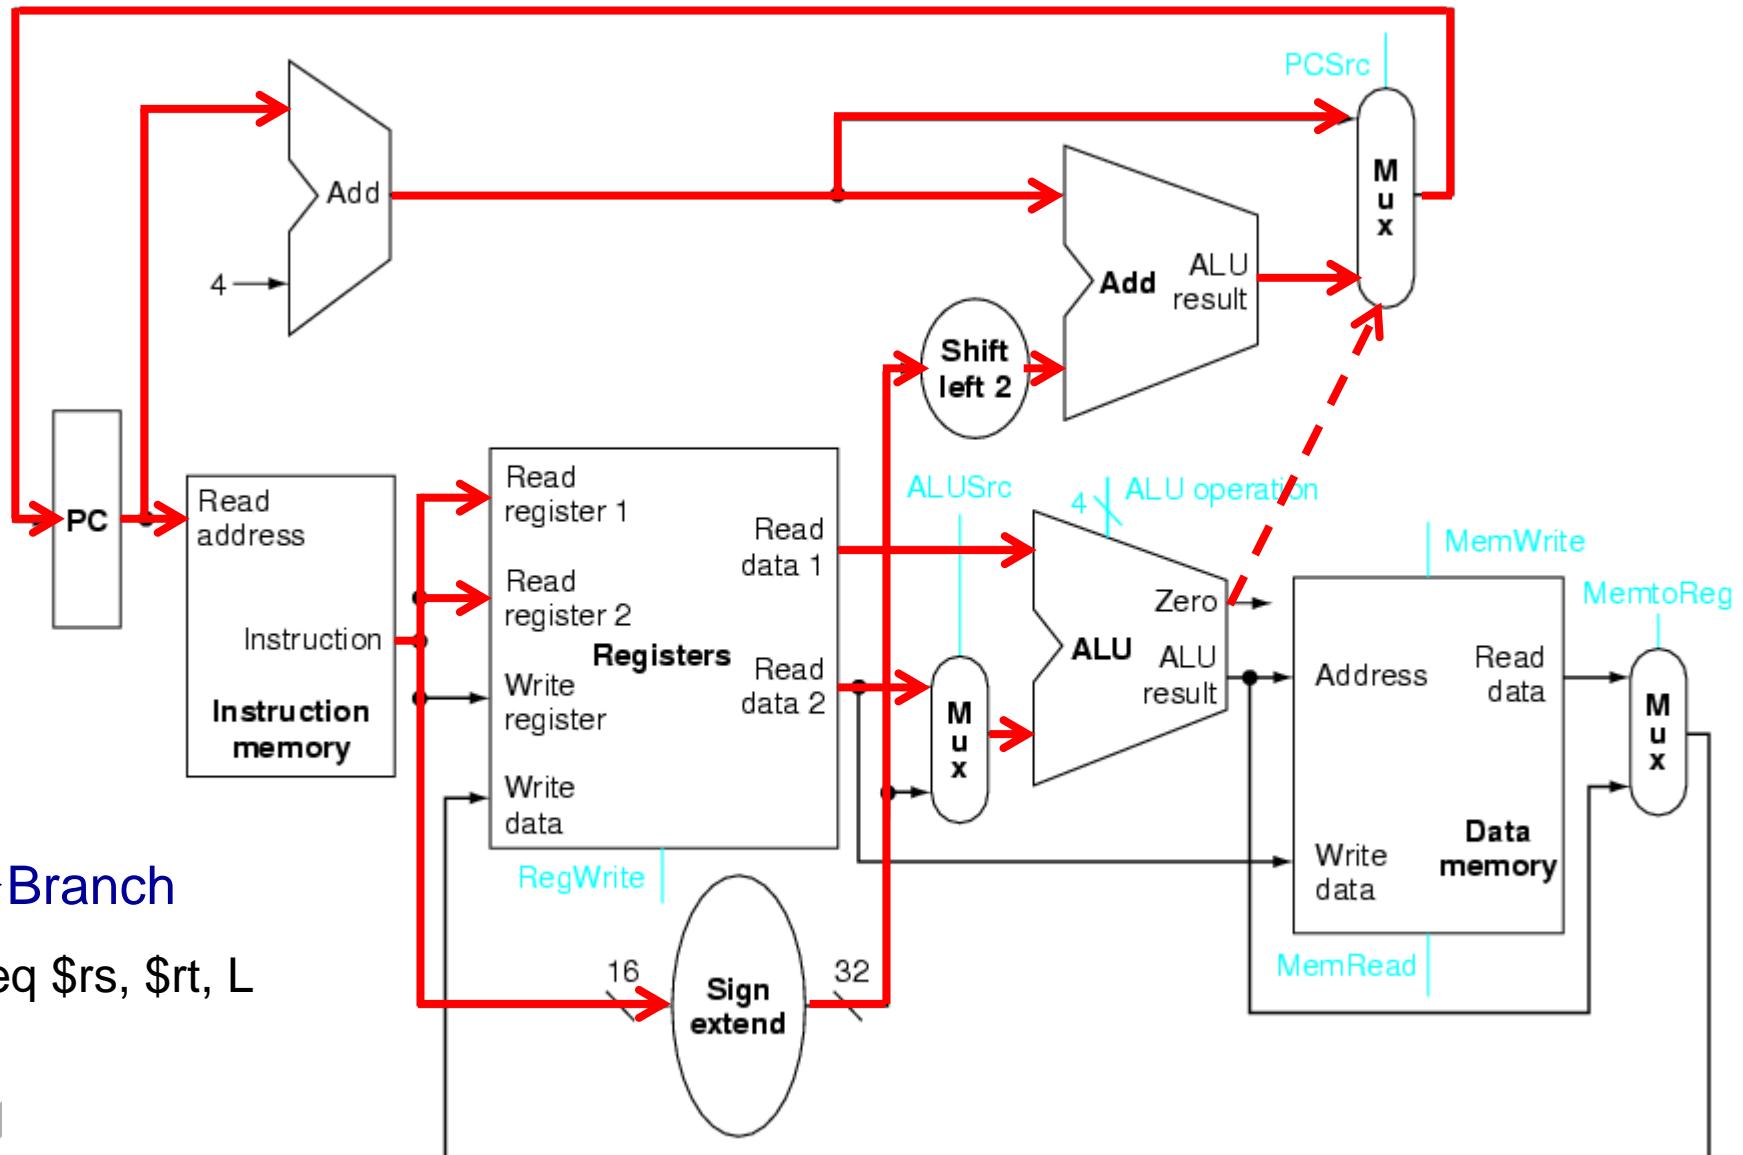
# A Simple Datapath for MIPS Inst. Set



# A Simple Datapath for MIPS Inst. Set



# A Simple Datapath for MIPS Inst. Set



➤ Branch

`beq $rs, $rt, L`



# ALU Control

- ALU used for
  - Load/Store:  $F = \text{add}$
  - Branch:  $F = \text{subtract}$
  - R-type:  $F$  depends on funct field

| ALU control | Function               |
|-------------|------------------------|
| 0000        | AND                    |
| 0001        | OR                     |
| 0010        | add                    |
| 0110        | subtract               |
| 0111        | set-on-less-than (slt) |
| 1100        | NOR                    |

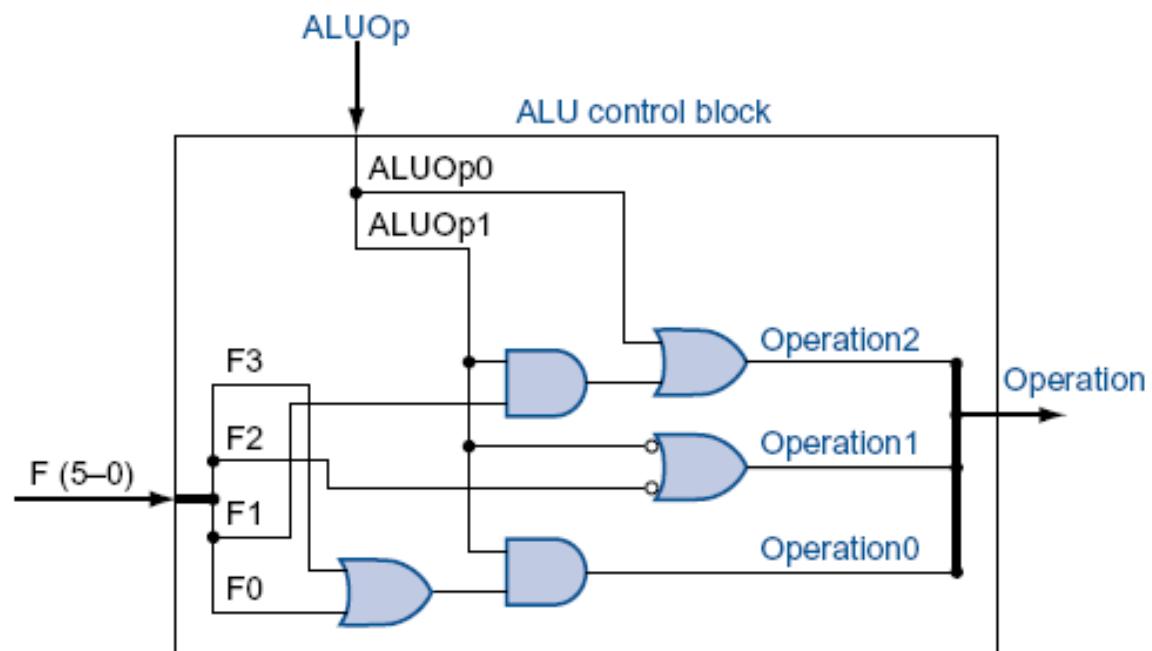
# ALU Control

- Assume 2-bit ALUOp derived from opcode
  - Combinational logic derives ALU control

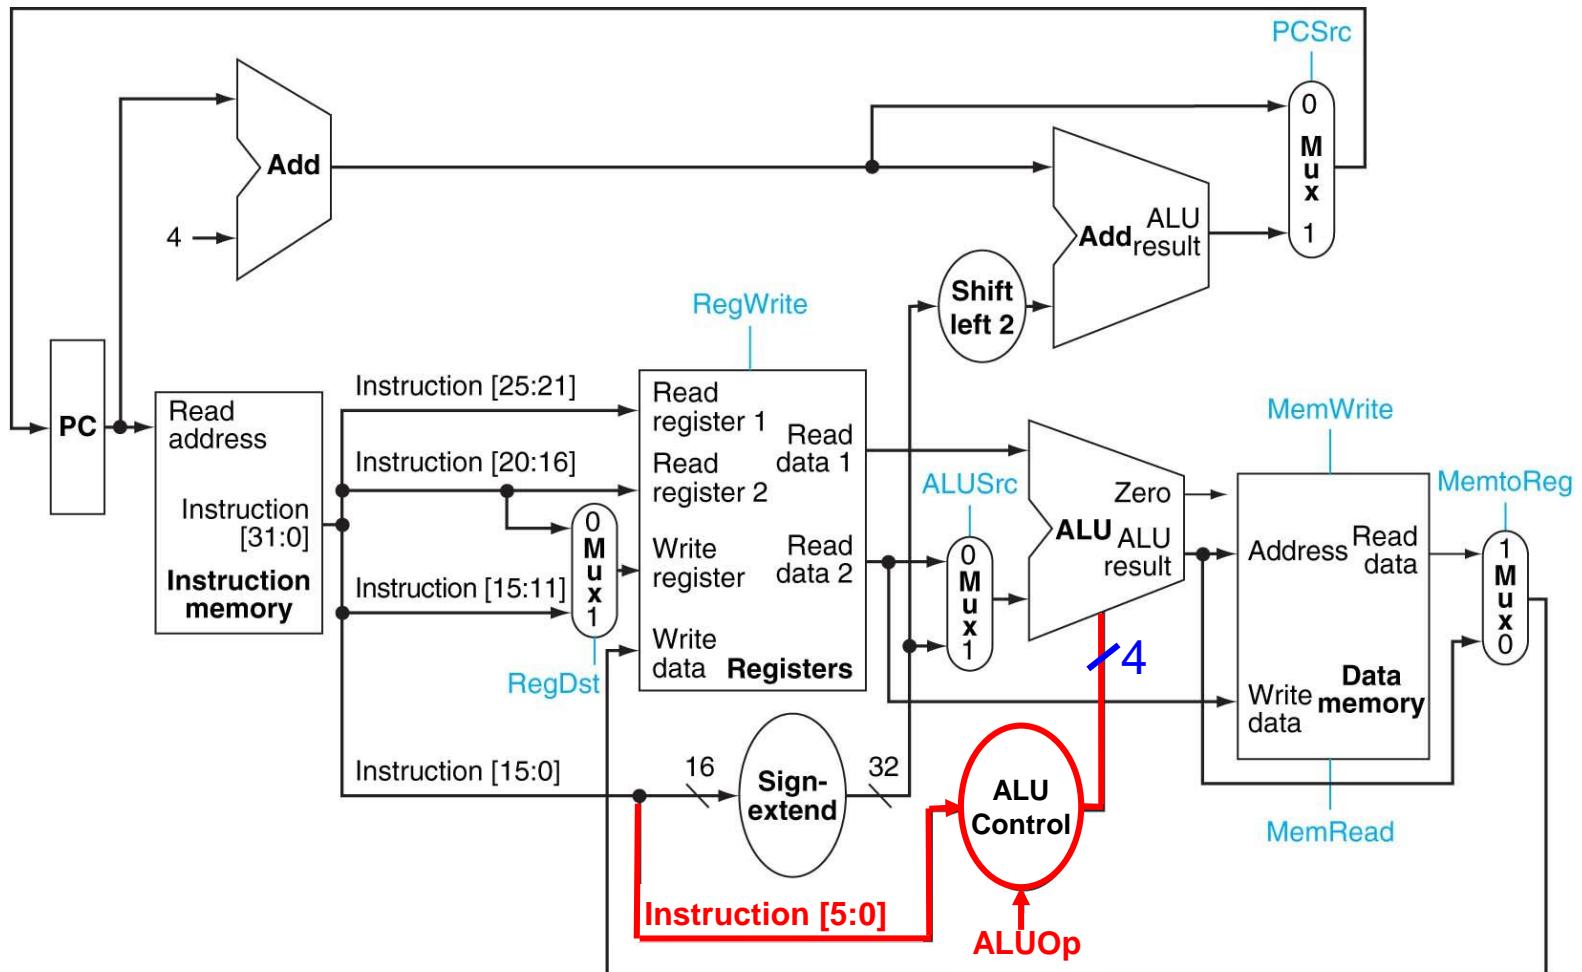
| opcode | ALUOp | Operation        | funct  | ALU function     | ALU control |
|--------|-------|------------------|--------|------------------|-------------|
| lw     | 00    | load word        | XXXXXX | add              | 0010        |
| sw     | 00    | store word       | XXXXXX | add              | 0010        |
| beq    | 01    | branch equal     | XXXXXX | subtract         | 0110        |
| R-type | 10    | add              | 100000 | add              | 0010        |
|        |       | subtract         | 100010 | subtract         | 0110        |
|        |       | AND              | 100100 | AND              | 0000        |
|        |       | OR               | 100101 | OR               | 0001        |
|        |       | set-on-less-than | 101010 | set-on-less-than | 0111        |

# ALU Control (Appendix)

| ALUOp  |        | Funct field |    |    |    |    |    |      | Operation<br>$C_3\ C_2\ C_1\ C_0$ |
|--------|--------|-------------|----|----|----|----|----|------|-----------------------------------|
| ALUOp1 | ALUOp0 | F5          | F4 | F3 | F2 | F1 | F0 |      |                                   |
| 0      | 0      | X           | X  | X  | X  | X  | X  | 0010 | (lw/sw)                           |
| X      | 1      | X           | X  | X  | X  | X  | X  | 0110 | (beq)                             |
| 1      | X      | X           | X  | 0  | 0  | 0  | 0  | 0010 | (add)                             |
| 1      | X      | X           | X  | 0  | 0  | 1  | 0  | 0110 | (sub)                             |
| 1      | X      | X           | X  | 0  | 1  | 0  | 0  | 0000 | (AND)                             |
| 1      | X      | X           | X  | 0  | 1  | 0  | 1  | 0001 | (OR)                              |
| 1      | X      | X           | X  | 1  | 0  | 1  | 0  | 0111 | (slt)                             |

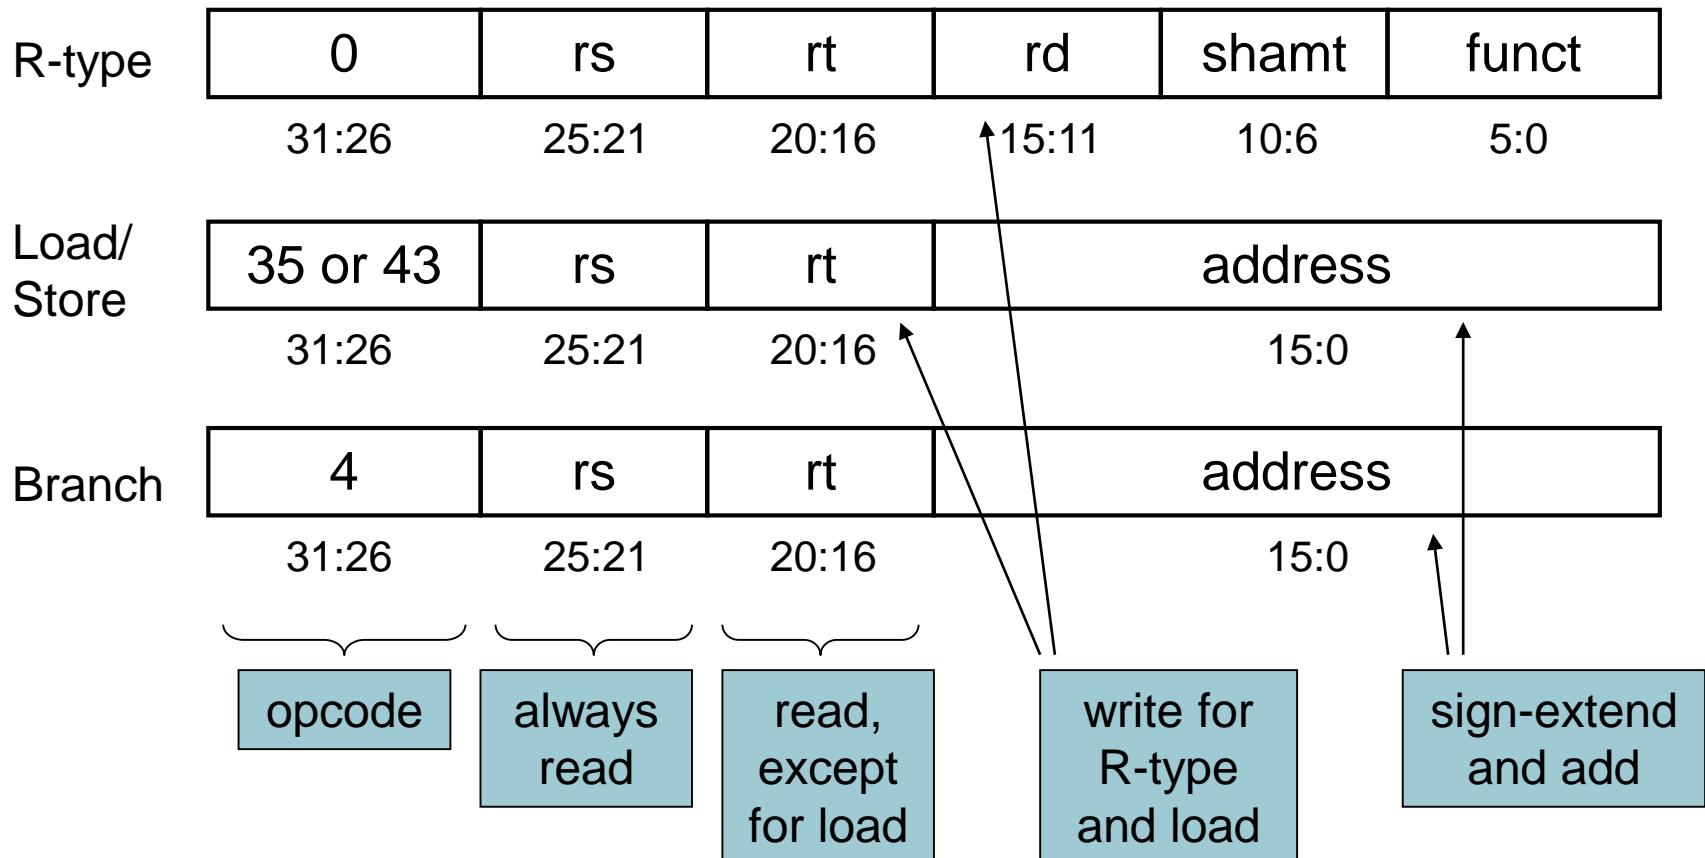


# The ALU Control



# The Main Control Unit

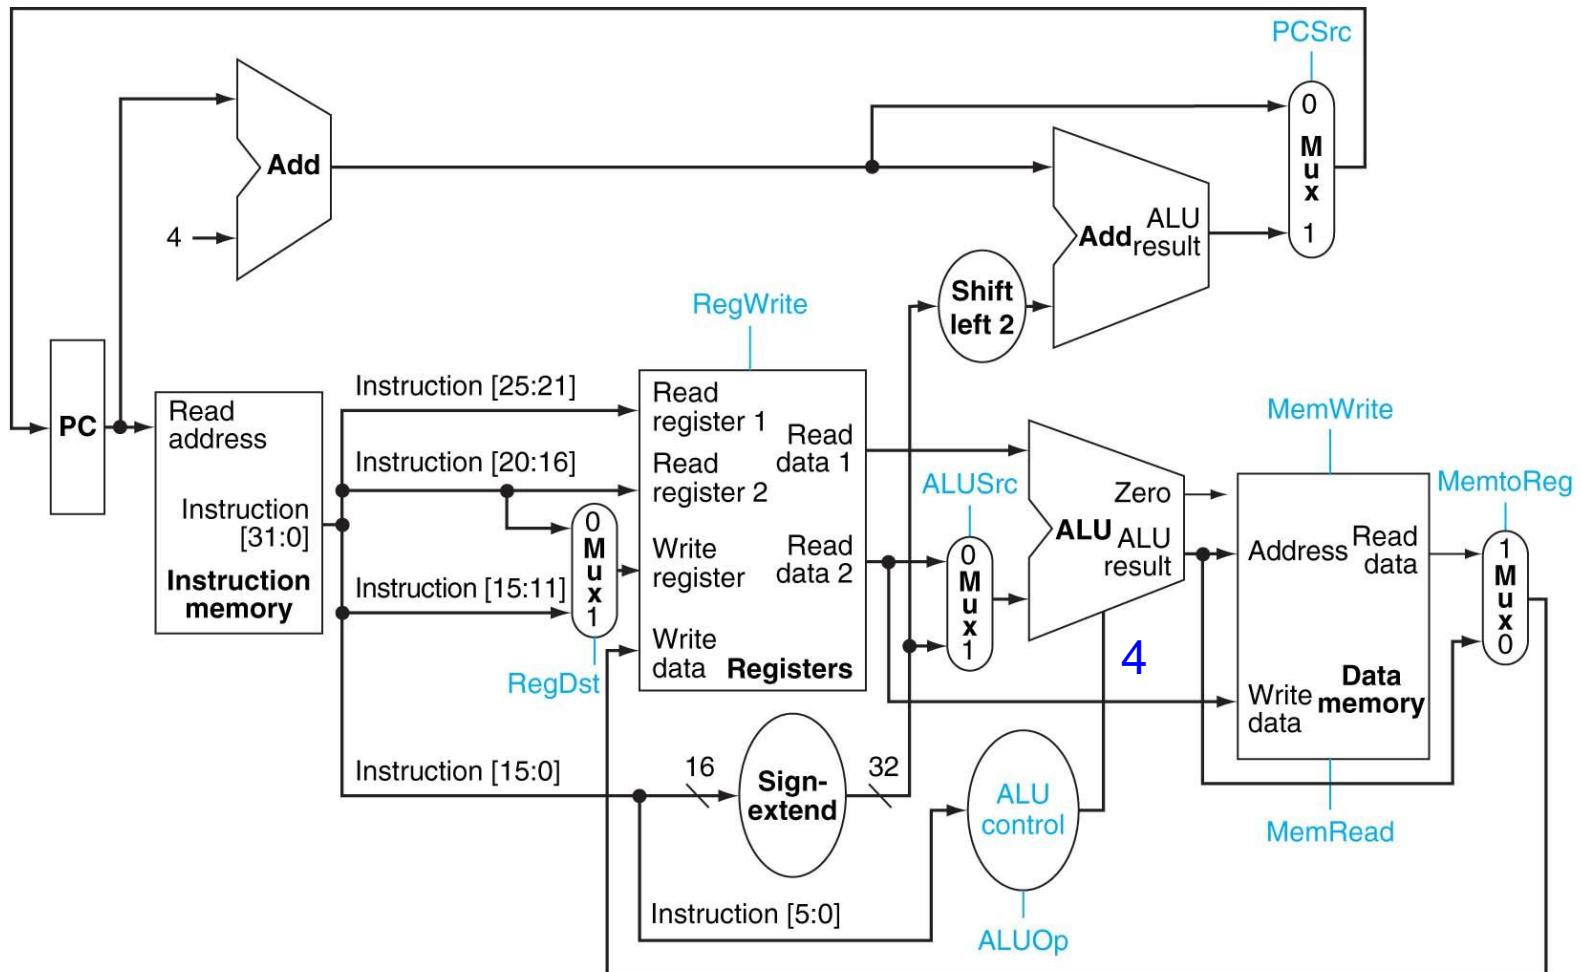
## Control signals derived from instruction



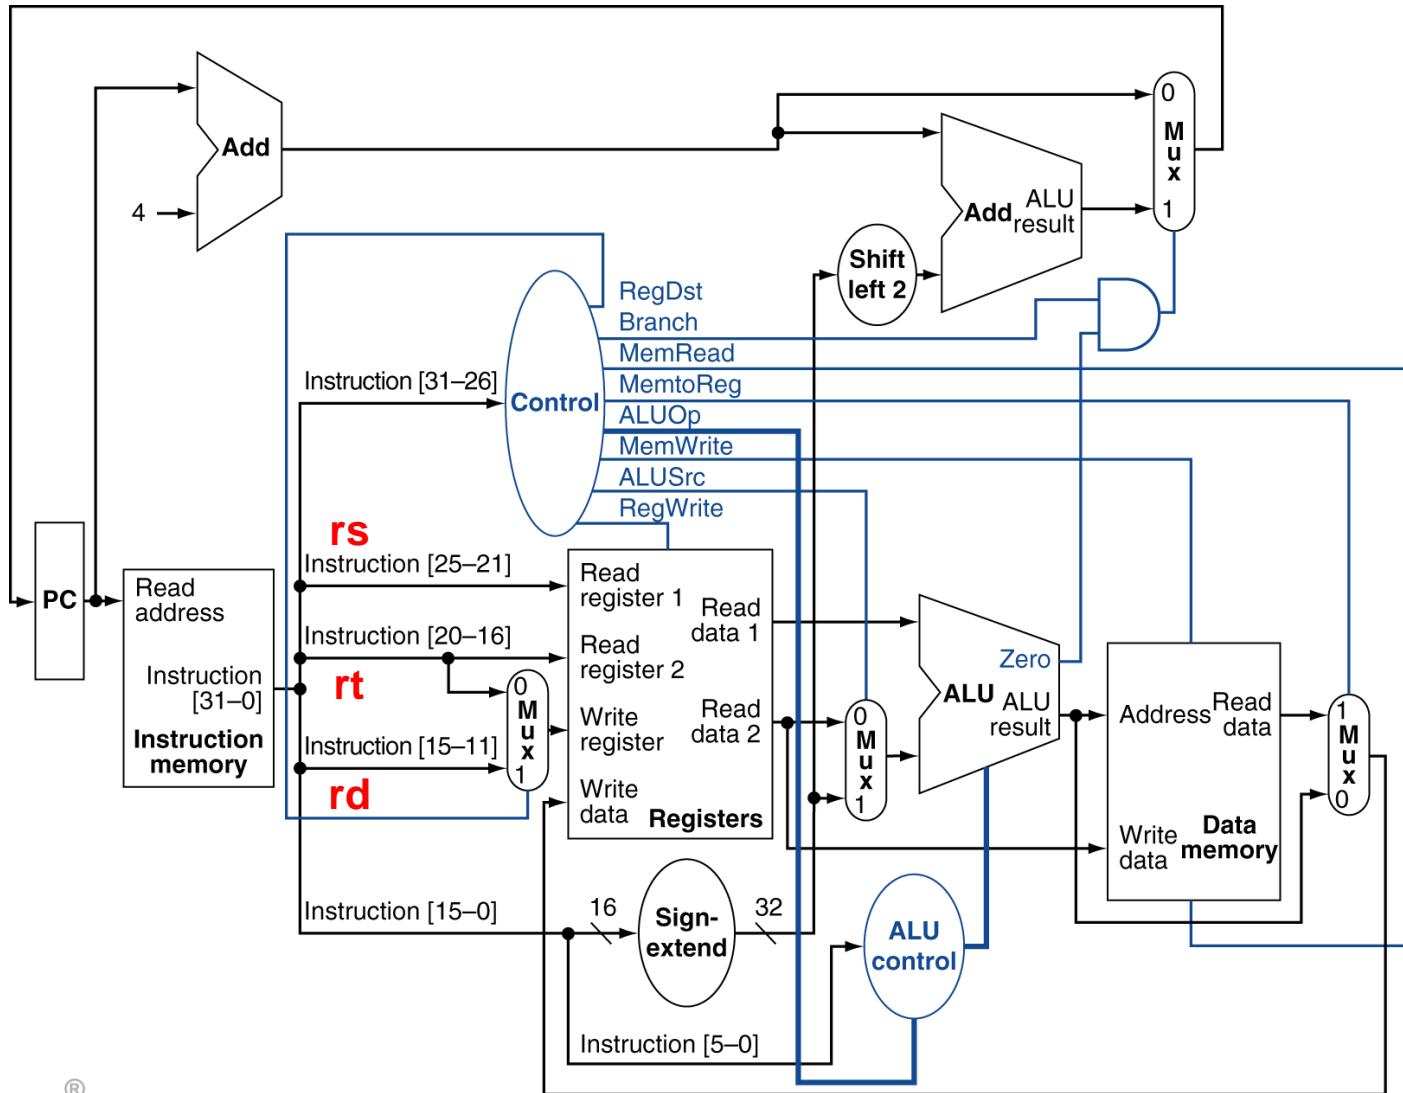
# The Main Control Unit

| Signal name | (set to 0)<br>Effect when deasserted  | (set to 1)<br>Effect when asserted  |
|-------------|---|---|
| RegDst      | The register destination number for the Write register comes from the <b>rt</b> field (bits 20:16). | The register destination number for the Write register comes from the <b>rd</b> field (bits 15:11).     |
| RegWrite    | None.   | The register on the Write register input is written with the value on the Write data input.             |
| ALUSrc      | The second ALU operand comes from the second register file output (Read <b>data 2</b> ).            | The second ALU operand is the sign-extended, lower 16 bits of the instruction.                          |
| PCSrc       | The PC is replaced by the output of the adder that computes the value of <b>PC + 4</b> .            | The PC is replaced by the output of the adder that computes the <b>branch target</b> .                  |
| MemRead     | None.   | Data memory contents designated by the address input are put on the Read data output.                   |
| MemWrite    | None.   | Data memory contents designated by the address input are replaced by the value on the Write data input. |
| MemtoReg    | The value fed to the register Write data input comes from the <b>ALU</b> .                          | The value fed to the register Write data input comes from the <b>data memory</b> .                      |

# The ALU Control

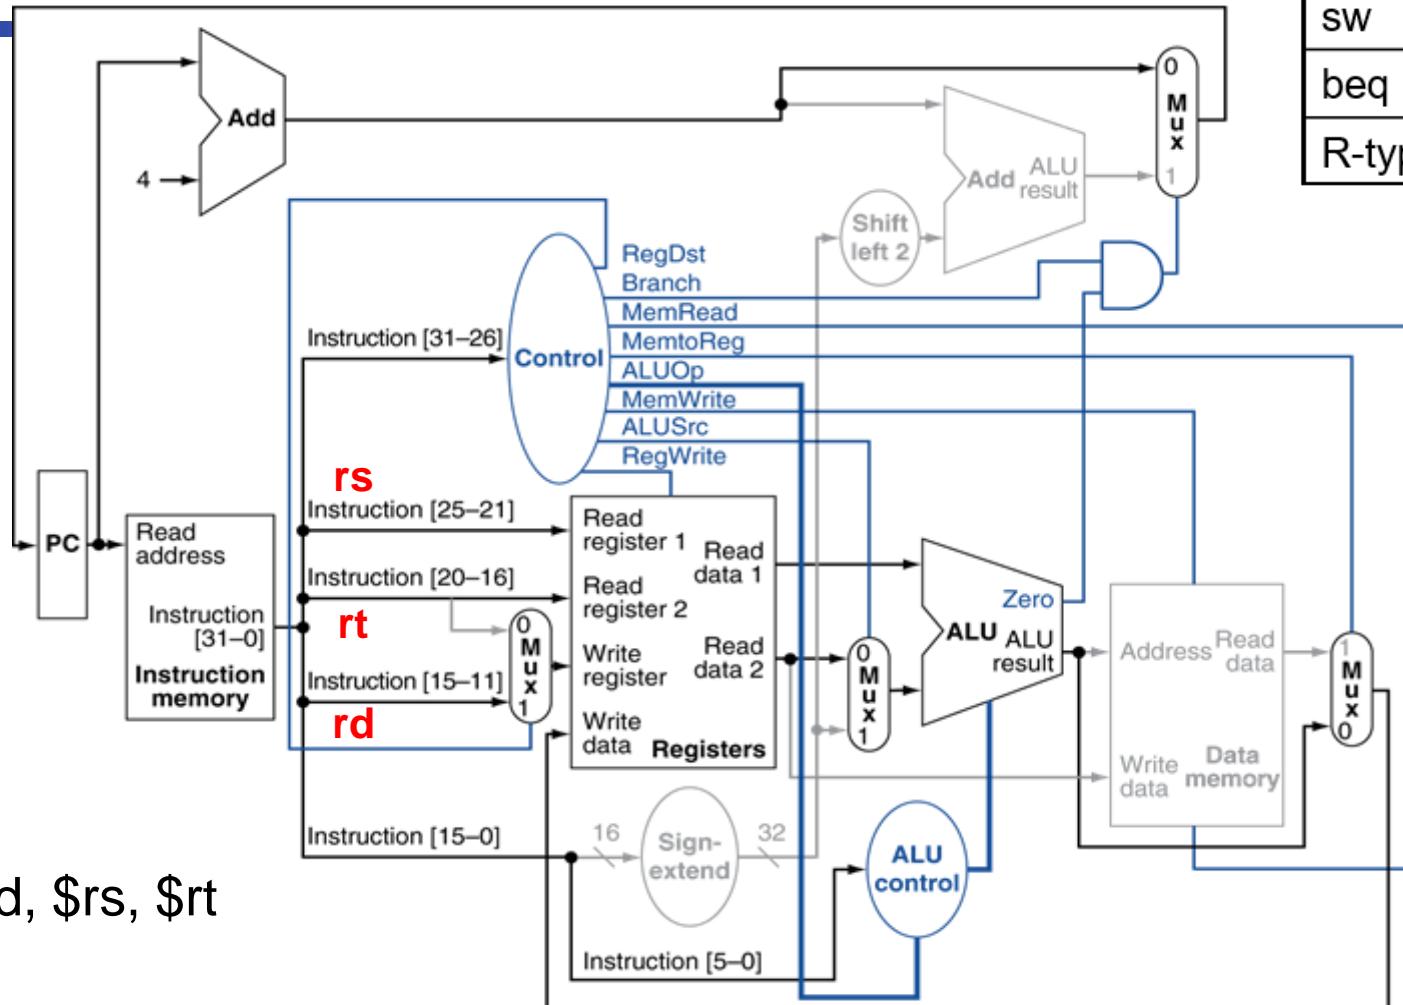


# Datapath With Control



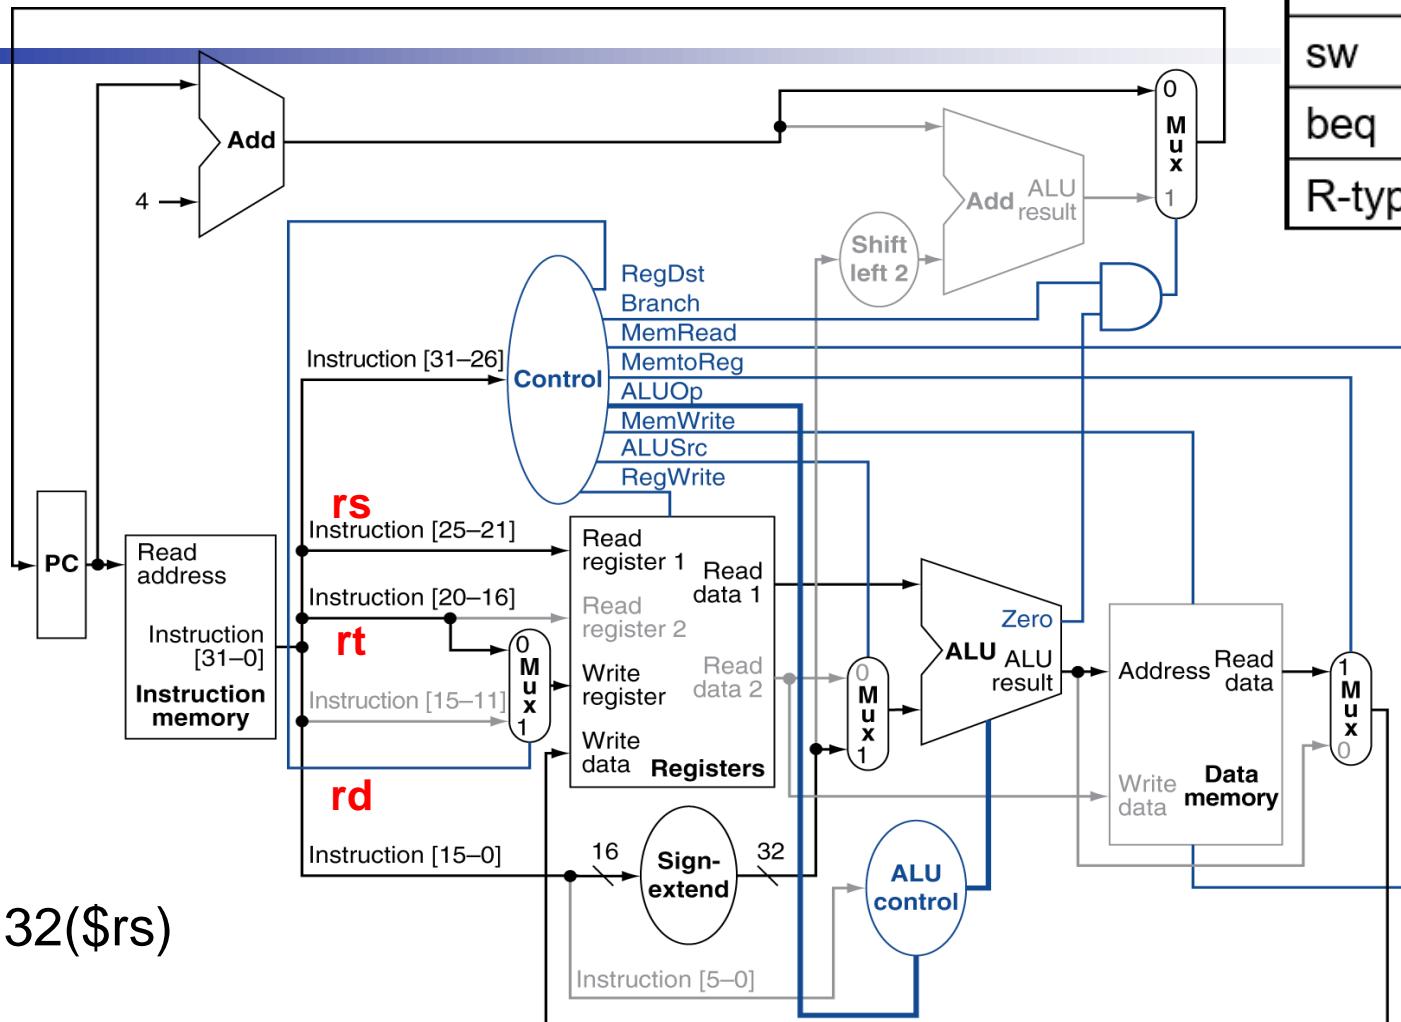
# R-Type Instruction

| opcode | ALUOp |
|--------|-------|
| lw     | 00    |
| sw     | 00    |
| beq    | 01    |
| R-type | 10    |



| Instruction | RegDst | ALUSrc | Mem-to-Reg | Reg Write | Mem Read | Mem Write | Branch | ALUOp1 | ALUOp0 |
|-------------|--------|--------|------------|-----------|----------|-----------|--------|--------|--------|
| R-format    | 1      | 0      | 0          | 1         | 0        | 0         | 0      | 1      | 0      |

# Load Instruction

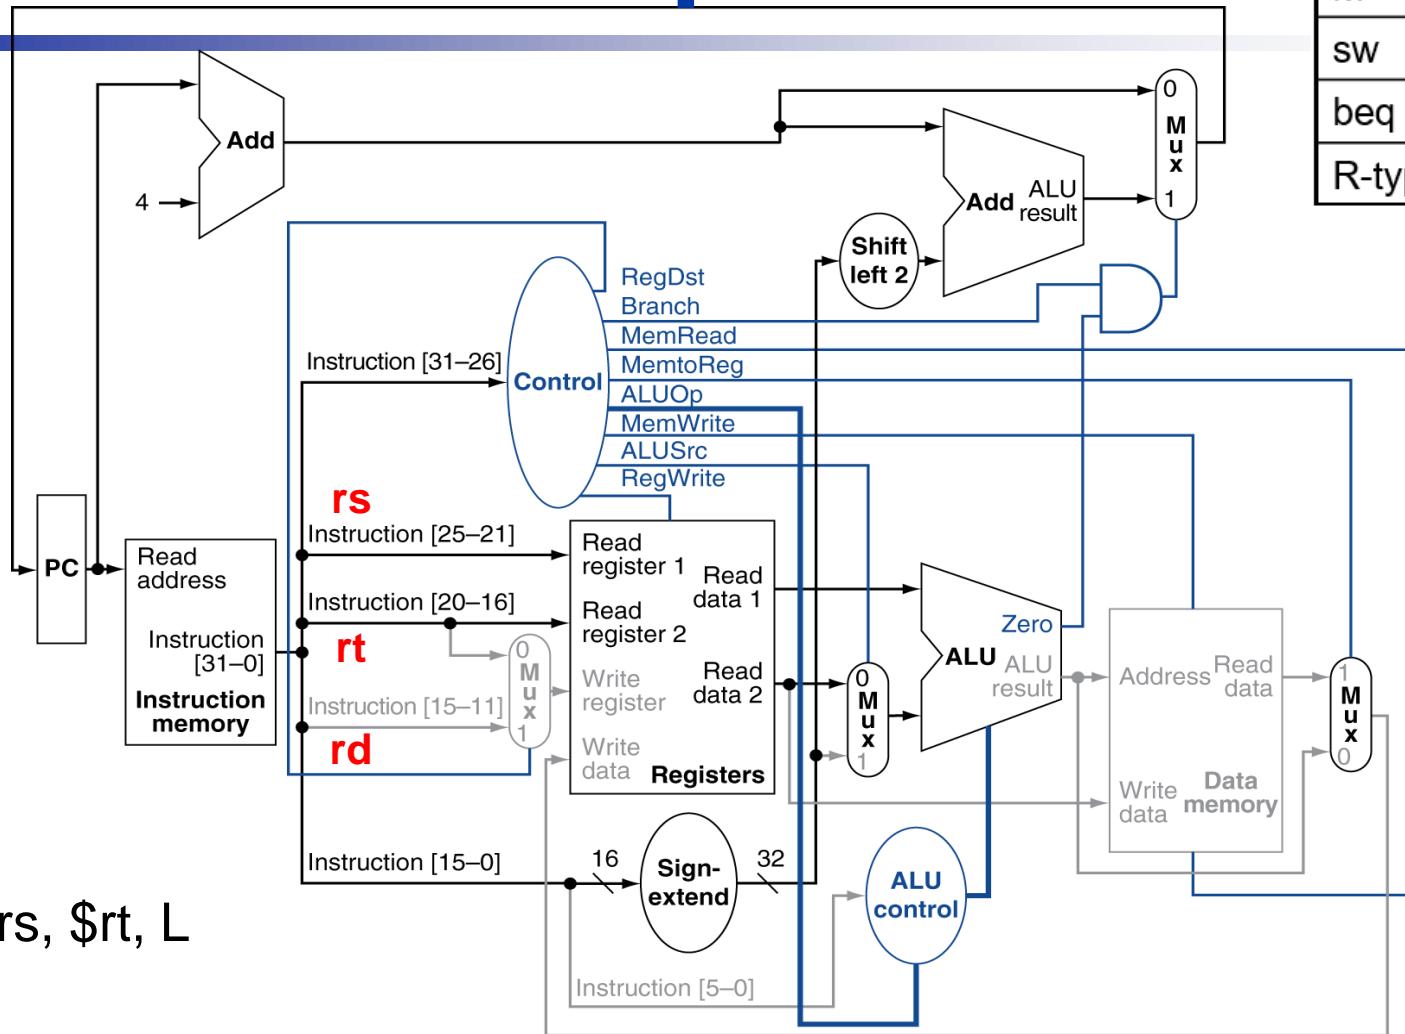


lw \$rt, 32(\$rs)

| Instruction | RegDst | ALUSrc | MemtoReg | RegWrite | MemRead | MemWrite | Branch | ALUOp1 | ALUOp0 |
|-------------|--------|--------|----------|----------|---------|----------|--------|--------|--------|
| lw          | 0      | 1      | 1        | 1        | 1       | 0        | 0      | 0      | 0      |

| opcode | ALUOp |
|--------|-------|
| lw     | 00    |
| sw     | 00    |
| beq    | 01    |
| R-type | 10    |

# Branch-on-Equal Instr.



beq \$rs, \$rt, L

| Instruction | RegDst | ALUSrc | MemtoReg | Reg Write | Mem Read | Mem Write | Branch | ALUOp1 | ALUOp0 |
|-------------|--------|--------|----------|-----------|----------|-----------|--------|--------|--------|
| beq         | X      | 0      | X        | 0         | 0        | 0         | 1      | 0      | 1      |



| opcode | ALUOp |
|--------|-------|
| lw     | 00    |
| sw     | 00    |
| beq    | 01    |
| R-type | 10    |

# The Main Control Unit

- The setting of the control lines:

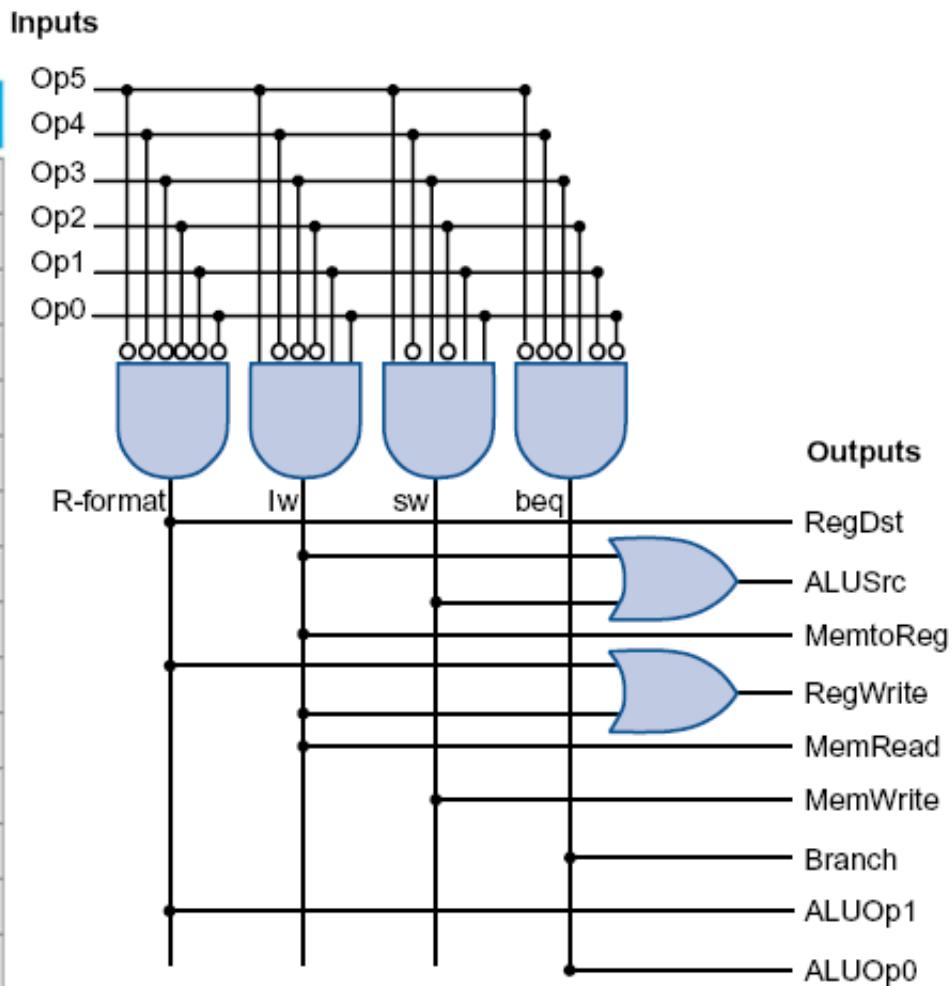
| Instruction | RegDst | ALUSrc | Memto-Reg | Reg-Write | Mem-Read | Mem-Write | Branch | ALUOp1 | ALUOp0 |
|-------------|--------|--------|-----------|-----------|----------|-----------|--------|--------|--------|
| R-format    | 1      | 0      | 0         | 1         | 0        | 0         | 0      | 1      | 0      |
| lw          | 0      | 1      | 1         | 1         | 1        | 0         | 0      | 0      | 0      |
| sw          | X      | 1      | X         | 0         | 0        | 1         | 0      | 0      | 0      |
| beq         | X      | 0      | X         | 0         | 0        | 0         | 1      | 0      | 1      |

- The encoding for each of the opcodes of interest:

| Name     | Opcode in decimal | Opcode in binary |     |     |     |     |     |  |
|----------|-------------------|------------------|-----|-----|-----|-----|-----|--|
|          |                   | Op5              | Op4 | Op3 | Op2 | Op1 | Op0 |  |
| R-format | 0 <sub>ten</sub>  | 0                | 0   | 0   | 0   | 0   | 0   |  |
| lw       | 35 <sub>ten</sub> | 1                | 0   | 0   | 0   | 1   | 1   |  |
| sw       | 43 <sub>ten</sub> | 1                | 0   | 1   | 0   | 1   | 1   |  |
| beq      | 4 <sub>ten</sub>  | 0                | 0   | 0   | 1   | 0   | 0   |  |

# Main Control Unit (Appendix)

|         | <b>Signal name</b> | <b>R-format</b> | <b>Iw</b> | <b>sw</b> | <b>beq</b> |
|---------|--------------------|-----------------|-----------|-----------|------------|
| Inputs  | Op5                | 0               | 1         | 1         | 0          |
|         | Op4                | 0               | 0         | 0         | 0          |
|         | Op3                | 0               | 0         | 1         | 0          |
|         | Op2                | 0               | 0         | 0         | 1          |
|         | Op1                | 0               | 1         | 1         | 0          |
|         | Op0                | 0               | 1         | 1         | 0          |
| Outputs | RegDst             | 1               | 0         | X         | X          |
|         | ALUSrc             | 0               | 1         | 1         | 0          |
|         | MemtoReg           | 0               | 1         | X         | X          |
|         | RegWrite           | 1               | 1         | 0         | 0          |
|         | MemRead            | 0               | 1         | 0         | 0          |
|         | MemWrite           | 0               | 0         | 1         | 0          |
|         | Branch             | 0               | 0         | 0         | 1          |
|         | ALUOp1             | 1               | 0         | 0         | 0          |
|         | ALUOp0             | 0               | 0         | 0         | 1          |



# Implementing Jumps

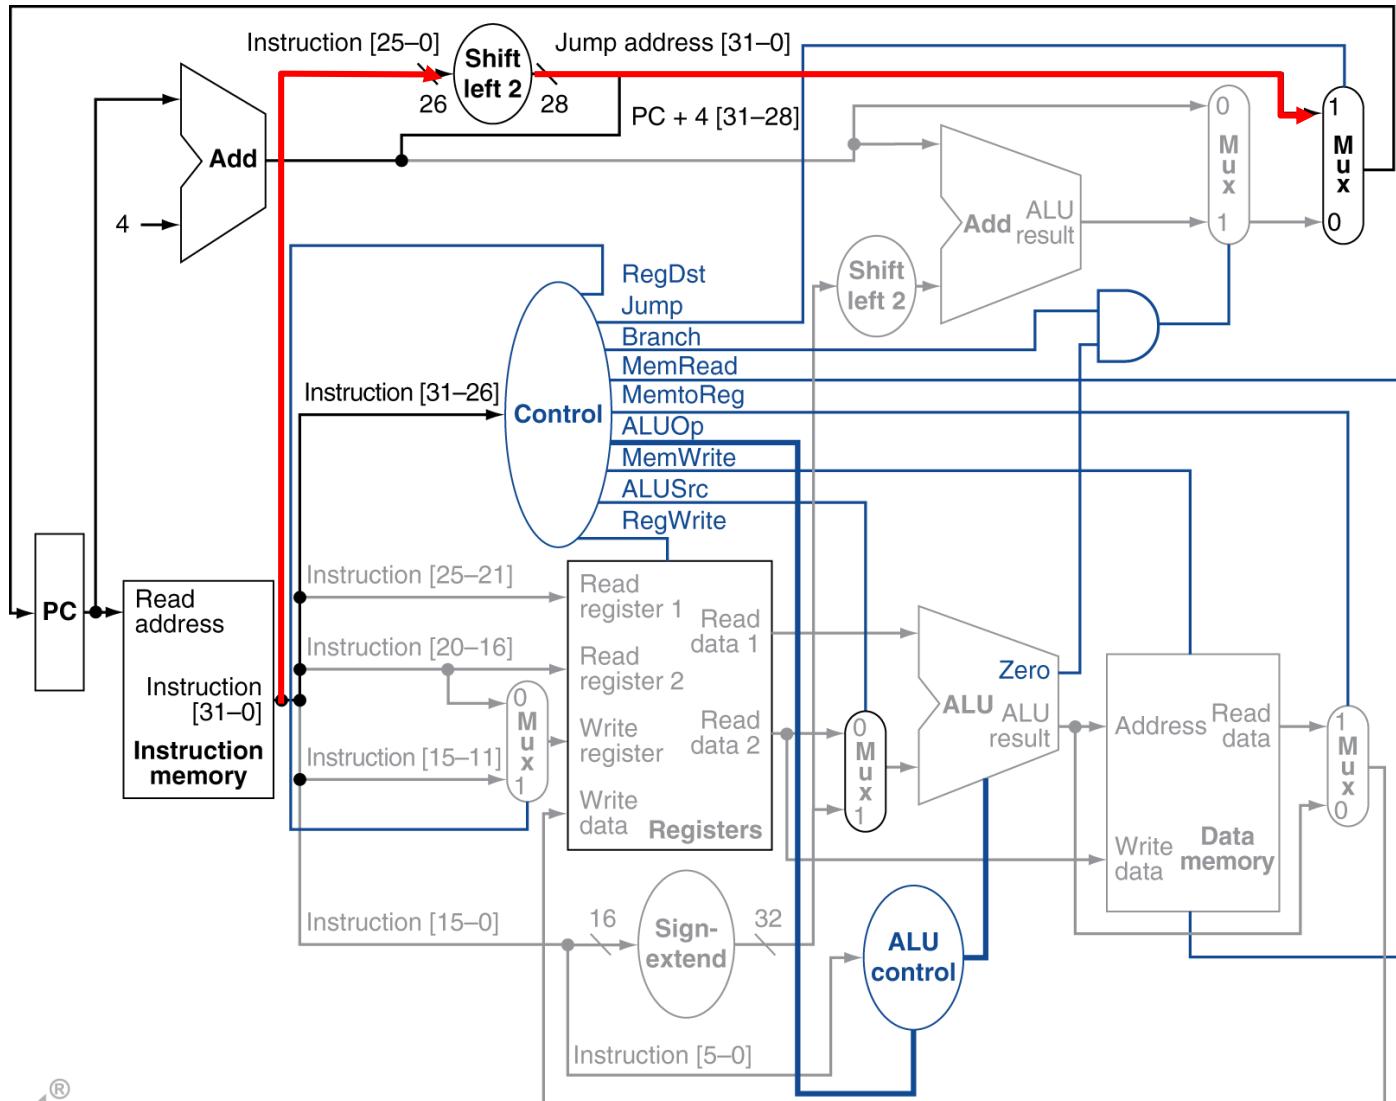
## Jump Instruction (opcode=2)

| Field         | 000010 | address |
|---------------|--------|---------|
| Bit positions | 31:26  | 25:0    |

- Jump uses word address
- Update PC with concatenation of
  - Top 4 bits of old PC
  - 26-bit jump address
  - 00
- Need an extra control signal decoded from opcode

| Target address |         |     |
|----------------|---------|-----|
| PC+4           | address | 00  |
| 31:28          | 27:2    | 1:0 |

# Datapath With Jumps Added

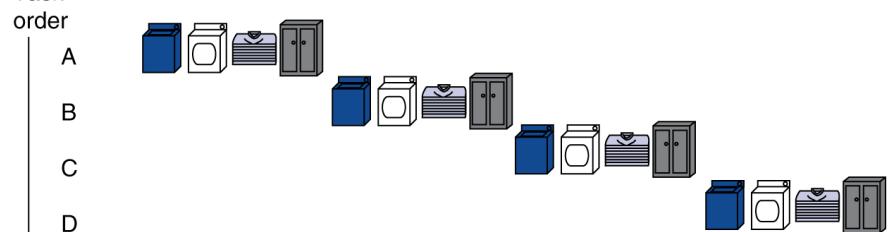
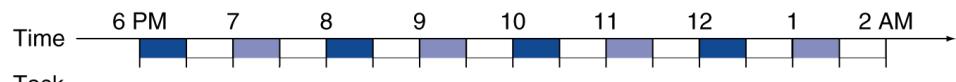


# Performance Issues

- Longest delay determines clock period
  - Critical path: load instruction
  - Instruction memory → register file → ALU → data memory → register file
- Not feasible to vary period for different instructions
- Violates design principle
  - Making the common case fast
- We will improve performance by pipelining

# Pipelining Analogy

- Pipelined laundry: overlapping execution
  - Parallelism improves performance



- Four loads:
  - Speedup  
 $= 8/3.5 = 2.3$
- Non-stop:
  - Speedup  
 $= 2n/(0.5n + 1.5) \approx 4$
  - number of stages

# MIPS Pipeline

- Five stages, one step per stage
  1. IF: Instruction fetch from memory
  2. ID: Instruction decode & register read
  3. EX: Execute operation or calculate address
  4. MEM: Access memory operand
  5. WB: Write result back to register

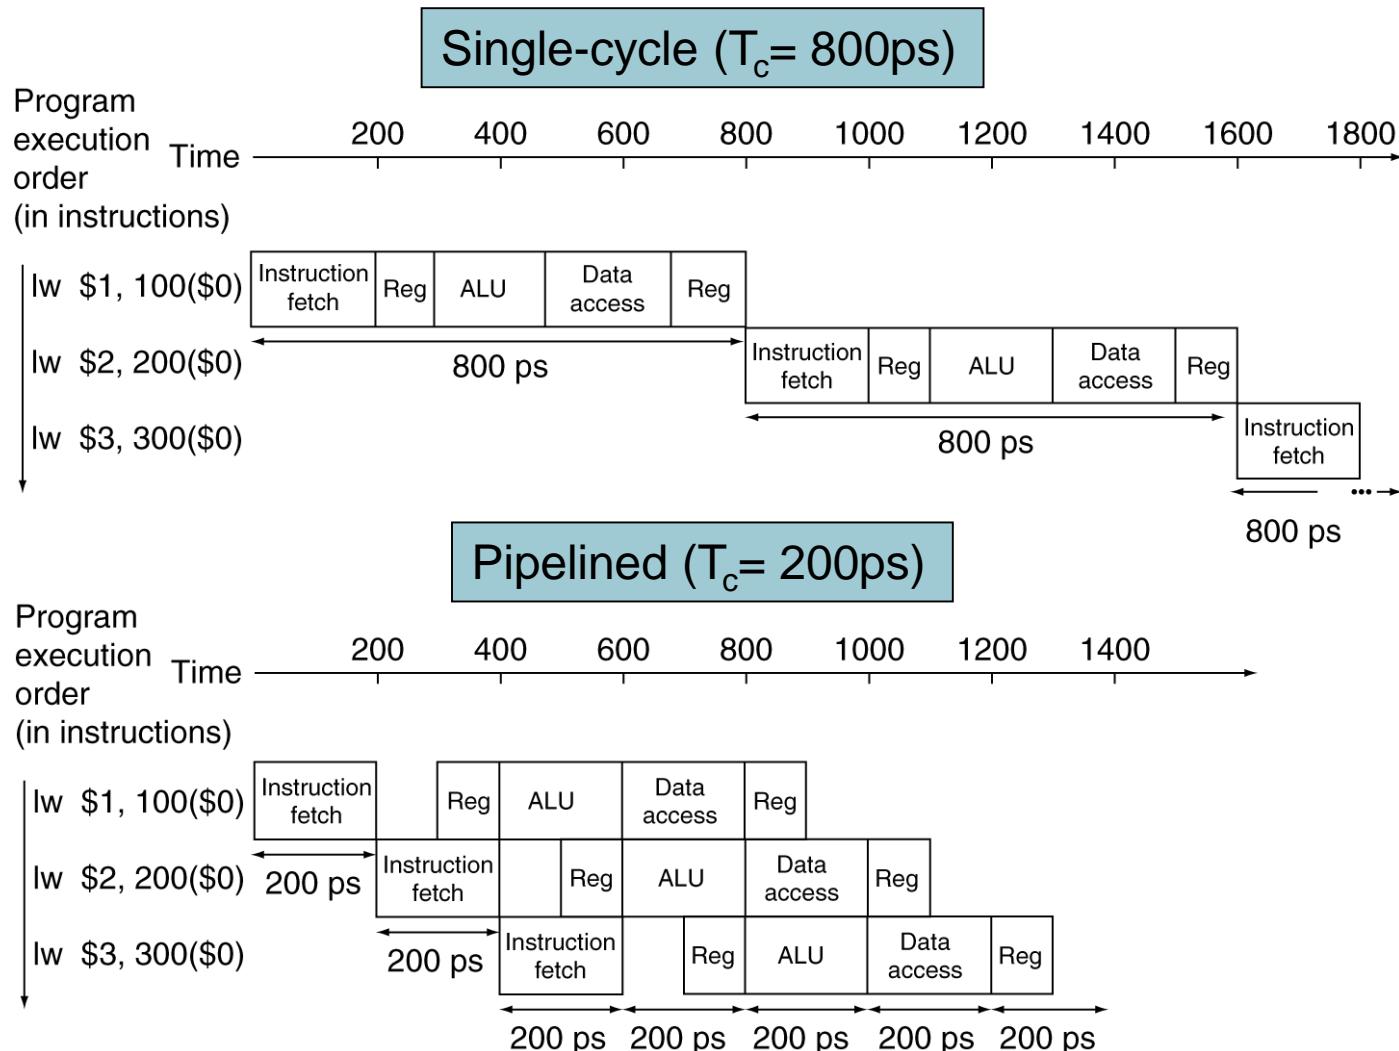
# Pipeline Performance

- Assume time for stages is
  - 100ps for register read or write
  - 200ps for other stages
- Compare pipelined datapath with single-cycle datapath

**IF            ID            EX            MEM            WB**

| Instr    | Instr fetch | Register read | ALU op | Memory access | Register write | Total time |
|----------|-------------|---------------|--------|---------------|----------------|------------|
| lw       | 200ps       | 100 ps        | 200ps  | 200ps         | 100 ps         | 800ps      |
| sw       | 200ps       | 100 ps        | 200ps  | 200ps         |                | 700ps      |
| R-format | 200ps       | 100 ps        | 200ps  |               | 100 ps         | 600ps      |
| beq      | 200ps       | 100 ps        | 200ps  |               |                | 500ps      |

# Pipeline Performance



# Pipeline Speedup

- If all stages are balanced
  - i.e., all take the same time
  - Time between instructions<sub>pipelined</sub>  
= Time between instructions<sub>nonpipelined</sub>  

---

Number of stages
- If not balanced, speedup is less
- Speedup due to increased throughput
  - Latency (time for each instruction) does not decrease

# Pipelining and ISA Design

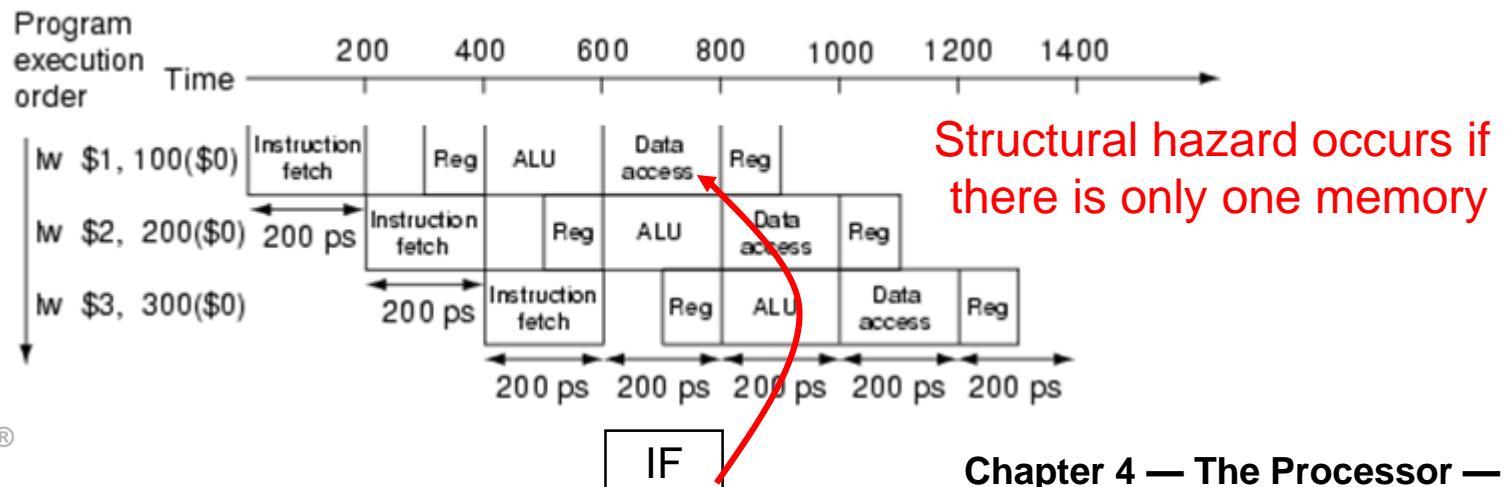
- MIPS ISA designed for pipelining
  - All instructions are 32-bits
    - Easier to fetch in one cycle
    - c.f. x86: 1- to 17-byte instructions
  - Few and regular instruction formats
    - Can decode and read registers in one step
  - Load/store addressing
    - Can calculate address in 3<sup>rd</sup> stage, access memory in 4<sup>th</sup> stage
  - Alignment of memory operands
    - Memory access takes only one cycle

# Hazards

- Situations that prevent starting the next instruction in the next cycle
- Structure hazards
  - A required resource is busy
- Data hazard
  - Need to wait for previous instruction to complete its data read/write
- Control hazard
  - Deciding on control action depends on previous instruction

# Structure Hazards

- Conflict for use of a resource
  - hardware cannot support the combination of instruction in the same clock.
- In MIPS pipeline with a single memory
  - Load/store requires data access
  - Instruction fetch would have to *stall* for that cycle
  - require separate instr./data memories (or caches)



# Data Hazards

- An instruction depends on completion of data access by a previous instruction
  - data that is needed to execute the instruction is still in the pipeline, not yet available.

- Ex. 1

add \$s0, \$t0, \$t1

sub \$t2, \$s0, \$t3



- Ex. 2 ( load-use data hazard )

lw \$s0, 4(\$t1)

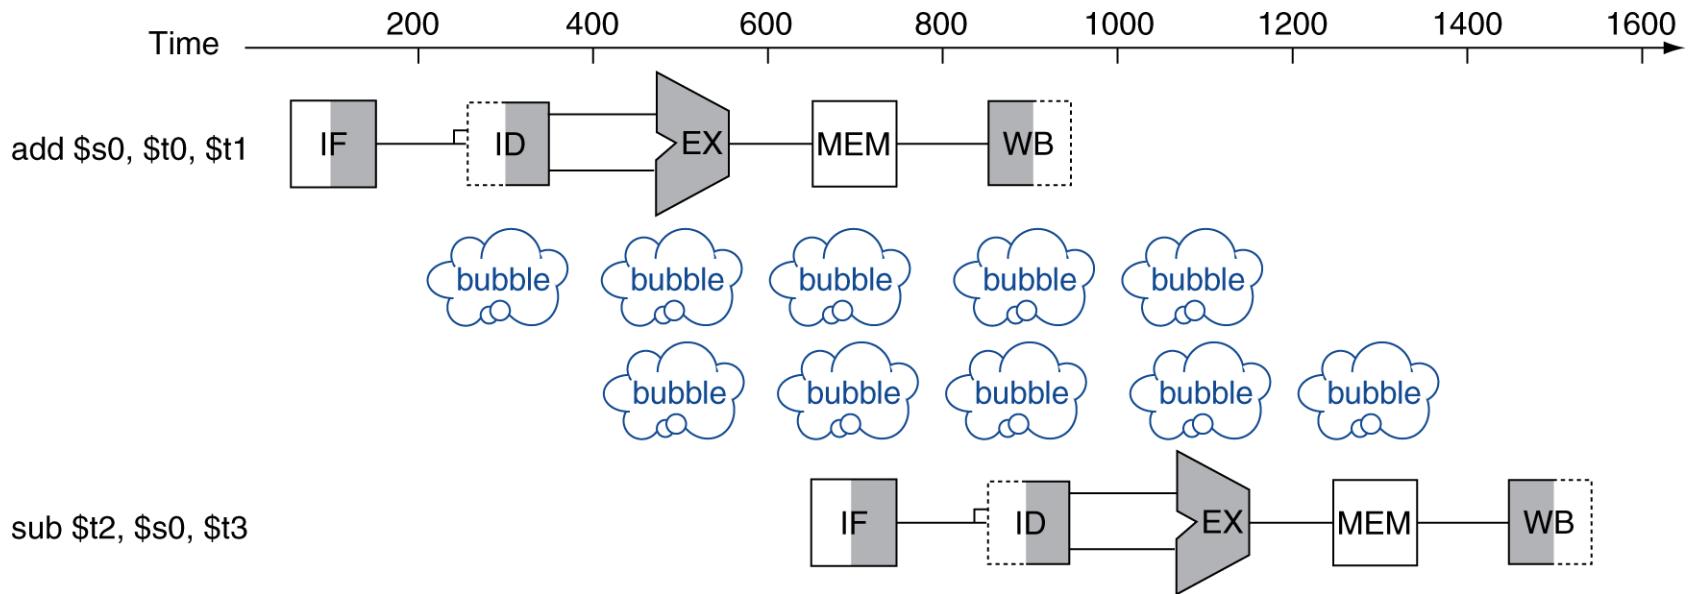
sub \$t2, \$s0, \$t3



# Data Hazards - stall

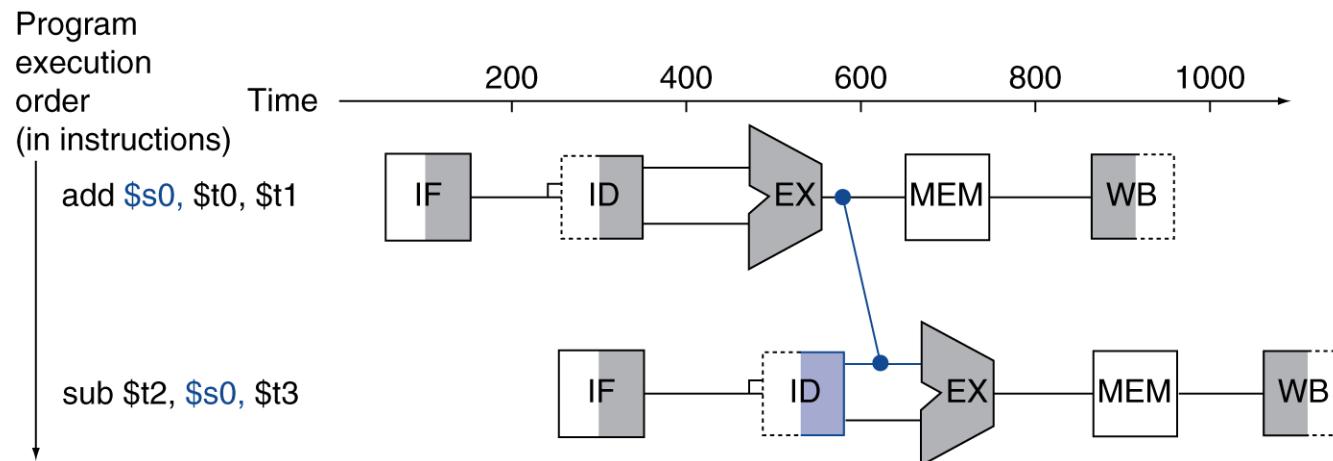
- Stall the instructions (bubble)

- add \$s0, \$t0, \$t1  
sub \$t2, \$s0, \$t3



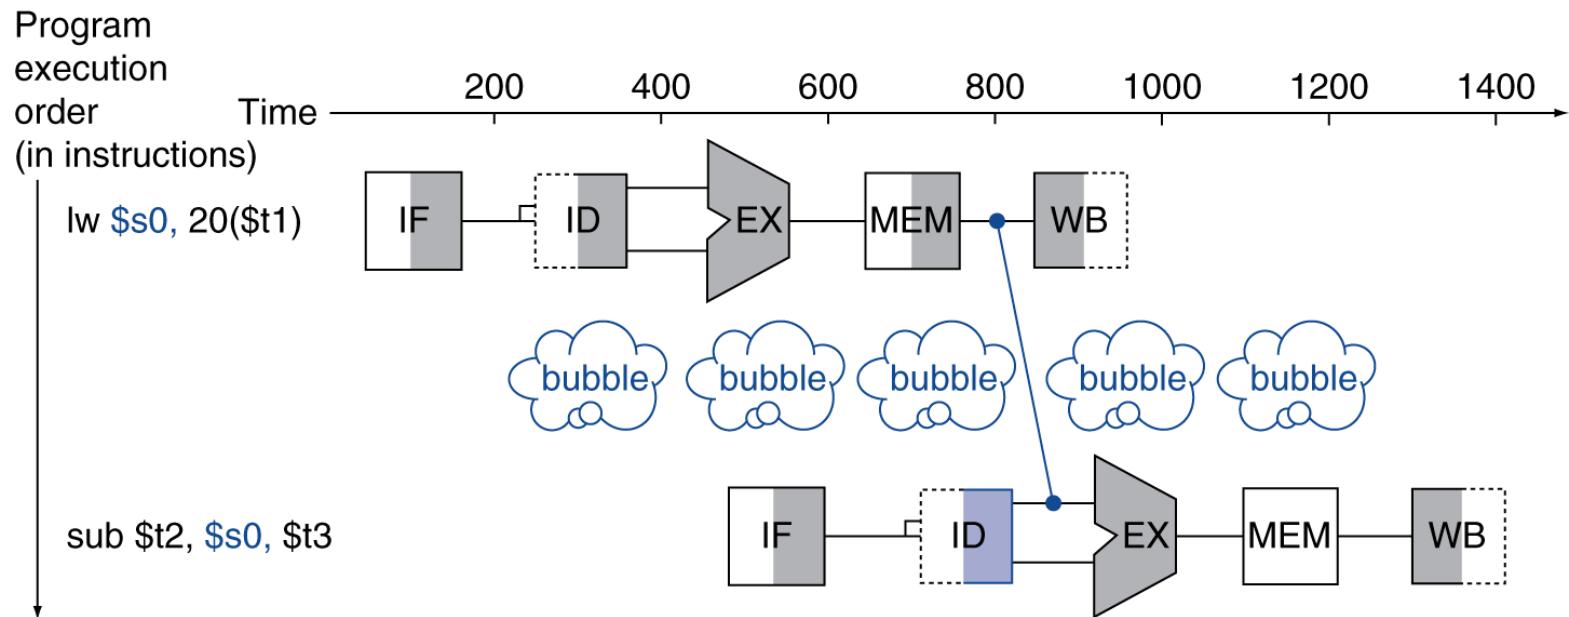
# Data Hazards - Forwarding

- Use result when it is computed
  - Don't wait for it to be stored in a register
  - Requires extra connections in the datapath



# Load-Use Data Hazard

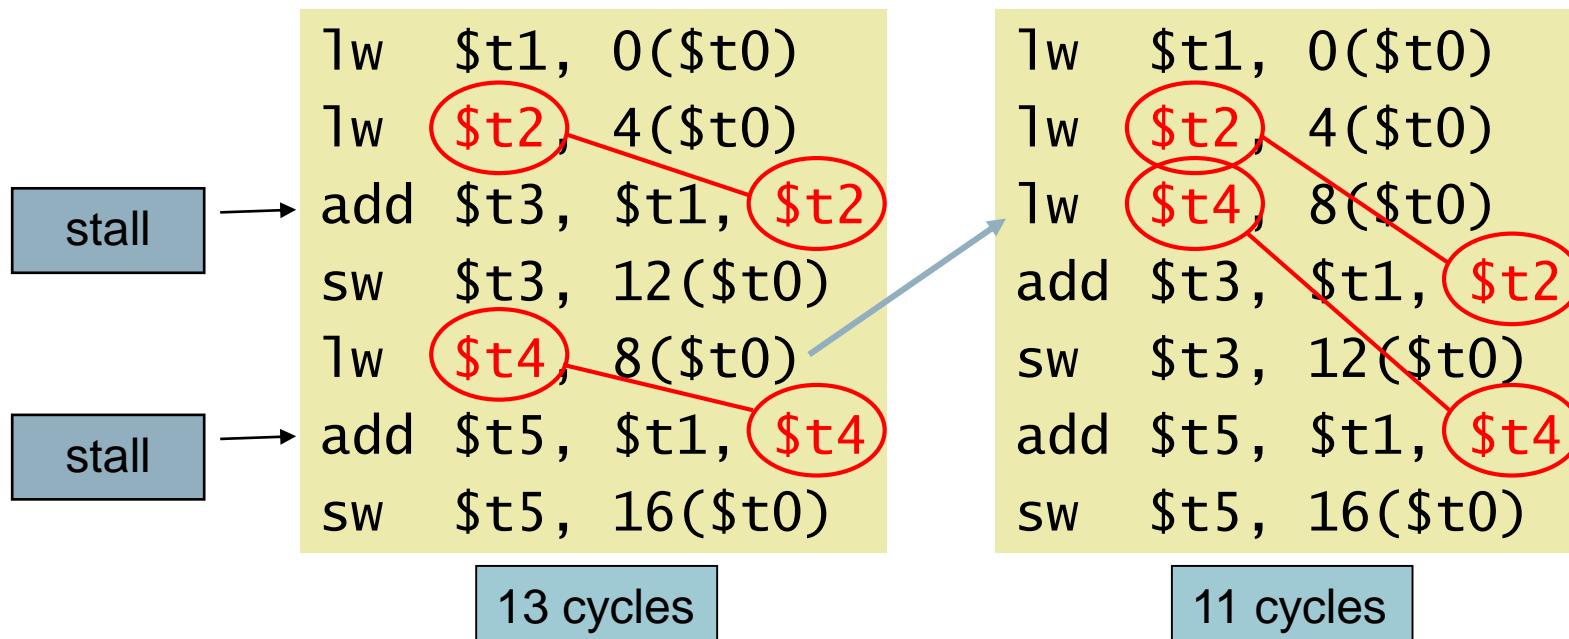
- Can't always avoid stalls by forwarding
  - If value not computed when needed
  - Can't forward backward in time!



# Code Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction
- C code for  $A = B + E; C = B + F;$

[3]      [0]      [1]      [4]      [0]      [2]



13 cycles

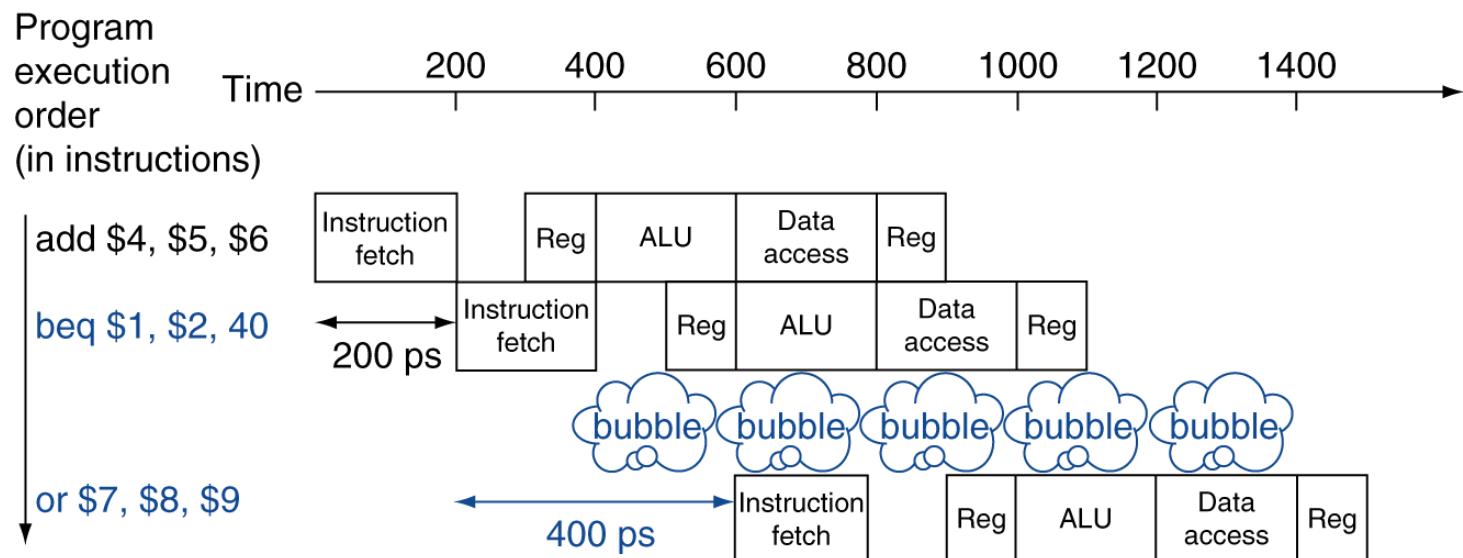
11 cycles

# Control Hazards

- Branch determines flow of control
  - Fetching next instruction depends on branch outcome
  - Pipeline can't always fetch correct instruction
    - Still working on ID stage of branch (however, the branch outcome is determined at EX stage.)
- In MIPS pipeline
  - Need to compare registers and compute target early in the pipeline
  - Add hardware to do it in ID stage

# Stall on Branch

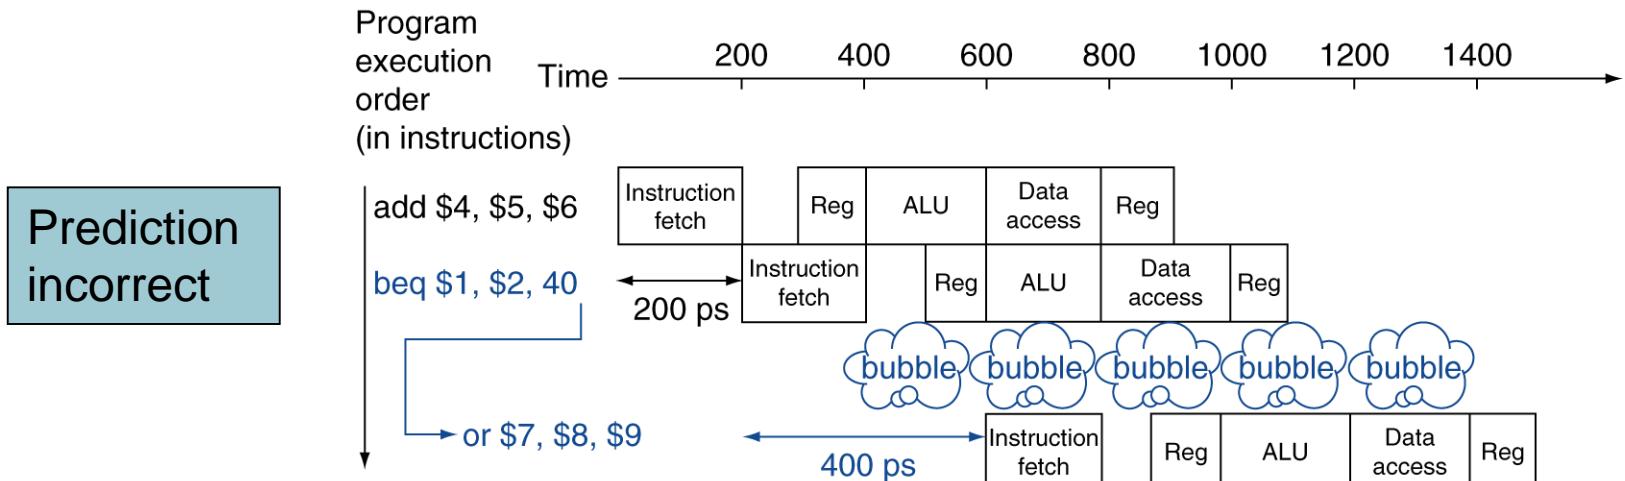
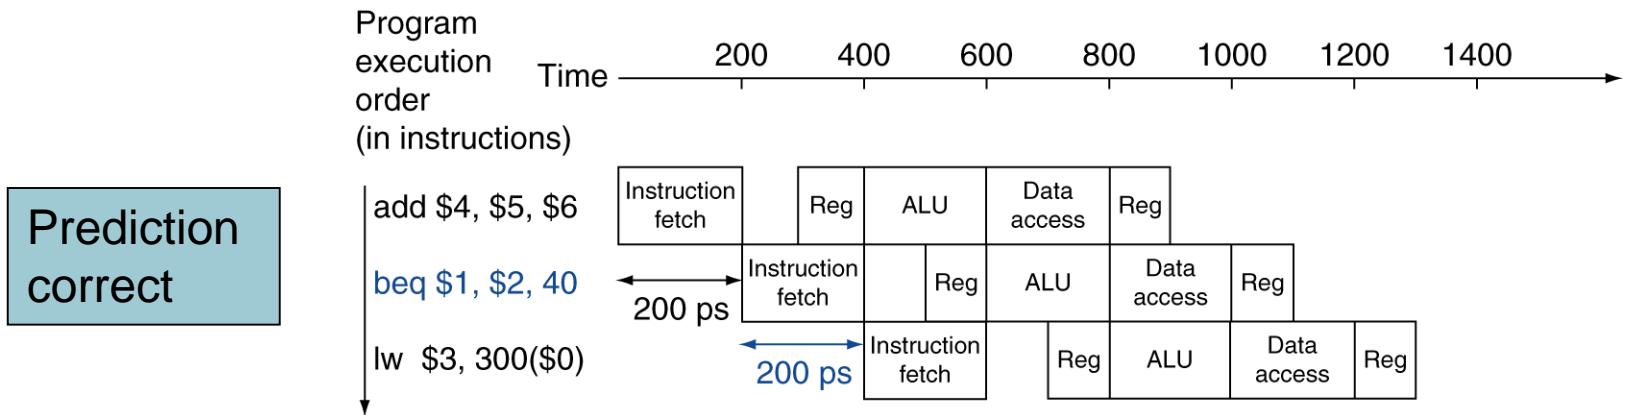
- Wait until branch outcome determined before fetching next instruction
  - Assume extra hardware has been added to test registers and calculate branch address at **ID stage**



# Branch Prediction

- Longer pipelines can't readily determine branch outcome early
  - Stall penalty becomes unacceptable
- Predict outcome of branch
  - Only stall if prediction is wrong
- In MIPS pipeline
  - Can predict branches not taken
  - Fetch instruction after branch, with no delay

# MIPS with Predict Not Taken



# More-Realistic Branch Prediction

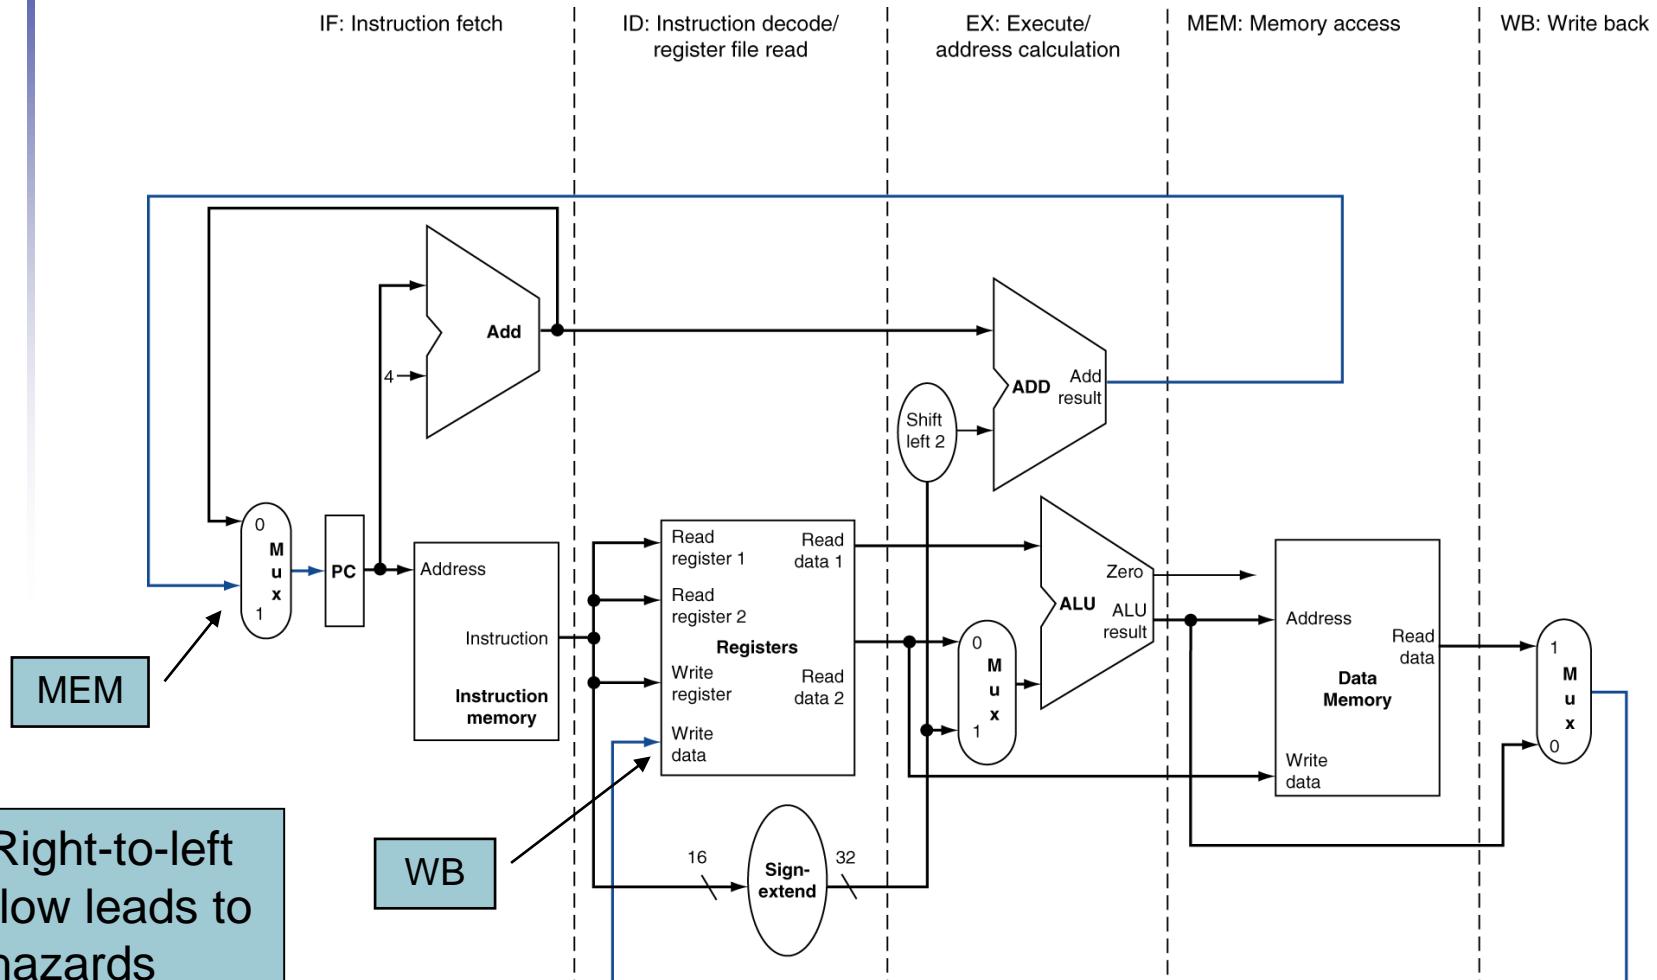
- Static branch prediction
  - Based on typical branch behavior
  - Example: loop and if-statement branches
    - Predict backward branches taken
    - Predict forward branches not taken
- Dynamic branch prediction
  - Hardware measures actual branch behavior
    - e.g., record recent history of each branch
  - Assume future behavior will continue the trend
    - When wrong, stall while re-fetching, and update history

# Pipeline Summary

## The BIG Picture

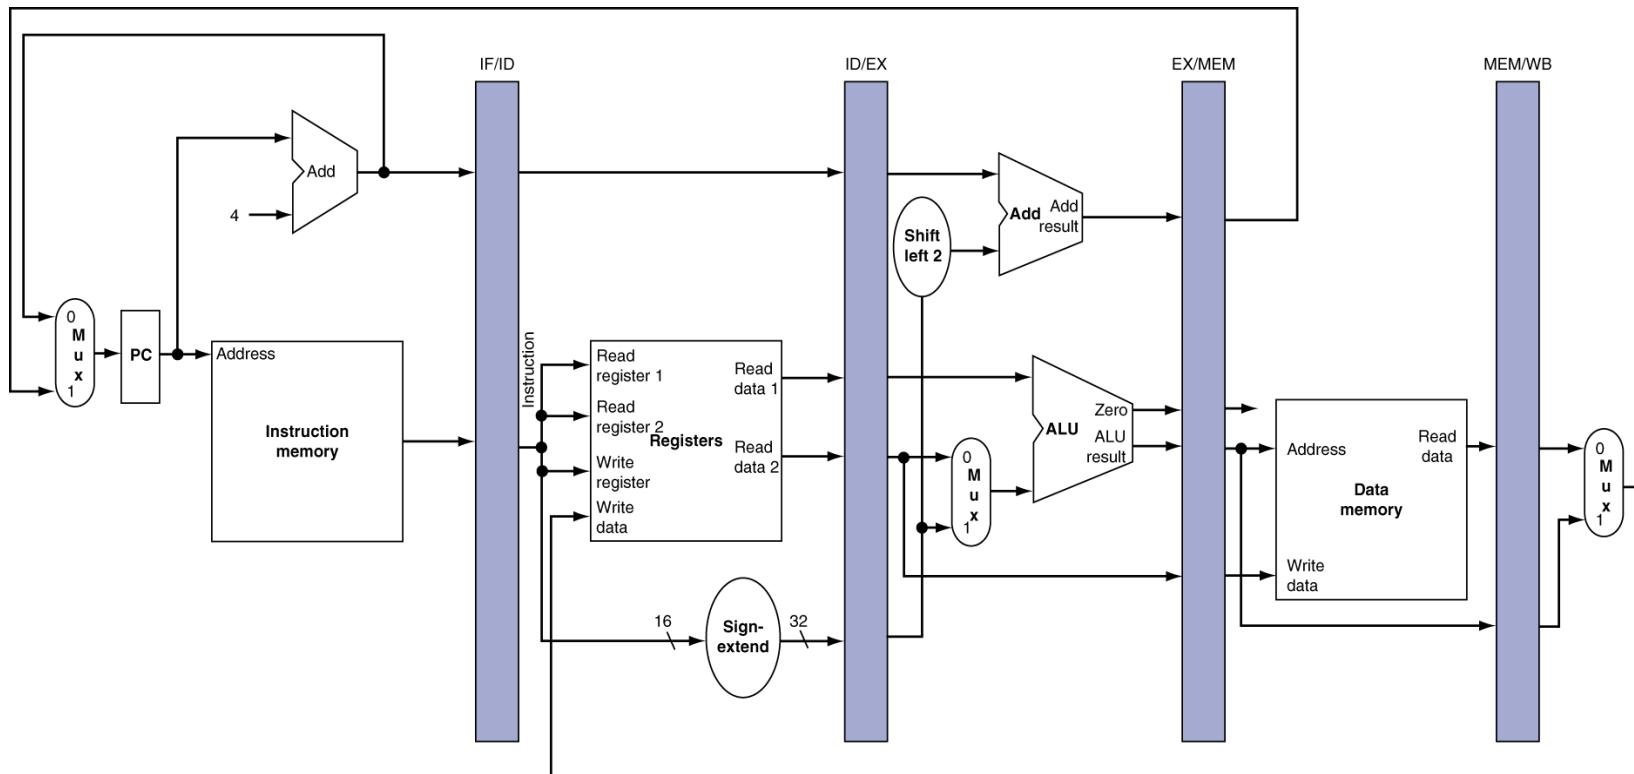
- Pipelining improves performance by increasing instruction throughput
  - Executes multiple instructions in parallel
  - Each instruction has the same latency
- Subject to hazards
  - Structure, data, control
- Instruction set design affects complexity of pipeline implementation

# MIPS Pipelined Datapath



# Pipeline registers

- Need registers between stages
  - To hold information produced in previous cycle

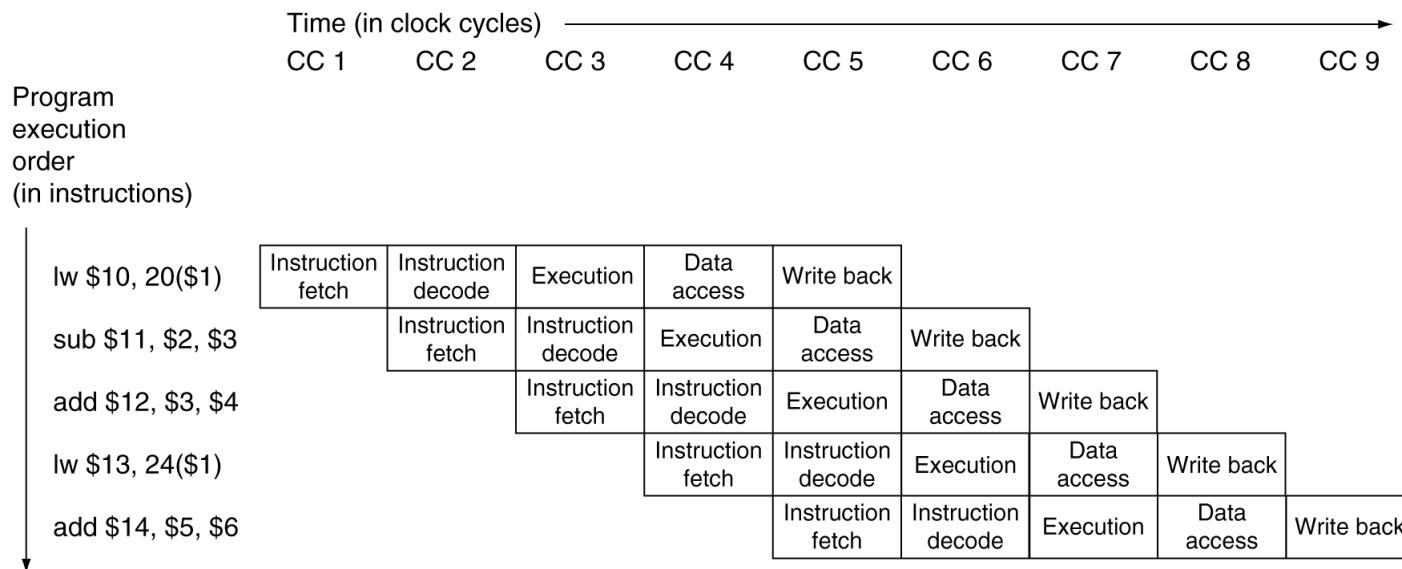


# Pipeline Operation Diagram

- Cycle-by-cycle flow of instructions through the pipelined datapath
  - “Single-clock-cycle” pipeline diagram
    - Shows pipeline usage in a single cycle
    - Highlight resources used
  - c.f. “multi-clock-cycle” diagram
    - Graph of operation over time
- We'll look at “single-clock-cycle” diagrams for load & store illustrations.

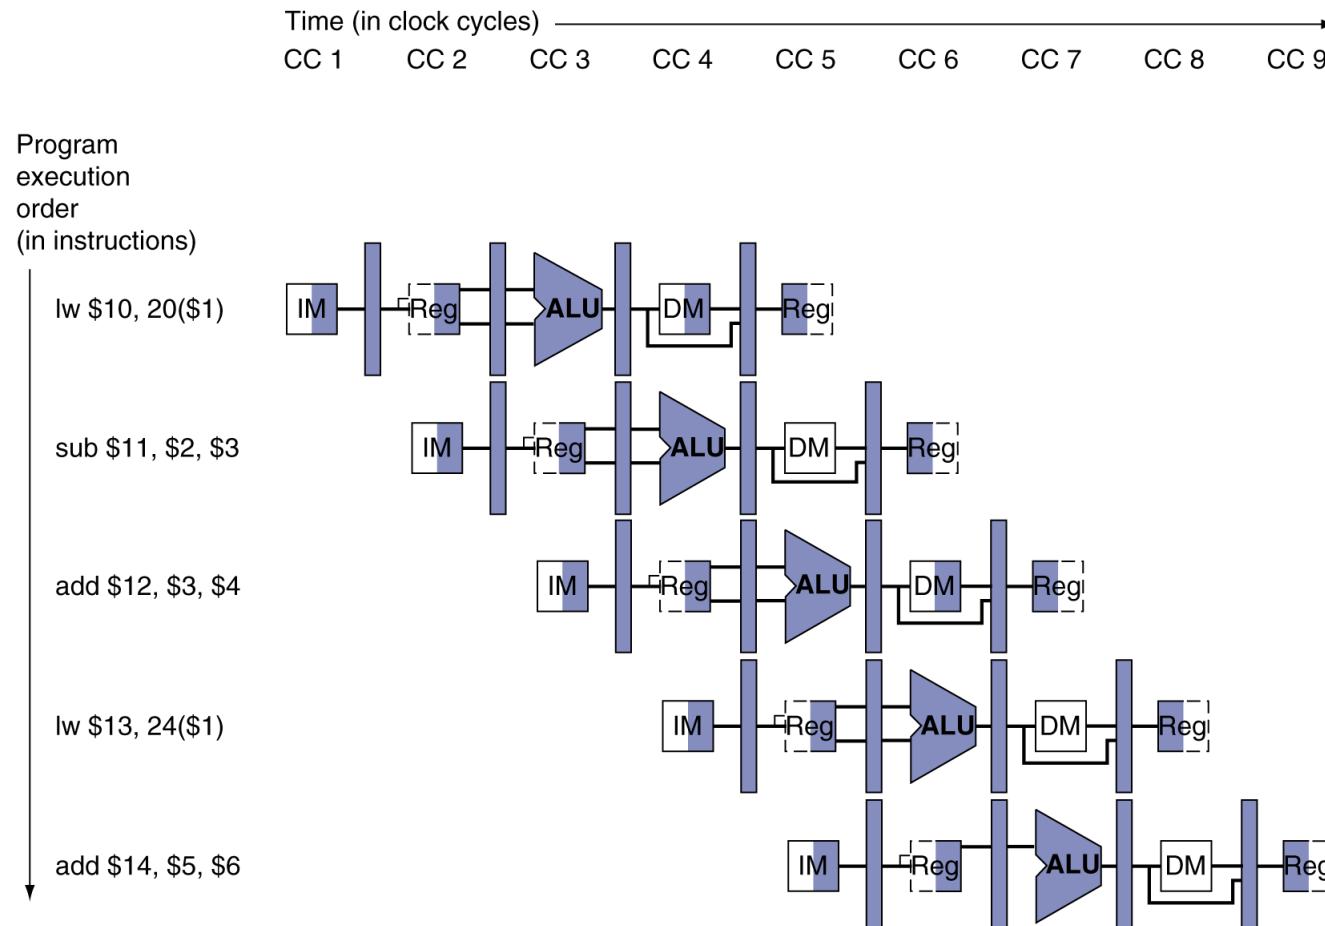
# Multi-Cycle Pipeline Diagram

## Traditional form



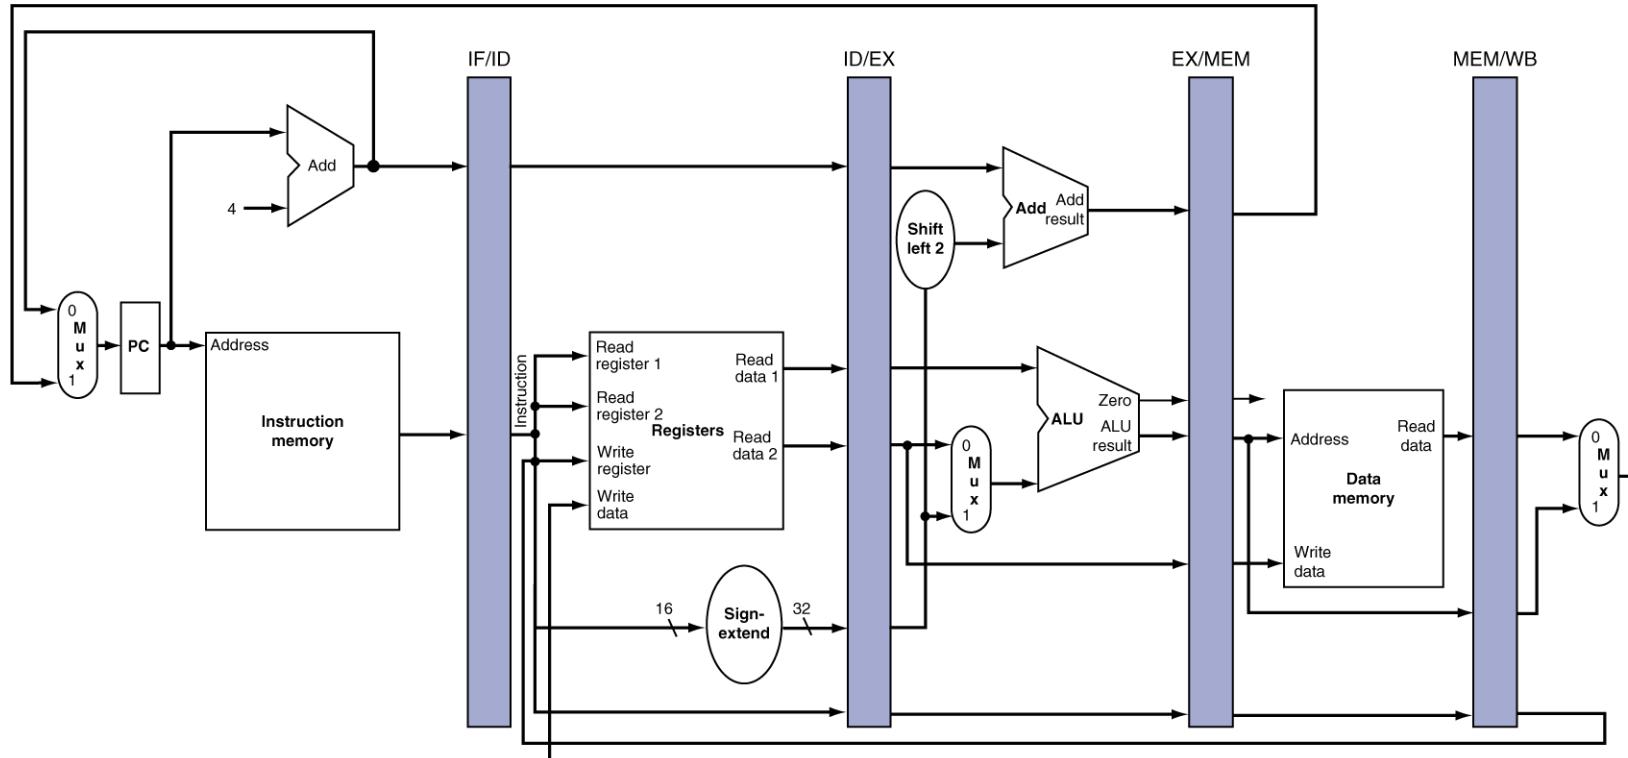
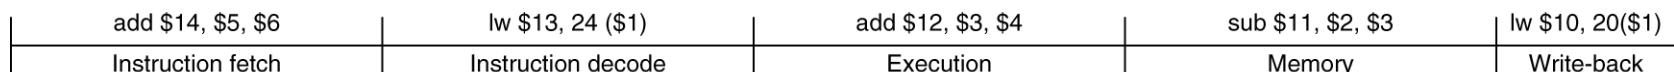
# Multi-Cycle Pipeline Diagram

- Form showing resource usage

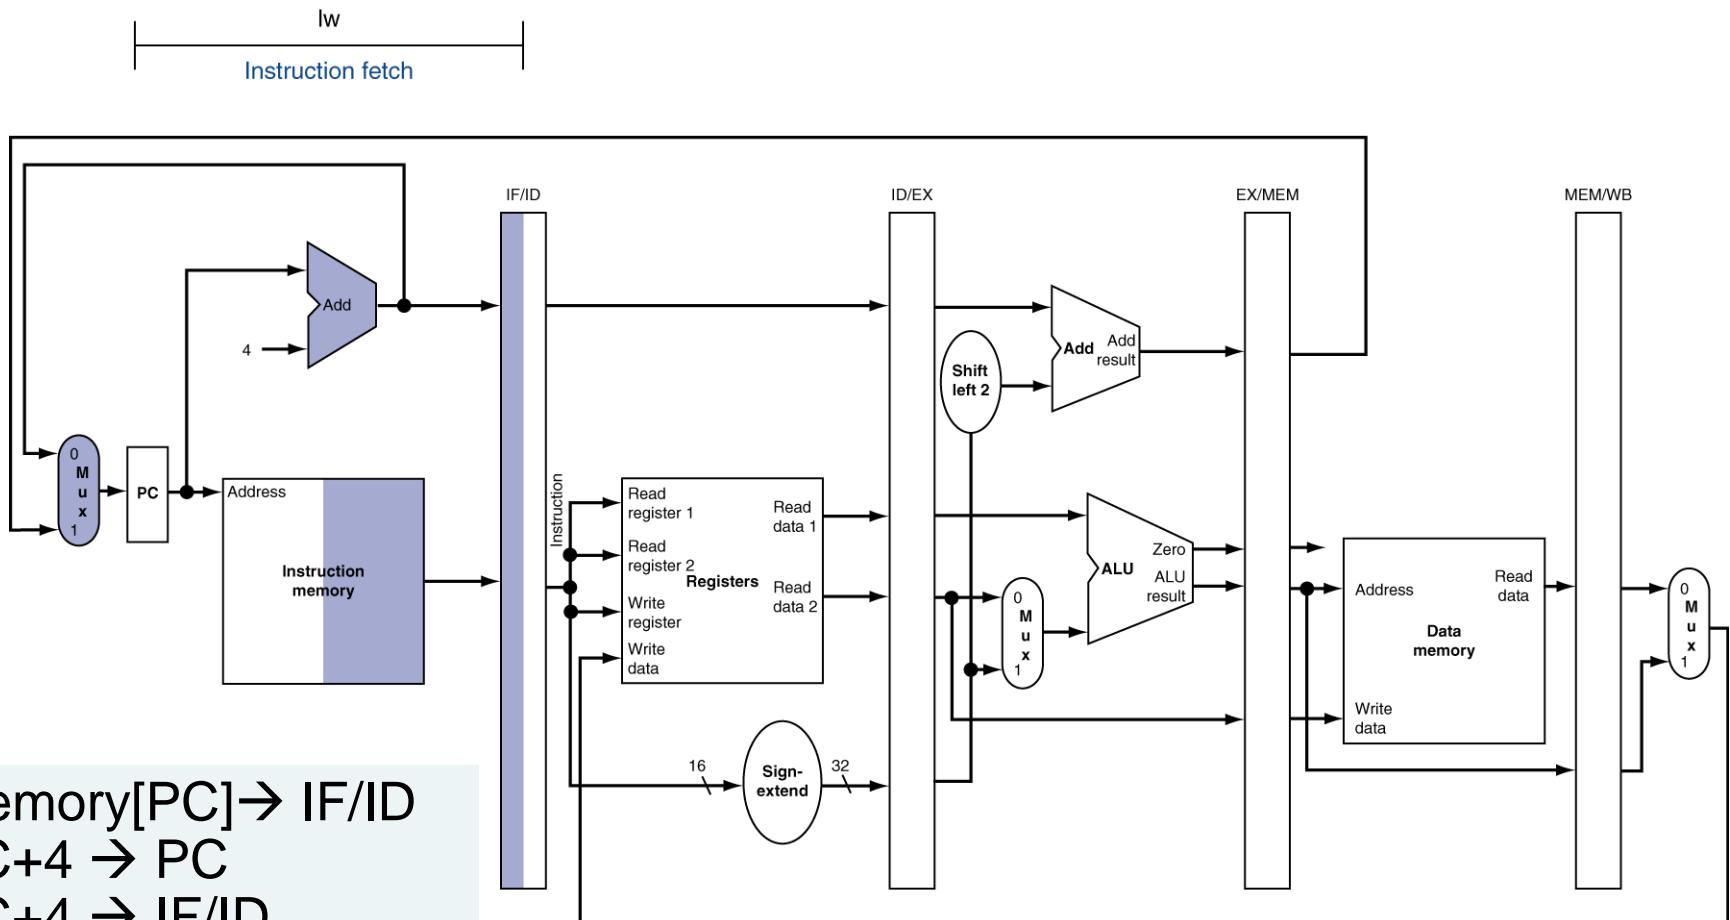


# Single-Cycle Pipeline Diagram

## State of pipeline in a given cycle



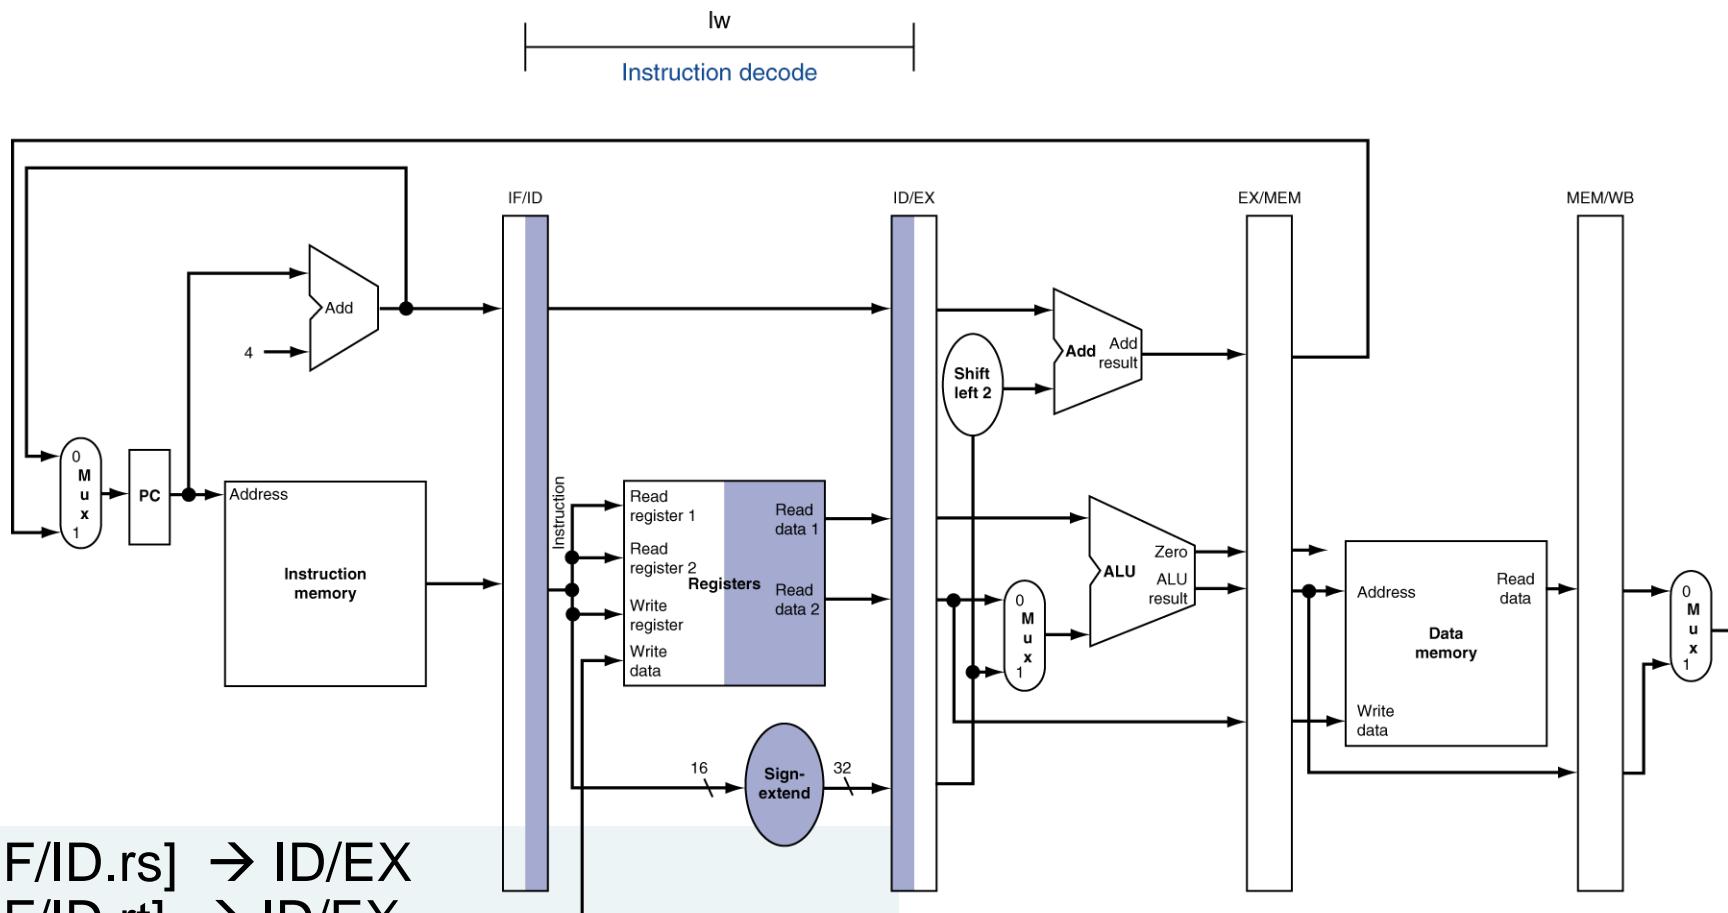
# IF for Load, Store, ...



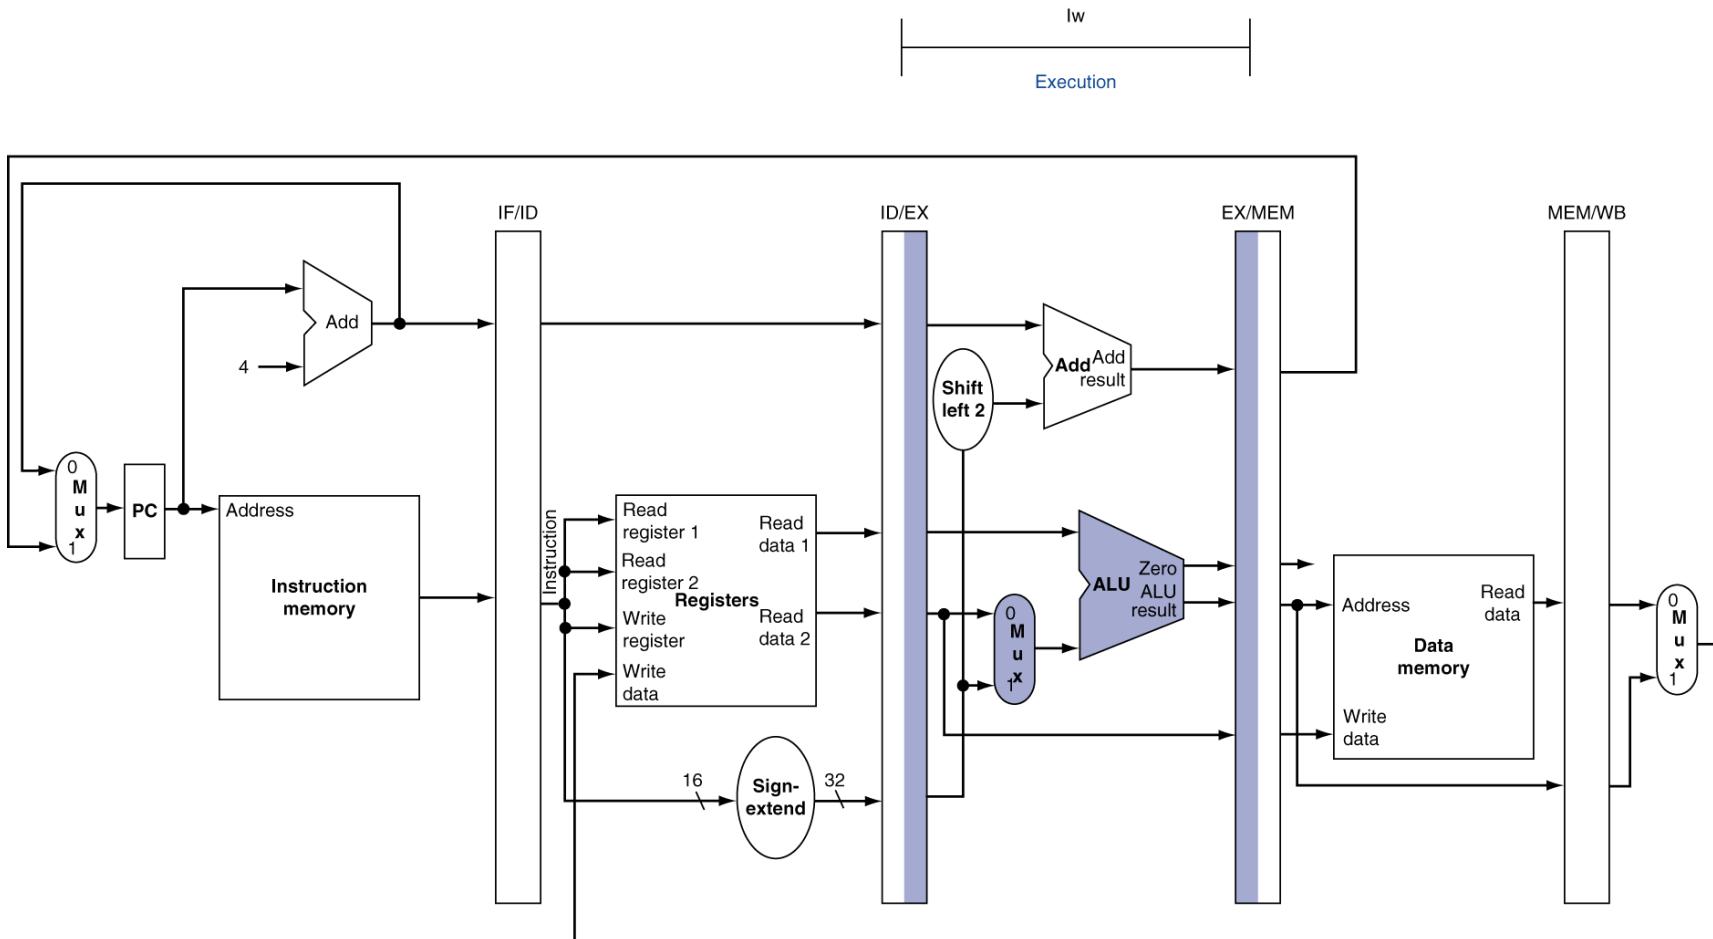
Note: Shaded left-half: Write

Shaded right-half: Read

# ID for Load, Store, ...



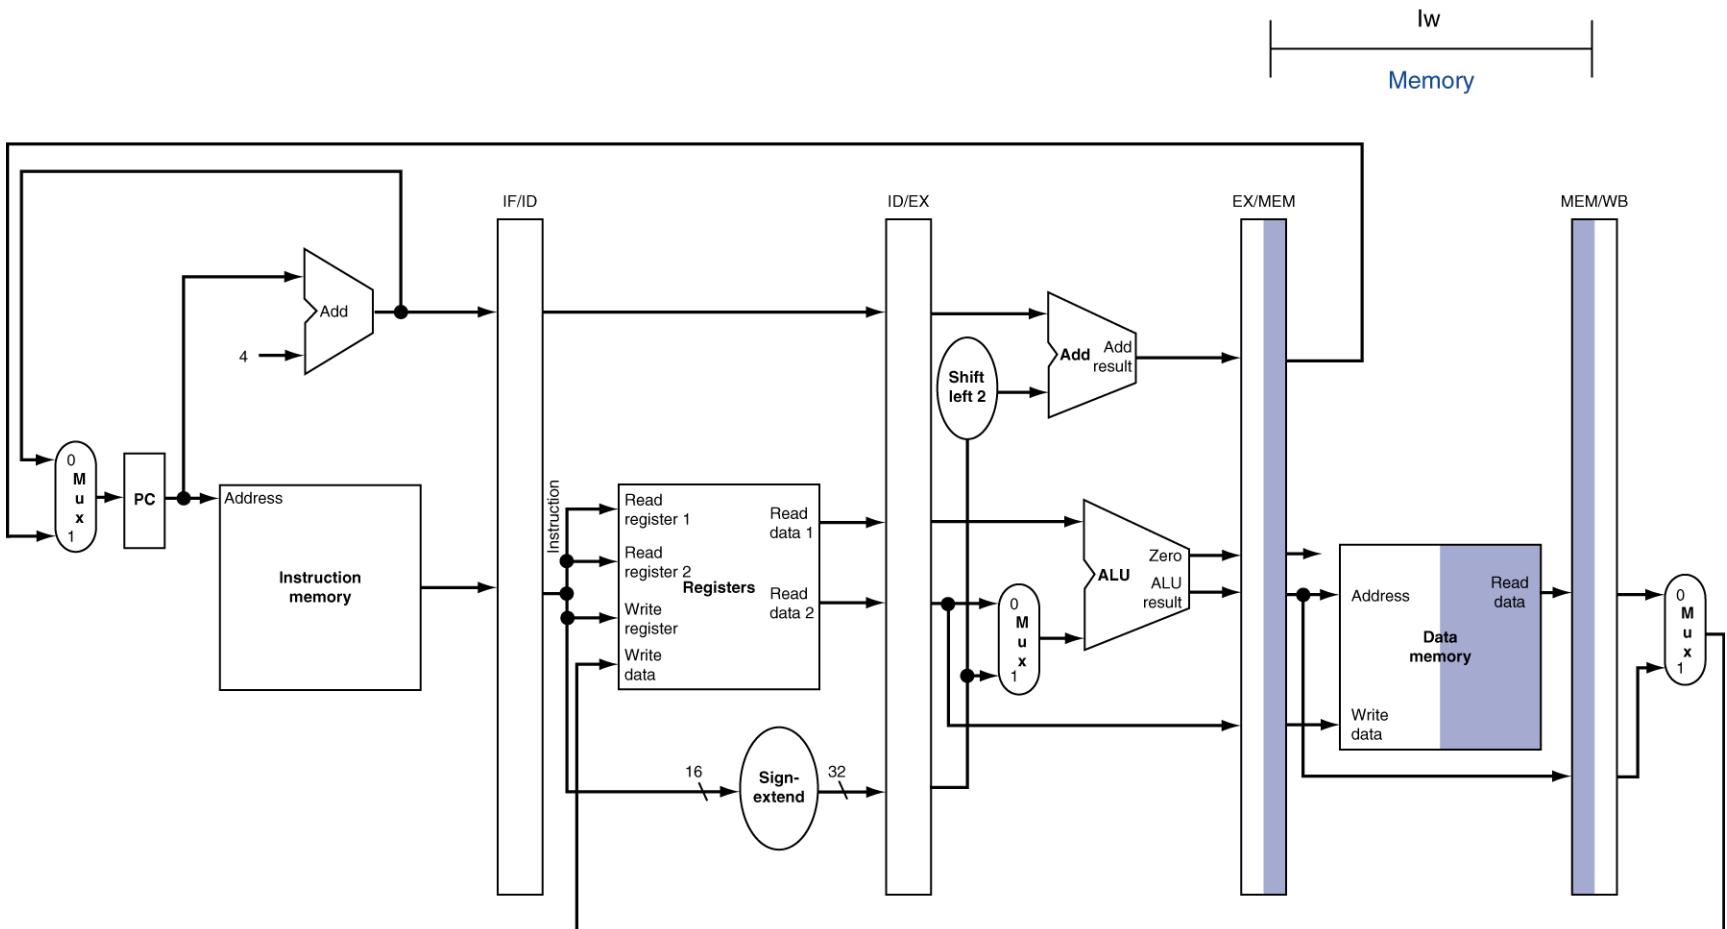
# EX for Load



$\text{mem-addr} = \text{ID/EX.reg[rs]} + \text{ID/EX.sign-ext32} \rightarrow \text{EX/MEM}$



# MEM for Load

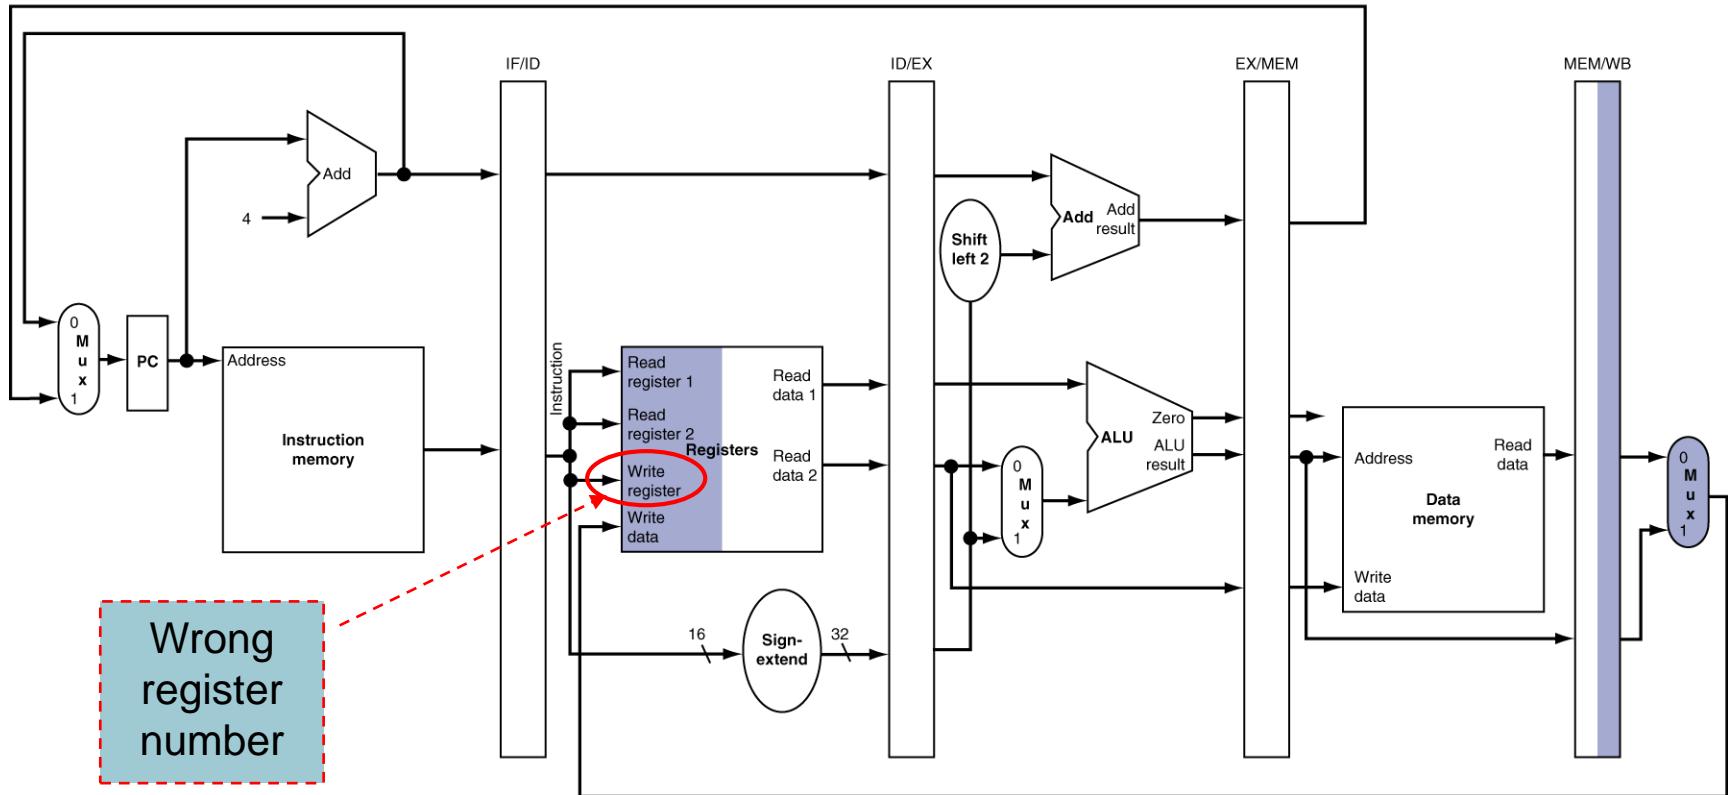


[mem-data= Memory[EX/EM. mem-addr] → MEM/WB]



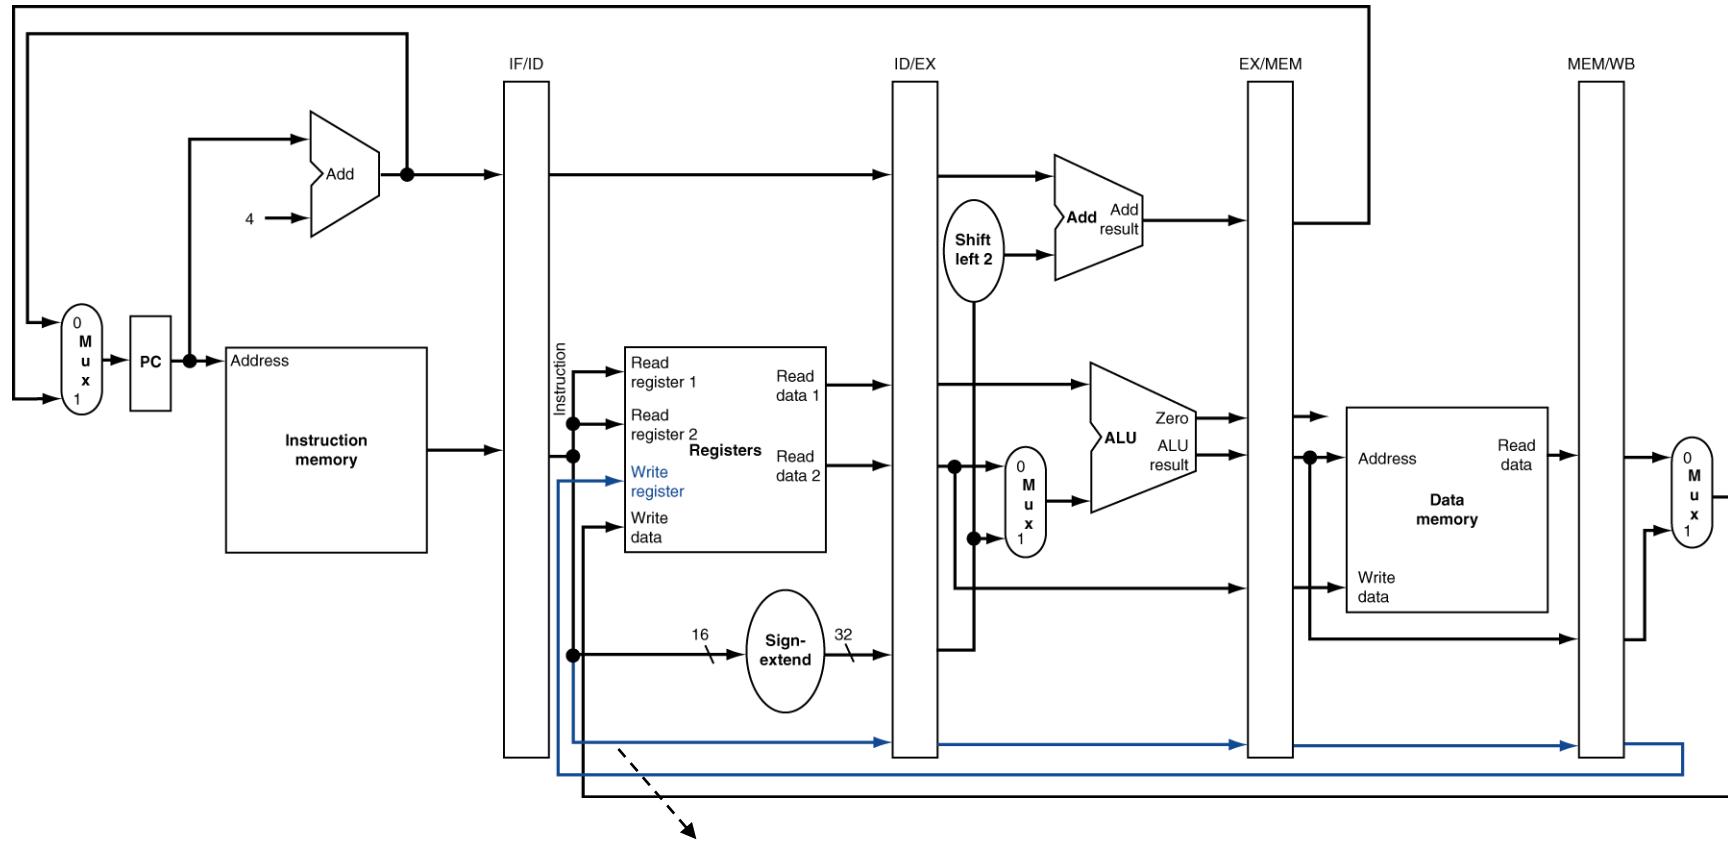
# WB for Load

lw  
Write back



MEM/WB.mem-data → Reg[ rt ]

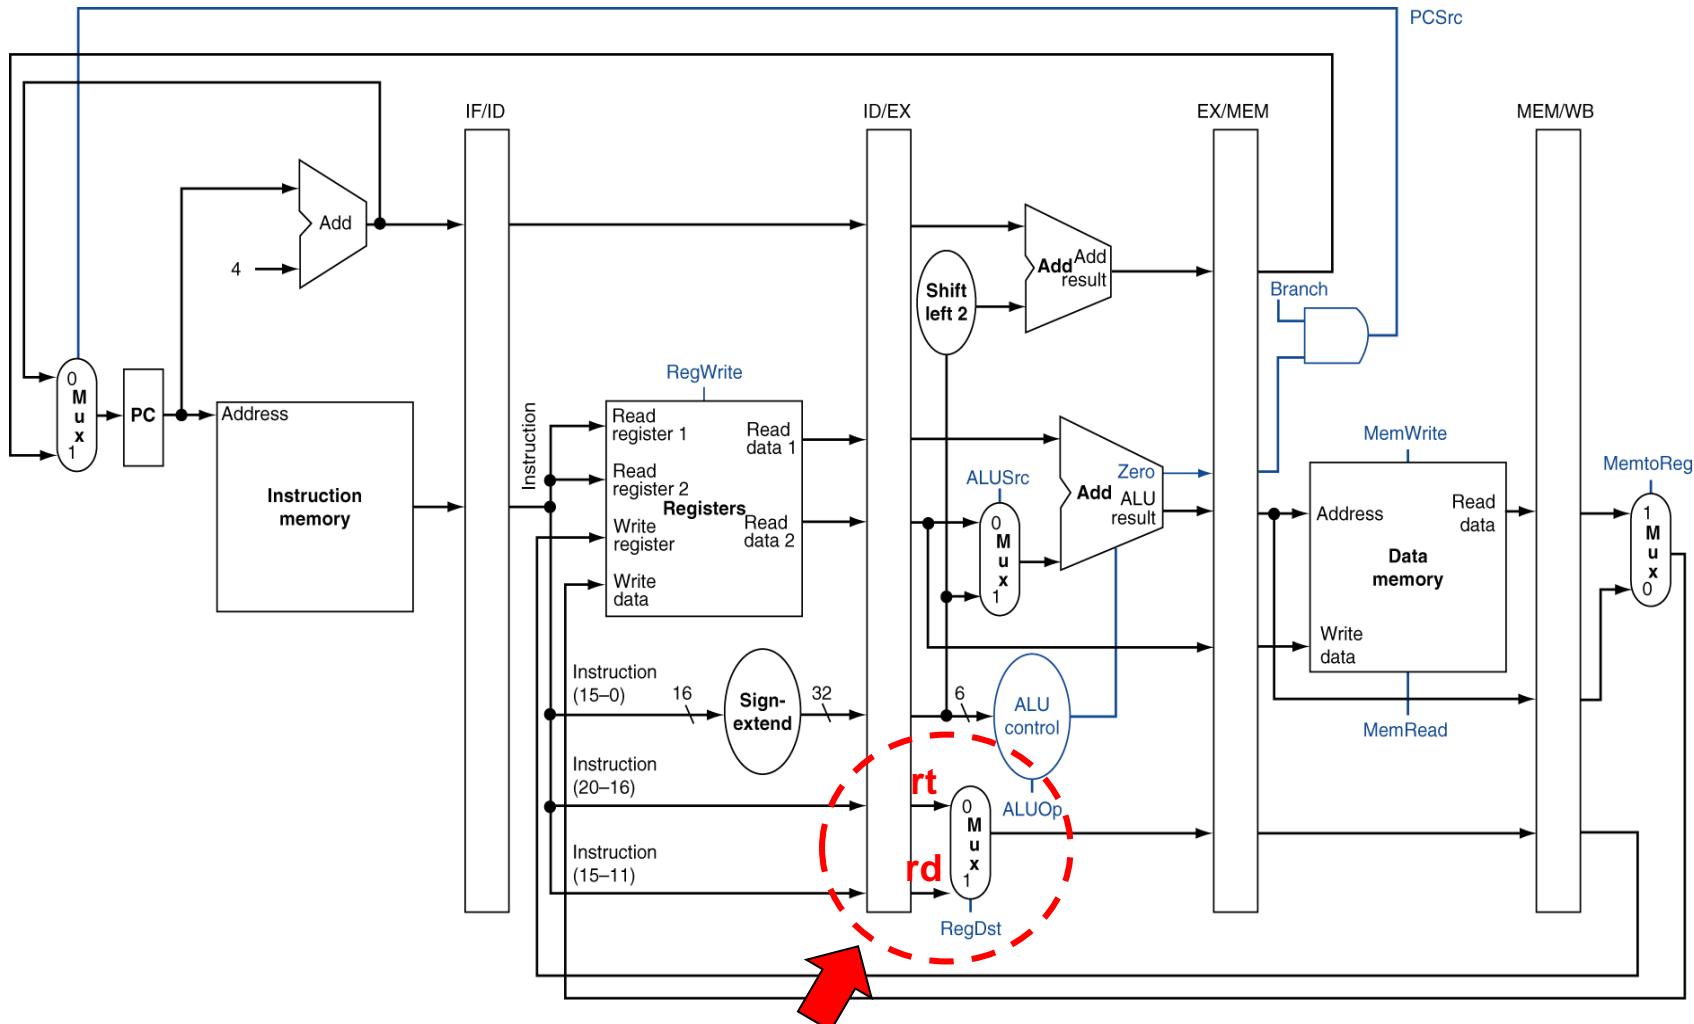
# Corrected Datapath for Load



IF/ID.rt → ID/EX → EX/MEM → MEM/WB for lw

How about R-type instruction ?

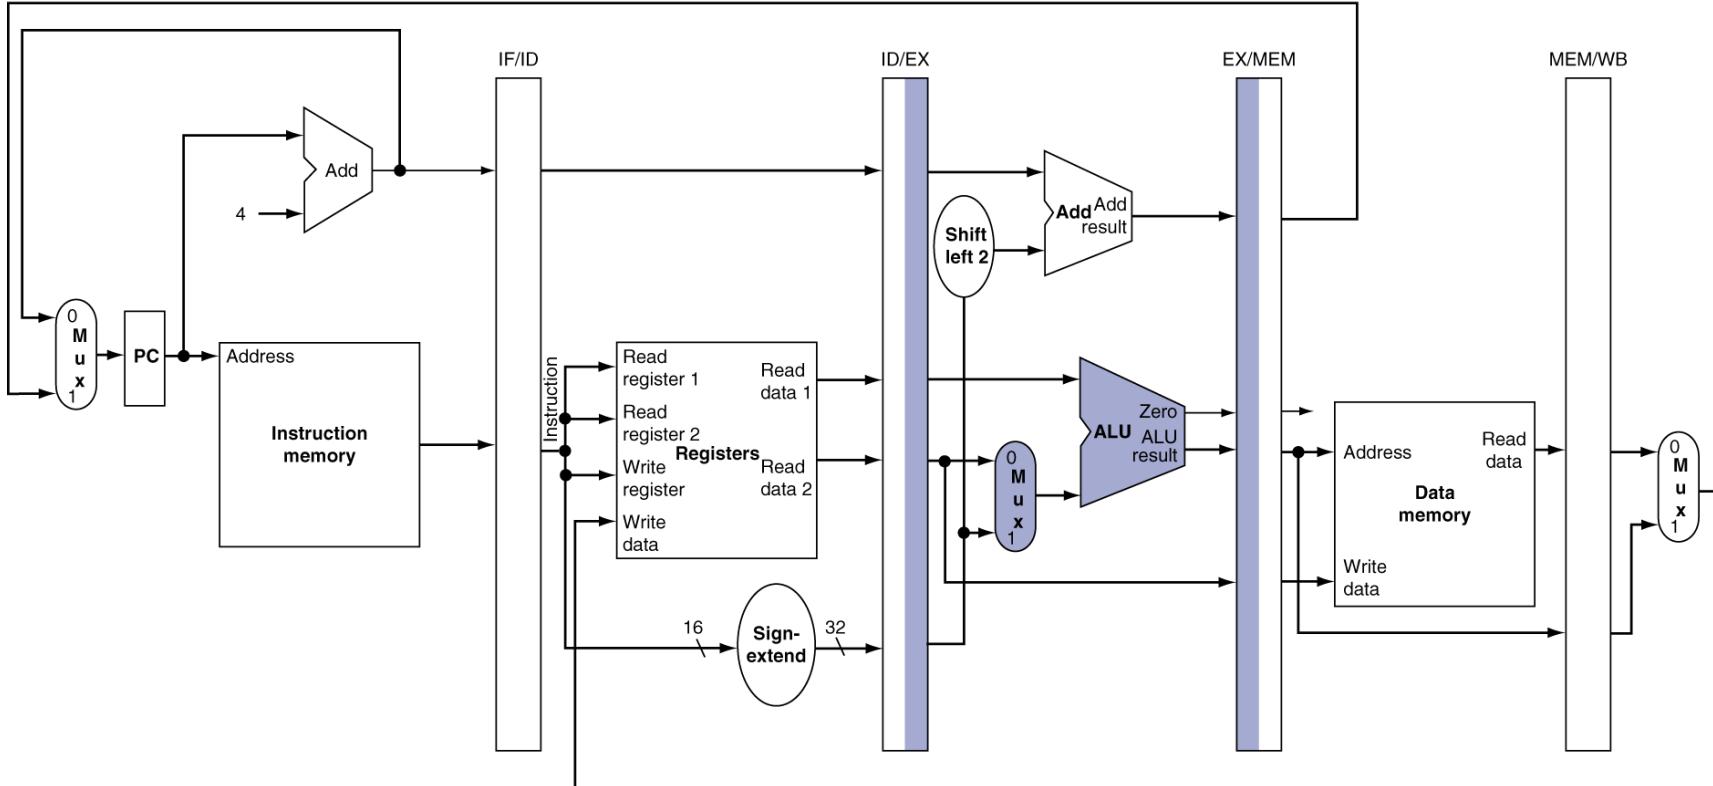
# Corrected Datapath for Load and R-type



# EX for Store

SW

Execution

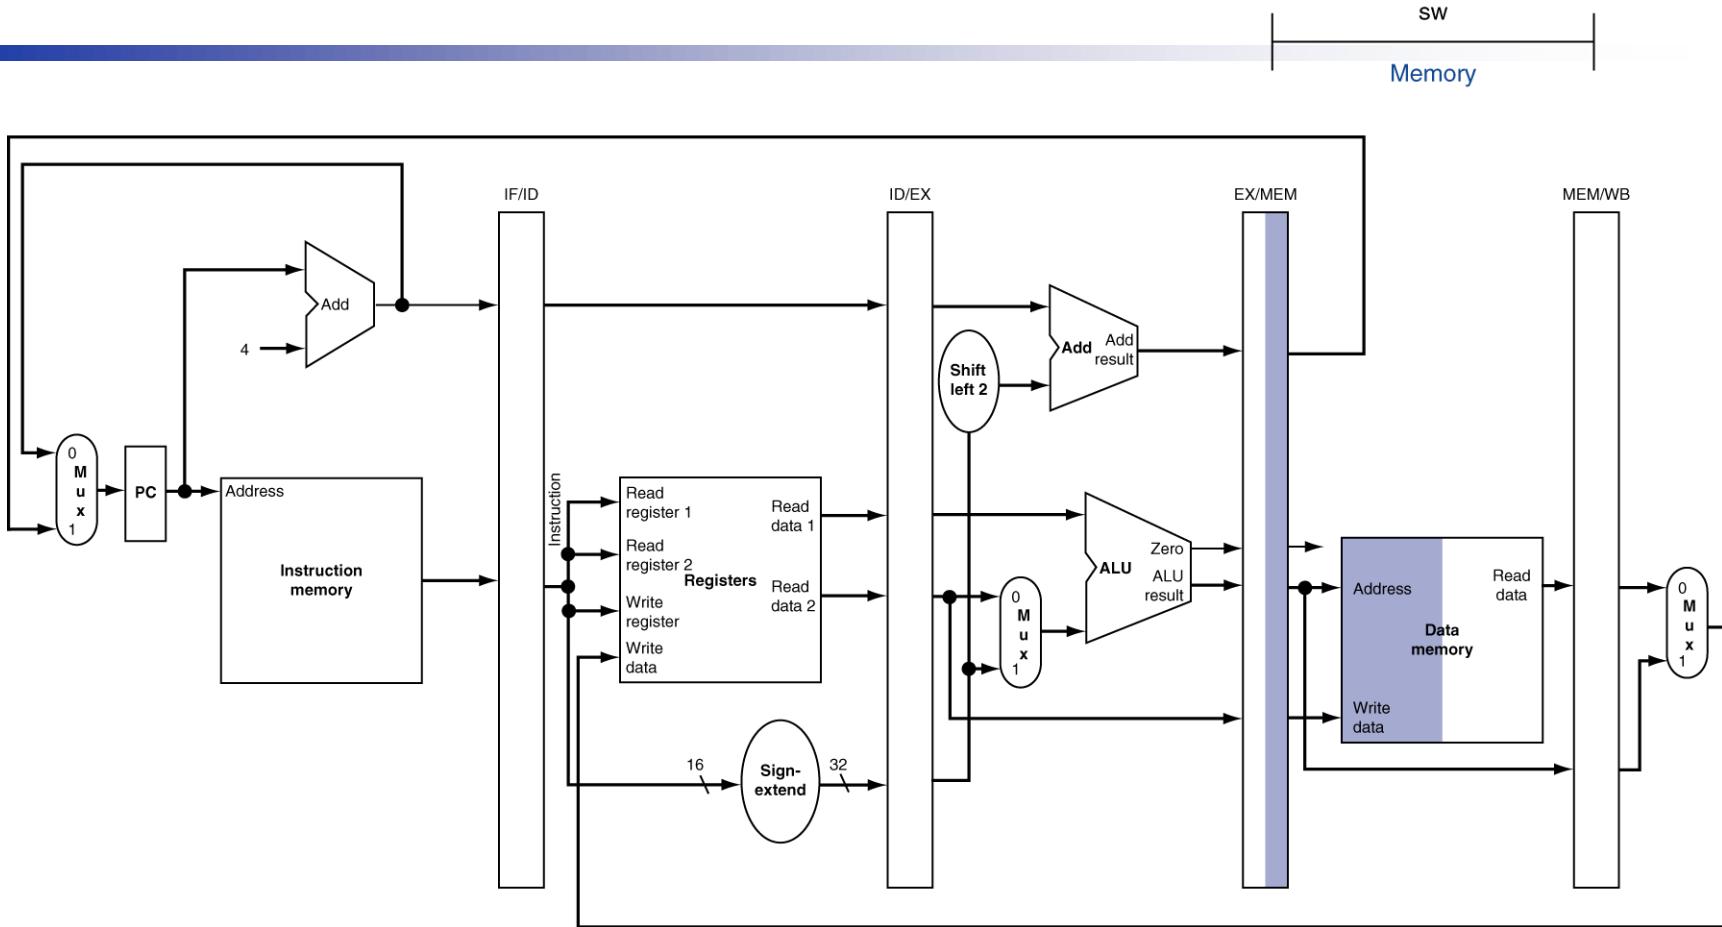


$$\text{mem-addr} = \text{ID/EX.reg[rs]} + \text{ID/EX.sign-ext32} \rightarrow \text{EX/MEM}$$

Anything else ? See next page



# MEM for Store

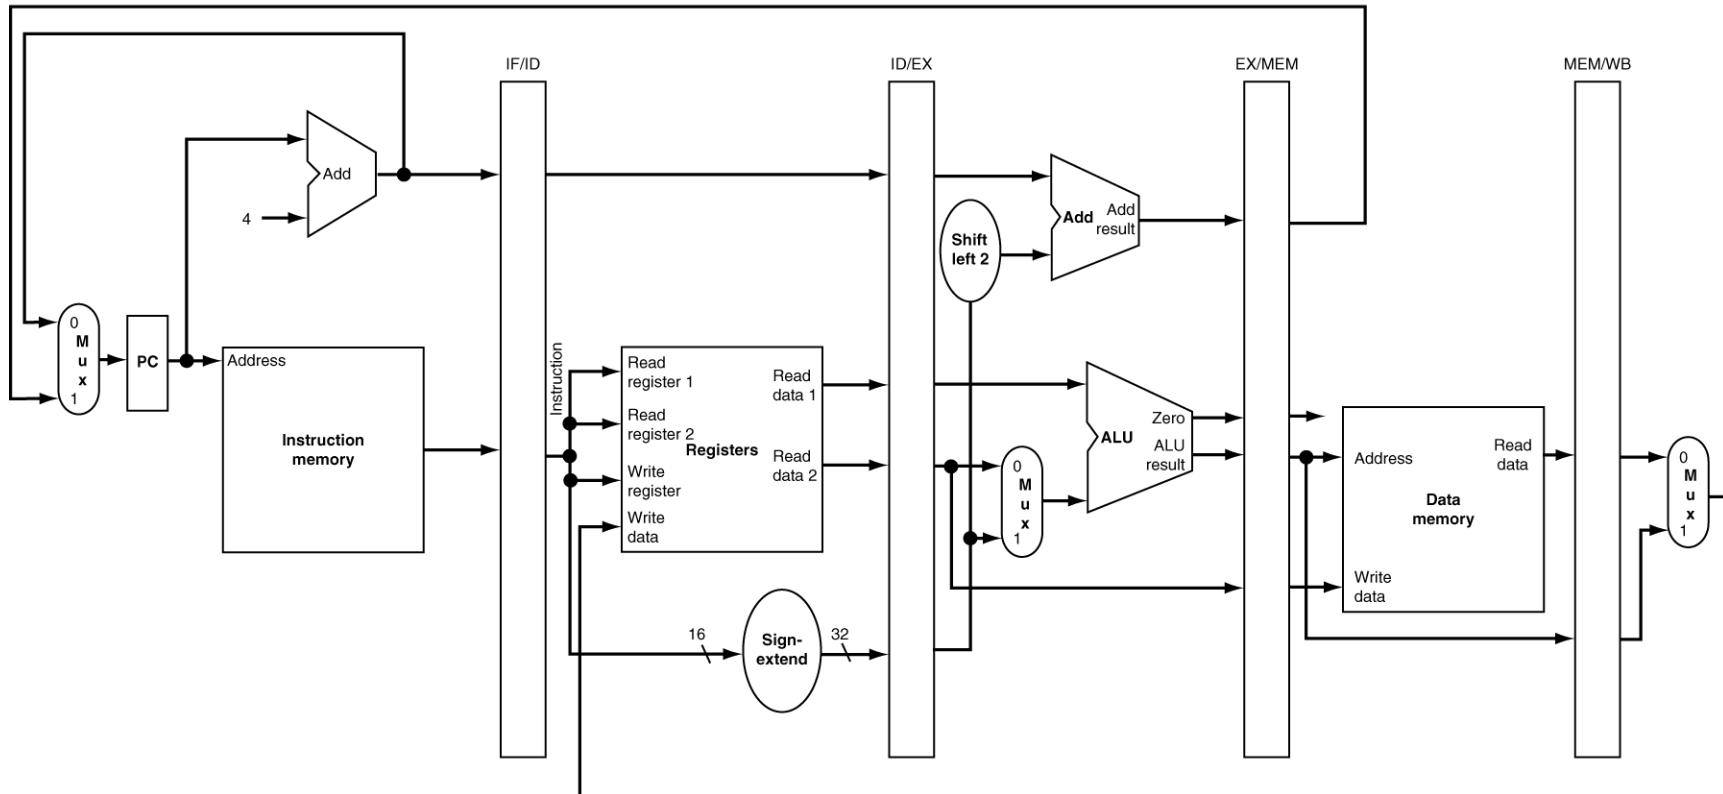


$\text{Reg}[ \text{rt} ] \rightarrow \text{MEM}[\text{EX/MEM. mem-addr}]$

How to get **Reg[rt]**? add “**ID/EX.Reg[rt] → EX/MEM**” to EX stage

# WB for Store

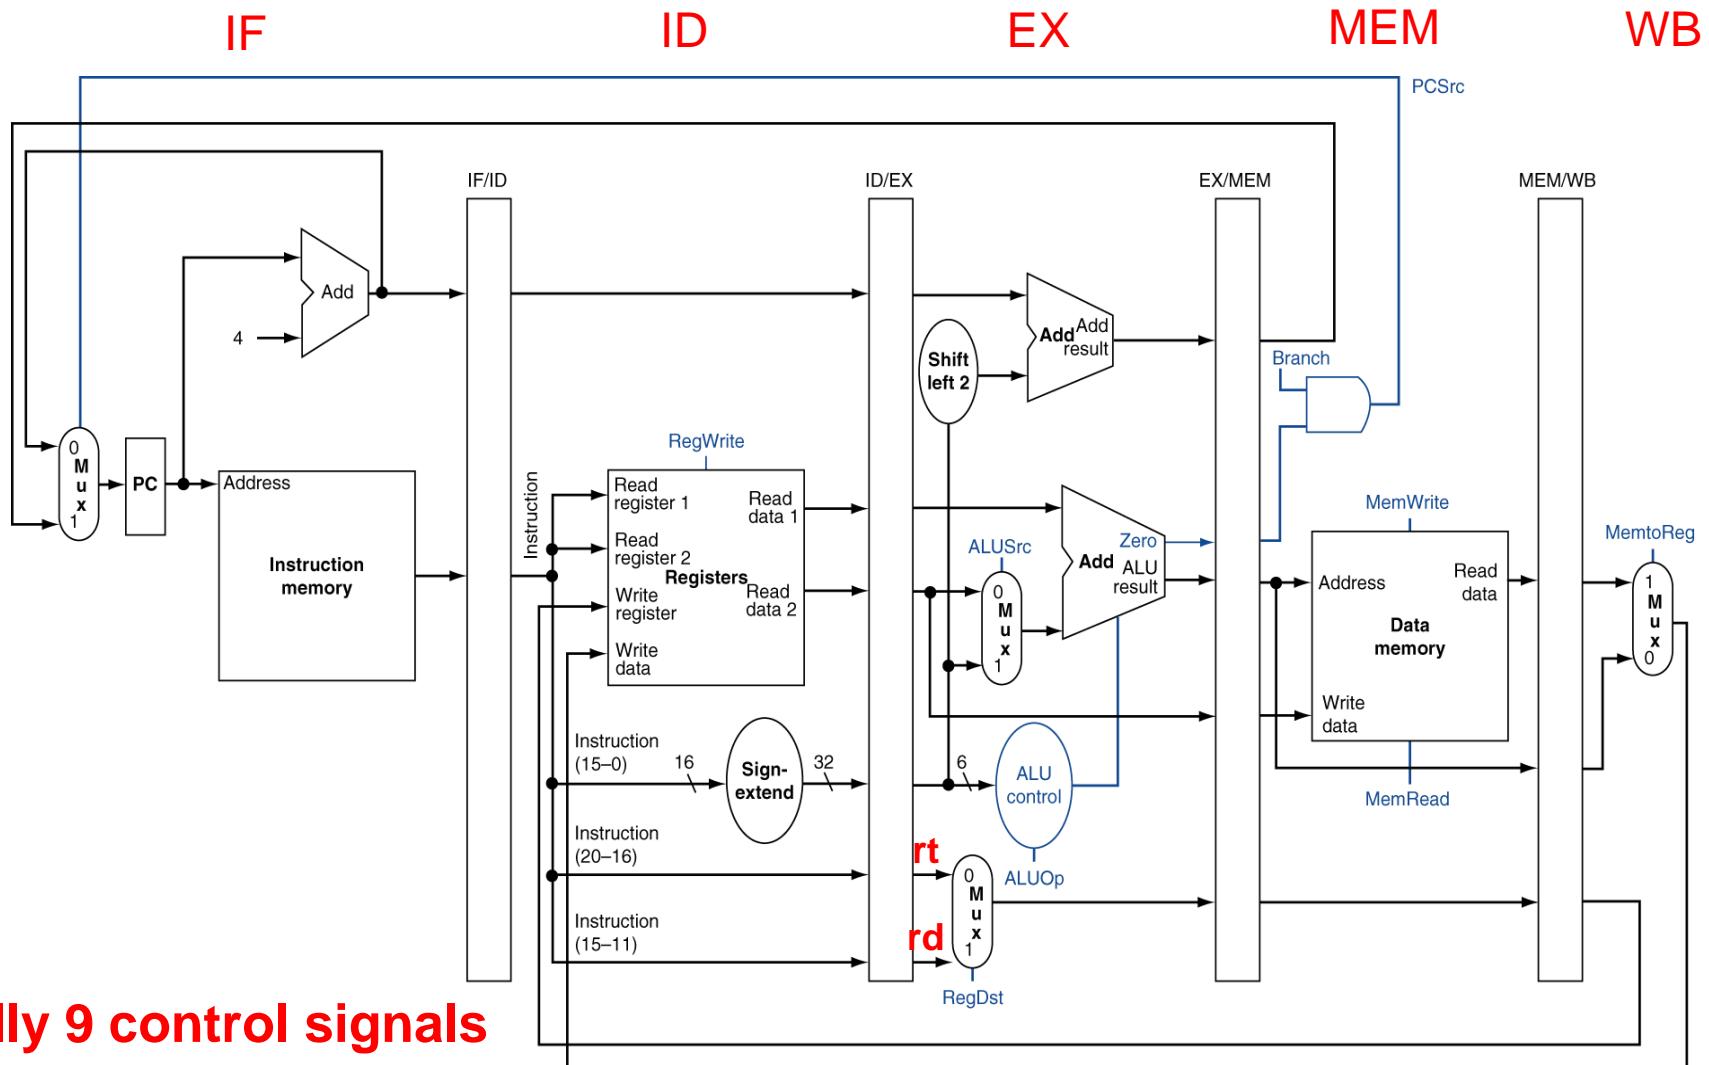
SW  
Write-back



Do nothing at this stage.

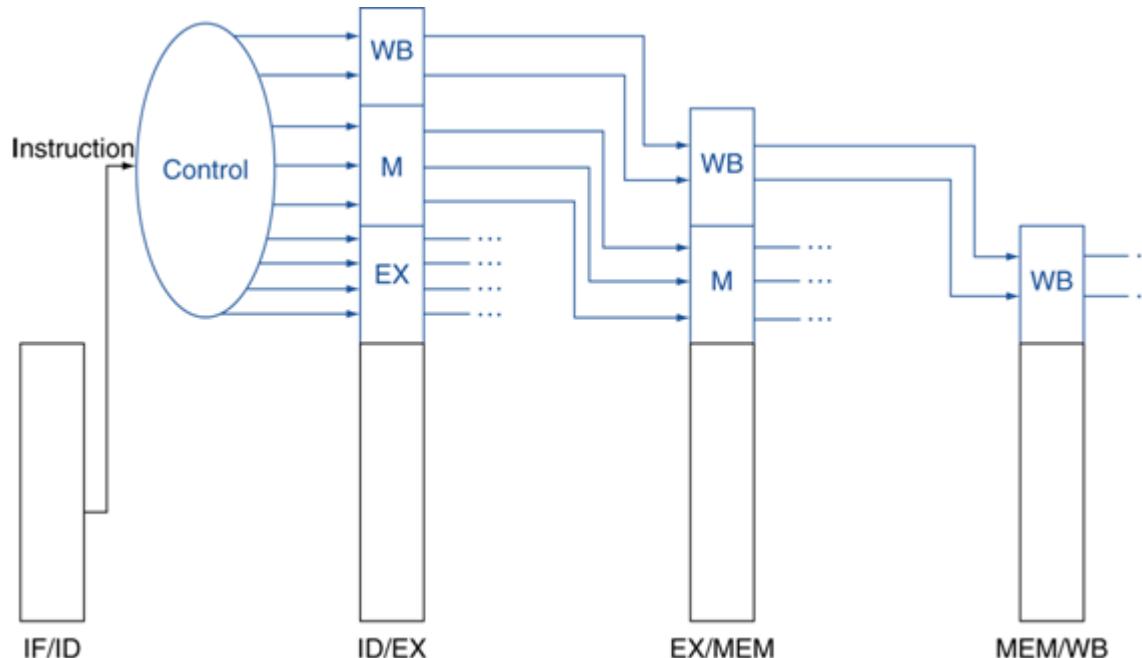


# Pipelined Control (Simplified)



# Pipelined Control

- Control signals derived from instruction
  - As in single-cycle implementation
- Pass control signals along just like the data
  - Create control info. at ID stage and then used in appropriate stage as pipeline move down

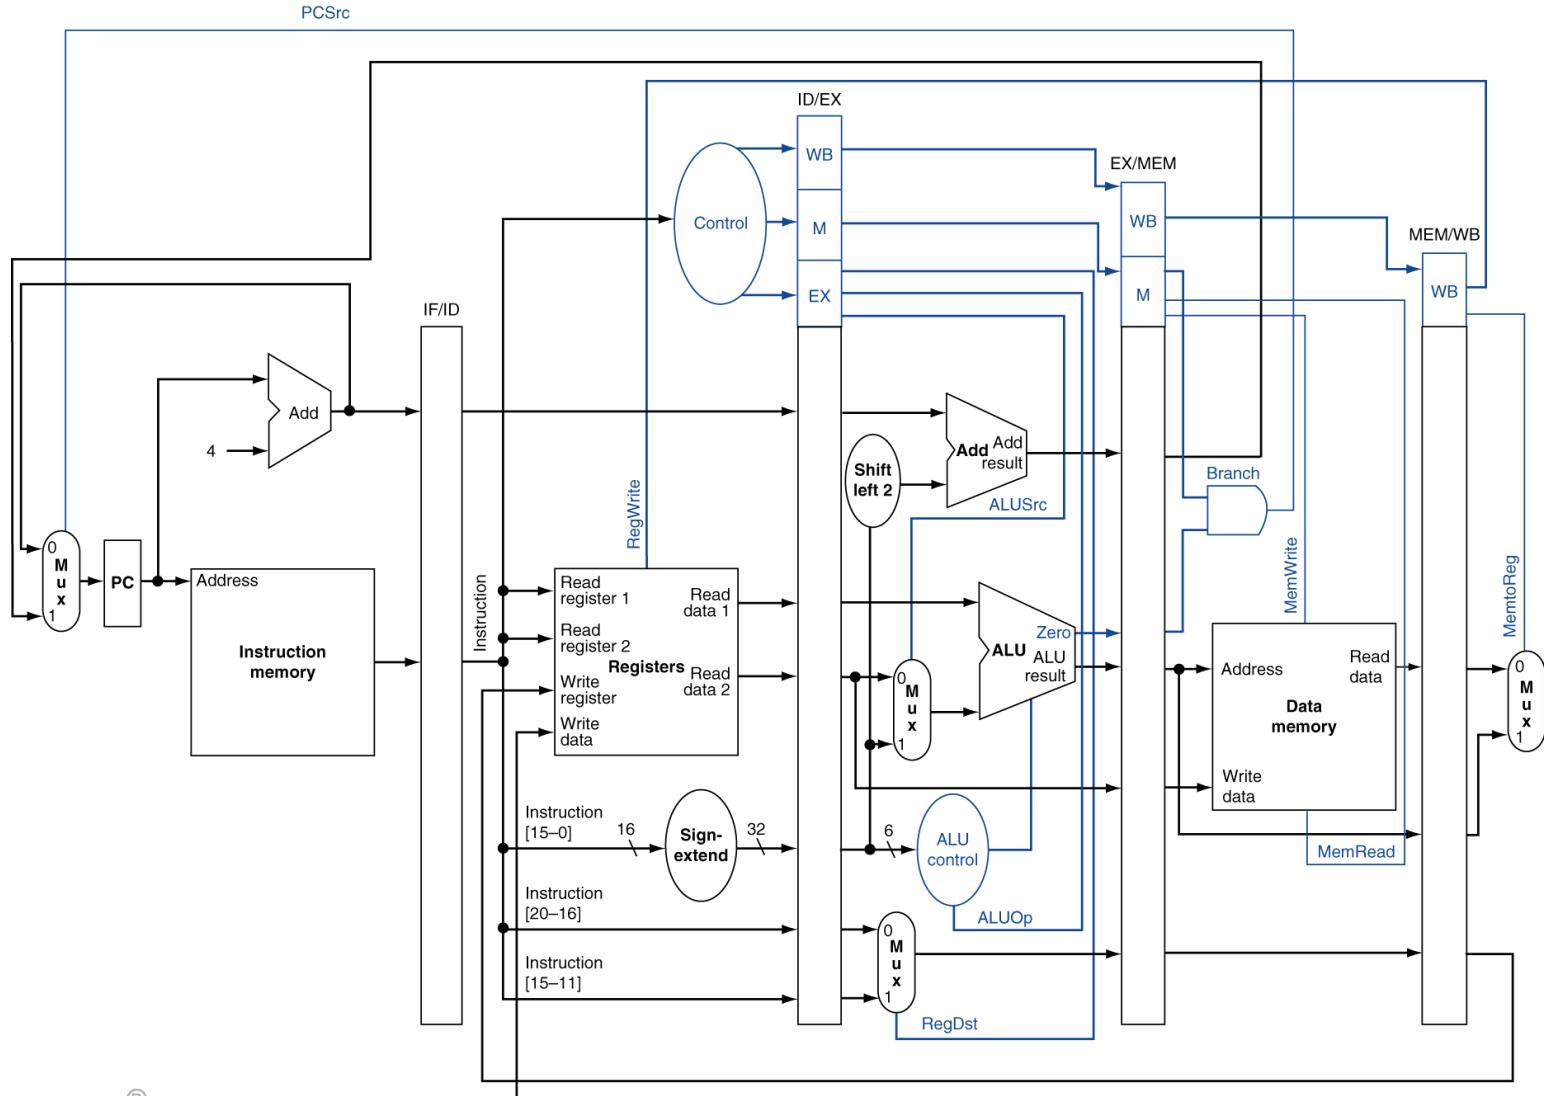


# Pipelined Control

| Instr.   | Execution/address calculation stage control lines |        |        |        | Memory access stage control lines |         |          | Write-back stage control lines |          |
|----------|---|--------|--------|--------|-----------------------------------|---------|----------|--------------------------------|----------|
|          | RegDst  | ALUOp1 | ALUOp0 | ALUSrc | Branch                            | MemRead | MemWrite | RegWrite                       | MemtoReg |
| R-format | 1   | 1      | 0      | 0      | 0                                 | 0       | 0        | 1                              | 0        |
| lw       | 0   | 0      | 0      | 1      | 0                                 | 1       | 0        | 1                              | 1        |
| sw       | X   | 0      | 0      | 1      | 0                                 | 0       | 1        | 0                              | X        |
| beq      | X   | 0      | 1      | 0      | 1                                 | 0       | 0        | 0                              | x        |

- **IF:** the control signals to read instruction memory and to write PC are always asserted for at each cycle, so there is nothing special to control at this stage.
- **ID:** As in IF, the same thing happens at every cycle, so no optional control signals to set.

# Pipelined Control



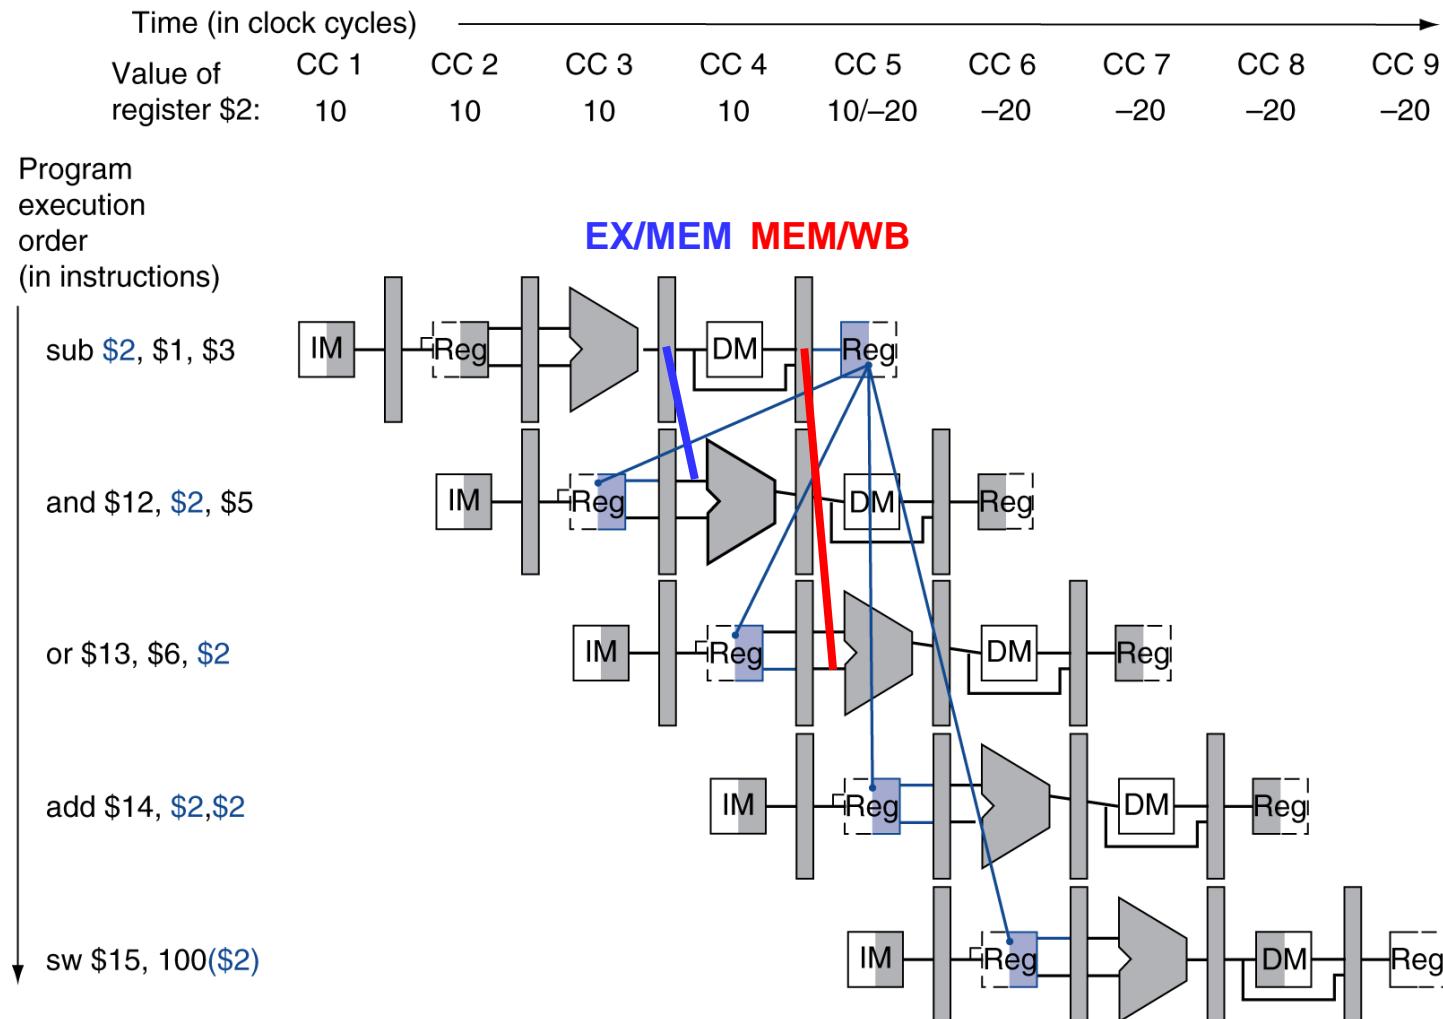
# Data Hazards in ALU Instructions

- Consider this sequence:

```
sub $2, $1, $3  
and $12, $2, $5  
or $13, $6, $2  
add $14, $2, $2  
sw $15, 100($2)
```

- We can resolve hazards with forwarding
  - How do we detect when to forward?

# Dependencies & Forwarding



# Detecting the Need to Forward

- Pass register numbers along pipeline
  - e.g., ID/EX.RegisterRs = register number for Rs sitting in ID/EX pipeline register
- ALU operand register numbers in EX stage are given by
  - ID/EX.RegisterRs, ID/EX.RegisterRt
- Data hazards when
  - 1a. EX/MEM.RegisterRd = ID/EX.RegisterRs
  - 1b. EX/MEM.RegisterRd = ID/EX.RegisterRt
  - 2a. MEM/WB.RegisterRd = ID/EX.RegisterRs
  - 2b. MEM/WB.RegisterRd = ID/EX.RegisterRt

Fwd from  
EX/MEM  
pipeline reg

Fwd from  
MEM/WB  
pipeline reg

# Detecting the Need to Forward

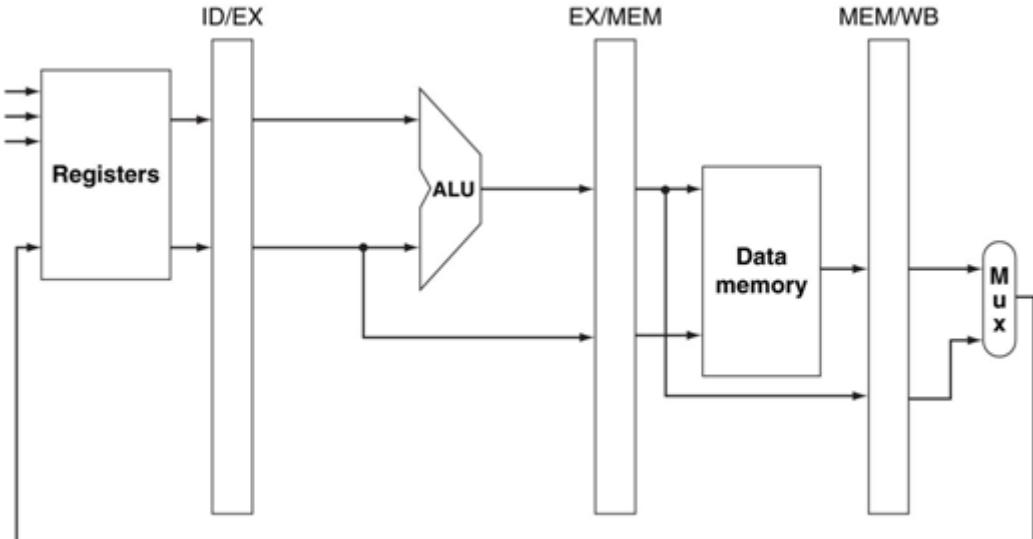
- But only if forwarding instruction will **write** to a register!
  - EX/MEM.RegWrite, MEM/WB.RegWrite
- And only if Rd for that instruction is **not \$zero**
  - EX/MEM.RegisterRd  $\neq 0$ ,  
MEM/WB.RegisterRd  $\neq 0$

# Forwarding Conditions

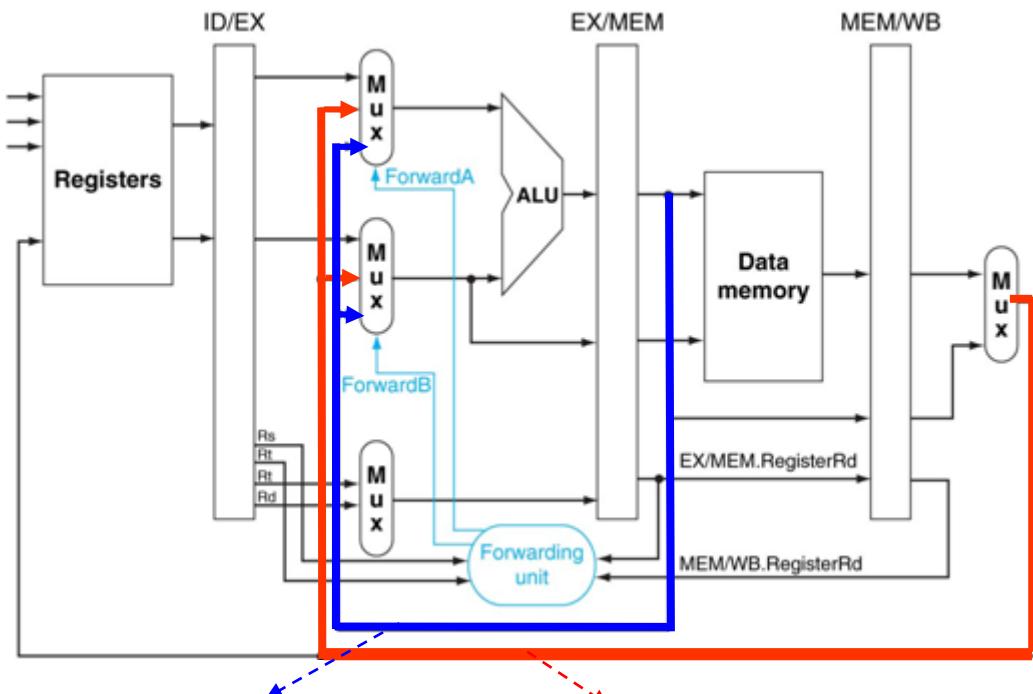
- EX hazard (*the sub-and* in p.83)
  - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0) and (EX/MEM.RegisterRd = ID/EX.RegisterRs))  
ForwardA = 10
  - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0) and (EX/MEM.RegisterRd = ID/EX.RegisterRt))  
ForwardB = 10
- MEM hazard (*the sub-or* in p.83)
  - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0) and (MEM/WB.RegisterRd = ID/EX.RegisterRs))  
ForwardA = 01
  - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0) and (MEM/WB.RegisterRd = ID/EX.RegisterRt))  
ForwardB = 01

# Forwarding

Without  
forwarding paths



With  
forwarding paths



1a. $\text{EX/MEM.RegRd} = \text{ID/EX.RegRs}$   
1b. $\text{EX/MEM.RegRd} = \text{ID/EX.RegRt}$

2a. $\text{MEM/WB.RegRd} = \text{ID/EX.RegRs}$   
2b. $\text{MEM/WB.RegRd} = \text{ID/EX.RegRt}$



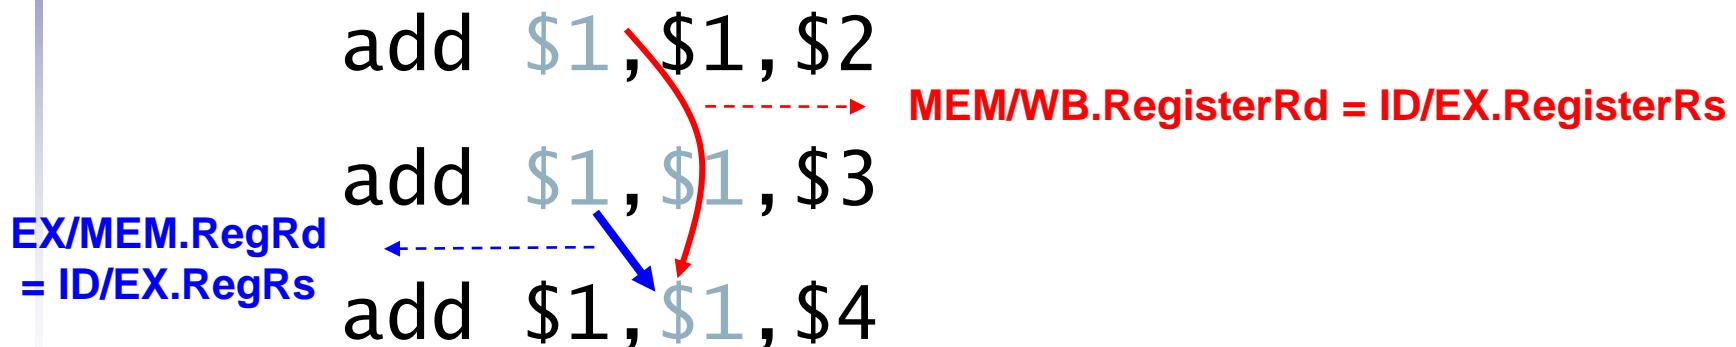
# Forwarding Control

The control values for the forwarding multiplexors

| Mux control   | Source | Explanation  |
|---------------|--------|--|
| ForwardA = 00 | ID/EX  | The first ALU operand comes from the register file.                            |
| ForwardA = 10 | EX/MEM | The first ALU operand is forwarded from the prior ALU result.                  |
| ForwardA = 01 | MEM/WB | The first ALU operand is forwarded from data memory or an earlier ALU result.  |
| ForwardB = 00 | ID/EX  | The second ALU operand comes from the register file.                           |
| ForwardB = 10 | EX/MEM | The second ALU operand is forwarded from the prior ALU result.                 |
| ForwardB = 01 | MEM/WB | The second ALU operand is forwarded from data memory or an earlier ALU result. |

# Double Data Hazard

- Consider the sequence:



- Both hazards occur
  - Want to use the most recent
- Revise MEM hazard condition
  - Only fwd if EX hazard condition isn't true

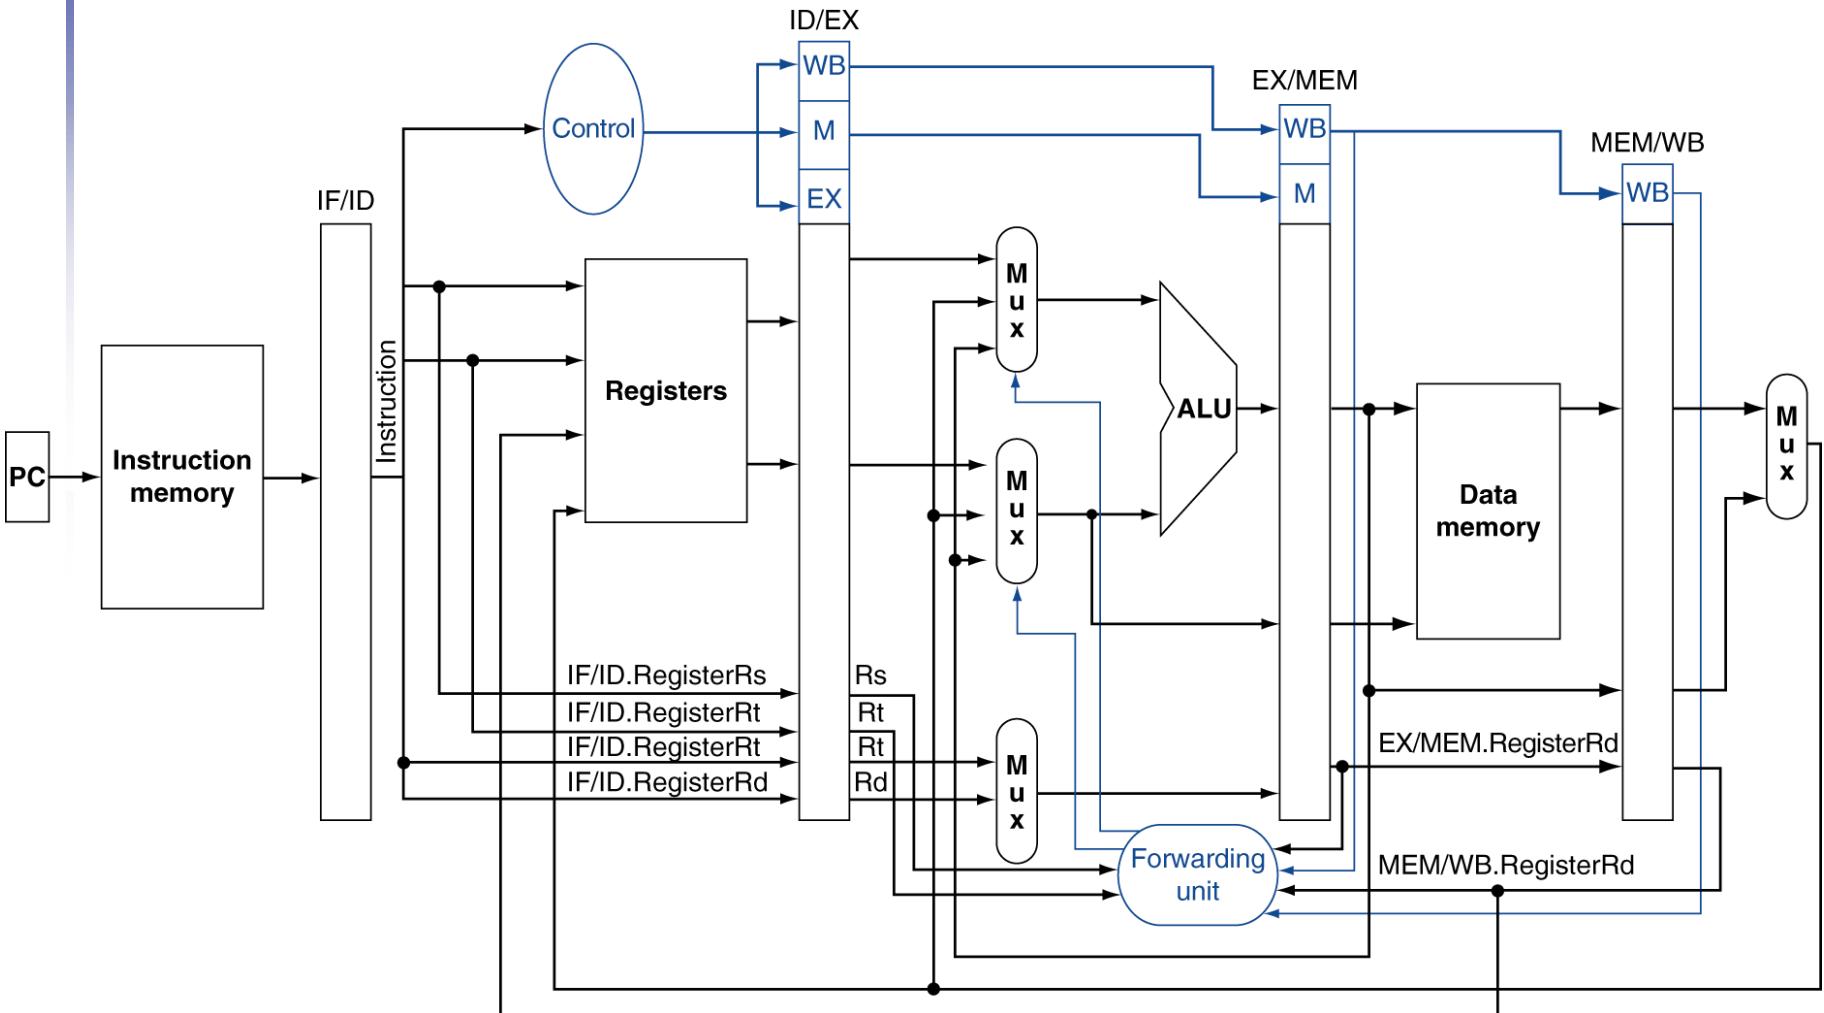
# Revised Forwarding Condition

## ■ MEM hazard

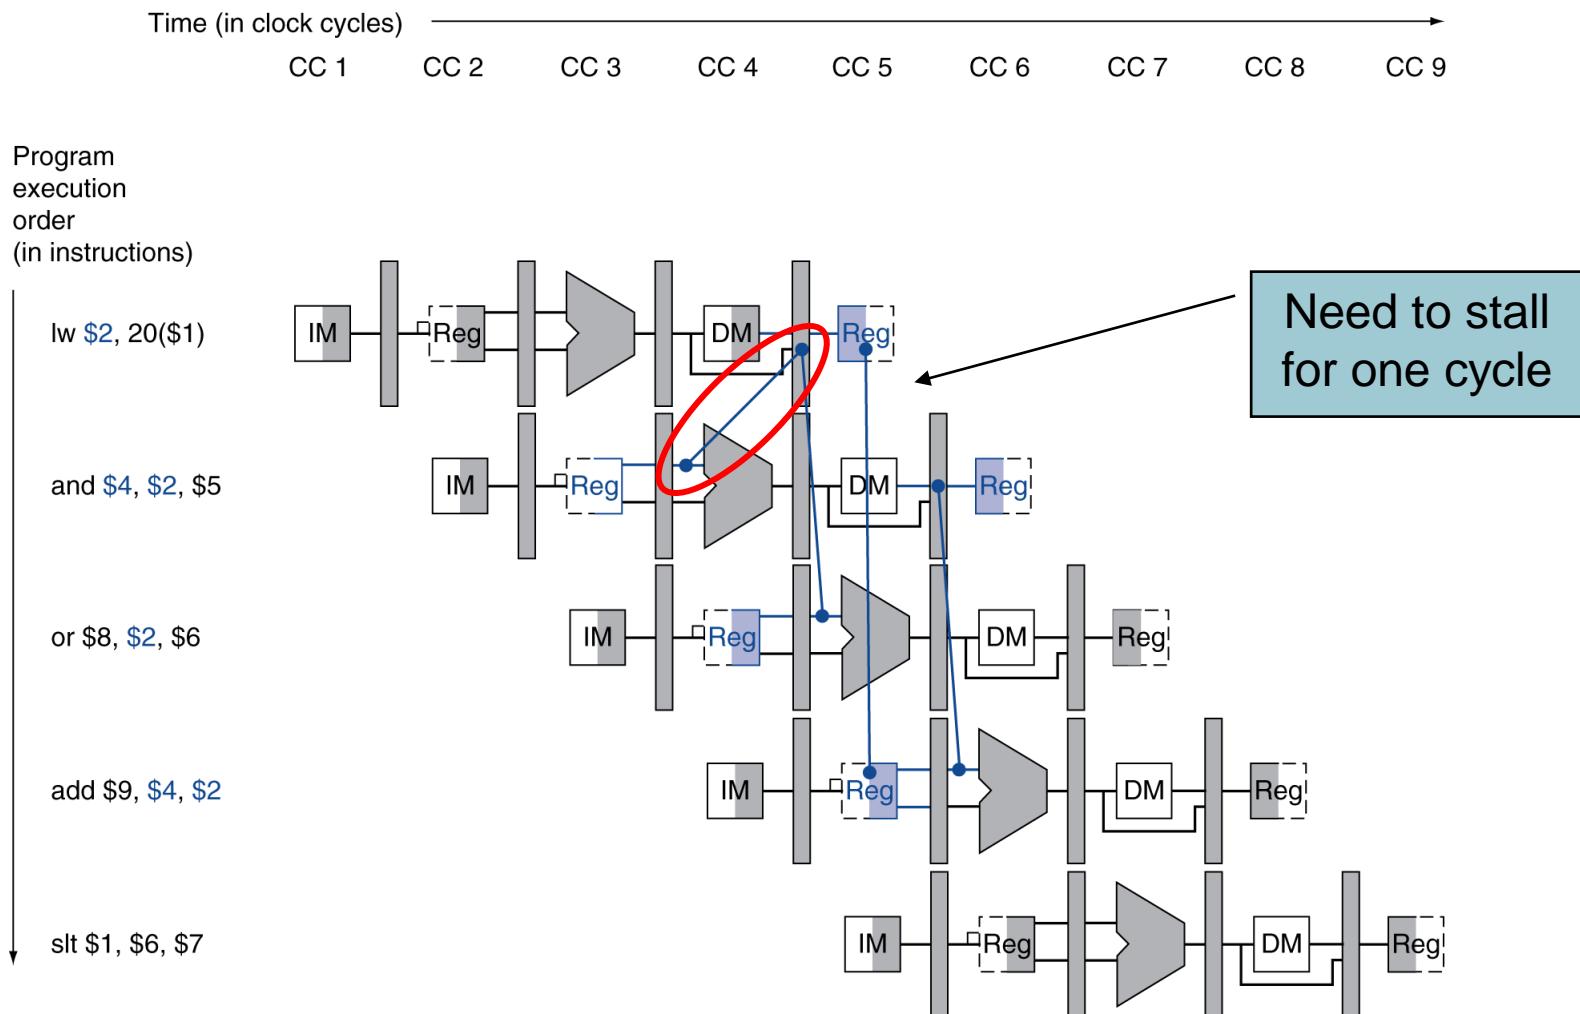
- if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)  
and **not** (EX/EM.RegWrite and (EX/EM.RegisterRd ≠ 0)  
and (EX/EM.RegisterRd = ID/EX.RegisterRs))  
and (MEM/WB.RegisterRd = ID/EX.RegisterRs))  
ForwardA = 01
- if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)  
and **not** (EX/EM.RegWrite and (EX/EM.RegisterRd ≠ 0)  
and (EX/EM.RegisterRd = ID/EX.RegisterRt))  
and (MEM/WB.RegisterRd = ID/EX.RegisterRt))  
ForwardB = 01

EX hazard

# Datapath with Forwarding



# Load-Use Data Hazard



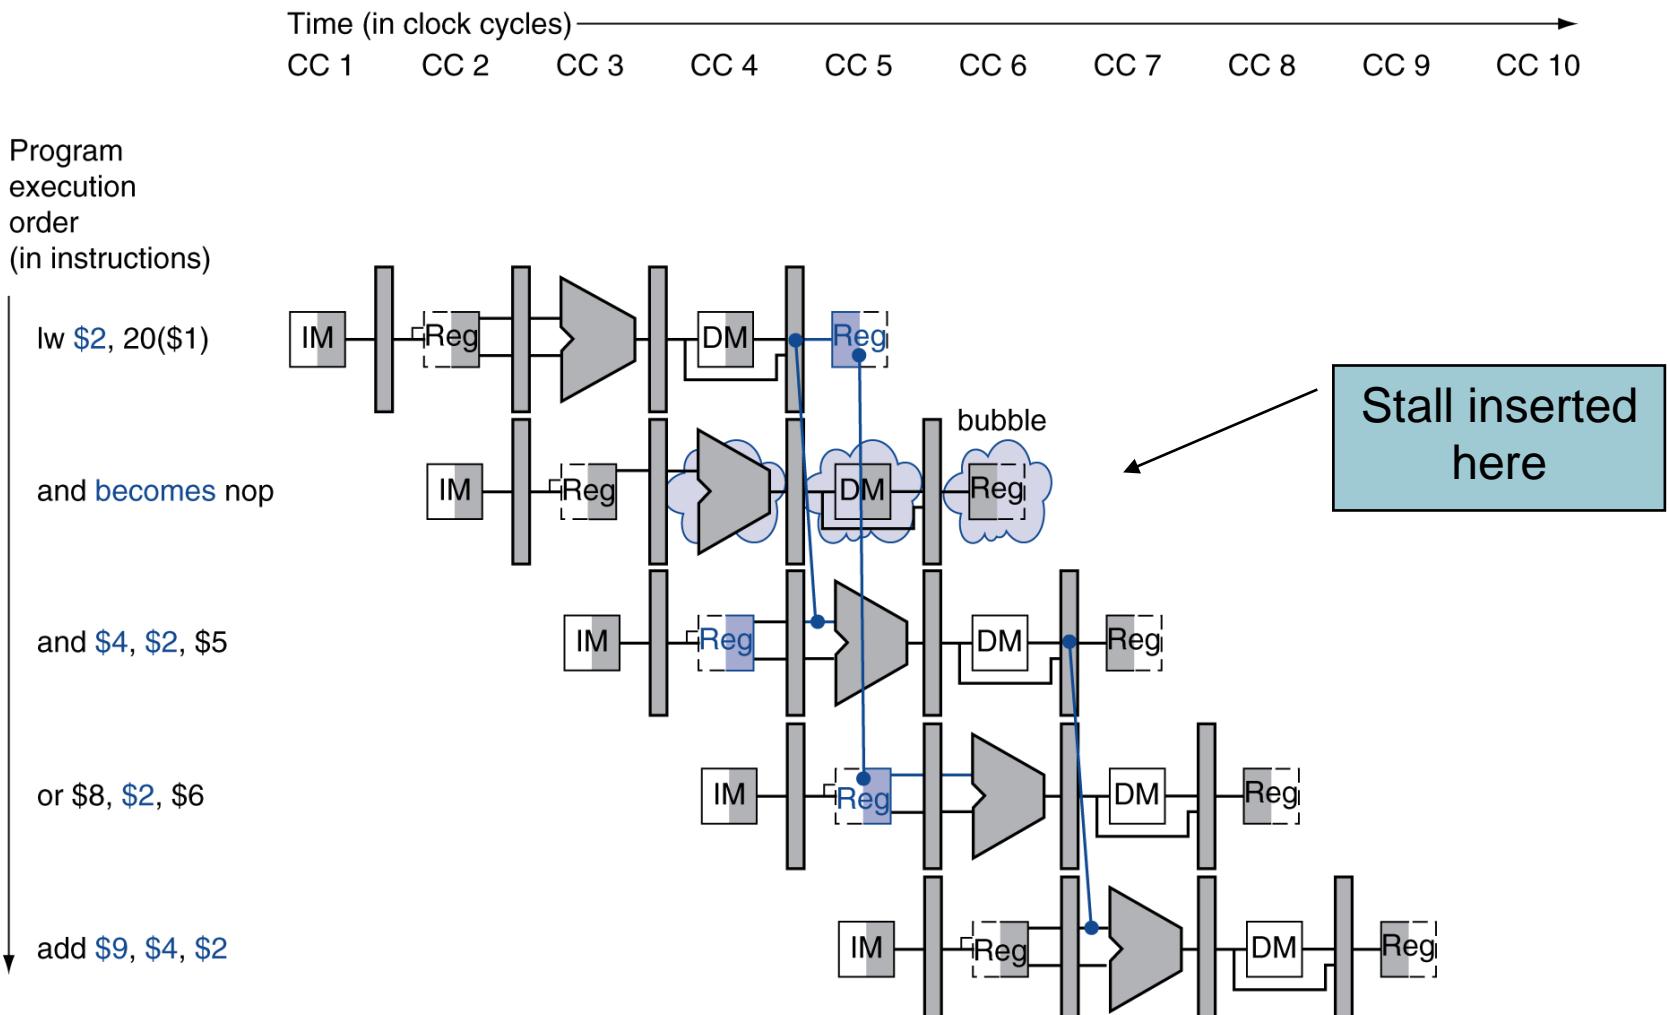
# Load-Use Hazard Detection

- Check when using instruction is decoded in **ID stage**
- ALU operand register numbers in ID stage are given by
  - IF/ID.RegisterRs, IF/ID.RegisterRt
- Load-use hazard when
  - ID/EX.MemRead and A load instr.  
 $((ID/EX.RegisterRt = IF/ID.RegisterRs) \text{ or } (ID/EX.RegisterRt = IF/ID.RegisterRt))$
- If detected, stall and insert bubble

# How to Stall the Pipeline

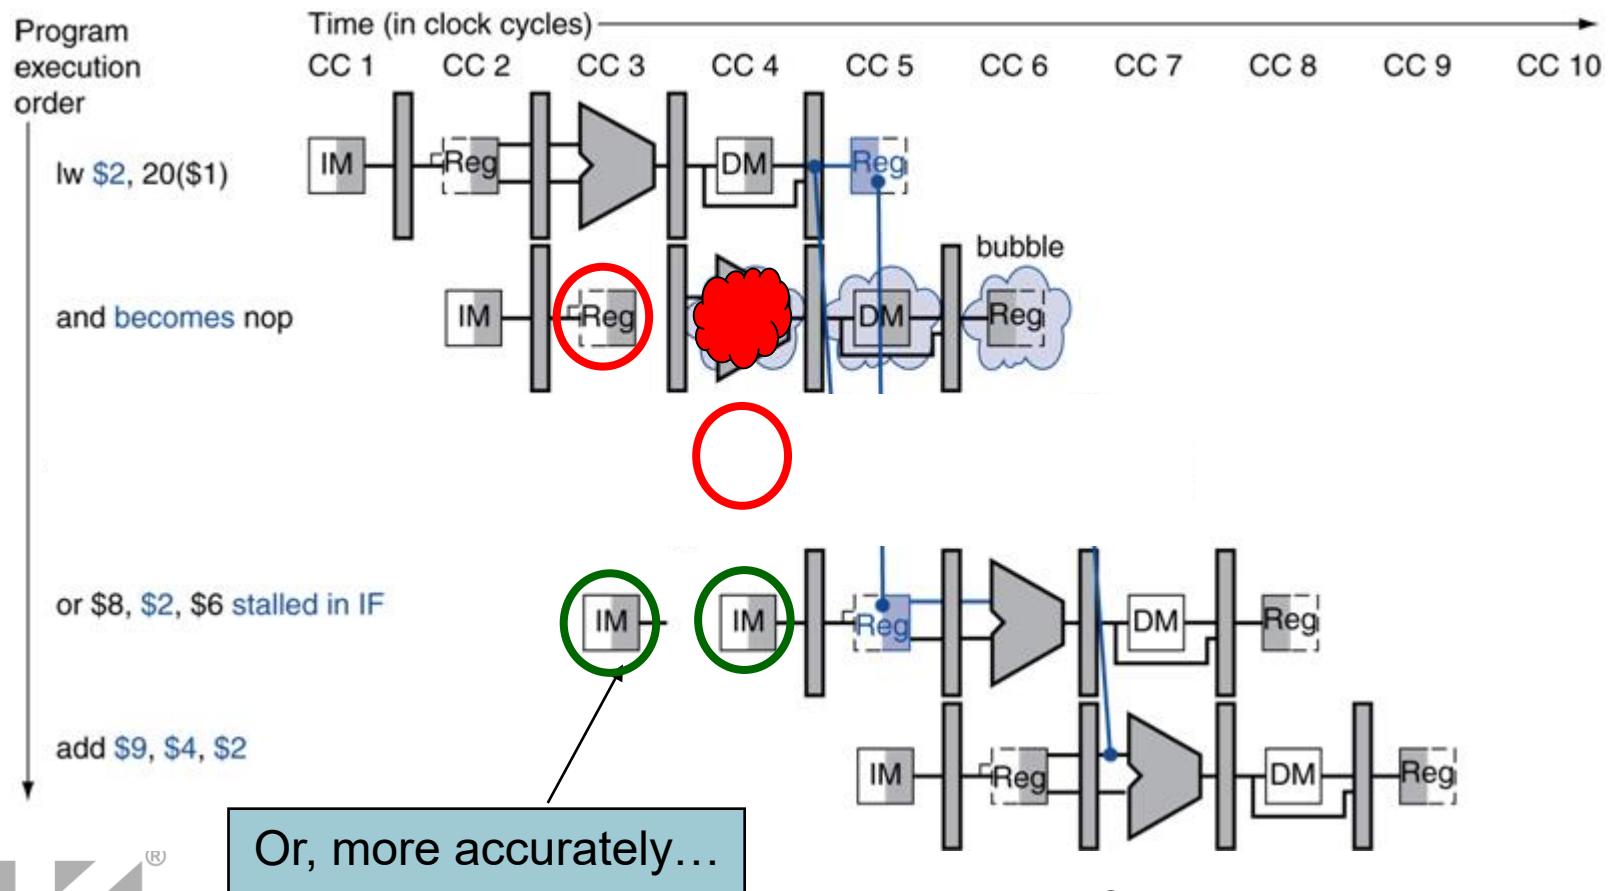
- Force control values in ID/EX register to 0
  - EX, MEM and WB do nop (no-operation)
- Prevent update of PC and IF/ID register
  - Using instruction is decoded again
  - Following instruction is fetched again
  - 1-cycle stall allows MEM to read data for 1w
    - Can subsequently forward to EX stage

# Stall/Bubble in the Pipeline



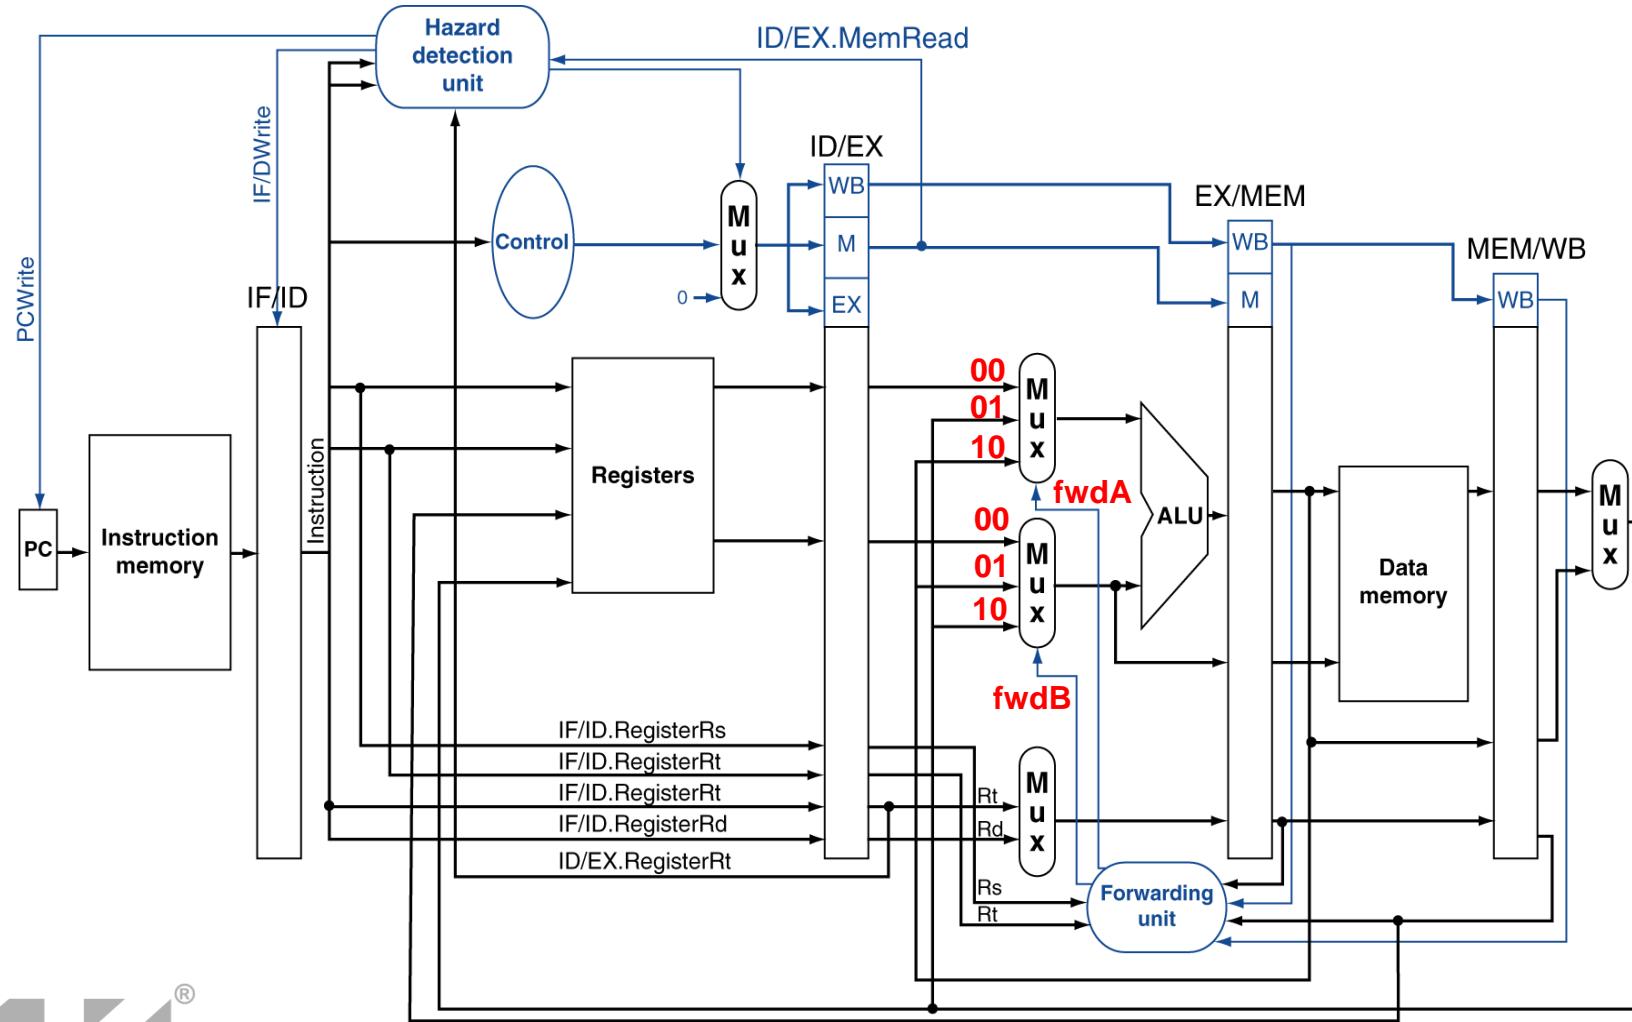
# Stall/Bubble in the Pipeline

- stall the pipeline: keep an instruction in the same stage
  - bubble – change the EX, MEM, WB controls fields of the ID/EX to 0
  - Keep the instr. in IF and ID for one more cycle (*and* and *or* below)



# Datapath with Hazard Detection

Hazard detected: 0 as mux input (and → bubble), PCWrite = 0, IF/Dwrite = 0



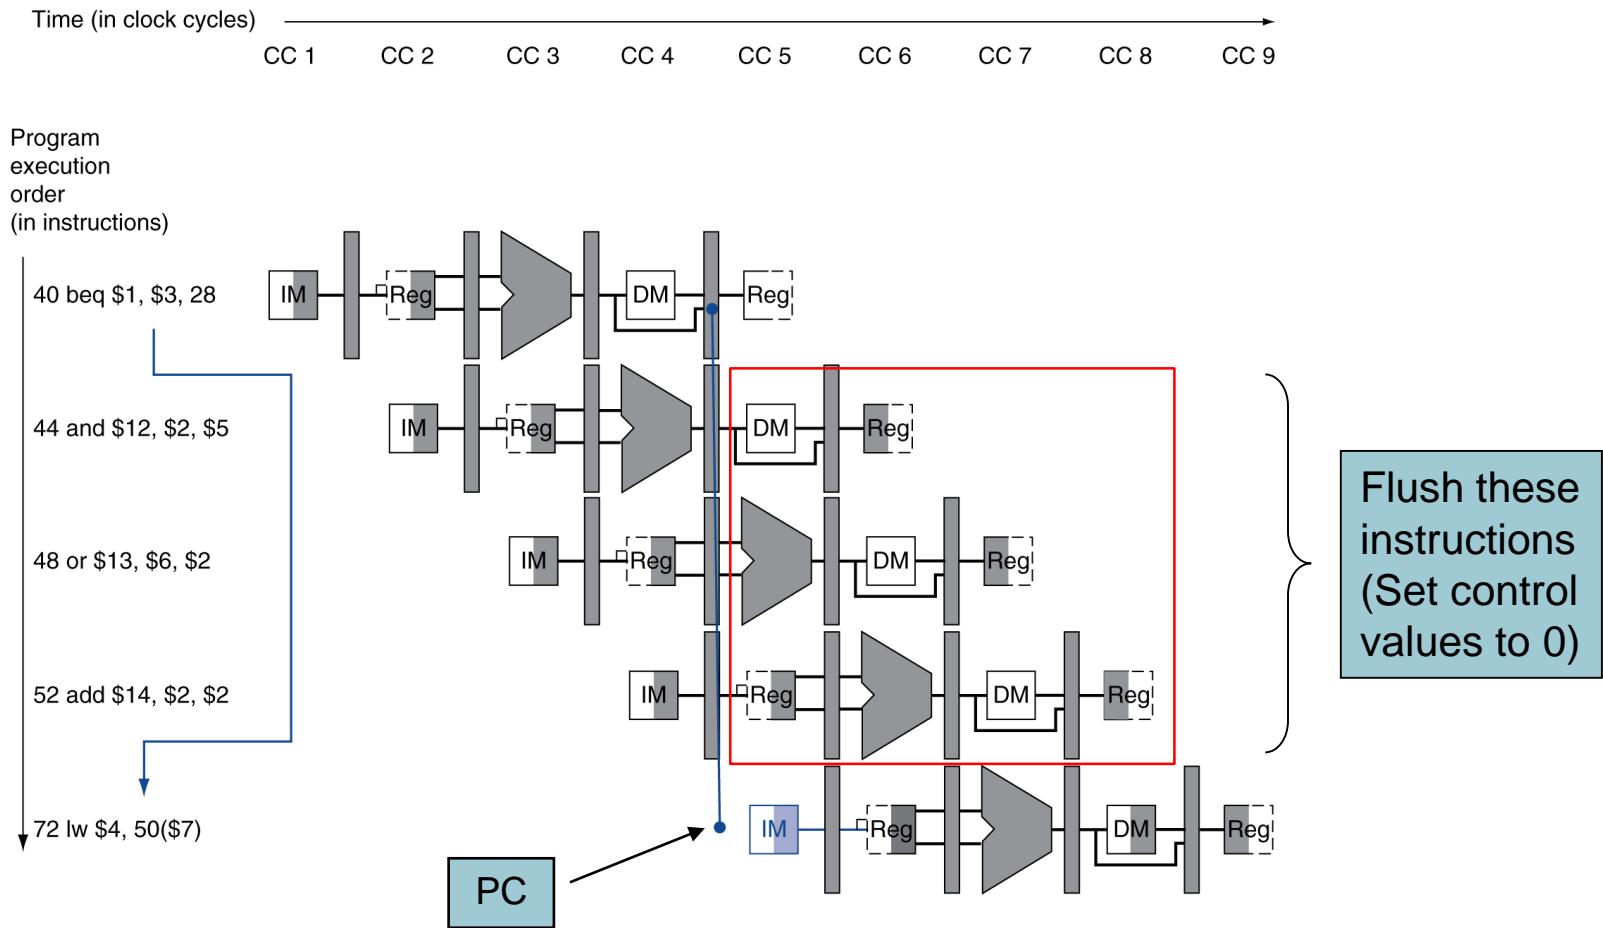
# Stalls and Performance

## The BIG Picture

- Stalls reduce performance
  - But are required to get correct results
- Compiler can arrange code to avoid hazards and stalls
  - Requires knowledge of the pipeline structure

# Branch Hazards

- If branch outcome determined in MEM



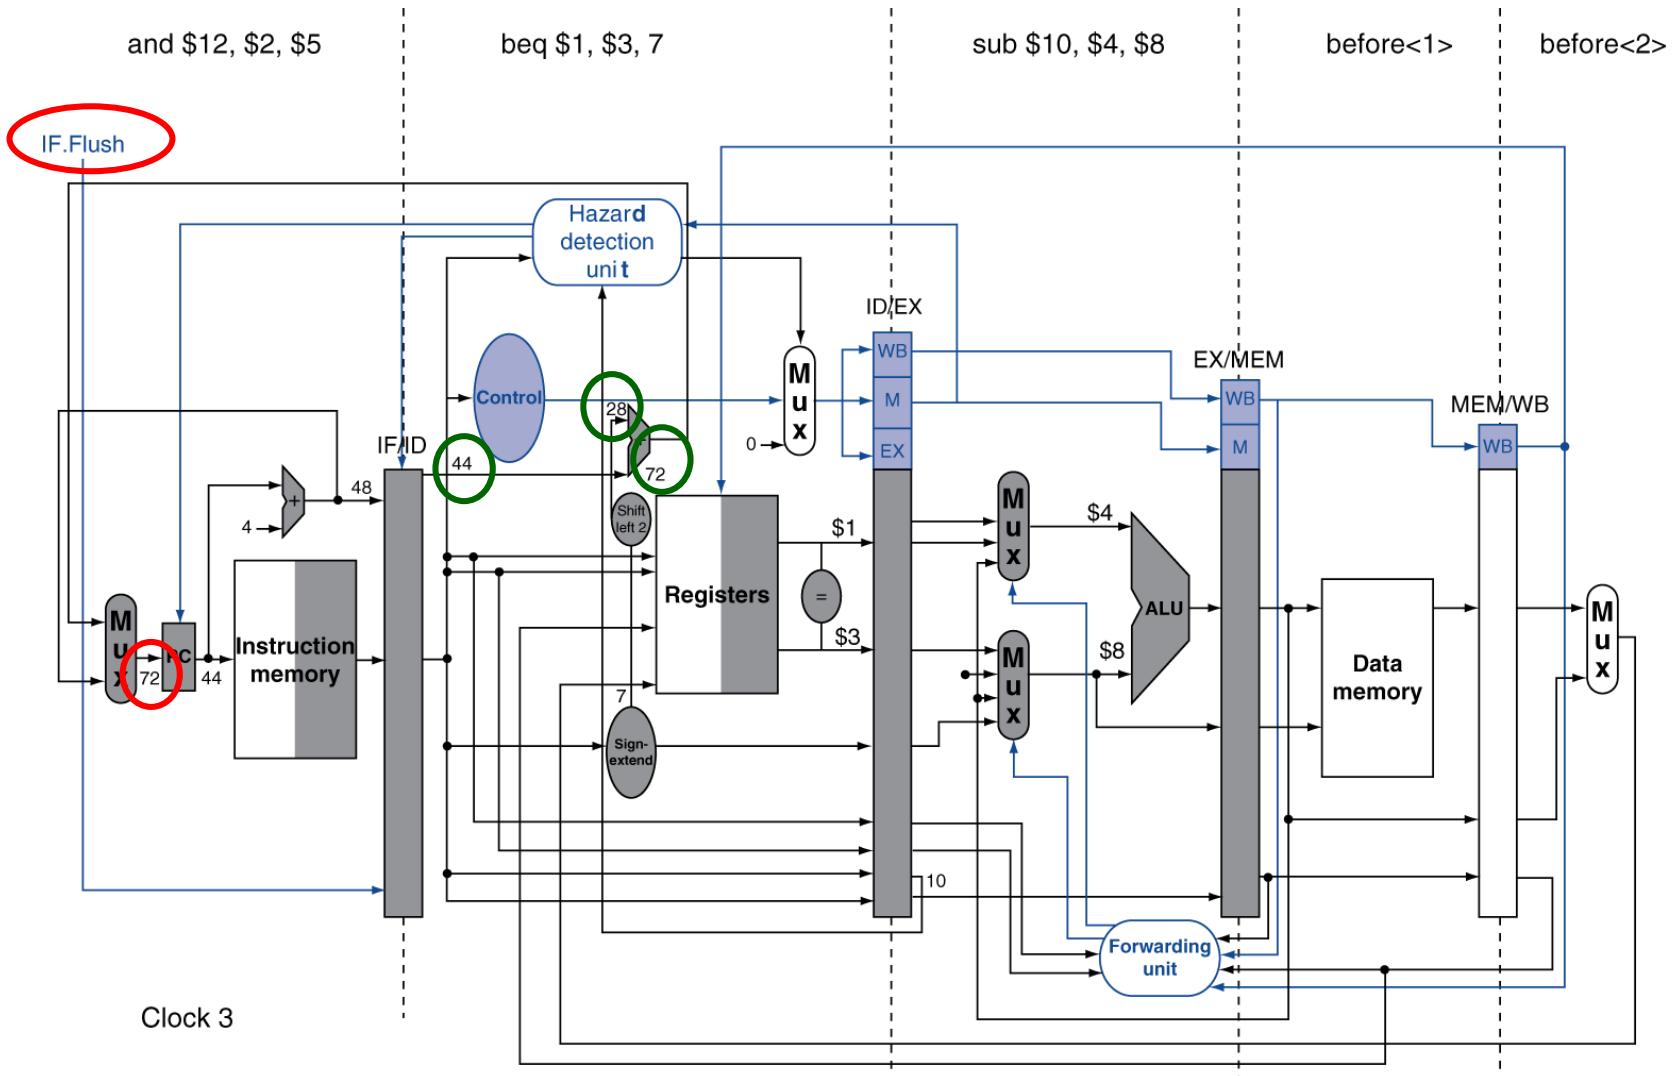
# Reducing Branch Delay

- Add hardware to determine outcome to **ID stage**
  - Target address adder
  - Register comparator (XOR two registers and OR the result)
- Example: branch taken

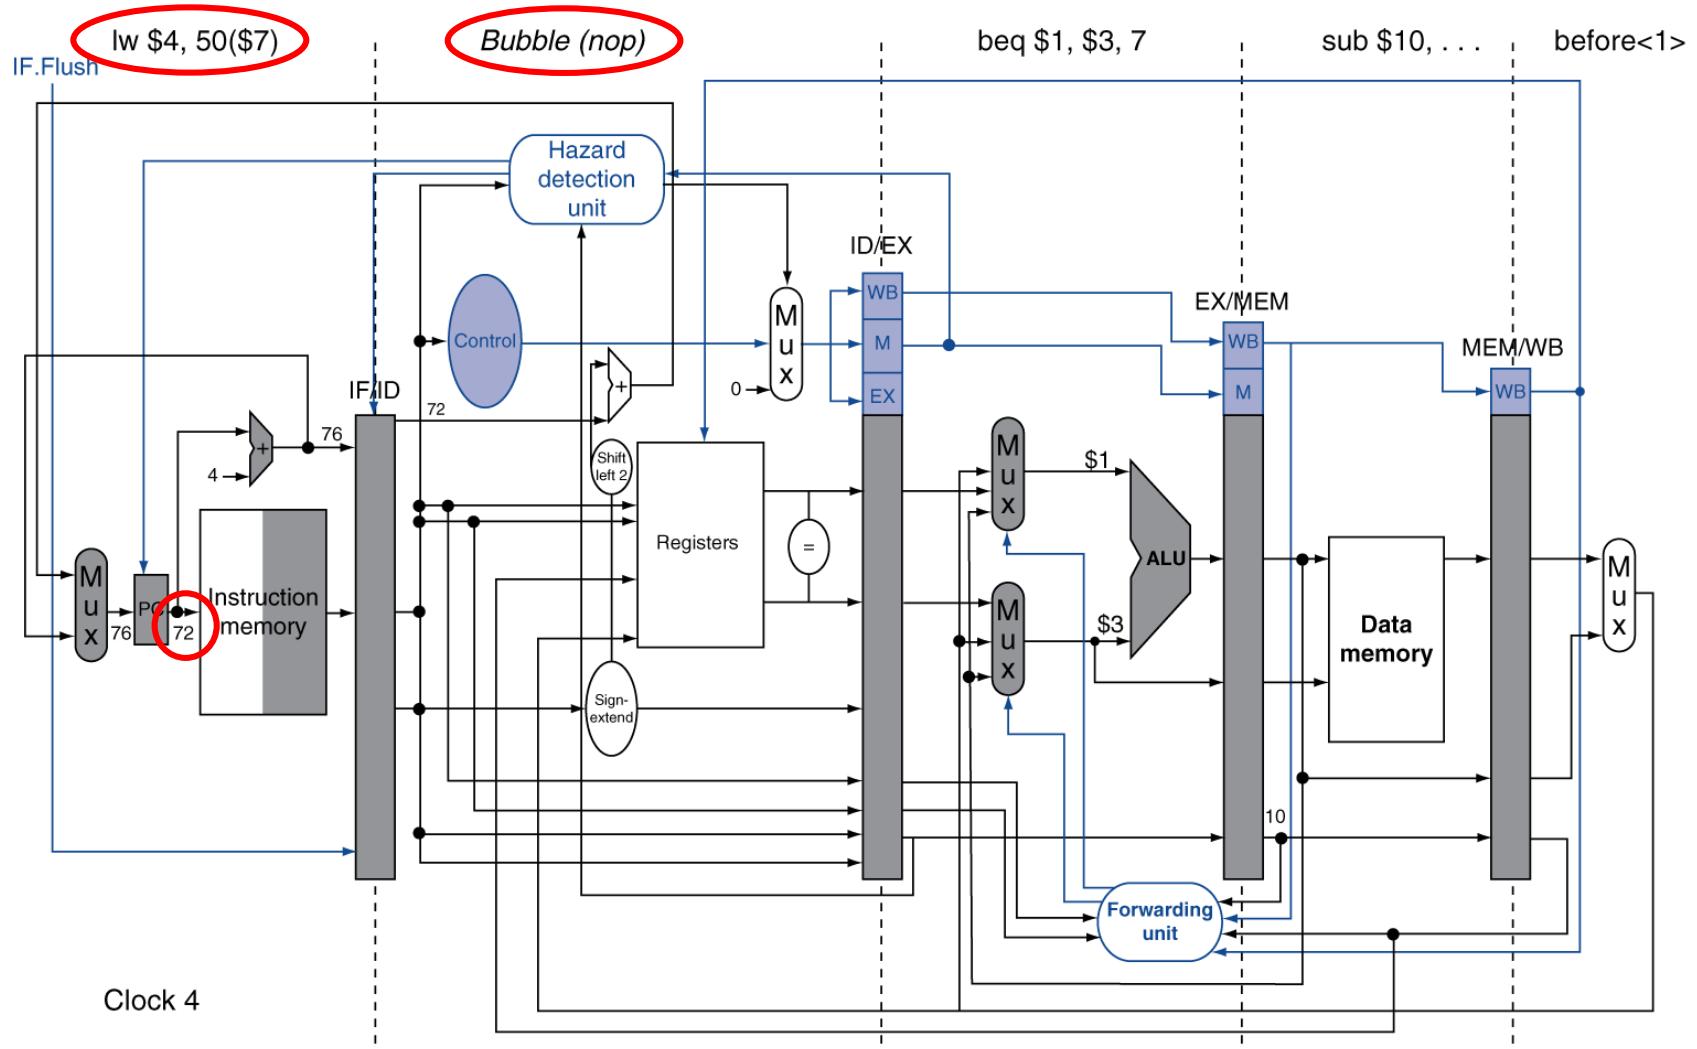
```
36: sub $10, $4, $8  
40: beq $1, $3, 7  
44: and $12, $2, $5  
48: or $13, $2, $6  
52: add $14, $4, $2  
56: slt $15, $6, $7  
...  
72: lw $4, 50($7)
```

For wrong branch decision:  
**flush** the instr. at IF stage  
(use **IF.Flush** to clear **IF/ID.instr.**)

# Example: Branch Taken

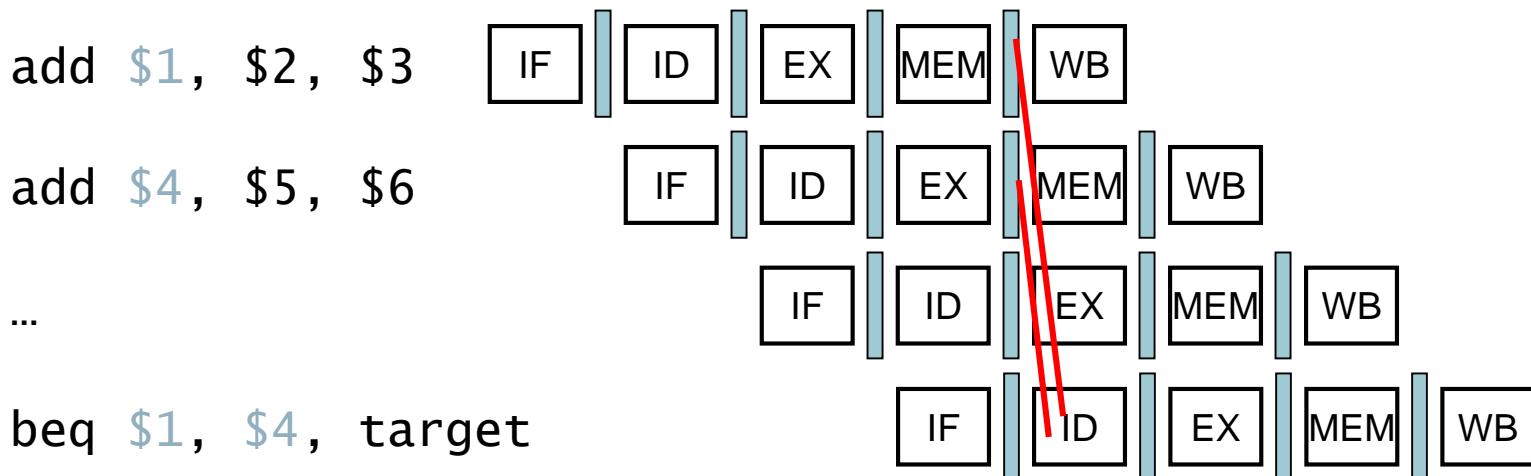


# Example: Branch Taken



# Data Hazards for Branches

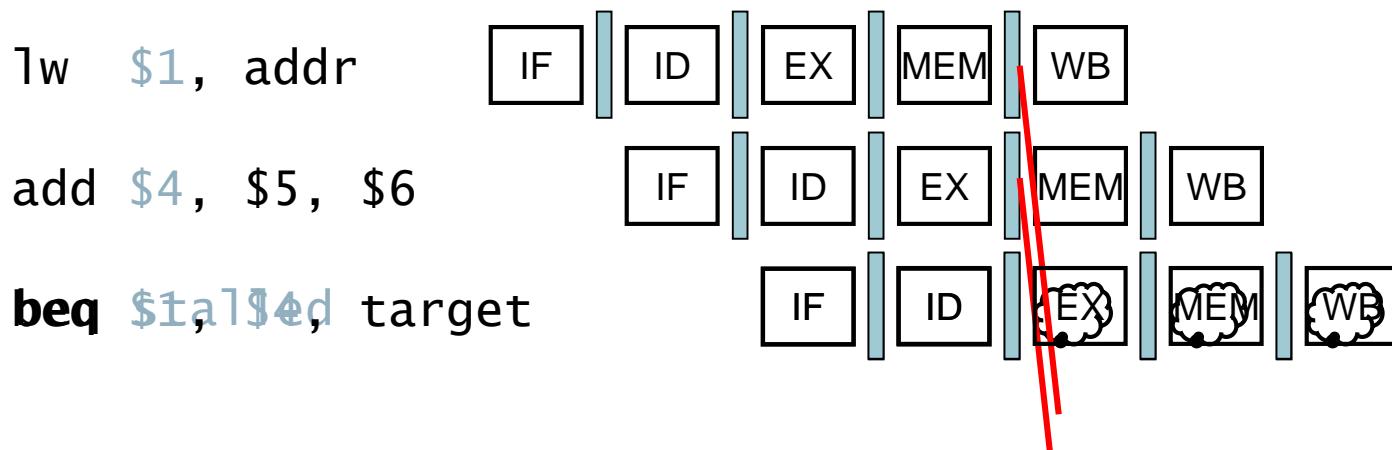
- Since we add hardware to determine branch outcome at **ID stage**, we need to compare two registers at ID stage
- If a comparison register is a destination of 2<sup>nd</sup> or 3<sup>rd</sup> preceding ALU instruction → resolve with forwarding



- Can resolve using forwarding

# Data Hazards for Branches

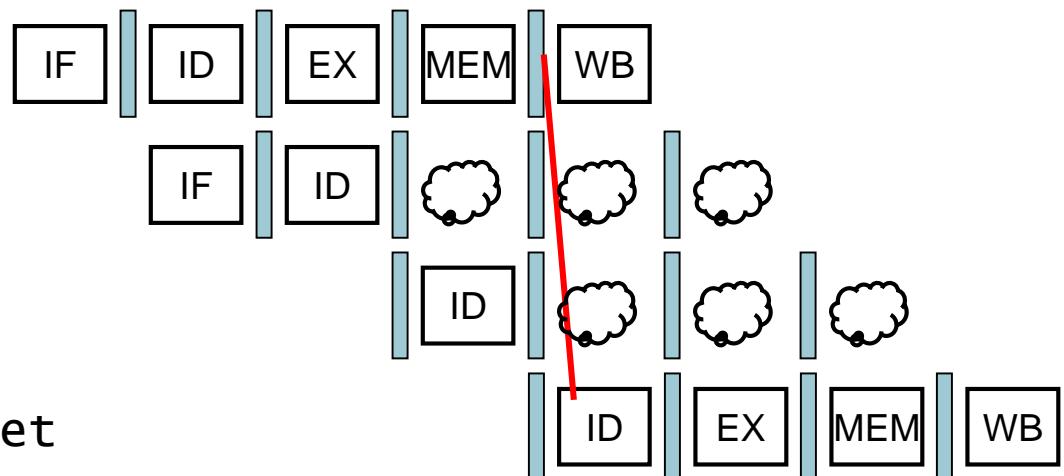
- If a comparison register is a destination of preceding ALU instruction or 2<sup>nd</sup> preceding load instruction
  - Need 1 stall cycle



# Data Hazards for Branches

- If a comparison register is a destination of immediately preceding load instruction
  - Need 2 stall cycles

lw \$1, addr



beq stalled

beq stalled

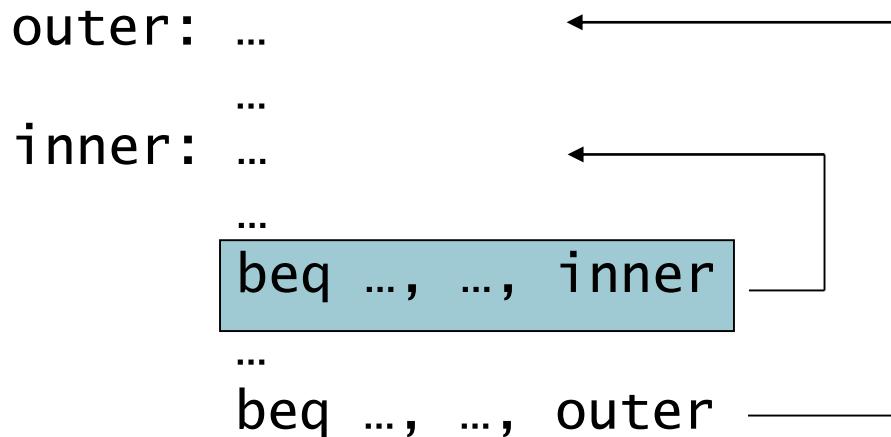
beq \$1, \$0, target

# Dynamic Branch Prediction

- In deeper and superscalar pipelines, branch penalty is more significant
- Use dynamic prediction
  - Branch prediction buffer (aka branch history table)
  - Indexed by address of branch instruction
  - Stores outcome (taken/not taken)
  - To execute a branch
    - Check table, expect the same outcome
    - Start fetching from fall-through or target
    - If wrong, flush pipeline and flip prediction

# 1-Bit Predictor: Shortcoming

- Inner loop branches mispredicted twice!

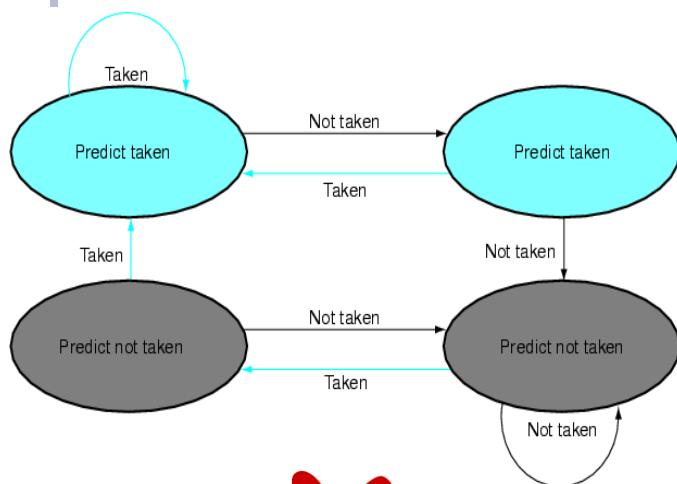


- Mispredict as taken on **last iteration** of inner loop
- Then mispredict as *not taken* on **first iteration** of inner loop next time around

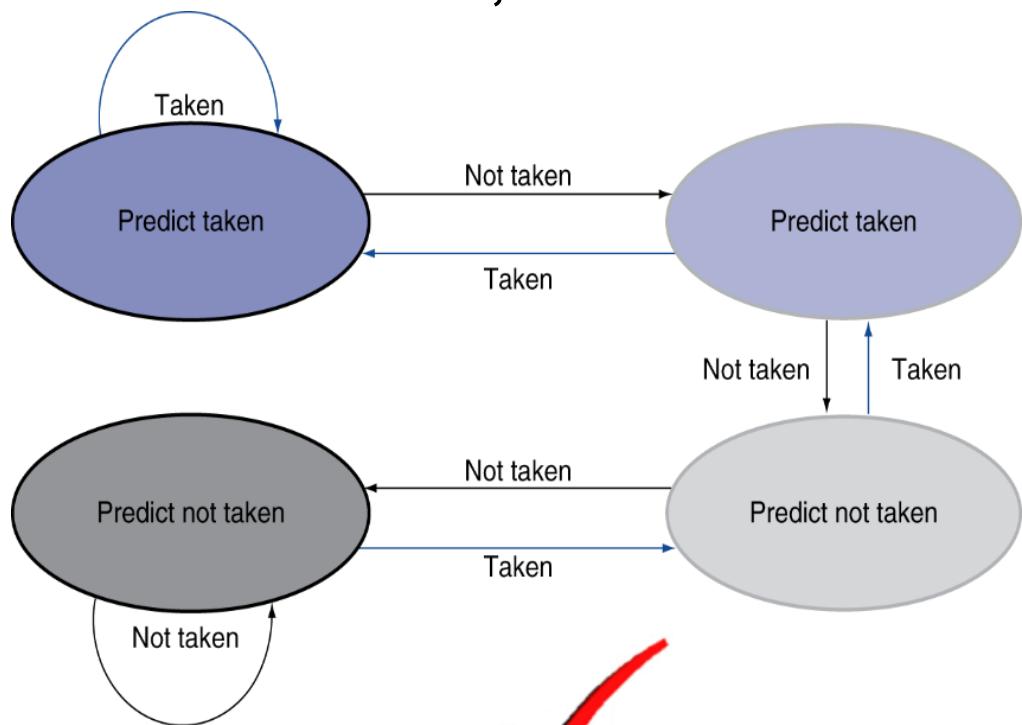
TTTTTN TTTTTN TTTTTN TTT  
TTTTT NTTTT NTTTT NTT

# 2-Bit Predictor

- Only change prediction on two successive mispredictions



The 4<sup>th</sup>, 5<sup>th</sup> Edition



# Branch predictor

- Branch: **T – N – T – N – N – T – N**
- Predictors are initialized to **taken**.
- 1-bit: **T – N – T – N – N – T – N**  
**T → T → N → T → N → N → T**
- 2-bit:  
**T – N – T – N – N – T – N**  
**T → T → T → T → T → N → T**
- 2-bit:  
**T – N – T – N – N – N – T – N**  
**T → T → T → T → T → N → N → N**

# Calculating the Branch Target

- Even with predictor, still need to calculate the target address
  - 1-cycle penalty for a taken branch
- **Branch target buffer**
  - Cache of target addresses
  - Indexed by PC when instruction fetched
    - If hit and instruction is branch predicted taken, can fetch target immediately

# Exceptions and Interrupts

- “Unexpected” events requiring change in flow of control
  - Different ISAs use the terms differently
- **Exception**
  - Arises within the CPU
    - e.g., undefined opcode, overflow, syscall, ...
- **Interrupt**
  - From an external I/O controller
- Dealing with them without sacrificing performance is hard

# Handling Exceptions

- In MIPS, managed by a System Control Coprocessor (CP0)
- Save PC of interrupted instruction
  - In MIPS: Exception Program Counter (EPC)
- Save indication of the problem
  - In MIPS: Cause register
  - We'll assume 1-bit: 0 for undefined opcode, 1 for overflow
- Jump to handler at 8000 0180
  - In MIPS, a single entry point for all exceptions
  - OS decodes the status register to find the cause
  - Alternative: vectored Interrupts
    - Handler address determined by the cause
    - E.g.: Undefined opcode: C000 0000
    - Overflow: C000 0020
    - .... C000 0040

# Handler Actions

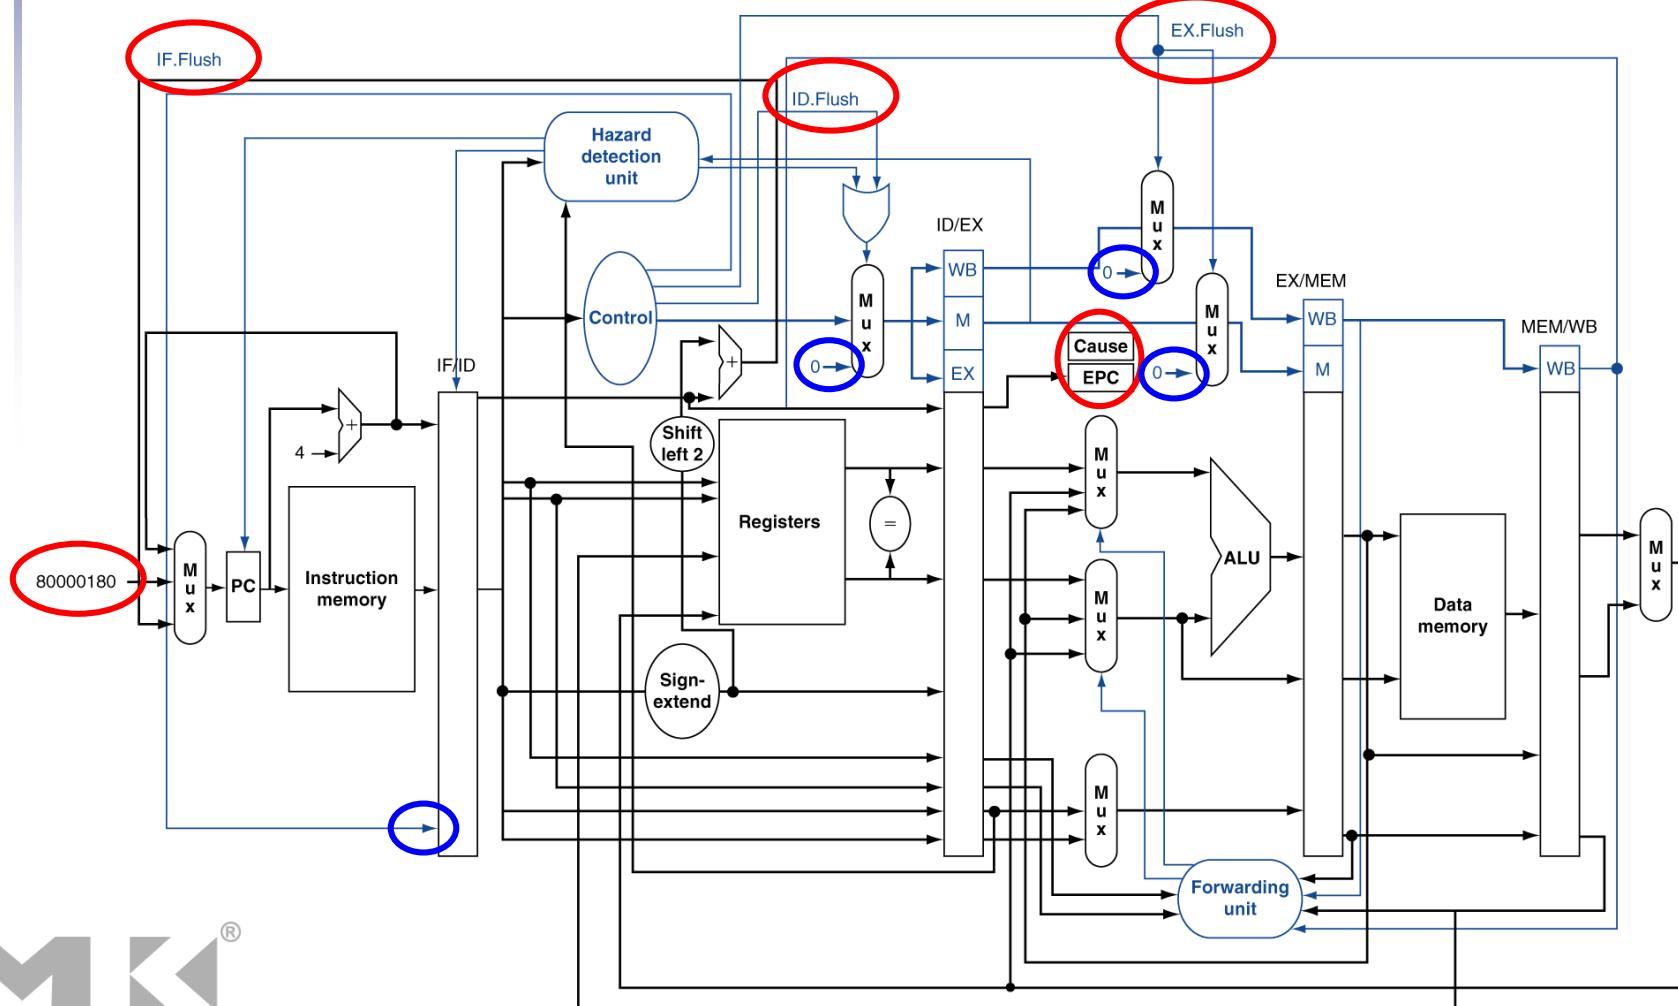
- Read cause, and transfer to relevant handler
- Determine action required
- If restartable
  - Take corrective action
  - use EPC to return to program
- Otherwise
  - Terminate program
  - Report error using EPC, cause, ...

# Exceptions in a Pipeline

- Another form of control hazard
- Consider overflow on add in EX stage
  - add \$1, \$2, \$1
    - Prevent \$1 from being clobbered
    - Complete previous instructions
    - Flush add and subsequent instructions
    - Set Cause and EPC register values
    - Transfer control to handler
- Similar to mispredicted branch
  - Use much of the same hardware

# Pipeline with Exceptions

- Flush instructions at IF, ID, EX when an exception is detected at EX.
- Supply 8000 0180 to PC. Add **Cause**, **EPC**, **IF.Flush**, **ID.Flush**, **EX.Flush**



# Exception Properties

- Restartable exceptions
  - Pipeline can flush the instruction
  - Handler executes, then returns to the instruction
    - Refetched and executed from scratch
- PC saved in EPC register
  - Identifies causing instruction
  - Actually  $PC + 4$  is saved
    - Handler must adjust

# Exception Example

- Exception on add in

```
40      sub    $11, $2, $4  
44      and    $12, $2, $5  
48      or     $13, $2, $6  
4C      add    $1,   $2,   $1  
50      slt    $15, $6, $7  
54      lw     $16, 50($7)
```

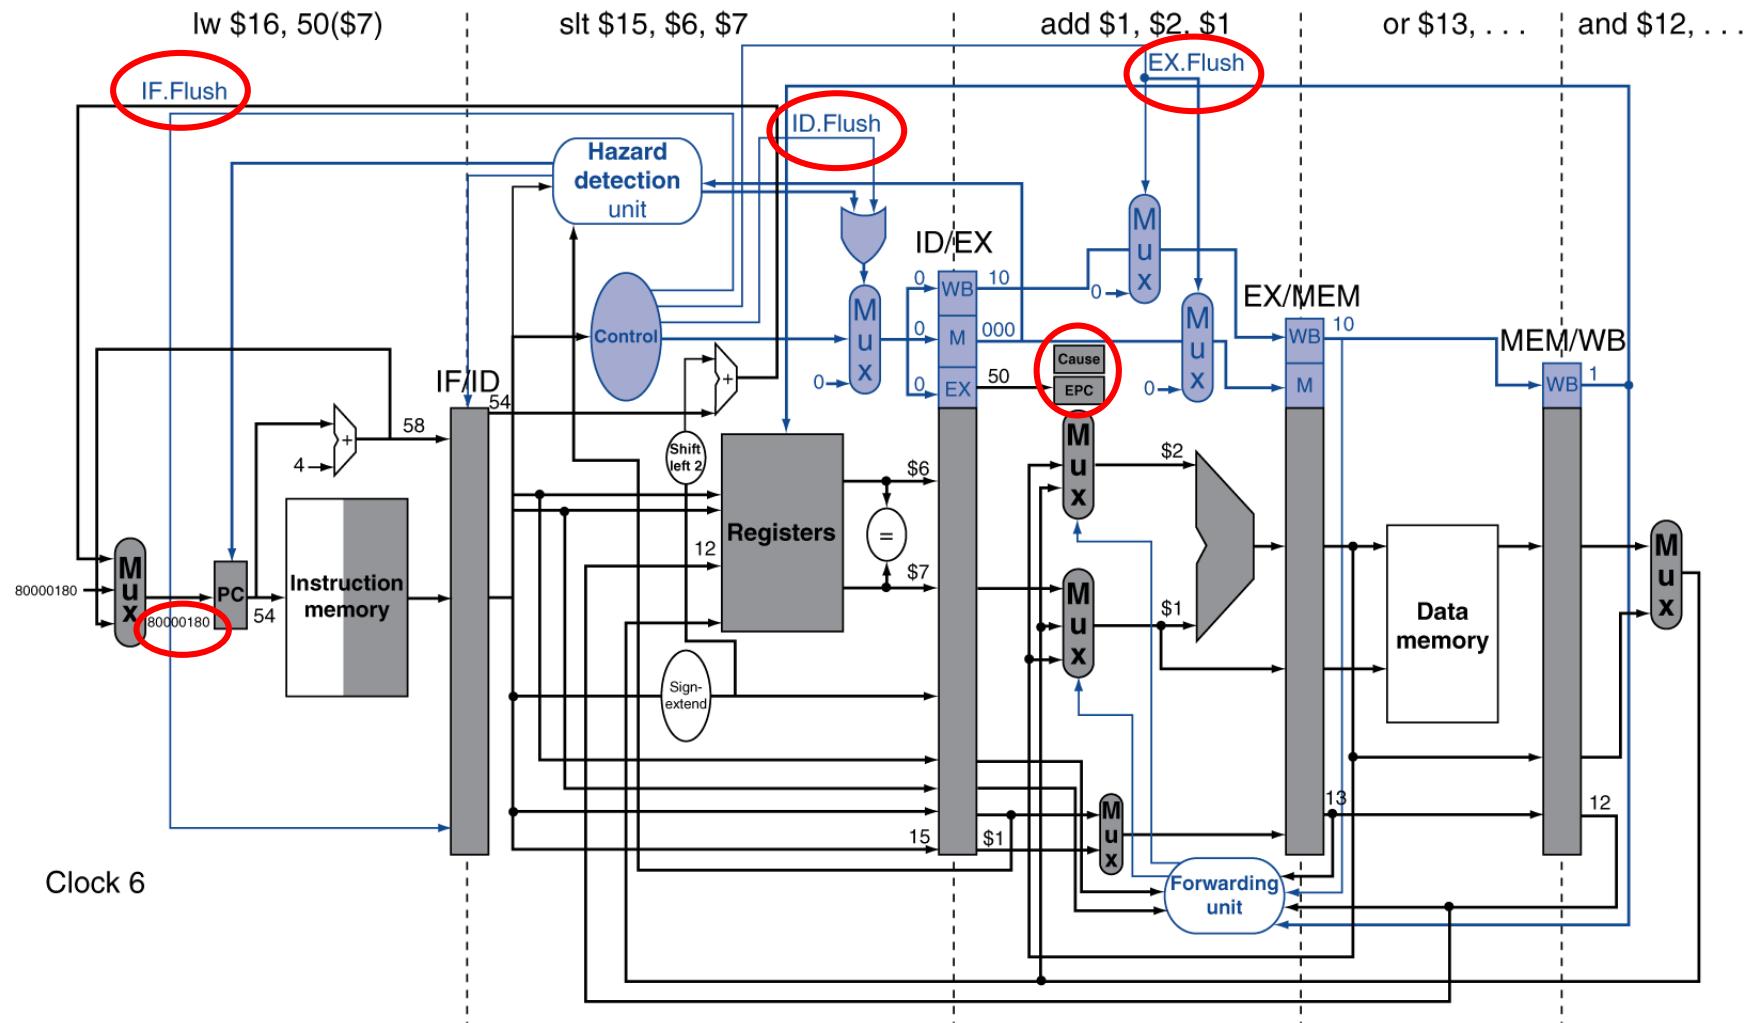
...

- Handler

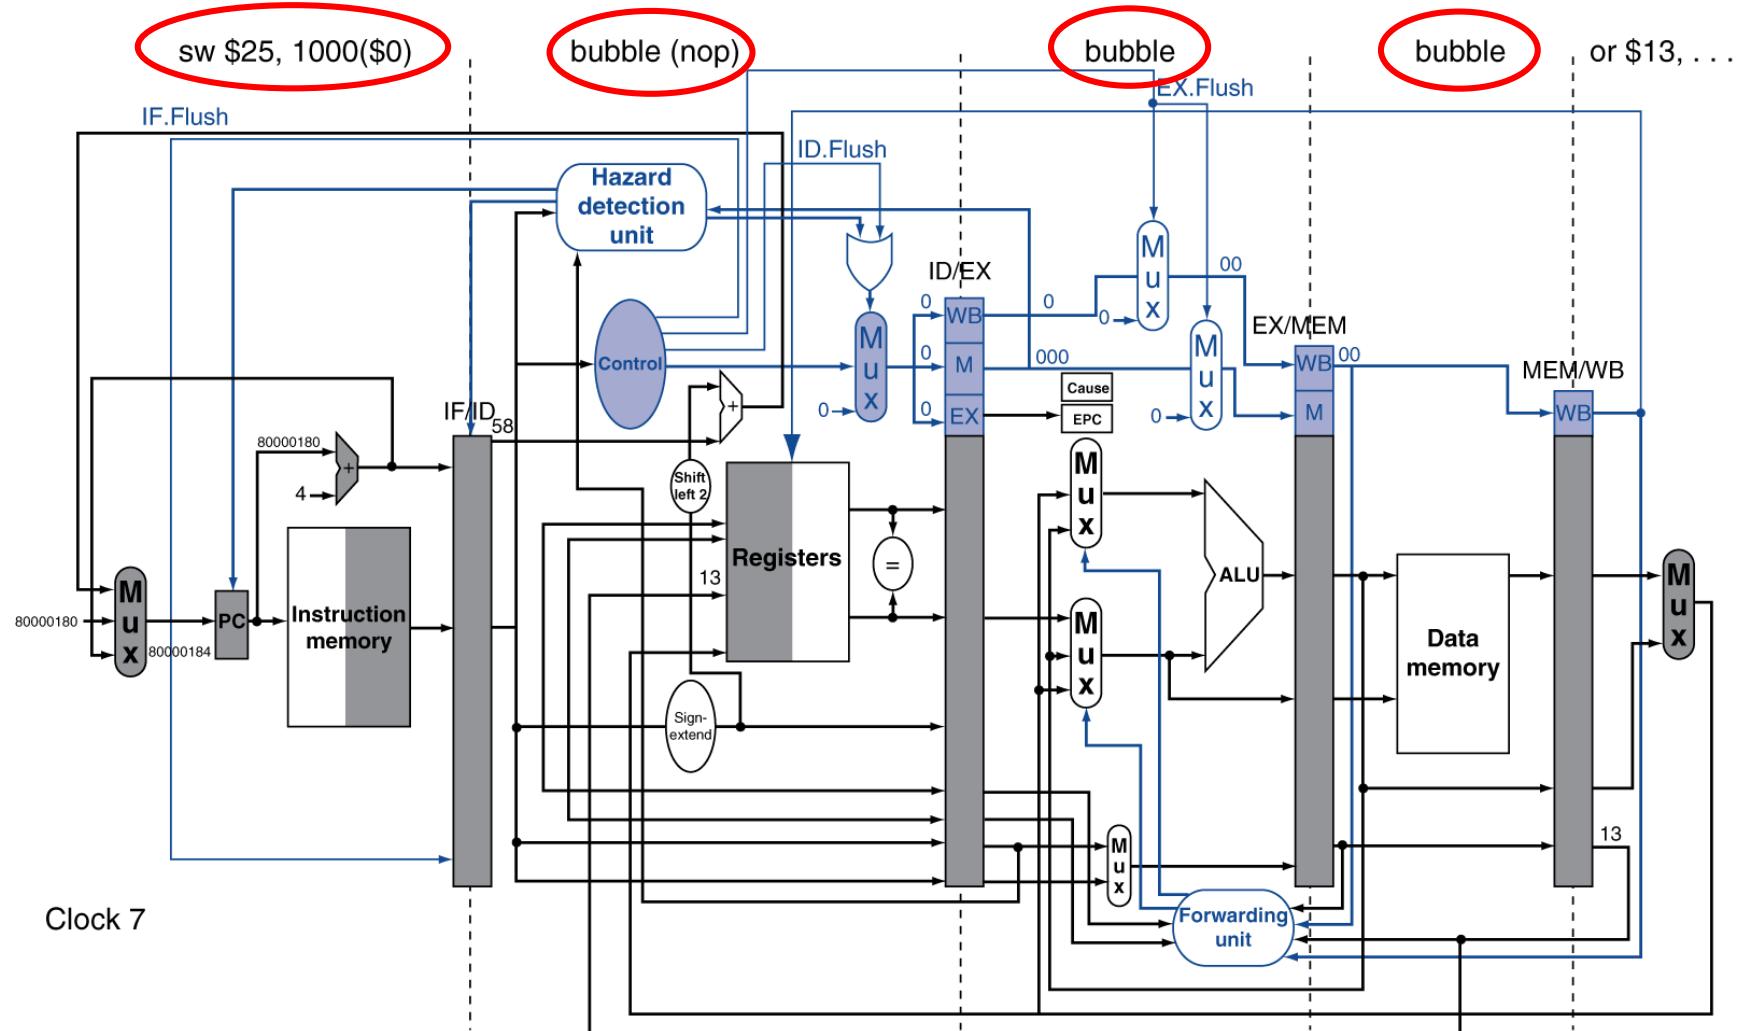
```
80000180    sw    $25, 1000($0)  
80000184    sw    $26, 1004($0)
```

...

# Exception Example



# Exception Example



# Multiple Exceptions

- Pipelining overlaps multiple instructions
  - Could have multiple exceptions at once
- Simple approach: deal with exception from earliest instruction
  - Flush subsequent instructions
  - “Precise” exceptions
- In complex pipelines
  - Multiple instructions issued per cycle
  - Out-of-order completion
  - Maintaining precise exceptions is difficult!

# Imprecise Exceptions

- Just stop pipeline and save state
  - Including exception cause(s)
- Let the handler work out
  - Which instruction(s) had exceptions
  - Which to complete or flush
    - May require “manual” completion
- Simplifies hardware, but more complex handler software
- Not feasible for complex multiple-issue out-of-order pipelines

# Instruction-Level Parallelism (ILP)

- Pipelining: executing multiple instructions in parallel
- To increase ILP
  - Deeper pipeline
    - Less work per stage  $\Rightarrow$  shorter clock cycle
  - Multiple issue
    - Replicate pipeline stages  $\Rightarrow$  multiple pipelines
    - Start multiple instructions per clock cycle
    - CPI < 1, so use Instructions Per Cycle (IPC)
    - E.g., 4GHz 4-way multiple-issue
      - 16 BIPS, peak CPI = 0.25, peak IPC = 4
    - But dependencies reduce this in practice

# Multiple Issue

- Static multiple issue
  - Compiler groups instructions to be issued together
  - Packages them into “issue slots”
  - Compiler detects and avoids hazards
- Dynamic multiple issue
  - CPU examines instruction stream and chooses instructions to issue each cycle
  - Compiler can help by reordering instructions
  - CPU resolves hazards using advanced techniques at runtime

# Speculation

- “Guess” what to do with an instruction
  - Start operation as soon as possible
  - Check whether guess was right
    - If so, complete the operation
    - If not, roll-back and do the right thing
- Common to static and dynamic multiple issue
- Examples
  - Speculate on branch outcome
    - Roll back if path taken is different
  - Speculate on load
    - Roll back if location is updated

# Static Multiple Issue

- Compiler groups instructions into “issue packets”
  - Group of instructions that can be issued on a single cycle
  - Determined by pipeline resources required
- Think of an issue packet as a very long instruction
  - Specifies multiple concurrent operations
  - ⇒ Very Long Instruction Word (VLIW)

# Scheduling Static Multiple Issue

- Compiler must remove some/all hazards
  - Reorder instructions into issue packets
  - No dependencies with a packet
  - Possibly some dependencies between packets
    - Varies between ISAs; compiler must know!
  - Pad with nop if necessary

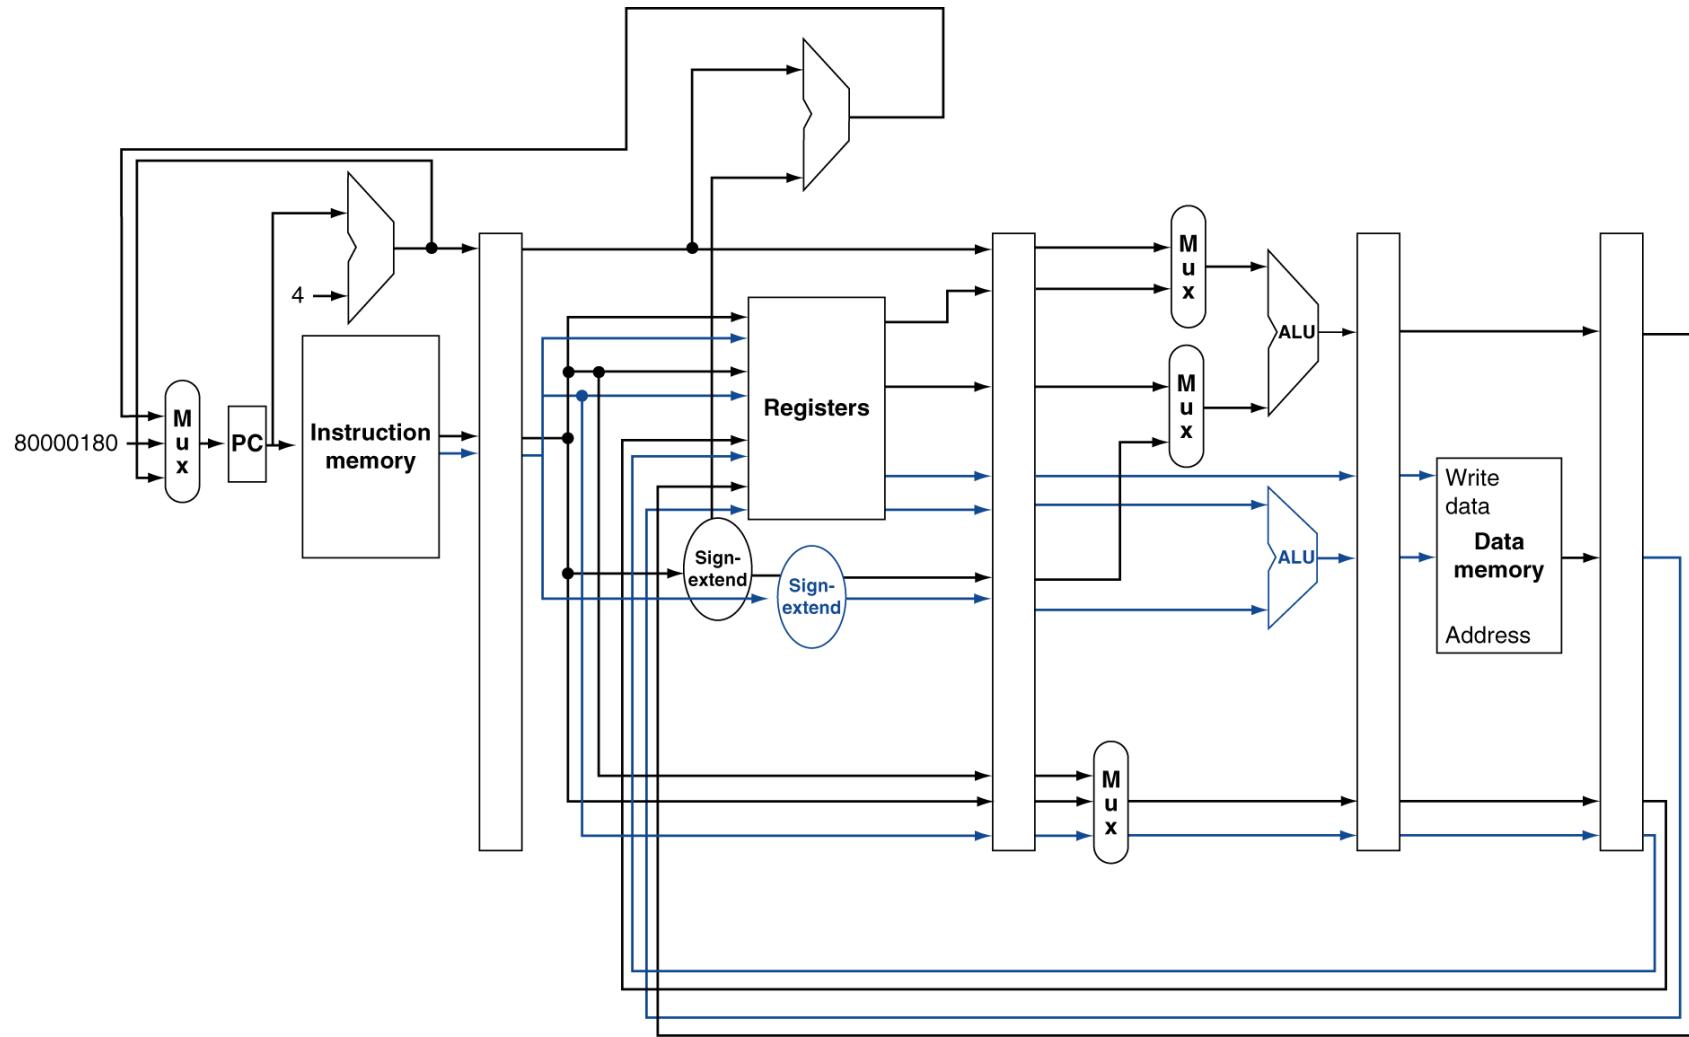
# MIPS with Static Dual Issue

- Two-issue packets

- One ALU/branch instruction
- One load/store instruction
- 64-bit aligned
  - ALU/branch, then load/store
  - Pad an unused instruction with nop

| Address | Instruction type | Pipeline Stages |    |    |     |     |     |    |
|---------|------------------|-----------------|----|----|-----|-----|-----|----|
| n       | ALU/branch       | IF              | ID | EX | MEM | WB  |     |    |
| n + 4   | Load/store       | IF              | ID | EX | MEM | WB  |     |    |
| n + 8   | ALU/branch       |                 | IF | ID | EX  | MEM | WB  |    |
| n + 12  | Load/store       |                 | IF | ID | EX  | MEM | WB  |    |
| n + 16  | ALU/branch       |                 |    | IF | ID  | EX  | MEM | WB |
| n + 20  | Load/store       |                 |    | IF | ID  | EX  | MEM | WB |

# MIPS with Static Dual Issue



# Hazards in the Dual-Issue MIPS

- More instructions executing in parallel
- EX data hazard
  - Forwarding avoided stalls with single-issue
  - Now can't use ALU result in load/store in same packet
    - add \$t0, \$s0, \$s1
    - load \$s2, 0(\$t0)
    - Split into two packets, effectively a stall
- Load-use hazard
  - Still one cycle use latency, but now two instructions
- More aggressive scheduling required

# Scheduling Example

- Schedule this for dual-issue MIPS

```
Loop: lw    $t0, 0($s1)      # $t0=array element
      addu $t0, $t0, $s2      # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi $s1, $s1,-4        # decrement pointer
      bne  $s1, $zero, Loop  # branch $s1!=0
```

|       | ALU/branch              | Load/store          | cycle |
|-------|-------------------------|---------------------|-------|
| Loop: | nop                     | lw    \$t0, 0(\$s1) | 1     |
|       | addi \$s1, \$s1,-4      | nop                 | 2     |
|       | addu \$t0, \$t0, \$s2   | nop                 | 3     |
|       | bne  \$s1, \$zero, Loop | sw    \$t0, 4(\$s1) | 4     |

- IPC = 5/4 = 1.25 (c.f. peak IPC = 2)

# Loop Unrolling

- Replicate loop body to expose more parallelism
  - Reduces loop-control overhead
- Use different registers per replication
  - Called “register renaming”
  - Avoid loop-carried “**anti-dependencies**”
    - Store followed by a load of the same register
    - Aka “name dependence”
      - Reuse of a register name

```

Loop: lw    $t0, 0($s1)
      addu $t0, $t0, $s2
      sw   $t0, 0($s1)
      addi $s1, $s1,-4
      bne  $s1, $zero, Loop

```

## Antidependence (name dependence)

- The ordering forced by the reuse of a name (register), **not** true data dependence.

## Register Renaming:

- to remove anti-dependence.
- e.g., \$t0 → \$t0, \$t1, \$t2, \$t3

```

Loop: lw    $t0, 0($s1)
      addu $t0, $t0, $s2
      sw   $t0, 0($s1)
      $t1
      lw    $t0, 0($s1)
      addu $t0, $t0, $s2
      sw   $t0, 0($s1)

```

```

      $t2
      lw    $t0, 0($s1)
      addu $t0, $t0, $s2
      sw   $t0, 0($s1)

```

```

      $t3
      lw    $t0, 0($s1)
      addu $t0, $t0, $s2
      sw   $t0, 0($s1)

```

```

      addi $s1, $s1,-4
      bne  $s1, $zero, Loop
      -16

```

|      |       |            |
|------|-------|------------|
| lw   | \$t0, | 0(\$s1)    |
| addu | \$t0, | \$t0, \$s2 |
| sw   | \$t0, | 0(\$s1)    |

|      |       |            |
|------|-------|------------|
| lw   | \$t1, | -4(\$s1)   |
| addu | \$t1, | \$t0, \$s2 |
| sw   | \$t1, | -4(\$s1)   |

|      |       |            |
|------|-------|------------|
| lw   | \$t2, | -8(\$s1)   |
| addu | \$t2, | \$t0, \$s2 |
| sw   | \$t2, | -8(\$s1)   |

|      |       |            |
|------|-------|------------|
| lw   | \$t3, | -12(\$s1)  |
| addu | \$t3, | \$t0, \$s2 |
| sw   | \$t3, | -12(\$s1)  |

addi \$s1, \$s1, -16  
 bne \$s1, \$zero, Loop

|       | ALU/branch            | Load/store       | cycle |
|-------|-----------------------|------------------|-------|
| Loop: | nop                   | lw \$t0, 0(\$s1) | 1     |
|       |                       | nop              | 2     |
|       | addu \$t0, \$t0, \$s2 | nop              | 3     |
|       |                       | sw \$t0, 0(\$s1) | 4     |
|       |                       |                  | 5     |



# Loop Unrolling Example

|       | ALU/branch             | Load/store        | cycle |
|-------|------------------------|-------------------|-------|
| Loop: | addi \$s1, \$s1, -16   | lw \$t0, 0(\$s1)  | 1     |
|       | nop                    | lw \$t1, 12(\$s1) | 2     |
|       | addu \$t0, \$t0, \$s2  | lw \$t2, 8(\$s1)  | 3     |
|       | addu \$t1, \$t1, \$s2  | lw \$t3, 4(\$s1)  | 4     |
|       | addu \$t2, \$t2, \$s2  | sw \$t0, 16(\$s1) | 5     |
|       | addu \$t3, \$t4, \$s2  | sw \$t1, 12(\$s1) | 6     |
|       | nop                    | sw \$t2, 8(\$s1)  | 7     |
|       | bne \$s1, \$zero, Loop | sw \$t3, 4(\$s1)  | 8     |

- IPC = 14/8 = 1.75
  - Closer to 2, but at cost of registers and code size

# Dynamic Multiple Issue

- “Superscalar” processors
- CPU decides whether to issue 0, 1, 2, ... instructions each cycle
  - Avoiding structural and data hazards
- Avoids the need for compiler scheduling
  - Though it may still help
  - Code semantics ensured by the CPU

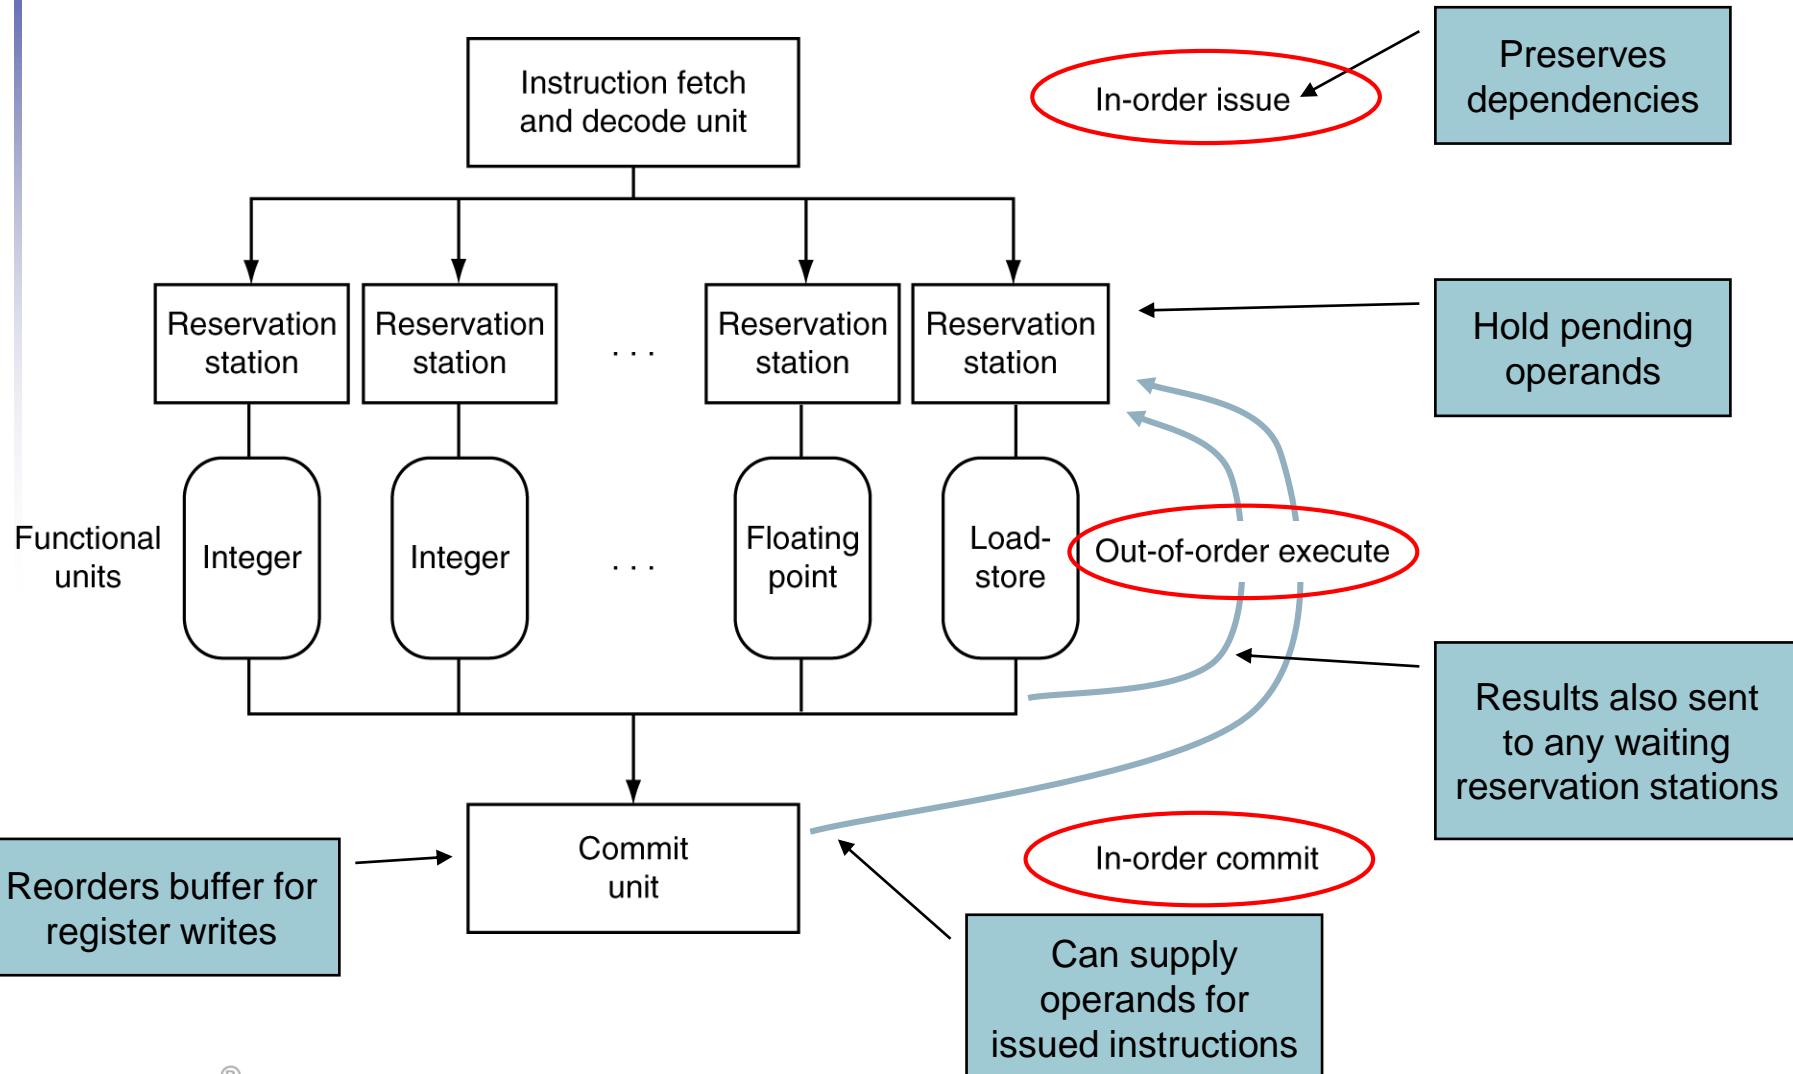
# Dynamic Pipeline Scheduling

- It is hardware support
- Allow the CPU to execute instructions out of order to avoid stalls (out-of-order execution)
  - But commit result to registers in order (in-order commit)

```
lw      $t0, 20($s2)
addu   $t1, $t0, $t2
sub    $s4, $s4, $t3
sltii $t5, $s4, 20
```

- Can start sub while addu is waiting for lw

# Dynamically Scheduled CPU



# Register Renaming

- Reservation stations and reorder buffer effectively provide register renaming
- On instruction issue to reservation station
  - If operand is available in register file or reorder buffer
    - Copied to reservation station
    - No longer required in the register; can be overwritten
  - If operand is not yet available
    - It will be provided to the reservation station by another function unit directly.
    - Register update may not be required

# Why Do Dynamic Scheduling?

- Why not just let the compiler schedule code? (Note: dynamic scheduling is hardware support)
- Not all stalls are predictable
  - e.g., cache misses
- Can't always schedule around branches
  - Branch outcome is dynamically determined
- Different implementations of an ISA have different latencies and hazards

# Does Multiple Issue Work?

## The BIG Picture

- Yes, but not as much as we'd like
- Programs have real dependencies that limit ILP
- Some dependencies are hard to eliminate
  - e.g., pointer aliasing
- Some parallelism is hard to expose
  - Limited window size during instruction issue
- Memory delays and limited bandwidth
  - Hard to keep pipelines full
- Speculation can help if done well



# Power Efficiency

- Complexity of dynamic scheduling and speculations requires power
- Multiple simpler cores may be better

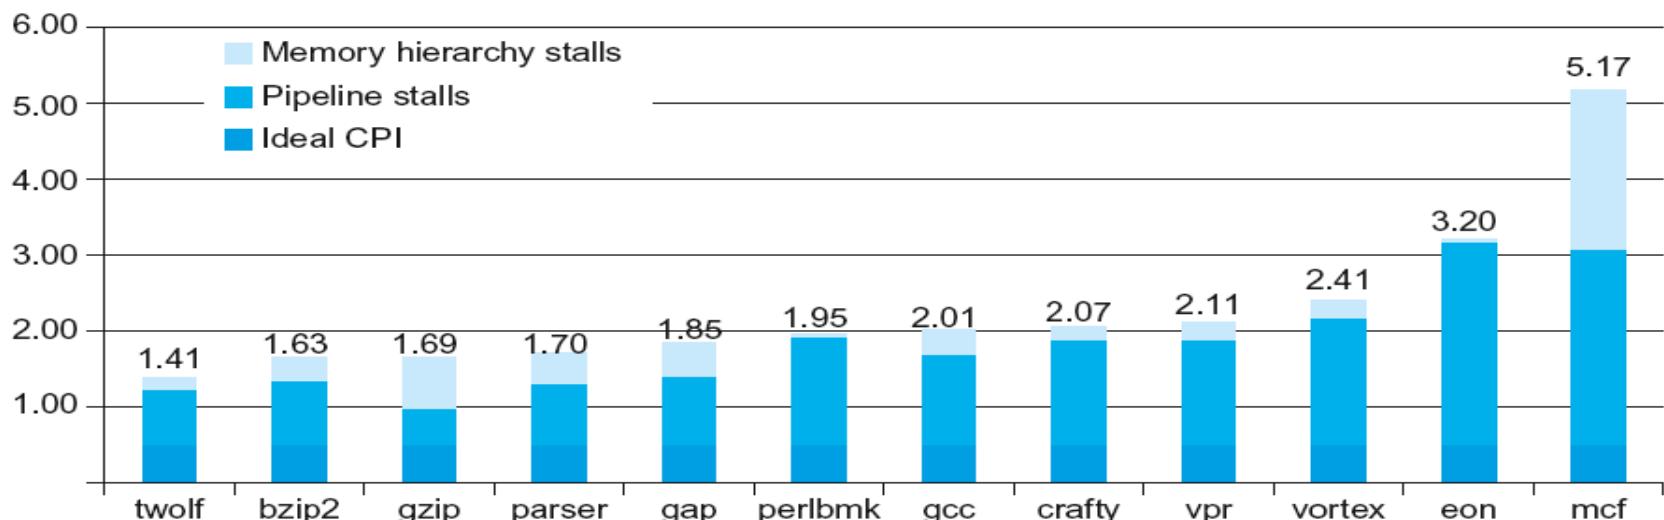
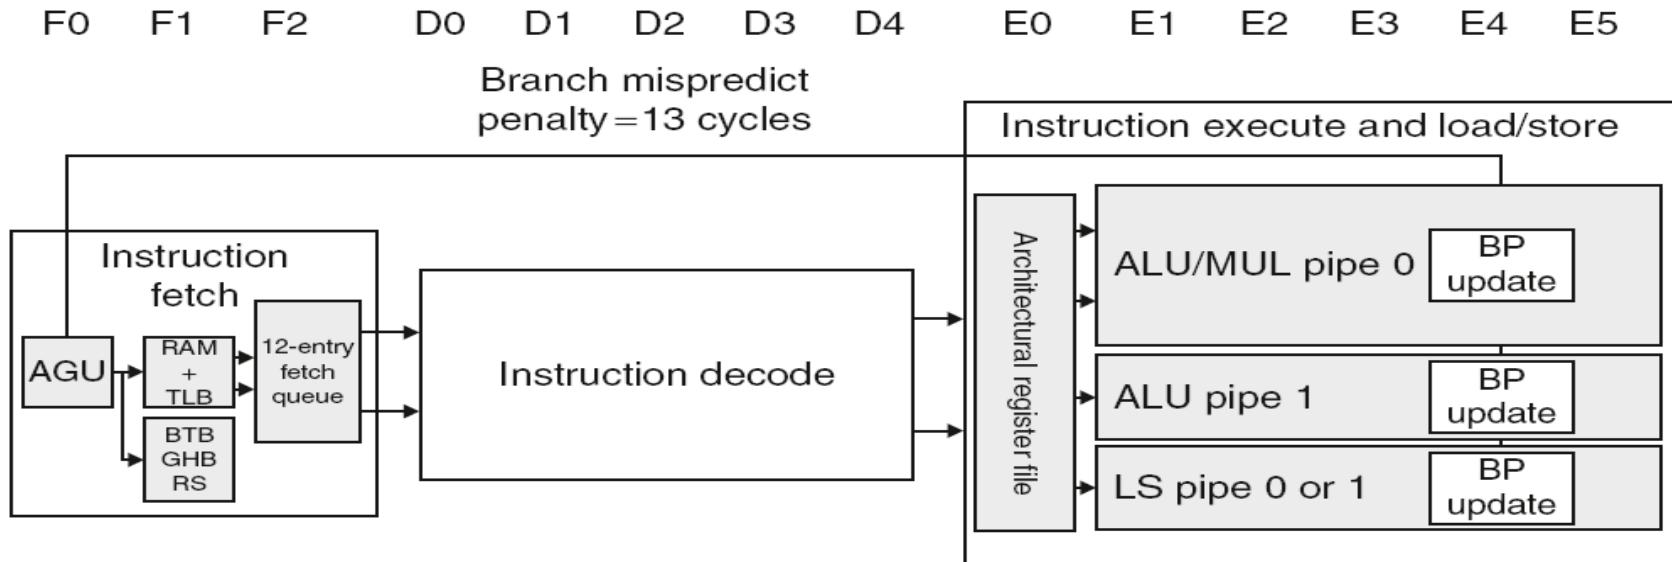
| Microprocessor | Year | Clock Rate | Pipeline Stages | Issue width | Out-of-order/Speculation | Cores | Power |
|----------------|------|------------|-----------------|-------------|--------------------------|-------|-------|
| i486           | 1989 | 25MHz      | 5               | 1           | No                       | 1     | 5W    |
| Pentium        | 1993 | 66MHz      | 5               | 2           | No                       | 1     | 10W   |
| Pentium Pro    | 1997 | 200MHz     | 10              | 3           | Yes                      | 1     | 29W   |
| P4 Willamette  | 2001 | 2000MHz    | 22              | 3           | Yes                      | 1     | 75W   |
| P4 Prescott    | 2004 | 3600MHz    | 31              | 3           | Yes                      | 1     | 103W  |
| Core           | 2006 | 2930MHz    | 14              | 4           | Yes                      | 2     | 75W   |
| UltraSparc III | 2003 | 1950MHz    | 14              | 4           | No                       | 1     | 90W   |
| UltraSparc T1  | 2005 | 1200MHz    | 6               | 1           | No                       | 8     | 70W   |

# Cortex A8 and Intel i7

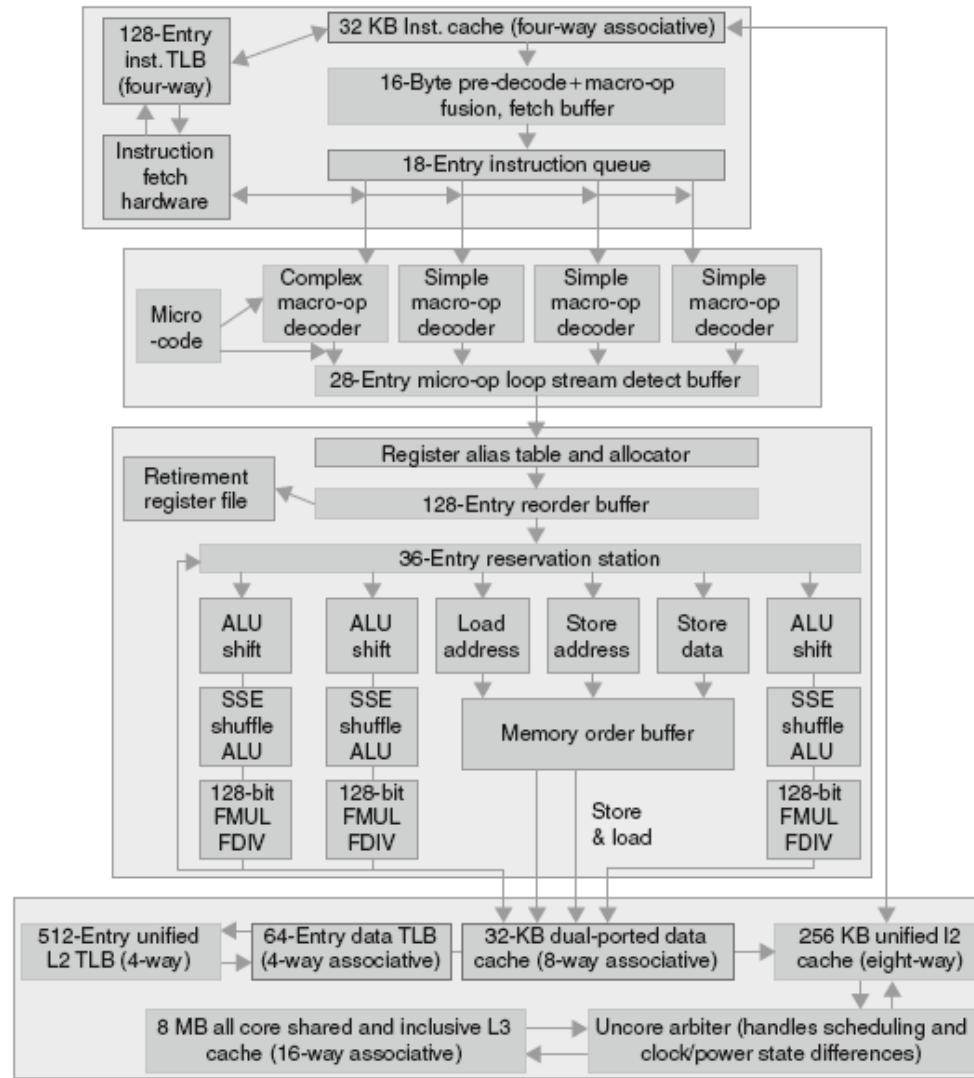
| Processor                             | ARM A8                 | Intel Core i7 920                     |
|---------------------------------------|------------------------|---------------------------------------|
| Market                                | Personal Mobile Device | Server, cloud                         |
| Thermal design power                  | 2 Watts                | 130 Watts                             |
| Clock rate                            | 1 GHz                  | 2.66 GHz                              |
| Cores/Chip                            | 1                      | 4                                     |
| Floating point?                       | No                     | Yes                                   |
| Multiple issue?                       | Dynamic                | Dynamic                               |
| Peak instructions/clock cycle         | 2                      | 4                                     |
| Pipeline stages                       | 14                     | 14                                    |
| Pipeline schedule                     | Static in-order        | Dynamic out-of-order with speculation |
| Branch prediction                     | 2-level                | 2-level                               |
| 1 <sup>st</sup> level caches/core     | 32 KiB I, 32 KiB D     | 32 KiB I, 32 KiB D                    |
| 2 <sup>nd</sup> level caches/core     | 128-1024 KiB           | 256 KiB                               |
| 3 <sup>rd</sup> level caches (shared) | -                      | 2- 8 MB                               |



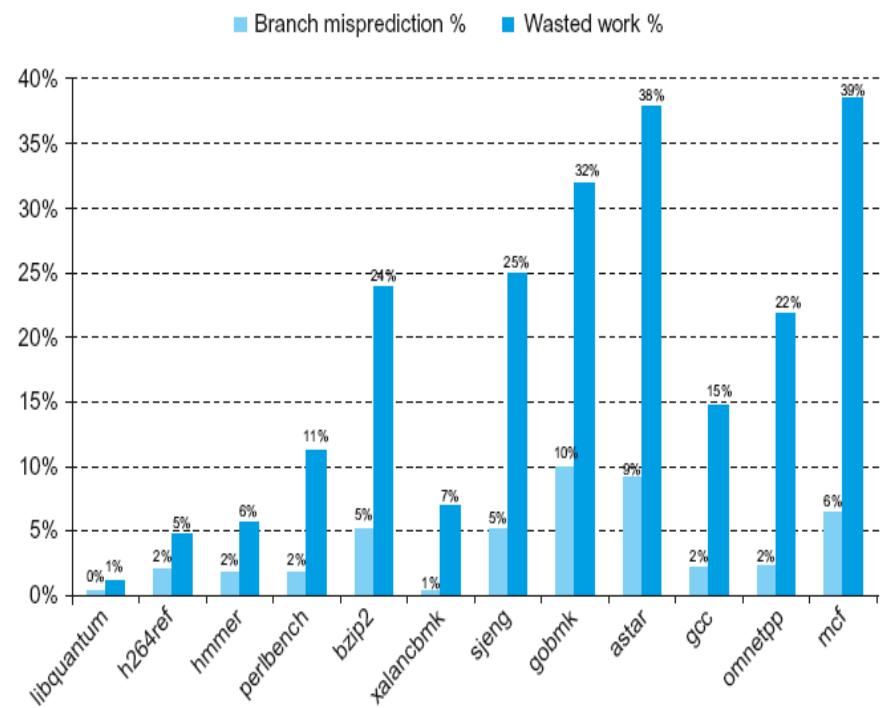
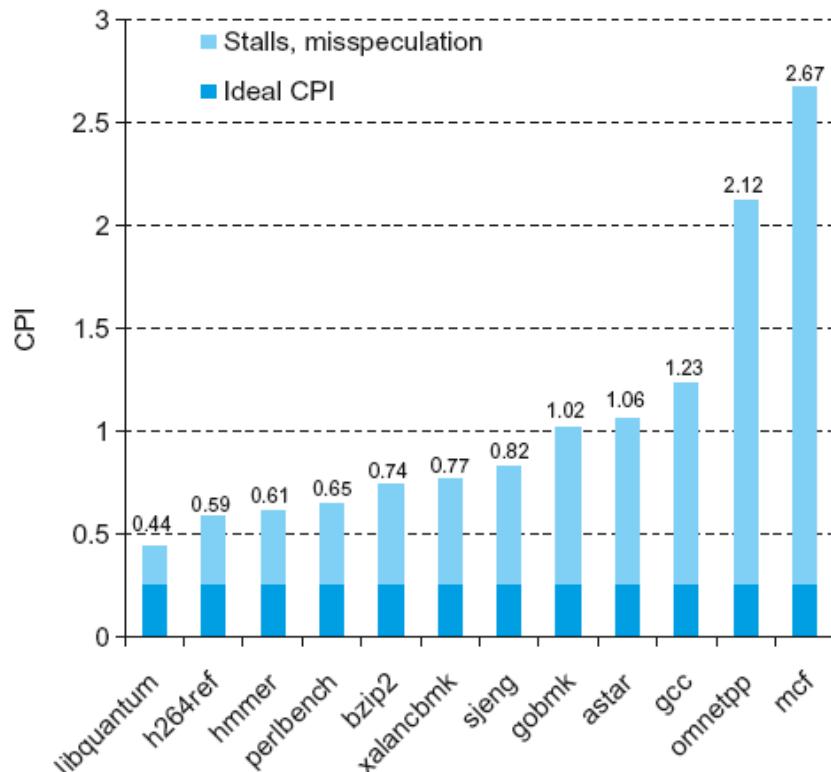
# ARM Cortex-A8 Pipeline & Performance



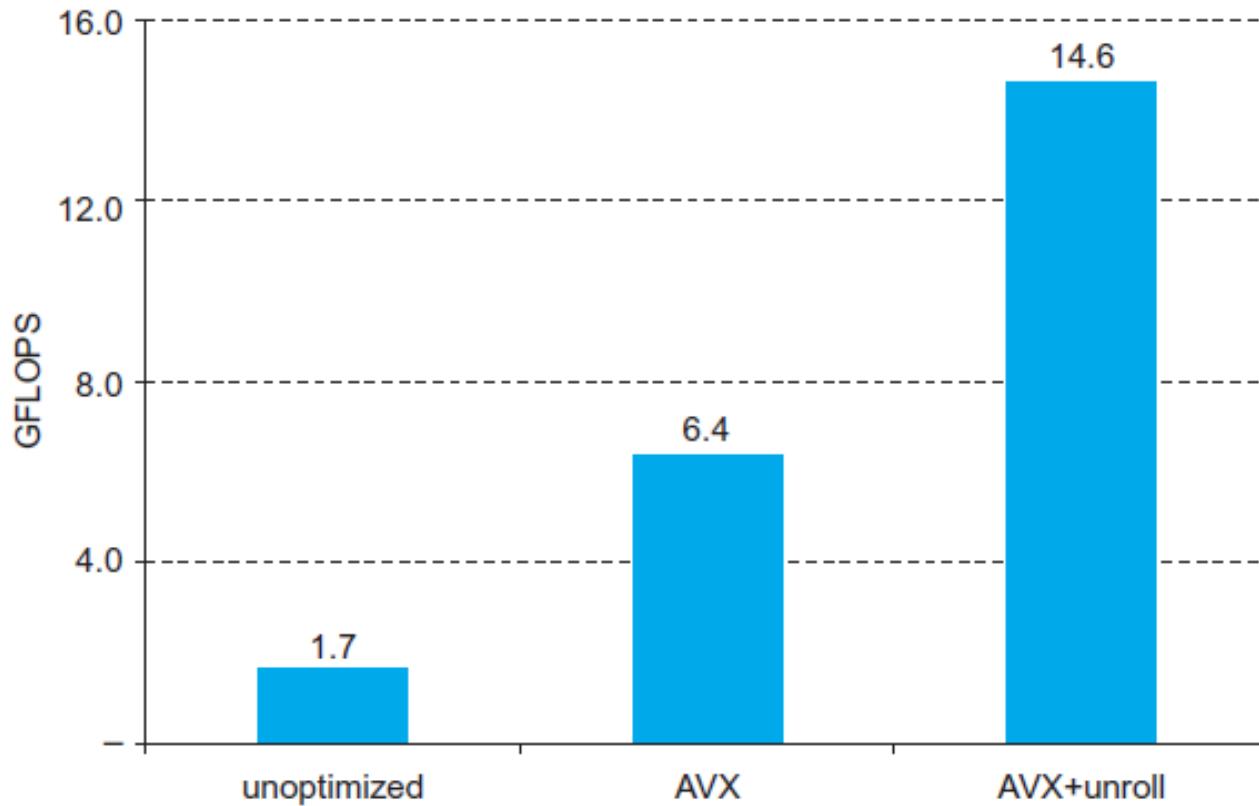
# Core i7 Pipeline



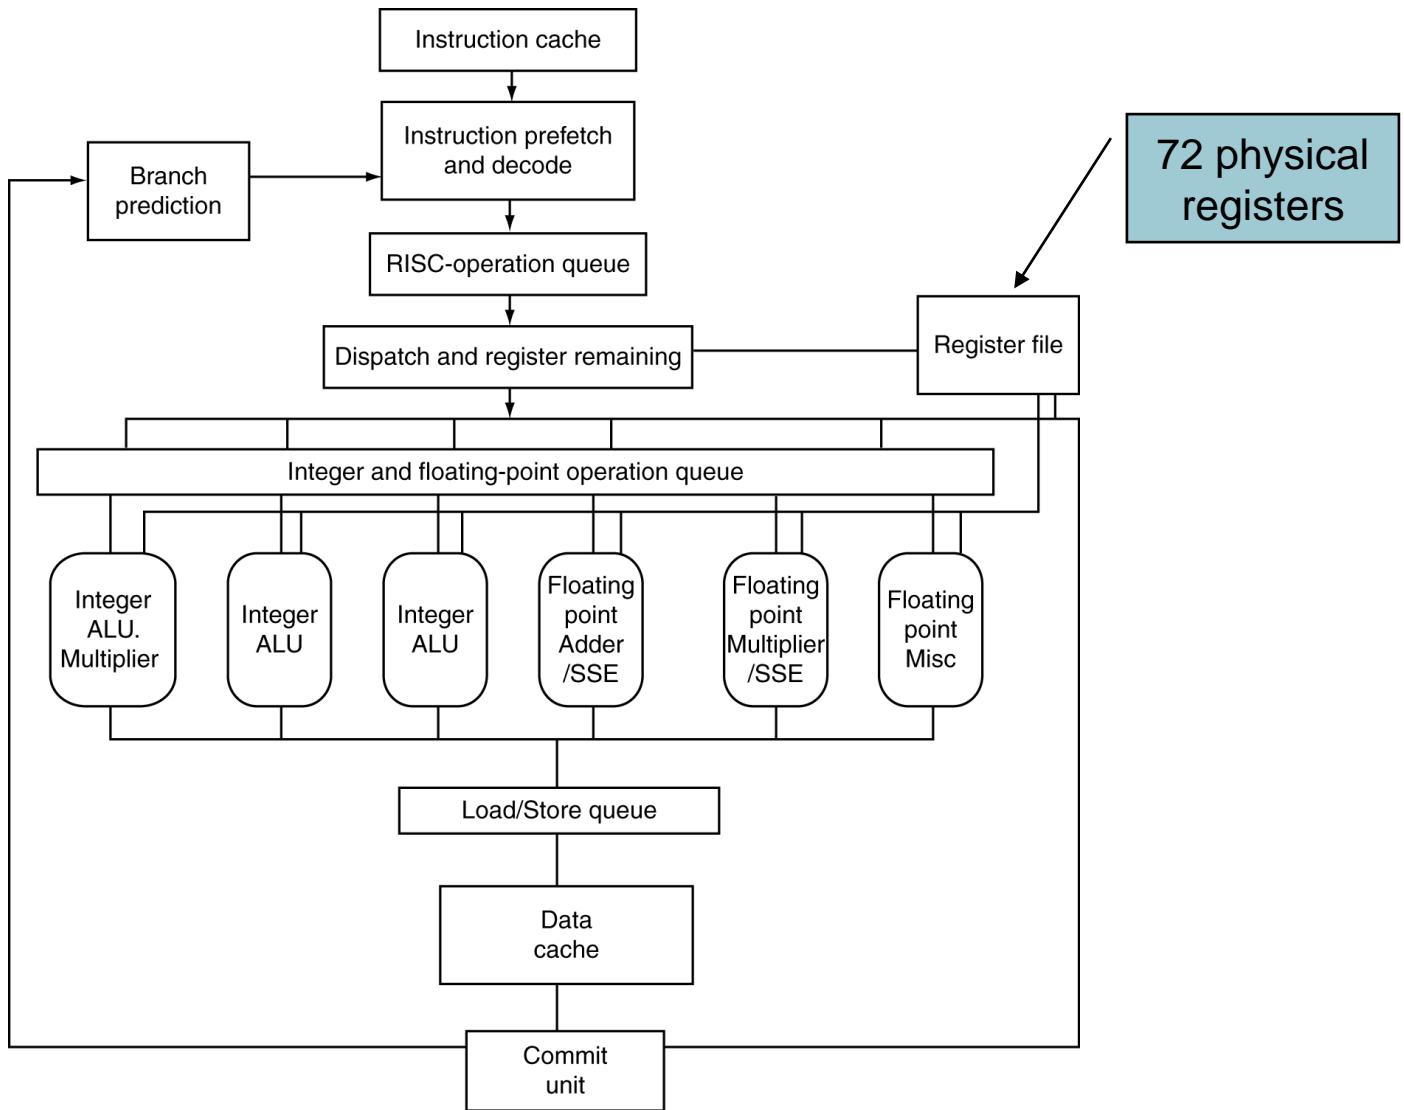
# Core i7 Performance



# Performance Impact

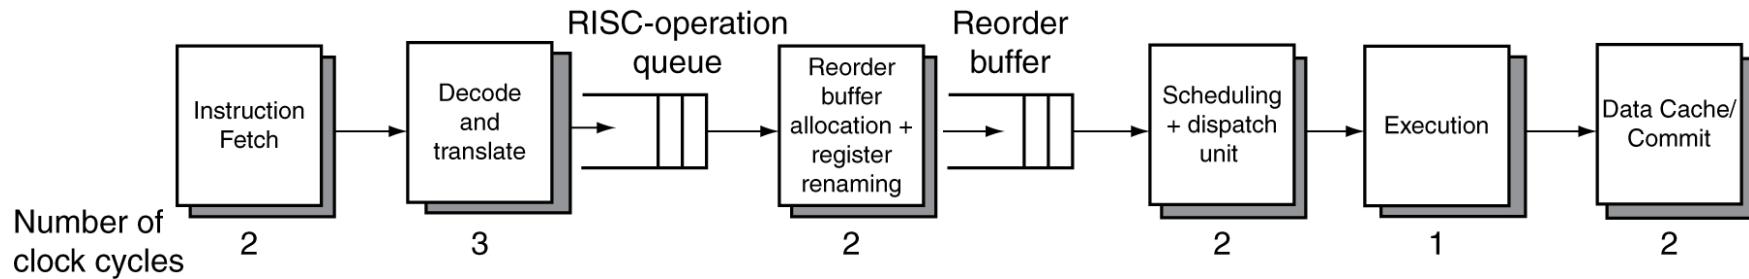


# The Opteron X4 Microarchitecture



# The Opteron X4 Pipeline Flow

## For integer operations



- FP is 5 stages longer
- Up to 106 RISC-ops in progress

## Bottlenecks

- Complex instructions with long dependencies
- Branch mispredictions
- Memory access delays

# Fallacies

- Pipelining is easy (!)
  - The basic idea is easy
  - The devil is in the details
    - e.g., detecting data hazards
- Pipelining is independent of technology
  - So why haven't we always done pipelining?
  - More transistors make more advanced techniques feasible
  - Pipeline-related ISA design needs to take account of technology trends
    - e.g., predicated instructions

# Pitfalls

- Poor ISA design can make pipelining harder
  - e.g., complex instruction sets (VAX, IA-32)
    - Significant overhead to make pipelining work
    - IA-32 micro-op approach
  - e.g., complex addressing modes
    - Register update side effects, memory indirection
  - e.g., delayed branches
    - Advanced pipelines have long delay slots

# Concluding Remarks

- ISA influences design of datapath and control
- Datapath and control influence design of ISA
- Pipelining improves instruction throughput using parallelism
  - More instructions completed per second
  - Latency for each instruction not reduced
- Hazards: structural, data, control
- Multiple issue and dynamic scheduling (ILP)
  - Dependencies limit achievable parallelism
  - Complexity leads to the power wall